


Version 2.14

UCSD CSE 30

Computer Organization and Systems Programming

C Part 2

Keith Muller



Review: Binary Numbering

- Binary is base 2
 - *adjective*: being in a **state of one of two mutually exclusive conditions** such as **on** or **off**, **true** or **false**, **molten** or **frozen**, **presence** or **absence** of a signal
 - From Late Latin *bīnārius* (“consisting of two”)
- **Two** symbols:
0 1
- Numbers in C that start with **0b** are binary
- Example: What is **0b110** in base 10?
 - $0b110 = 110_2 = (1 \times 2^2) + (1 \times 2^1) + (0 \times 2^0) = 6_{10}$
- **A bit** is a **single binary digit**
- **A byte** is an **8-bit** value

powers of two



$$\text{Unsigned binary Number} = \sum_{i=0}^{i=n-1} b_i \times 2^i = b_{n-1}2^{N-1} + b_{n-2}2^{N-2} + \dots + b_12^1 + b_02^0$$

Review: Hexadecimal Numbering

- hexadecimal is base 16
 - From “hexa” (Ancient Greek ἑξά-) \Rightarrow six
 - and from “decem” (Latin) \Rightarrow ten

- Sixteen** symbols

0 1 2 3 4 5 6 7 8 9 a b c d e f



- Numbers in C that start with **0x** are hexadecimal numbers
 - 16**₁₀ = **0x**10₁₆
- Example: What is **0xa5** in base 10?
 - 0xa5** = **a5**₁₆ = (**10** \times 16¹) + (**5** \times 16⁰) = 165₁₀
- Hexadecimal numbers are **very commonly used** in programming to express binary values
 - Imagine the difficulty in correctly expressing a 64-bit binary value in your code

$$\text{Unsigned Hex Number} = \sum_{i=0}^{n-1} b_i \times 16^i = b_{n-1}16^{n-1} + b_{n-2}16^{n-2} + \dots + b_116^1 + b_016^0$$

Binary <---> Hexadecimal Equivalences

- Hex → Binary: $16^1 = 2^4$ 1 digit hex = 4 digits binary
 1. Replace hex digits with binary digits
 2. Drop **leading zeros**
 - Example: 0x2d to binary
 - 0x2 is 0b0010, 0xd is 0b1101
 - Drop two leading zeros, answer is 0b101101
- Binary → Hex: $2^4 = 16^1$
 1. **Pad** with enough **leading zeros** until number of digits is a multiple of 4
 2. **Replace** each **group of 4** with the **HEX equivalent**
 - Example: 0b101101
 - **Pad on the left** to: 0b 0010 1101
 - Replace to get: 0x2d

Number Base Overview (as written in C)

- Decimal is base 10 and Hexadecimal is base 16,
- **Hex digits** have 16 values 0 - 9 a - f (written in C as 0x0 – 0xf)
- No standard prefix in C for binary (most use **hex**)
 - gcc (compiler) allows **0b** prefix **others might not**

Hex digit	0x0	0x1	0x2	0x3	0x4	0x5	0x6	0x7
Decimal value	0	1	2	3	4	5	6	7
Binary value	0b0000	0b0001	0b0010	0b0011	0b0100	0b0101	0b0110	0b0111

Hex digit	0x8	0x9	0xa	0xb	0xc	0xd	0xe	0xf
Decimal value	8	9	10	11	12	13	14	15
Binary value	0b1000	0b1001	0b1010	0b1011	0b1100	0b1101	0b1110	0b1111

Hex to Binary (group 4 bits per digit from the right)

- Each Hex digit is 4 bits in base 2 $16^1 = 2^4$

0x f a 5 3

1111 1010 0101 0011

0b1111101001010011

↑ binary start with a 0b in C

Binary to Hex (group 4 bits per digit from the right)

- 4 binary bits is one Hex digit $2^4 = 16^1$

0b 0110 1010 0011 1111

6 a 3 f

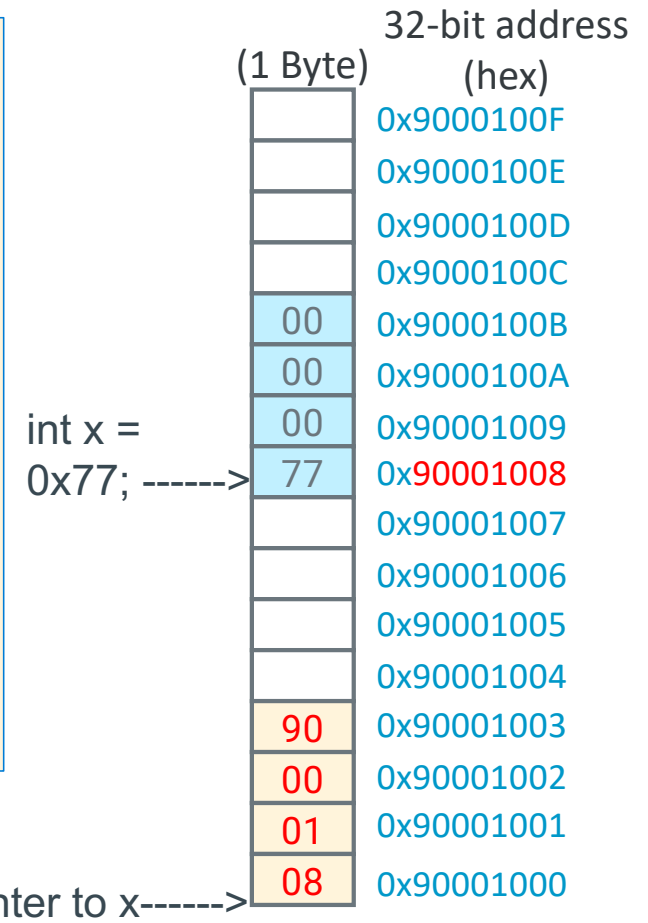
0x6a3f

hex start with 0x in C



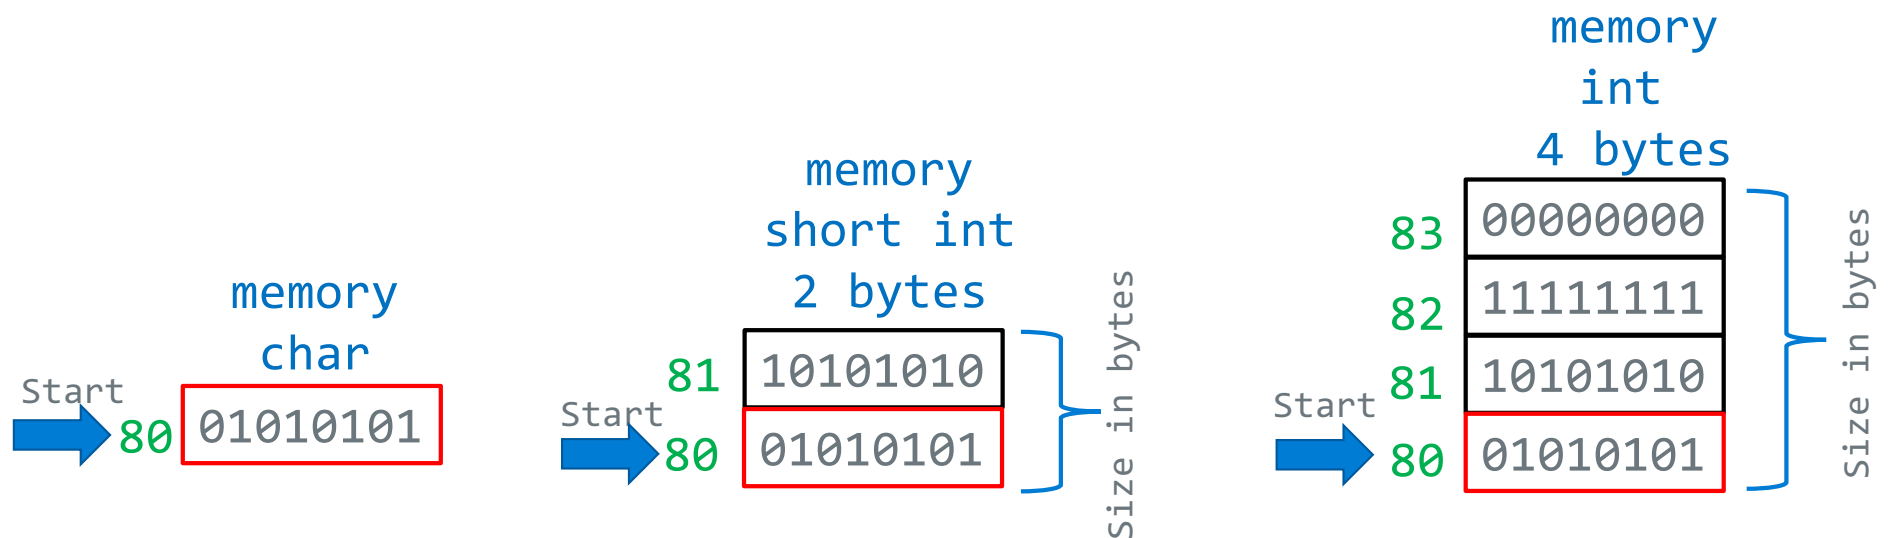
Address and Pointers

- An **address** refers to a location in memory, the **lowest** or **first byte** in a **contiguous sequence of bytes**
- A **pointer** is a **variable** whose **contents** (or value) can be properly used as an **address**
 - The **value in a pointer** *should* be a **valid address allocated to the process by the operating system**
- The **variable x** is at **memory address 0x90001008**
- The **variable pt** is at **memory location 0x90001000**
- The **contents** of **pt** is the **address of x 0x90001008**



Variables in Memory: Size and Address

- The **number of contiguous bytes** a variable uses is based on the *type* of the variable
 - Different **variable types** require different numbers of **contiguous bytes**
- **Variable names** map to a starting address in memory
- **Example Below:** Variables all starting at address 0x80, each box is a byte



Variables: Size

- Integer types

- `char`, `int`

- Floating Point

- `float`, `double`

- Modifiers for each base type

- `short` [int]
- `long` [int, double]
- `signed` [char, int]
- `unsigned` [char, int]
- `const`: variable read only

- `char` type

- One byte in a byte addressable memory
- **Signed** vs **Unsigned** Char implementations
- **Be careful** `char` is unsigned on arm and signed on other HW like intel

C Data Type	AArch-32 contiguous Bytes	AArch-64 contiguous Bytes	printf specification
<code>char</code> (arm unsigned)	1	1	%c
<code>short int</code>	2	2	%hd
<code>unsigned short int</code>	2	2	%hu
<code>int</code>	4	4	%d / %i
<code>unsigned int</code>	4	4	%u
<code>long int</code>	4	8	%ld
<code>long long int</code>	8	8	%lld
<code>float</code>	4	4	%f
<code>double</code>	8	8	%lf
<code>long double</code>	8	16	%Lf
<code>pointer *</code>	4	8	%p

size of a pointer is the word size

sizeof(): Variable Size (number of bytes) Operator

```
#include <stddef.h>
/* size_t type may vary by system but is always unsigned */
```

sizeof() operator returns a value of type **size_t**:

the number of bytes used to store a variable or variable type

```
size_t size = sizeof(variable_type);
```

or

```
size_t size = sizeof(variable_name); // preferred!
```

- The argument to sizeof() is often an expression:

```
size = sizeof(int * 10);
```

- reads as:

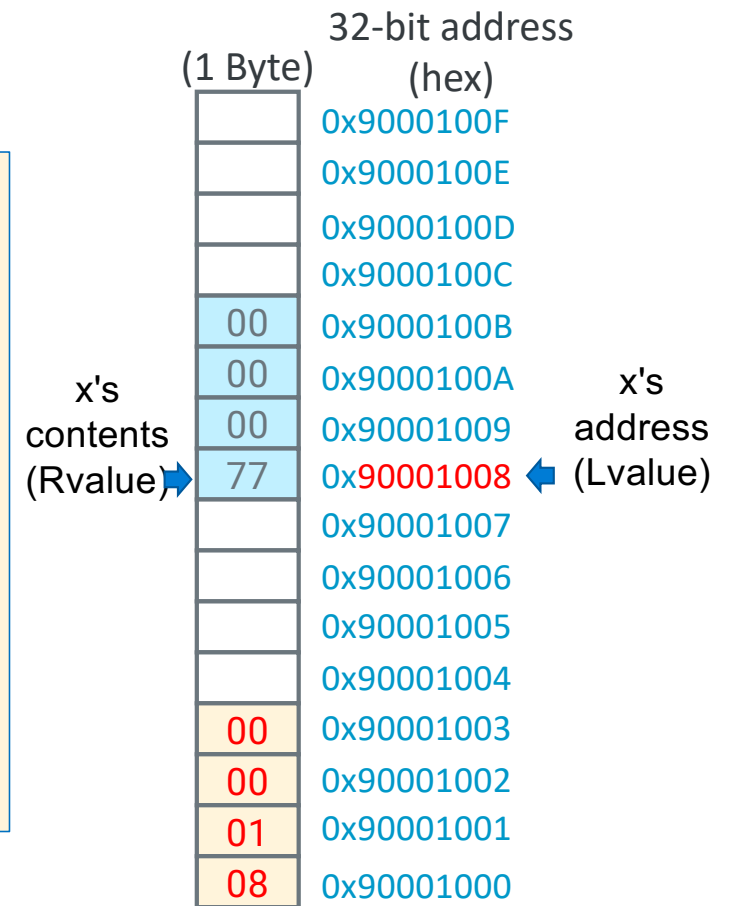
- number of bytes required to store **10 integers (an array of [10])**

Memory Addresses & Memory Content

Variable names in a C statement evaluation

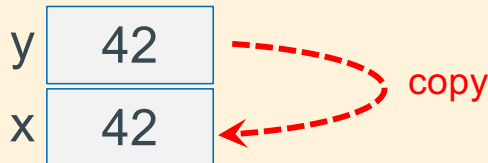
```
x = x + 1;    // Lvalue = Rvalue
```

- **Lvalue:** when on the left side (Lside or Left value) of the = sign
 - address where it is stored in memory – a constant
 - **Address assigned to a variable cannot be changed at runtime**
 - Does not require a memory read
 - **Lside Must evaluate to an address**
- **Rvalue:** when on the right side (Rside or Right value) of an = sign
 - contents or value stored in the variable (at its memory address)
 - requires a memory read to obtain contents



Memory Addresses & Memory Content

`y = 42;` One memory write required
`x = y;` One memory read required
 // Lvalue = Rvalue



- `x` on left side (**Lside**) of the assignment operator = evaluates to:
 - **Address** of the memory assigned to the `x` – this is `x`'s **Lvalue**
- `y` on right side (**Rside**) of the assignment operator = evaluates to:
 - **Contents** of the memory assigned to the variable `y` (type determines length – number of bytes) - this is `y`'s **Rvalue**
- So, `x = y;` is:

Read memory at `y` (**Rvalue**); write it to memory at `x`'s address (**Lvalue**)

Introduction: Address Operator: &

- Unary **address operator** (&) produces the **address** of where an **identifier** is in memory
 - Print assigned address to **g**
- **Example** this might print:
value of g is: 42
address of g is: 0x71a0a0
(the address will vary)
- **Tip:** printf() format specifier to display an address/pointer (in hex) is "%p"

```
int main(void)
{
    int g = 42;

    printf("value of g is: %d\n", g);
    printf("address of g is: %p\n", &g);
    return EXIT_SUCCESS;
}
```


Introduction: Address Operator: &

- Requirement: **identifier must have a Lvalue**
 - Cannot be used with **constants** (e.g., 12) or **expressions** (e.g., x + y)
 - Example: **&12** does not have an *Lvalue*,
 - so, **&12** is not a legal expression
- How can I get an **address for use on the Rside**?
 - **&var** (any variable identifier or name)
 - **function_name** (name of a **function**, not func());
 - **&func_name** is equivalent
 - **array_name** (name of the **array** like array_name[5]);
 - **&array_name** is equivalent

Pointer Variables

- In C, there is a *variable type* for **storing an address**: a *pointer*
 - **Contents** of a pointer is an unsigned (positive numbers) memory address

```
type *name; // defines a pointer; name contains address of a variable of type
```

- A *pointer* is defined by placing a *star* (or *asterisk*) (*) before the identifier (name)
- You also must specify the *type of variable* to which the pointer points
- **Pointers are typed!** Why?
 - The compiler needs to know the *size* (`sizeof()`) of the data **you are pointing at** (number of consecutive bytes to access) to use (dereference) the pointer
- When the **Rside** of a **variable** contains a **memory address**, (it **evaluates** to an **address**) the variable is called a **pointer variable**

Pointer Variables - 2

- A pointer cannot point at itself, why?

```
int *p = &p; /* is not legal - type mismatch */
```

- `p` is defined as `(int *)`, a pointer to an int, but
- the type of `&p` is `(int **)`, a pointer to a pointer to an int
- Pointer variables all use the **same amount of memory** no matter what they point at (in all but very tiny, often old design, cpu's)

```
int *iptr;  
char *cptr;  
  
printf("iptr(%u) cptr(%u)\n", sizeof(iptr), sizeof(cptr));
```

- Above prints on a 32-raspberry pi

```
% ./example  
iptr(4) cptr(4)
```

Defining Pointer Variables

- Assigning a value to a pointer:

```
int *p = &i;  /* p points at i (assign address i to p) */
```

- Is the same as writing the following definition and assignment statements

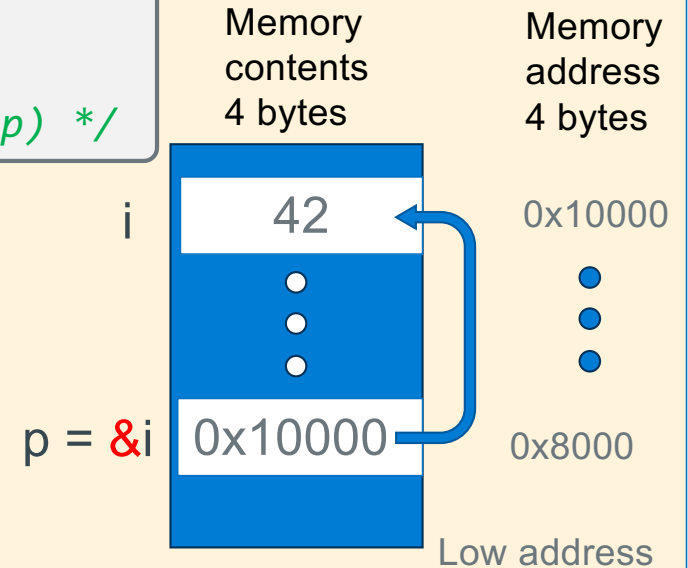
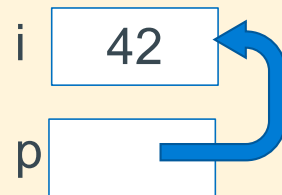
```
int *p;      /* p is defined (not initialized) */  
p = &i;      /* p points at i (assign address of i to p) */
```

- The ***** is part of the definition of **p** and is not part of the variable name
 - The name of the variable is simply **p**, not ***p**
- C mostly ignores whitespace, so these three definitions are equivalent

```
int  *p = &i;      /* Style A */  
int * p = &i;      /* Style B */  
int*  p = &i;      /* Style C */
```

Using Pointer Variables and the Address Operator & - 1

```
int i = 42;  
int *p; /* p contains the address of an integer */  
p = &i; /* p "points at" i (assign address of i to p) */
```



- **Warning:** be careful when defining multiple pointers on the same line:

`int *p1, p2;` is not the same as: `int *p1, *p2;`

Some find this clearer instead:


```
int *p1;  
int *p2;
```

Using Pointer Variables and the Address Operator & - 2

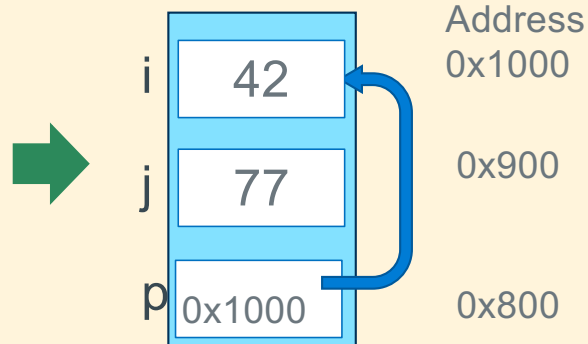
- As with any variable, you can change a pointers value (contents)

```
p = &j;      /* p now points at j */
```

See that p's value (contents) is the address of i



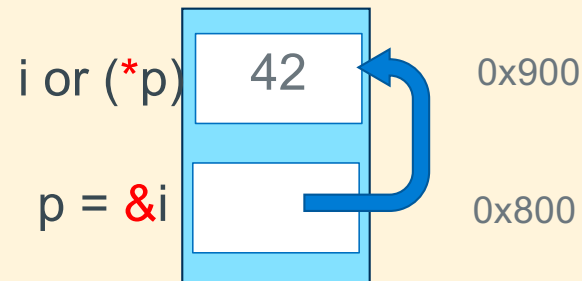
```
p = &i;      /* p now points at i */
```



Indirection (or dereference) Operator: *

- The **indirection operator** (*) or the *dereference operator to a variable* is the **inverse** of the *address operator* (&)
- **address operator** (&) can be thought of as:

"get the address of this box"




- **indirection operator** (*) can be thought of as:
"follow the arrow to the next box and get its contents"
- **Indirection operator causes an additional read to occur**, when on either the Rside or Lside of a statement

Rside Indirection (or dereference) Operator: *

- Performs the following steps when the * is on the Rside:
 1. read the contents of the variable to get an address
 2. read and return the contents at that address
 - (requires two reads of memory on the Rside)

```
z = *x; // copy the contents of memory pointed at by x to z
```



Two reads here
(1) read to get an address
(2) read the address to get the value

Rside Indirection (or dereference) Operator: *

*Contents of **p** is the **address** of **i**
(**p** points at **i**)*

```
int i = 42;  
int *p;
```


```
p = &i;
```

No reads here

```
printf("*p is %d\n", *p);
```

```
% ./a.out  
*p is 42
```

Two reads here
(1) read to get an address
(2) read the address to get
the value

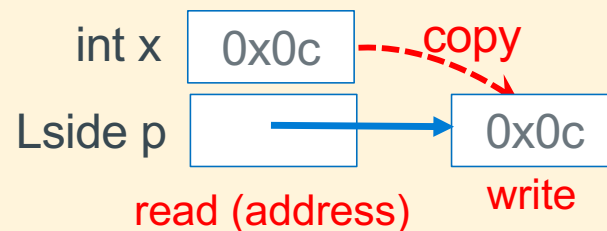


Lside Indirection Operator

Performs the following steps when the ***** is on the Lside:

1. **read** the **contents** of the **variable** to get **an address**
2. **write** the evaluation of the Rside expression to that address
 - (requires **one read of memory and one write of memory on the Lside**)

```
*p = x; // copy the value of x to the memory pointed at by p
```



Left Side Indirection (or dereference) Operator: *

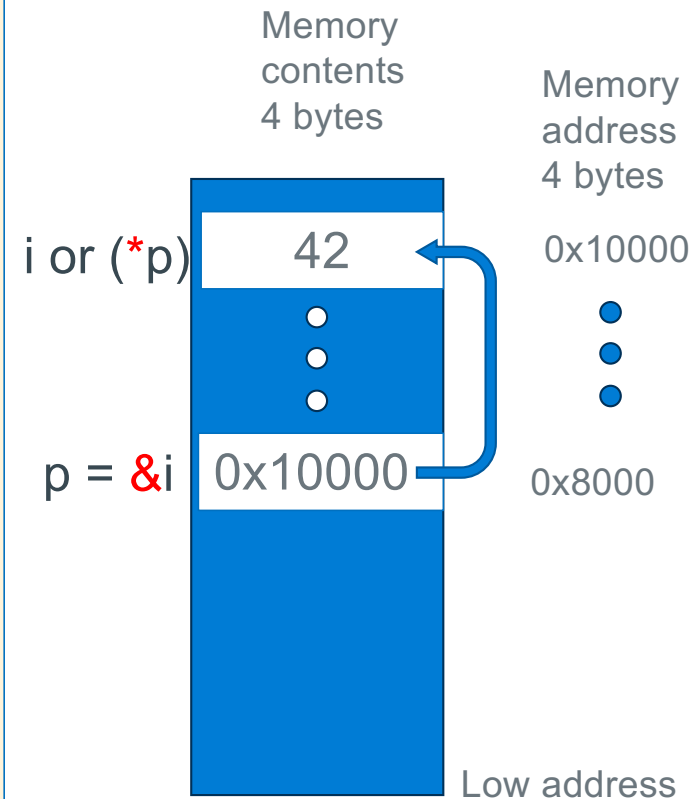
*Contents of **p** is the address of **i**
(**p** points at **i**)*

```
int i;  
int *p;  
  
p = &i;  
*p = 42;  
  
printf("*p is %d\n", *p);
```

No reads here
one write

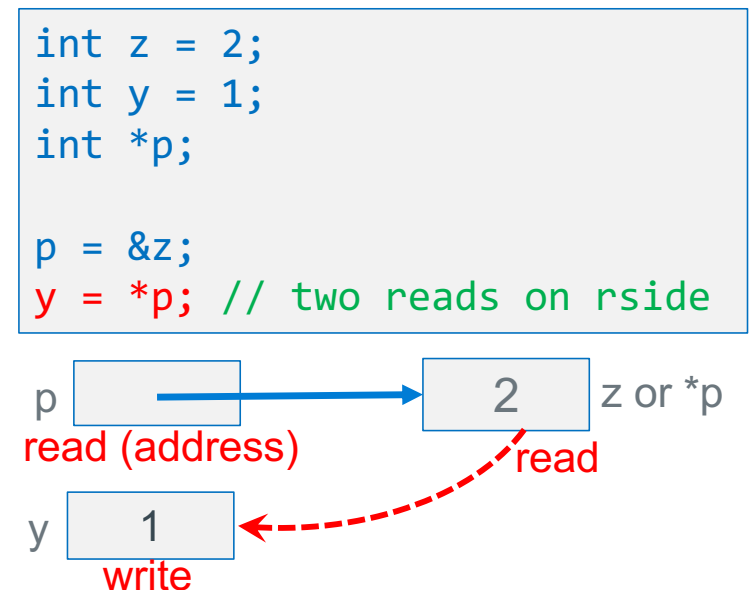
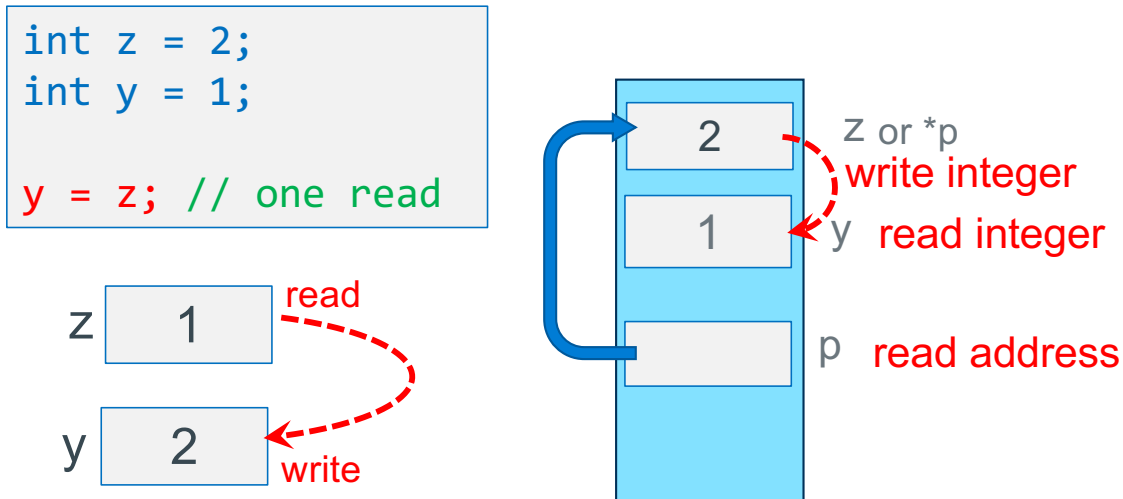
One read here
(1) read to get
an address
(2) one write

```
% ./a.out  
*p is 42
```



Each use of a * operator results in one additional read: Rside

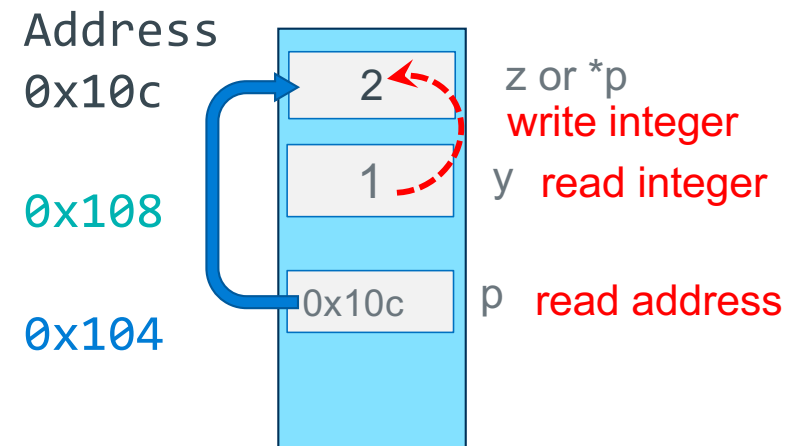
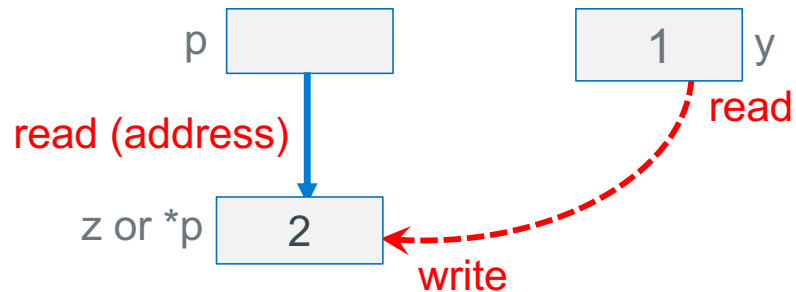
RULE: Each * when used as a dereference operator in a **statement** (either **Lside** or **Rside**) it causes an additional read to be performed



Aside: `y = *(&z);` // same as `y = z`

Each use of a * operator results in one additional read: Lside

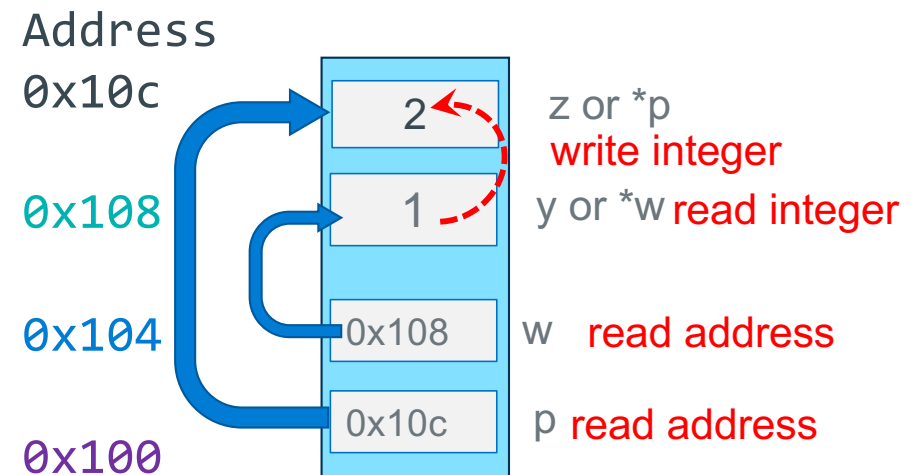

```
int z = 2;  
int y = 1;  
int *p;  
  
p = &z;  
*p = y;    // one read on lside
```



Each use of a * operator results in one additional read : both sides

```
int z = 2;  
int y = 1;  
int *w;  
int *p;
```

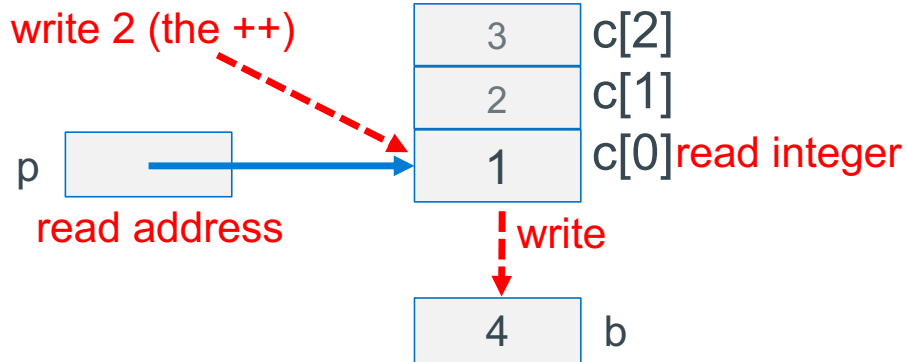
```
p = &z;  
w = &y;  
*p = *w;
```



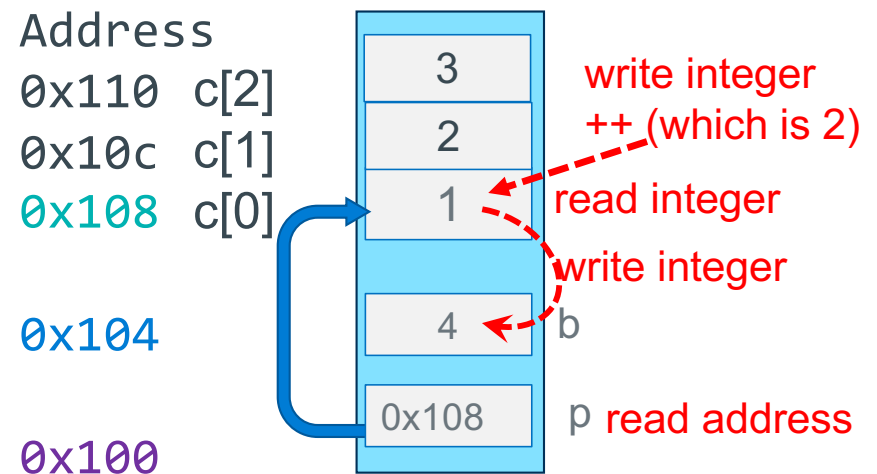
Each use of a * operator results in one additional read : both sides

```
int c[] = {1, 2, 3};  
int b = 4;  
int *p;
```

```
p = c;  
b = (*p)++;
```




2 reads and 2 writes




Each use of a * operator results in one additional read : both sides

```
int c[] = {1, 2, 3};  
int b = 4;  
int *p;
```

```
p = c;  
*(++p) = b;
```



2 reads and 2 writes




Pointer to Pointers (Double Indirection)

- Define a pointer to a pointer (p2 below)


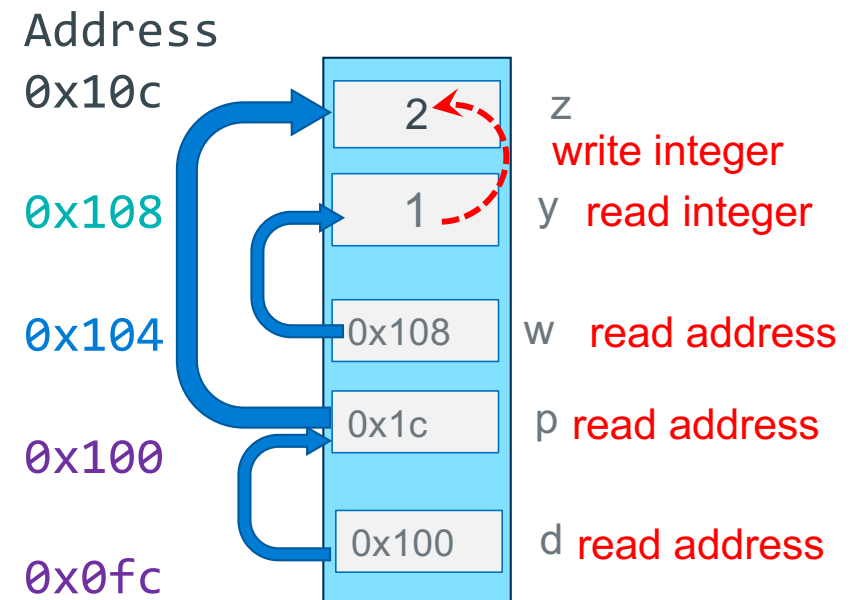
```
int i = 2;  
int *p1;  
int **p2; // pointer to a pointer to an int  
  
p1 = &i;  
p2 = &p1;  
printf("%d\n", (**p2) * (**p2));
```

- C allows any number of pointer indirections
 - more than two levels is very uncommon in real applications as it reduces readability and generates a lot of memory reads
- RULE (important):** number of ***** in the **variable definition** tells you **how many reads** it takes to get to the **base type**
 $\text{\#reads to base type} = \text{number of } * \text{ (in the definition)} + 1$
- Example:
`int **p2;` // requires 3 reads to get to the int



Double Indirection: Lside

```
int z = 2;  
int y = 1;  
int *w;  
int *p;  
int **d;  
  
p = &z;  
w = &y;  
d = &p;  
**d = *w;
```

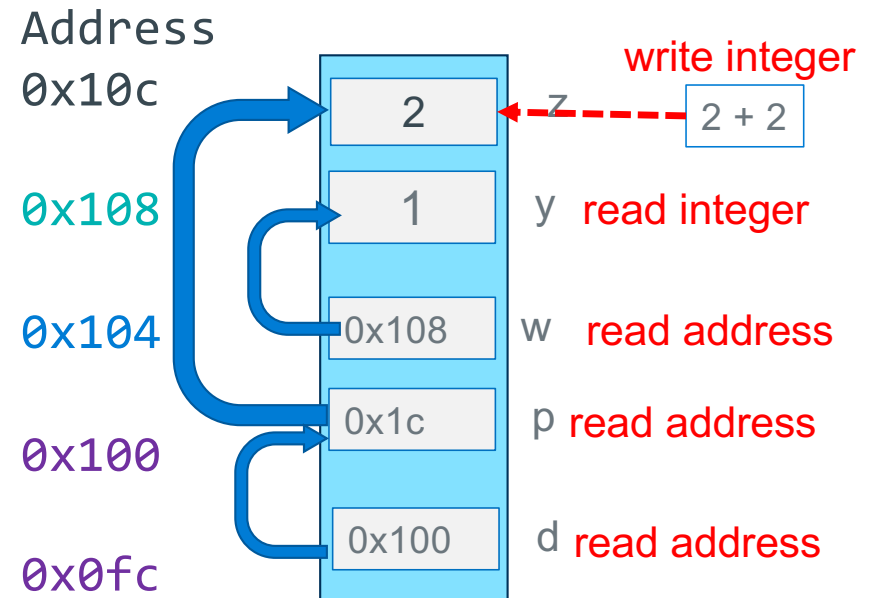
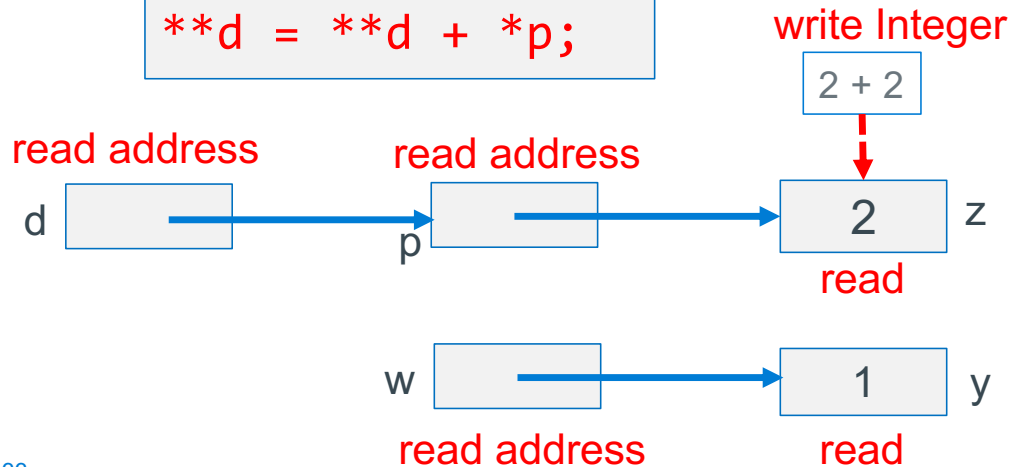


Double Indirection: Rside

```

int z = 2;
int y = 1;
int *w;
int *p;
int **d;

p = &z;
w = &y;
d = &p;
**d = **d + *p;
    
```



Important Observation


****d on Lside is two reads**
****d on Rside is three reads**

What is Aliasing?

- **Two or more** variables are **aliases** of each other when they all **reference the same memory** (so different names, same memory location)
- **Example:** When one pointer is copied to another pointer it **creates an alias**
- **Side effect:** Changing one variables value (content) changes the value for other variables
 - **Multiple variables** all **read and write** the **same** memory location
 - Aliases occur either by **accident** (coding errors) or **deliberate** (**careful: readability**)

```
int i = 5;
int *p;
int *q;

p = &i;
q = p;    // *p & *q now aliases
*q = 4;   // changes i and *p
```



Result *p, *q and i all have the value of 4

Defining Arrays

Definition: `type name[count]`

- **"Compound"** data type where each value in an array is an element of `type`
- Allocates **name** with a *fixed* `count` array elements of type `type`
- Allocates (`count * sizeof(type)`) bytes of *contiguous memory*
- Common usage is to specify a compile-time constant for `count`


```
#define BSZ 6  
int b[BSZ];
```

BSZ is a macro replaced by the C preprocessor

- Array **names are constants** and **cannot be assigned** (the name cannot appear on the Lside by itself)

```
int a [BSZ];  
a = b;          // invalid does not copy the array  
                // must copy arrays element by element
```

```
int b[6];
```




Array Initialization

- Initialization: `type name[count] = {val0,...,valN};`
 - `{ }` (*optional*) initialization list can only be used at **time of definition**
 - If no `count` supplied, `count` is determined by compiler using the number of array initializers
- `int block[20] = {};` *//only works with constant size arrays*
 - defines an **array of 20 integers** each element filled with zeros
 - Performance comment: do not zero automatic arrays unless really needed!
- When a `count` is given:
 - extra initialization values** are **ignored**
 - missing initialization values** are set to **zero**

```
int block[5] = {2, 3, 5, 6, 11, 13};
```

not needed and if used **may** truncate initialization list

6 initialization values given, **only 5 are used**



Accessing Arrays Using Indexing


- **name** [**index**] selects the **index** element of the array
 - **index should be unsigned**
 - **Elements range from: 0 to count – 1** (`int x[count];`)
- **name** [**index**] can be used as an **assignment target** or as a **value in an expression**
- **Array name** (by itself with no []) on the **Rside** evaluates to the **address of the first element of the array**

```
int a[2] = {1, 2};  
a[0] = a[1];
```

```
int b[5];  
int *p = b;
```

p 9000

1 word
(int = 4 bytes)




How many elements are in an array?

- The number of elements of space allocated to an array (called **element count**) and indirectly the total size in bytes of an array **is not stored anywhere!!!!!!**
- An **array name** is just the **address of the first element in a block of contiguous memory**
 - So, an array does not know its own size!

```
#define SZ 6
int block[SZ];    // you specify the array has SZ elements
int indx;        // use when SZ is defined

for (indx = 0; indx < SZ; indx++)
    block[indx] = 0;
```

```
int b[6];
```



Determining Element Count: compile time calculation

- Programmatically determining the element count in a compiler calculated array
`sizeof(array) / sizeof(of just one element in the array)`
- `sizeof(array)` **only works** when used in the SAME **scope** where the array variable was defined

```
#include <stddef.h>
int main()
{
    int block[] =
        {2, 3, 5, 6, 11, 13};    // automatic: compiler calculates array size

    int cnt = (int)(sizeof(block) / sizeof(block[0]));    // in this case cnt = 6

    for (int indx = 0; indx < cnt; indx++)
        block[indx] = 0;
```


Pointers and Arrays - 1

- A few slides back we stated: **Array name** (by itself) on the **Rside** evaluates to the **address of the first element of the array**

```
int buf[] = {2, 3, 5, 6, 11};
```

- Array indexing syntax (`[]`) an operator that performs *pointer arithmetic*
- buf** and **&buf[0]** on the **Rside** are **equivalent**, **both evaluate** to the address of the first array element

```
int *p = buf;           // or int *p = &buf[0];  
int *p1 = &buf[1];  
int *p2 = &buf[2];  
int *p3 = &buf[3];
```




Pointers and Arrays - 2

- When `p` is a pointer, the actual evaluation of the address:
 - `(p+1)` depends on the base type the pointer `p` points at
- `(p+1)` adds `1 x sizeof(what p points at)` bytes to `p`
 - `++p` is equivalent to `p = p + 1`
- Using pointer arithmetic to find array elements:
 - Address of the second element `&buf[1]` is `(buf + 1)`
 - It can be referenced as `*(buf + 1)` or `buf[1]`


```
int buf[] = {2, 3, 5, 6, 11};
int *p;
p = buf;

*p = *p + 10; // {12, 3, 5, 6, 11}
*(p + 1) = *(buf + 1) + 10; // {12, 13, 5, 6, 11}
```

1 byte Memory Content
One byte per row



Pointer Arithmetic In Use – C's Performance Focus




- **Alert!:** C performance focus **does not** perform any array “bounds checking”
- **Performance by Design:** *bound checking **slows down execution** of a properly written program*
- **Example:** array **a** of length **i**, C **does not verify** that **a[j]** or ***(a + j)** is valid (does not check: $0 \leq j < i$)
 - C simply “*translates*” and accesses the memory specified from: **a[j]** to be ***(a + j)** which may be *outside the bounds* of the array
 - OS only ***“faults”*** for an incorrect access to memory (read-only or not assigned to your process)
 - It does not fault for out of bound indexes or out of scope
- **Lack of bound checking** is a **common source of errors and bugs** and is a common criticism of C

Pointer Arithmetic

- You cannot add two pointers (*what is the reason?*)
- A pointer *q* can be subtracted from another pointer *p* when the pointers are the same type – **best done only within arrays!**
- The value of $(p - q)$ is the number of **elements between** the two pointers
 - Using memory address arithmetic (*p* and *q* *R*side are both **byte addresses**):
 - Notice that it is `sizeof(*p)` below: it is what ***p* points at** and not **`sizeof(p)` which is the size of the pointer!**

distance in elements = $(p - q) / \text{sizeof}(*p)$

$$(p + 3) - p = 3 = (0x08c - 0x080) / 4 = 3$$




Pointer Comparisons

- Pointers (**same type**) can be compared with the comparison operators:

<, <=, ==, !=, >=, >

```
int numb[] = {9, 8, 1, 9, 5};
int *end;
int *a;
end = numb + (int) (sizeof(numb)/sizeof(*numb));
a = numb;
while (a < end) // compares two pointers (address)
    /* rest of code including doing an a++ */
```

size of what
numb points at!



- Invalid, Undefined, or **risky** pointer arithmetic (some examples)
 - Add, multiply, divide on two pointers
 - Subtract two pointers of different types or pointing at different arrays
 - Compare two pointers of different types
 - Subtract a pointer from an integer

Using Pointers to Traverse an array

```
int x[] = {0xd4c3b2a1, 0xd4c3b200, 0x12345684};
int cnt = (int)(sizeof(x) / sizeof(*x));

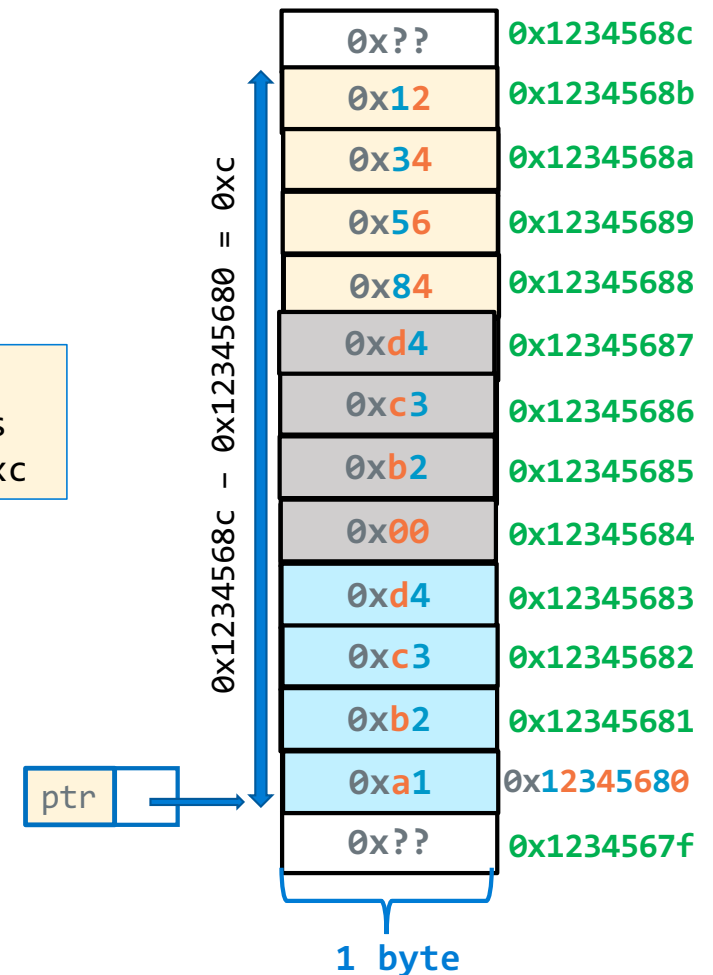
for (int j = 0; j < cnt; j++)
    printf("%#x\n", x[j]);
}
```

cnt = 3;
 actual space used by x is cnt * sizeof(*x); = 12 bytes
 Calculating by addresses: 0x1234568c - 0x12345680 = 0xc

```
int x[] = {0xd4c3b2a1, 0xd4c3b200, 0x12345684};
int cnt = (int)(sizeof(x) / sizeof(*x));
int *ptr = x;           // or &x[0]

for (int j = 0; j < cnt; j++)
    printf("%#x\n", *(ptr + j));
}
```

Brute force translation to pointers



Fast Ways to Traverse an Array: Use a Limit Pointer


```
int x[] = {0xd4c3b2a1, 0xd4c3b200, 0x12345684};
int cnt = (int)(sizeof(x) / sizeof(*x));
```

```
int *ptr;
int *xptr;
ptr = x;
xptr = ptr + cnt; //or &x[0]
```

xptr is a **loop limit pointer**
it points **1 element past**
the end of the array

```
while (ptr < xptr) {
    printf("%#x\n", *ptr);
    ptr++;
}
```

```
% ./a.out
0xd4c3b2a1
0xd4c3b200
0x12345684
```



C Precedence and Pointers

- ++ -- pre and post increment combined with pointers can create code that is complex, hard to read and difficult to maintain
- Use () to help readability

common	With Parentheses	Meaning
*p++	*(p++)	(1)The Rvalue is the object that p points at (2)increment pointer p to next element ++ is higher than *
(*p)++		(1)Rvalue is the object that p points at (2)increment the object
*++p	*(++p)	(1)Increment pointer p first to the next element (2)Rvalue is the object that the incremented pointer points at
++*p	++(*p)	Rvalue is the incremented value of the object that p points at

Operator	Description	Associativity
() [] . -> ++ --	Parentheses or function call Brackets or array subscript Dot or Member selection operator Arrow operator Postfix increment/decrement	left to right
++ -- + - ! ~ (type) * & sizeof	Prefix increment/decrement Unary plus and minus not operator and bitwise complement type cast Indirection or dereference operator Address of operator Determine size in bytes	right to left
* / %	Multiplication, division and modulus	left to right
+ -	Addition and subtraction	left to right
<< >>	Bitwise left shift and right shift	left to right
< <= > >=	relational less than/less than equal to relational greater than/greater than or equal to	left to right
== !=	Relational equal to or not equal to	left to right
&&	Bitwise AND	left to right
^	Bitwise exclusive OR	left to right
	Bitwise inclusive OR	left to right
&&	Logical AND	left to right
	Logical OR	left to right
? :	Ternary operator	right to left
= += -= *= /= %= &= ^= = <<= >>=	Assignment operator Addition/subtraction assignment Multiplication/division assignment Modulus and bitwise assignment Bitwise exclusive/inclusive OR assignment	right to left
,	comma operator	left to right

Example of a hard-to-understand pointer statement

```
int array[] = {2, 5, 7, 9, 11, 13};
int *ptr = array;
int x;
```

common	With Parentheses	Meaning
*p++	*(p++)	(1) The Rvalue is the object that p points at (2) increment pointer p to next element ++ is higher than *
(*p)++		(1) Rvalue is the object that p points at (2) increment the object
*++p	*(++p)	(1) Increment pointer p first to the next element (2) Rvalue is the object that the incremented pointer points at
++*p	++(*p)	Rvalue is the incremented value of the object that p points at

```
x = 1 + (*ptr++)++; // yuck!!
```

```
/* Same as the one line above */
```


```
x = 1 + *ptr;           // x = 1 + 2 = 3;
*ptr = *ptr + 1;        // (*ptr)++ is array[0]=2+1=3;
ptr = 1 + ptr;          // ptr = &array[1] = ptr points at 5
```

```
x = 1 + ++(*ptr++);
```

```
*ptr = *ptr + 1;        // (*ptr)++ is array[0]=2+1=3
x = 1 + *ptr;           // x = 1 + 3 = 4;
ptr = 1 + ptr;          // ptr = &array[1]; ptr -> 5
```



C Strings - 1

- C does not have a **dedicated type** for strings
- Strings are an **array of characters** terminated by a **sentinel termination character**
- `'\0'` is the **Null termination character**; has the **value of zero** (do not confuse with `'0'`)
- An **array of chars** contains **a string only when** it is terminated by a `'\0'`
- **Length of a string** is the number of characters in it, not including the `'\0'`
- Strings in C are **not** objects
 - **No embedded information about them**, you just have a **name** and a **memory location**
 - You **cannot use +** or **+=** to concatenate strings in C
 - For example, you must **calculate string length** using code at runtime looking for the sentinel



C Strings - 2

- First '`\0`' encountered from the start of the string always indicates the end of a string
- The '`\0`' **does not have to be** in the **last element in the space allocated to the array**
 - But String length is always less than the size of the array it is contained in
- In the example below, the array `buf` contains two strings (but only `cat` is seen as the string)
 - One string starts at `&(buf[0])` is `"cat"` with a string length of 3
 - The other string starts at `&(b[4])` is `"o"` with a string length of 1
 - `"o"` has two bytes: `'o'` and `\0`



Defining Strings: Initialization

- When you combine the automatic length definition for arrays with double quote(") **initialization**
 - Compiler automatically adds the null terminator '\0' for you

```
char a[4] = {'c', 'a', 't', '\0'};  
char b[] = "cat";  
char c[] = {'c', 'a', 't', '\0', 'a', 'b'};  
char empty[] = "";  
// compiler calculates size, adds '\0'  
// array size 6, first string length 3  
// empty string - contains '\0'  
// string length = 0
```

Background: Different Ways to Pass Parameters

- **Call-by-reference (or pass by reference)**

- Parameter in the called function is an alias (references the same memory location) for the supplied argument
- Modifying the parameter modifies the calling argument

Call-by-value (or pass by value) (C)


- What **Called** Function Does
 - Passed Parameters are used like local variables
 - Modifying the passed parameter in the function is allowed just like a local variable
 - So, writing to the parameter, only changes the copy
- The return value from a function in C is **by value**

Passing Parameters – Call by Value Example

```
int main(void)
{
    int x = 5;
    inc(x); // makes a copy of x
    printf("%d\n", x); // 5 or 6 ?
}

void inc(int i) // i is local to inc
{
    ++i;
}
```

if this was an expression like `inc(x+1)` it evaluates and stores the result in the memory allocated for the copy

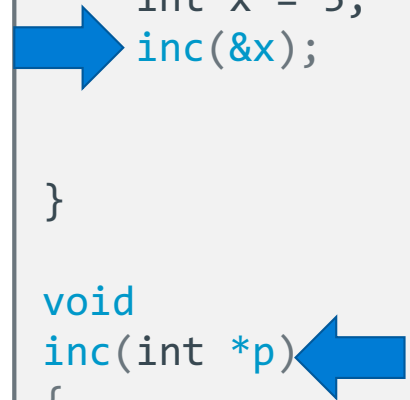


- when `inc(x)` is called, a copy of `x` is made to another memory location
 - `inc()` cannot change the variable `x` since `inc()` does not have the address of `x`, it is local to `main()` so, 5 is printed
- The `inc()` function is free to change its copy of the argument (just like any local variable) remember it does NOT change the parameter in `main()`

Output Parameters (Mimics Call by Reference)

- Passing a **pointer parameter** with the **intent** that the called function will use the address it to store values for use by the **calling function**, then **pointer parameter** is called an **output parameter**
- To pass the address of a variable `x` use the **address operator** (`&x`) or the contents of a pointer variable that points at `x`, or the name of an array (the arrays address)
- To be receive an address in the called function, define the **corresponding parameter type** to be a **pointer** (add `*`)
 - It is common to describe this method as: "pass a **pointer to x**"
- C is still using "*pass by value*"
 - we pass the **value** of the **address/pointer** in a **parameter copy**
 - **The called routine** uses the address to change a variable in the caller's scope

```
void inc(int *p);  
int  
main(void)  
{  
    int x = 5;  
    inc(&x);  
}  
  
void  
inc(int *p)  
{  
  
}
```



Example Using Output Parameters

```
void inc(int *p);  
int  
main(void)  
{  
    int x = 5;  
    inc(&x);  
    printf("%d\n", x);  
    return EXIT_SUCCESS;  
}
```

Pass the
address of x (&x)


```
void  
inc(int *p)  
{  
    if (p != NULL)  
        *p += 1;    // or (*p)++  
}
```

Receive an
address copy in
the variable p
(int *p)

Write to the output variable (*p)

At the Call to inc() in main()

1. Allocate space for p
2. Copy x's address into p



With a pointer to X,

inc() can change x in main()
this is called a side effect

p just like any other local variable

Array Parameters: Call-By-Value or Call-By-Reference?

- `Type []` array parameter is automatically “promoted” to a pointer of type `Type *`, and a copy of the *pointer* is *passed by value*

the name is the address, so this is passing a pointer to the start of the array

```
void passa(int []);
int main(void)
{
    int numbers[] = {9, 8, 1, 9, 5};

    passa(numbers);
    printf("numbers size:%lu\n", sizeof(numbers)); // 20
    return EXIT_SUCCESS;
}
```

```
void passa(int a[])
{
    printf("a size:%lu\n", sizeof(a)); // 4
    return;
}
```

IMPORTANT:

See the size difference 20 in `main()` in `passa()` is 4 bytes (size of a pointer) on `pi-cluster` and 8 on `ieng6`


- Call-by-value pointer (callee can change the pointer parameter to point to something else!)
- Acts like call-by-reference (called function can change the contents caller's array)

Arrays As Parameters: What is the size of the array?


- It's tricky to use arrays as parameters, as **they are passed as pointers to the start of the array**
 - In C, Arrays do not know their own size and at runtime there is no “bounds” checking on indexes

1 word content
(int = 4 bytes)

Observe numb is the name of an array (whose Rvalue is its starting address)



Remember **a** is parameter copy so is a separate variable that contains a pointer to num



```
int sumAll(int *);
```

```
int main(void)
```

```
{
```

```
    int numb[] = {9, 8, 1, 9, 5};
```

```
    int sum = sumAll(numb);
```

```
    return EXIT_SUCCESS;
```

```
}
```

```
int sumAll(int *a)
```

```
{
```

```
    int i, sum = 0;
```

```
    int sz = (int) (sizeof(a)/sizeof(*a));
```

```
    for (i = 0; i < sz; i++) // this does not work
```

```
        sum += a[i];
```

```
}
```

```
}
```

this is a POINTER to the first element.....
so sizeof(a) is the size of a pointer, not the array it points at
Net result:
sz is 4/4 = 1 on picluster

Arrays As Parameters, Approach 1: Pass the size


Two ways to pass array size

1. pass the **count** as an additional argument
2. add a **sentinel element** as the last element

remember you can only use `sizeof()` to calculate element count where the array is defined

```
int sumAll(int *a, int size);
int main(void)
{
    int numb[] = {9, 8, 1, 9, 5};
    int cnt = (int)(sizeof(numb)/sizeof(numb[0]));

    printf("sum is: %d\n", sumAll(numb, cnt));
    return EXIT_SUCCESS;
}
```



```
int sumAll(int *a, int size)
{
    int sum = 0;
    int *end;
    end = a + size;

    while (a < end)
        sum += *a++;
    return sum;
}
```

same as:
sum = sum + *a;
a++;

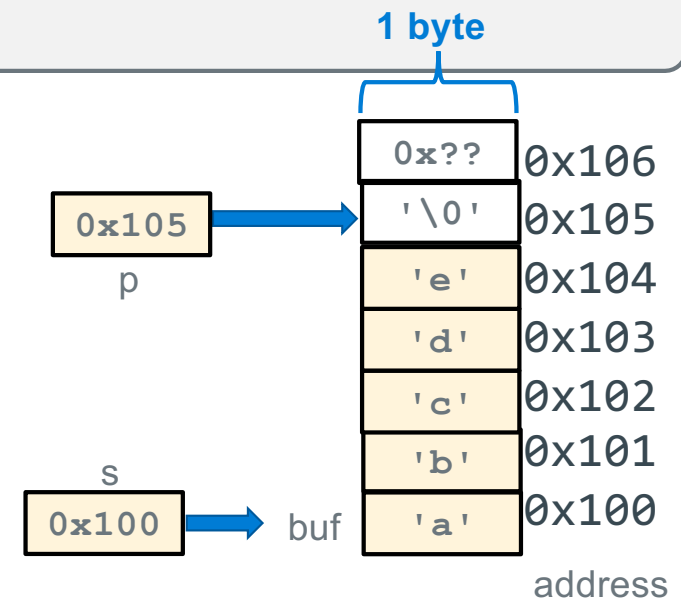
Arrays As Parameters, Approach 2: Use a sentinel element

- A **sentinel** is an element that contains a value that is not part of the normal data range
 - Forms of 0 are often used (like with strings). Examples: '\0', NULL

```
int strlen(char *a); // returns number of chars in string, not counting \0
int main(void)
{
    char buf[] = {'a', 'b', 'c', 'd', 'e', '\0'}; // or buf[] = "abcde";

    printf("Number of chars is: %d\n", strlen(buf));
    return EXIT_SUCCESS;
}
```

```
/* Assumes parameter is a terminated string */
int strlen(char *s)
{
    char *p = s;
    if (p == NULL)
        return 0;
    while (*p != '\0')
        p++;
    return (p - s);
}
```



Reference: Some String Routines in libc (#include <string.h>)

Function	Description
strlen(<i>str</i>)	returns the # of chars in a C string (before null-terminating character).
strcmp(<i>str1</i> , <i>str2</i>), strncmp(<i>str1</i> , <i>str2</i> , <i>n</i>)	compares two strings; returns 0 if identical, <0 if <i>str1</i> comes before <i>str2</i> in alphabet, >0 if <i>str1</i> comes after <i>str2</i> in alphabet. <i>strncmp</i> stops comparing after at most <i>n</i> characters.
strchr(<i>str</i> , <i>ch</i>) strrchr(<i>str</i> , <i>ch</i>)	character search: returns a pointer to the first occurrence of <i>ch</i> in <i>str</i> , or NULL if <i>ch</i> was not found in <i>str</i> . <i>strrchr</i> find the last occurrence.
strstr(<i>haystack</i> , <i>needle</i>)	string search: returns a pointer to the start of the first occurrence of <i>needle</i> in <i>haystack</i> , or NULL if <i>needle</i> was not found in <i>haystack</i> .
strcpy(<i>dst</i> , <i>src</i>), strncpy(<i>dst</i> , <i>src</i> , <i>n</i>)	copies characters in <i>src</i> to <i>dst</i> , including null-terminating character. Assumes enough space in <i>dst</i> . Strings must not overlap. <i>strncpy</i> stops after at most <i>n</i> chars, and <u>does not</u> add null-terminating char.
strcat(<i>dst</i> , <i>src</i>), strncat(<i>dst</i> , <i>src</i> , <i>n</i>)	concatenate <i>src</i> onto the end of <i>dst</i> . <i>strncat</i> stops concatenating after at most <i>n</i> characters. <u>Always</u> adds a null-terminating character.
strspn(<i>str</i> , <i>accept</i>), strcspn(<i>str</i> , <i>reject</i>)	strspn returns the length of the initial part of <i>str</i> which contains <u>only</u> characters in <i>accept</i> . strcspn returns the length of the initial part of <i>str</i> which does <u>not</u> contain any characters in <i>reject</i> .

Do not overuse strlen()

- C string library function `strlen()` calculates string length **at runtime**
- **Do not overuse `strlen()`, as it walks the array each time called**

```
int count_e(char *s) //  $O(n^2)$  !!!
{
    int count = 0;
    if (s == NULL)
        return 0;
    for (int j = 0; j < strlen(s); j++) {
        if (s[j] == 'e')
            count++;
    }
    return count ;
}
```



```
int count_e(char *s) //  $O(n)$  !!!
{
    int count = 0;
    if (s == NULL)
        return 0;
    while (*s) {
        if (*s++ == 'e')
            count++;
    }
    return count ;
}
```

The NULL Constant and Pointers

- **NULL is a constant** that **evaluates to zero (0)**
- You **assign a pointer variable to contain NULL** to **indicate that the pointer does not point at anything**
- A **pointer variable** with a **value of NULL** is called a “**NULL pointer**” (invalid address!)
- Memory location 0 (address is 0) is not a valid memory address in any C program
- Dereferencing NULL at runtime will cause a program fault (segmentation fault)!

```
p = NULL;  
i = *p;           /* segmentation fault! */  
*(int *)900000 = 25; /* cast 900000 to a pointer */  
                  /* if writeable address space, it works */  
                  /* that memory location just changed */
```

Using the NULL Pointer

- Many functions return NULL to indicate an error has occurred

```
/* these are all equivalent */  
int *p = NULL;  
int *p = (int *)0;    // cast 0 to a pointer type  
int *p = (void *)0;   // automatically gets converted to the correct type
```

- NULL is considered “false” when used in a Boolean context
 - Remember: false expressions** in C are defined to be zero or NULL
- The following two are equivalent (the second one is preferred for readability):

```
if (p) ...  
if (p != NULL) ...
```

Simple String IO - Reading

Task	Example Function Calls
Read a string	<pre>#include <stdio.h> char *strptr; char myStr[BFSZ]; strptr = fgets(myStr, BFSZ, stdin);</pre> <div>must pass the size of the array so fgets() knows how much space there is</div>

`char *fgets(char array[], int size, FILE *stream)`

- `char *` is a pointer (address) to an **array of char**
- reads in at most **one less than size** characters from **stream** and stores them into **array**
- Reading stops after an **EOF** or a newline '\n'
 - If a newline ('\n') is read, it is stored into the buffer
 - **A terminating null byte ('\0') is always stored after the last character in the buffer**

t	h	i	s		i	s		a		s		t		r		i		n		g		\n	\0
---	---	---	---	--	---	---	--	---	--	---	--	---	--	---	--	---	--	---	--	---	--	----	----

- Returns a **NULL at end of file** (or a read failure), otherwise a pointer to array (pointers later...)
- See `man 3 fgets`

Pointer returns from a function call (NULL Examples)

This function returns a pointer to the character that follows the first comma ','

```
char *next(char *ptr)
{
    if (ptr == NULL)
        return NULL;

    while ((*ptr != '\0') && (*ptr != ','))
        ptr++;

    if (*ptr == ',')
        return ++ptr;
    return NULL;
}
```

```
#include <stdlib.h>
#include <stdio.h>
#define BUFSZ 512
char *next(char *);

int main()
{
    char buf[BUFSZ];
    char *ptr;

    while (fgets(buf, BUFSZ, stdin) != NULL) {
        printf("buf: %s\n", buf);

        if ((ptr = next(buf)) != NULL)
            printf("after: %s\n", ptr);
        else
            printf("no comma found\n");
    }
    return EXIT_SUCCESS;
}
```

Returning a Pointer To a Local Variable (Dangling Pointer)

- There are many situations where a function will return a pointer, but a function must never return a pointer to a memory location that is **no longer valid** such as:
 - Address of a **passed parameter copy** as the caller may or will deallocate it after the call
 - Address of a **local variable (automatic)** that is invalid on function return
- These errors are called a **dangling pointer**

n is a parameter with the scope of bad_idea it is no longer valid after the function returns

```
int *bad_idea(int n)
{
    return &n; // NEVER do this
}
```

a is an automatic (local) with a scope and **lifetime** within bad_idea2 a is no longer a valid location after the function returns

```
int *bad_idea2(int n)
{
    int a = n * n;
    return &a; // NEVER do this
}
```

```
/*
 * this is ok to do
 * it is NOT a dangling
 * pointer
 */

int *ok(int n)
{
    static int a = n * n;
    return &a; // ok
}
```

String Literals (Read-Only) in Expressions

- When strings in quotations (e.g., "string") are **part of** an **expression** (i.e., *not part of an array initialization*) they are called **string literals**

```
printf("literal\n");  
printf("literal %s\n", "another literal");
```

- What is a **string literal**:
 - Is a **null-terminated string** in a **const char array**
 - Located in the **read-only data segment of memory**
 - Is **not assigned a variable name** by the compiler, so it is only accessible by the location in memory where it is stored
- **String literals** are a type of **anonymous variable**
 - Memory containing **data without a name bound** to them (**only the address is known**)
- The **string literal in the printf()'s**, are replaced with the **starting address of the corresponding array** (first or [0] element) when the code is compiled

String Literals, Mutable and Immutable arrays - 1

- **mess1** is a **mutable array** (type is `char []`) with enough space to hold the string + `'\0'`

```
char mess1[] = "Hello World";  
*(mess1 + 5) = '\0'; // shortens string to "Hello"
```

`mess1[]` Hello World\0

- **mess2** is a **pointer** to an **immutable array** with space to hold the string + `'\0'`

```
char *mess2 = "Hello World"; // "Hello World" read only string literal  
// mess2 is a pointer NOT an array!  
*(mess2 + 1) = '\0'; // Not OK (bus error)
```

`mess2` → Hello World\0 ← read only string literal

- **mess3** is a **pointer** to a mutable array

```
char *mess3 = (char []) {"Hello World"}; // mutable string  
*(mess3 + 1) = '\0'; // ok
```

using the cast `(char [])`
makes it mutable

`mess3` → Hello World\0 ← mutable string

2D Array of Char (where elements may contain strings)

- 2D array of chars (where rows may include strings)
- Each row has the same fixed number of memory allocated
- All the rows are the same length regardless of the actual string length
- **The column size must be large enough for the longest string (fills rest with zeros '\0')**

```
char aos[3][14] = {"my", "two dimension", "char array"};
```

high
memory

aos[2]	c	h	a	r		a	r	r	a	y	'\0'			
aos[1]	t	w	o		d	i	m	e	n	s	i	o	n	'\0'
aos[0]	m	y	'\0'											

low
memory


```
#define ROWS 3
char aos[ROWS][14] = { "my", "two dimensional", "char array"};
char (*ptc)[14] = aos; // ptc points at a row of 14 chars

for (int i = 0; i < ROWS; i++)
    printf("%s\n", *(ptc + i));
```

high
memory

Array of Pointers to Strings (This is NOT a 2D array)


- 2D char arrays are an inefficient way to store strings (wastes memory) unless all the strings are similar lengths, so 2D char arrays *are rarely used* with string elements
- **An array of pointers** is common for strings as "rows" can vary in length
- `char *aos[3];` `aos[2][0]` `aos[2][11]`



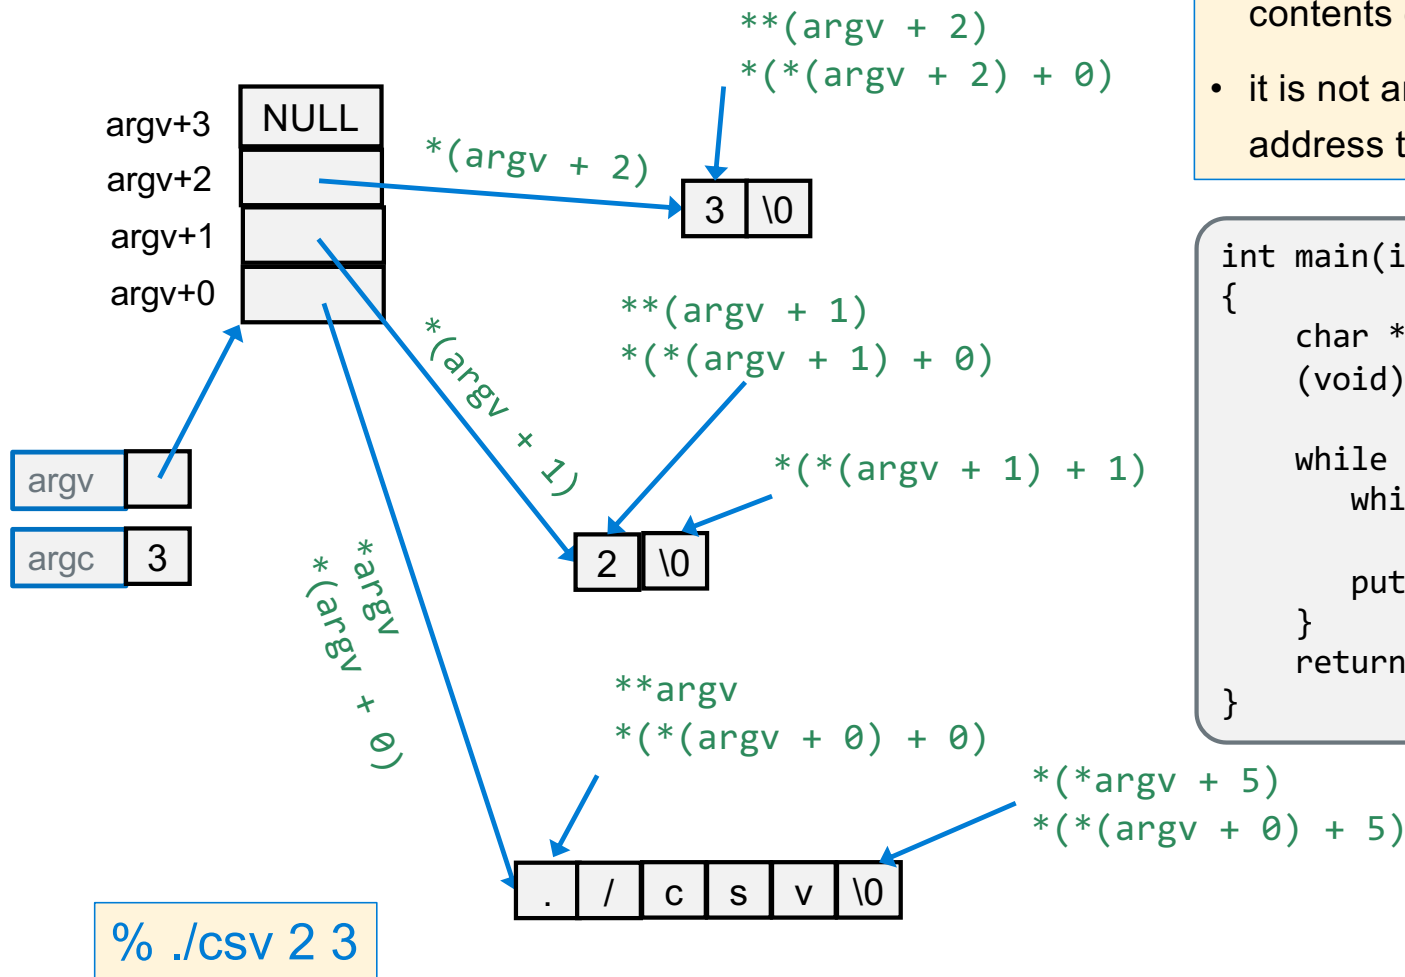
- `aos` is an **array of pointers**; each pointer points at a **character array (also a string here)**
- **Not a 2D array**, but any char can be accessed as if it was in a 2D array of chars
 - When I was learning, this was the most confusing syntax aspects of C

main() Command line arguments: argc, argv

- Arguments are passed to main() as a pointer to an array of pointers to char arrays (strings)(****argv**)
Conceptually: % *argv[0] *argv[1] *argv[2] ...
- argc** is the number of VALID elements (they point at something)
- *argv** (**argv[0]**) is **usually** is the **name** of the executable file (% **./vim** file.c)
- argv[argc]** or ***(argv + argc)** always contains a NULL (0) sentinel
- argv** elements point at **mutable** (fixed size) **strings!**



Accessing argv char at a time



- `argv` is a pointer variable, whose contents can be changed
- it is not an array name, which is just an address that cannot be changed

```
int main(int argc, char **argv)
{
    char *pt;
    (void)argc; // shut up the compiler


    while ((pt = *argv++) != NULL) {
        while (*pt != '\0') {
            putchar(*pt++);
            putchar('\n');
        }
        return EXIT_SUCCESS;
    }
}
```


Defining an Array of Pointer to Strings

- `mess4` is an array of pointers to immutable arrays


```
char *mess4[] = {"Hello", "World"}; // immutable string  
*(mess4 + 1) = '\0';                // bus error
```

Bus error: writing
read only memory
Seg fault: writing
unallocated memory



- `mess5` is an array of pointers to mutable arrays

```
char *mess5[] = { (char []){"Hello"}, (char []){"World"}};  
*(mess5 + 1) = '\0';                // OK!
```

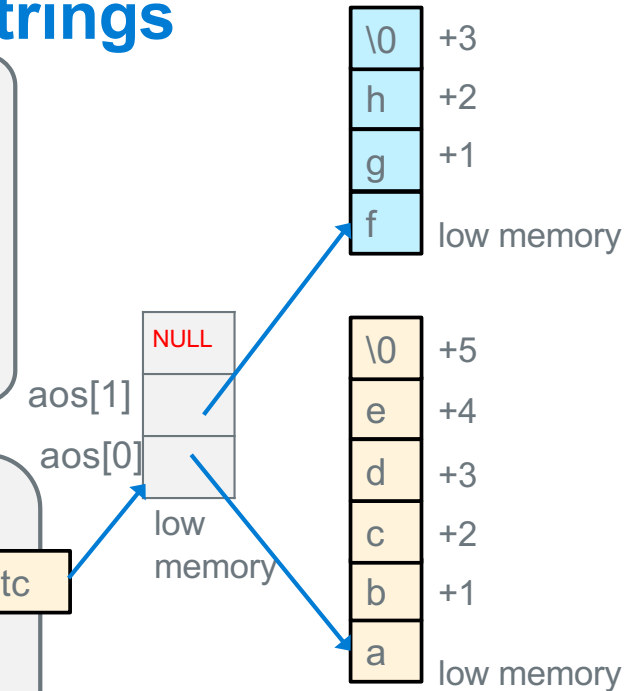


Defining an Array of Pointers to Mutable Strings

- Make an **array of pointers** to **mutable strings** requires using a **cast to an array (char [])**
- Add a NULL sentinel at the end to indicate the end of the array

```
char *aos[] = {  
    (char []) {"abcde"},  
    (char []) {"fgh"},  
    (char *) {NULL}  
};  
char **ptc = aos;
```

```
printf("%c\n", (*(aos + 1) + 1));  
  
while (*ptc != NULL) {  
    printf("%s\n", *ptc);    // prints string  
  
    for (int j = 0; *(*ptc + j); j++)  
        putchar(*(*ptc + j)); // char in string  
  
    putchar('\n');  
    ptc++;  
}
```



```
%. /a.out  
abcde  
abcde  
fgh  
fgh
```

Pointers to Functions (Function Pointers)

- Similar in concept to an array name, a function name ends up being the address of the first instruction in a function

- A function pointer variable contains the address of a function

- Generic format: `returnType (*name)(type1, ..., typeN)`

- Looks like a function prototype with extra * in front of name
- Why are parentheses around (*name) needed?

`returnType *name(type1, ..., typeN) //wrong`

- Above says name is a function returning a pointer to returnType

- Using the function:

`(*name)(arg1, ..., argN)`

`name(arg1, ..., argN)`

Calls the pointed-to function with the given arguments and returns the return value

Pointers to Function Example

```
int add1(int);
int sqr(int);
void array_update(int (*)(int), int *, int);
void print_array(int *, int);

int main(void)
{
    int array[] = {4, 8, 15, 16, 23, 42};
    int cnt = sizeof(array)/sizeof(array[0]);

    print_array(array, cnt);
    array_update(add1, array, cnt);
    print_array(array, cnt);
    array_update(sqr, array, cnt);
    print_array(array, cnt);
    return EXIT_SUCCESS;
}
```

```
void array_update(int (*f)(int), int *a, int cnt)
{
    while (a < endpt) {
        *a = f(*a);
        a++;
    }
}
```

```
void print_array(int *a, int cnt)
{
    int *endpt = a + cnt;

    while (a < endpt)
        printf("%d ", *a++);
    printf("\n");
}
```


```
int add1(int i)
{
    return i + 1;
}

int sqr(int i)
{
    return i * i;
}
```

```
./a.out
4 8 15 16 23 42
5 9 16 17 24 43
25 81 256 289 576 1849
```

string buffer overflow: common security flaw

- A **buffer overflow** occurs when data is written **outside the boundaries** of the memory allocated to **target variable** (or target buffer)
- **strcpy()** is a very *common source of buffer overrun security flaws*:
 - always ensure that the **destination array is large enough** (and don't forget the null terminator)
- **strcpy()** can cause problems when the **destination** and **source regions overlap**



strcpy() buffer overflow: over-write of an adjacent variable

```
int main(void)
{
    char s1[] = "before";
    char r2[] = "xyz";
    char s2[] = "after";

    printf("s2: %s\nr2: %s\nr2: %s\n", s2, r2, s1);

    strcpy(r2, "hello");

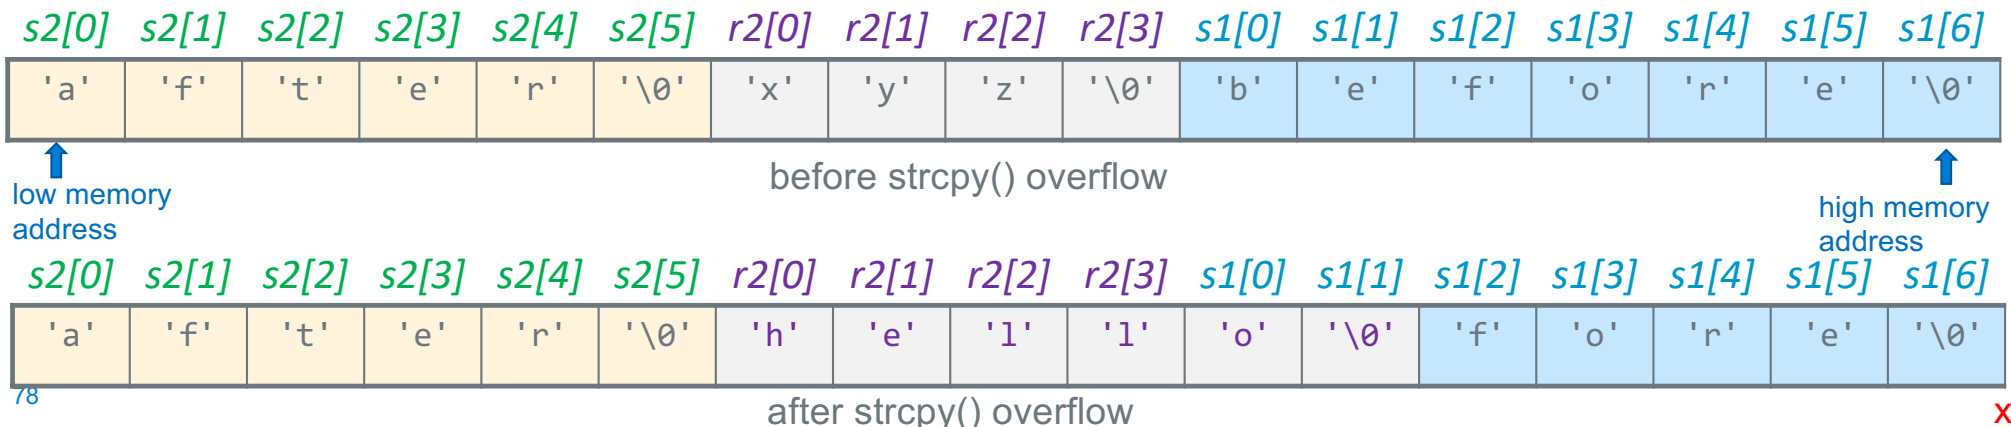
    printf("\ns2: %s\nr2: %s\nr2: %s\n", s2, r2, s1);
    return EXIT_SUCCESS;
}
```

these are mutable
arrays, not literals

compile on pi-cluster with
gcc test.c

```
./a.out
s2: after
r2: xyz
s1: before

s2: after
r2: hello
s1: o
```




Extra Slides

-

Memory Size

- Since memory addresses are implemented in hardware using binary
 - The **Size (number of byte sized cells)** of Memory is specified in **powers of 2**
- Memory size/capacity in **bytes** is specified by the “**Number of bits**” in an address
 - 32 bits of address = $2^{32} = 4,294,967,296$
 - Address Range is 0 to $2^{32} - 1$ (unsigned)
- Shorthand notation for address size (Memory Capacity)
 - KB = 2^{10} (K=1024) kilobyte
 - MB = 2^{20} megabyte
 - GB = 2^{30} gigabyte
 - TB = 2^{40} terabyte
 - PB = 2^{50} petabyte



Fixed size types in C (later addition to C)

- Sometimes programs need to be written for a particular range of integers or for a particular size of storage, regardless of what machine the program runs on
- In the file `<stdint.h>` the following fixed size types are defined for use in these situations:

Signed Data types	Unsigned Data types	Exact Size
<code>int8_t</code>	<code>uint8_t</code>	8 bits (1 byte)
<code>int16_t</code>	<code>uint16_t</code>	16 bits (2 bytes)
<code>int32_t</code>	<code>uint32_t</code>	32 bits (4 bytes)
<code>int64_t</code>	<code>uint64_t</code>	64 bits (8 bytes)

Defining Strings: Initialization Equivalents

- Following definitions create **equivalent** 4-character arrays
 - These are all strings as they all include a null ('\0') terminator

```
char a[4] = {'c', 'a', 't', '\0'};
char b[4] = {'c', 'a', 't', 0};
char c[4] = {'c', 'a', 't'};           // missing initial value defaults to 0
char d[4] = { 99, 97, 116, 0};         // 99 = 'c', 97 = 'a', 116 = 't'
char e[4] = "cat";
char f[4] = "cat\0";                   // literal has 5 chars; array f string
                                        // length is 3
```

Pointer Practice

```
int *ptr;
```

Declares a variable, `ptr`, which is a pointer to (it contains the address of) an `int` in memory

```
int x = 5;
```

Declares two variables, `x` and `y`, that contain ints, and *initializes* them to 5 and 2, respectively

```
int y = 2;
```

```
ptr = &x;
```

Sets `ptr` to contain the address of `x` ("`ptr` points to `x`")

```
y = 1 + *ptr;
```

"Dereference `ptr`"

Sets `y` to "1 plus the value stored at the address held by `ptr`. Because `ptr` points to `x`, this is equivalent to `y = 1 + x`;


```
x = *(&y);
```

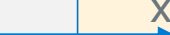
Sets `x = y`; The `*` and `&` cancel each other. get the address of `y` and then get the contents pointed by that address


`ptr`

`x` write

`y` write

`ptr`  `x`
`y`

`ptr`  `x` read
`y` write

`ptr`  `x` write
`y` read

strtol() and strtoul() examples of passing a pointer to a pointer

```
long int strtol(const char *str, char **endptr, int base);
```

```
unsigned long int strtoul(const char *str, char **endptr, int base);
```

reruns the string converted to a long or unsigned long

str pointer to the string to convert

endptr pass the address of a variable that is a char pointer (output variable)

base: number base used by the string

- **Example**: string is to contain just positive numbers ≥ 0 (in ascii) with no extra stuff
- If the string is not valid, then
 - ***endptr** **!=** **'\0'** then string contains more than just numbers (bad input)
 - ***endptr** stores the address of the first invalid character found in the buffer pointed (**str**)
- How to use **endptr** when it does not contain **NULL**:
 - If there are other conversion errors (you can read the man page) then **errno** **!=** 0
 - When conversion is ok, **errno** is unaltered (always clear it before calling these routines)

strtol() and strtoul() examples of passing a pointer to a pointer

```
#include <stdlib.h>
#include <errno.h>
char *endptr;
char buf[] = "33"; // test buffer string
int number;

errno = 0; // set errno to 0 (zero) before each call
number = (int)strtol(buf, &endptr, 10)
// check if the string was a proper number
// *entpr should be at the end of the string == '\0'

if ((*endptr != '\0') || (errno != 0)) {
    // handle the error
}
printf("%d\n", number);
```

Copying Strings: Use the Sentinel; libc: strcpy()

- To copy an array, you must copy each character from source to destination array
- Watch overwrites: strcpy assumes the target array size is equal or larger than source array



<i>index</i>	0	1	2	3	4	5
<i>char</i>	'H'	'e'	'l'	'l'	'o'	'\0'

```
char str1[80];  
strcpy(str1, "hello");
```

```
char *strcpy(char *s0, char *s1)  
{  
    char *str = s0;  
  
    if ((s0 == NULL) || (s1 == NULL))  
        return NULL;  
    while (*s0++ = *s1++)  
        ;  
    return str; // address of dest string  
}
```

Accessing Characters in an Array of Pointers to Strings

`aos[1][1]` is `*(*aos + 1) + 1` which contains 'n'; address is `(*(aos + 1) + 1)`



```
char *ptr;

ptr = *aos;

if (*ptr == ',')
    or
if (**aos) == ','
```

You cannot write to an immutable literal

- You can use & to get the address of an anonymous variable as shown
 - Though the Rvalue of "Hello" is the address
- You cannot write to an immutable literal

```
char *pt;  
  
// pt = "Hello";  
pt = (char *) &"Hello";  
*pt = 'a'; // bus error  
printf("%s\n", pt);
```