

Version 2.11

UCSD CSE 30 Section B

Computer Organization and Systems Programming

Lecture 1

April 2, 2024

Keith Muller

DEC PDP 11/45 - 1973



CSE30 Section B Spring 2024

Instructor: Keith Muller

- **I highly encourage feedback**
 - Please bring any issues to my attention, I will promptly address them
- How to **contact me directly**:
 - kmuller@ucsd.edu
 - Please do not use canvas messages

- In Person Office Hours: CSE 2109
 - Tue, Thu: 2:00 PM to 3:00 PM
 - These office hours are group meetings
 - Ask questions, review material, or just come to listen
 - Students who attend office hours tend to do better in the course

- Zoom Office Hours
<https://ucsd.zoom.us/j/94331007124>
 - Friday: 4:00 PM to 4:45 PM
 - These office hours can be individual or for a group if you like
 - Additional office times By Appointment
 - Send me email to schedule

CSE 30 Spring 2024 – Staff Covers Both Sections A & B

Section A (Cao) and B (Muller) share the same pool of TA's and Tutors

TA's

- Nitya Agarwal
- Mihir Kekkar
- Yuchen Jing
- Liam Fernandez

Tutors

Ali Alabiad
Bryan Cho
Charlotte Dong
Vivian Liu
Kate Romero
Kevin Shen
Charvi Sukla
Fong Vachirathanusorn
Joseph Edmonston
Thanh-Nhan Lam

Tutors

Christian Lee
Jessie Ouyang
Brandon Reponte
Adrian Rosing
Luffy Saito
Leica Shen
Shijie Wang
Alex Simonyan
Reese Whitlock

Overview of Grading - See Syllabus (Canvas) for More Details

70 pts – Attending Lecture in person

- 5 points per section B lecture

120 pts total – Canvas Quizzes

240 pts total – Programming Assignments

190 pts - Midterm – In Person

380 pts - Final – In Person

1000 pts total for graded assignments

- **Special grading circumstances** (e.g., extended absence, illness, other issues, etc.)
 - **PROMPTLY** Contact me directly (kmuller@ucsd.edu)

Lecture 1 QR Code

- Class attendance points: To encourage you to attend lecture
 - Over the years we have found that students that attend lectures in CSE30 get better grades
- Section B has 20 lectures, attend 14 to get the 70 points
 - Attending more than 14 gets you up to 30 more points
- Attendance is taken at the start of class using google forms that is accessed with a lecture QR code in the slides
 - **For the first lecture only, the form will be open until 9 PM**
 - bring a device that can use QR codes to access google forms and allows you to sign into UCSD SSO
 - You will be required to supply a code word announced in class
 - You will eventually get an email acknowledgement from google that your attendance was recorded
- **ONLY If you cannot access the google form**, send me email (kmuller@ucsd.edu) with the code word in the subject line
 - The email must be timestamped within the first 20 minutes of lecture



If you have issues with this method,
please contact me

CSE30 Spring 2024 Section B Specific

- There are two sections: Section A (Cao) and **Section B (Muller)**
- **What is the same in the two sections**
 - **study topics** (roughly in-sync by the end of each week)
 - **quizzes**
 - **Programming Assignments**
- **What is different between the two sections**
 - **lecture materials**
 - **midterm questions (from Sect B lecture)**
 - **final questions (from Sect B lecture)**
- **In-person lecture attendance is strongly encouraged** (attendance points)
 - Lectures are podcast recorded
- **Discussion section attendance is optional but strongly encouraged**
 - You may attend either discussion section and still be enrolled in Sect B
 - Section B sections are podcast recorded
- **See the syllabus for grading details**

CSE30 Class Resources

- **Section B Lecture Slides:** <https://github.com/cse30-sp24/Muller-Slides>
 - Located on class github in both **pptx** and **pdf** format
 - Slides **are updated constantly** to correct errors and to improve content
 - Version is at the upper left on the title slide
 - **Always check** you have the current version the morning before lecture
- **Class github:** <https://github.com/cse30-sp24>
- **Piazza:** https://piazza.com/ucsd/spring2024/cse30_sp24_a0/home
 - **First Place to go to** for **Q/A** and **important announcements**
 - **Public piazza posts are for:** general questions on PA's and lectures
 - **Do not post publicly** any parts of an assignment, quiz or exam solution
 - **Private posts are for:** specific situation relating to just you or you are not sure
- **Tutor Lab hour schedule:** <https://autograder.ucsd.edu>
 - For getting help from the tutors
- **Canvas:** <https://canvas.ucsd.edu/courses/54650>
 - Links to quizzes, textbooks, programming assignments, exams
- **Gradescope:** <https://www.gradescope.com>
 - Quizzes and Submitting programming assignments

Surviving Section B Lectures (In-person)

- **Make sure you bring your copy of lecture slides to class, it helps**
- **How to get my attention in class**
 - I **never intentionally ignore questions; I just may not see you**
 - **Raise your hand**, or just **call out** if I appear to **ignore you by accident**
- **You must SLOW ME DOWN:** Otherwise, **I tend to speed up**
 - Please do not be shy, **speak up** and **remind me to slow down**
- If you have questions, or I went too fast, or the material is not clear, etc.
 - Please ask me to go over it again (do this right away, not 5 slides later)
 - Just don't sit there and waste your time
 - **my responsibility:** help you learn the material
 - **your responsibility:** **ask questions** (I love questions, they also slow me down!)

How to do well in CSE30 - 1

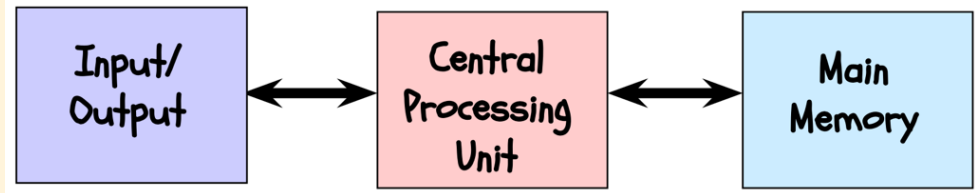
- Go to lecture
 - Before lecture go over the class slides
 - Lecture slides are posted the day before class (last minute updates that morning)
 - Keep your lecture slides up to date (I update them to fix errors and address questions)
- Go to Discussion Sessions
 - ask the TA's and Tutors for help
- Studying for exams
 - All the exam question topics are found in my slides and the PA writeups
 - Try to write the exam yourself, with practice you will be able to guess the questions
- Post to piazza when you have questions
- Do the readings on time
- Review the material: watch the podcasts and occasional special topic videos

How to do well in CSE30 - 2

- **Most important:** **Keep up, do not procrastinate as it is hard to catch up**
 - The class material starts easy and gets much harder over the quarter
 - Do not expect you can do later programming assignments in less than 5 days
 - Do not expect to learn the material by binge watching podcasts, this never ends well
- **Please be careful when using web resources** for this class
 - a lot of the material you will find is either not correct or does not apply to our programming environment
 - this is especially true with assembly language programming topics
- **Are you struggling?**
 - Do not wait, **ask for help as soon as possible** – do not fall behind
 - **Best advice: Come to my office hours** (or schedule a zoom meeting)
 - Give me a chance to help you
 - I will spend as much time as necessary to help you understand the material

A General-Purpose Computer – Von Neuman Architecture

- Since the middle of the 20th century, many **architectural approaches** to the **general-purpose computer** have been tried
- The **architecture** which **nearly all modern computers** are based was proposed by John Von Neuman in the late 1940's
- The **major components** are:



- **Central Processing Unit (CPU)**: a device which **fetches**, **interprets (decodes)**, and **executes** a **specified set of operations** called **instructions**
- **Memory**: **Storage** of **N words** of **W bits**, where **W** is a fixed architectural parameter, and **N** can be expanded to meet **workload** (the programs running on the CPU) and **cost requirements**
- **I/O**: **Devices for communication with the outside world** (including external persistent storage)
 - External connections (from CPU to memory and I/O) typically use industry **"standards"**
 - **Standards** enable technologies from **different companies to interoperate**

What is Computer Architecture?

Instruction Set Architecture (ISA)

- Functional behavior of a computer system as viewed by a programmer
 - describes how the CPU is controlled by software programs
 - specifies both what the processor can do as well as how it gets it done
- Architectural Characteristics (partial list):
 - supported data types (how data is encoded)
 - CPU registers (number, size, use, etc.)
 - how the hardware manages main memory
 - instructions a microprocessor can execute
 - Operations they perform
 - Instruction "format" (bit patterns) in memory
 - input/output model

Machine Organization


- Physical (design) realization of what is specified by the instruction set architecture
- It deals with how the hardware components are linked together to meet the requirements specified by instruction set architecture
 - The ISA allows variability in the physical design implementation to address different workload needs (cost, scalability, etc.)
- Machine organizational characteristics (partial)
 - Hardware component choices to achieve:
 - Expandability
 - Configurability
 - etc.
 - Physical layout
 - Number and type of peripherals (I/O devices)

Von Neuman Architecture

- **Distinguishing feature:** Memory contains **both** program CPU instructions and data
- **CPU Instructions** are encoded in memory with patterns of ones and zeros (similar to binary numbers)
 - **Encoded CPU instructions** are called **machine code (or machine language)**
- **Example:** three 32-bit instructions (shown in hexadecimal format below)

```
81 fe 89 32
81 54 22 af
81 22 10 9A
```

- **Instructions** operate on **data** that is stored in a **small capacity volatile (fast) memory** in the CPU
 - This memory is called **registers**
- CPU **reads/writes data** from **memory** from these **data registers** to **operate on the data**




C, Assembly and Executable Programs

- **Assembly language** is a **symbolic version** of the **machine code (language)**
 - **Instructions** describe operations the hardware can perform (e.g., =, +, -, *)
 - **Unique to a specific ISA**: e.g., ARM-32 versus IA-64
 - May be stored in a **human readable text file**
 - You can write in assembly language just like C or Java
 - Assembly is much easier to program than machine code
- A **high-level language** (like C) is **compiled** into an **assembly language equivalent**
 - A statement in C is represented by a sequence of one or more assembly language instructions (why do you think it is a sequence?)
- **Assembly language program**
 - assembly language program is **translated (assembled)** into **machine code**
- An **executable program** contains
 - **series of instructions in machine code** (the program)
 - (maybe some) **data** to operate on

Assembly & Machine Code Example: ARM-32 (32-bits)

Consider an addition statement
R0 = R1 + R3;



Simple R-type instructions follow the following template:

OP Rd, Rn, Rm

Assembly Language (human readable)
ADD R0, R1, R3

machine code pattern in memory

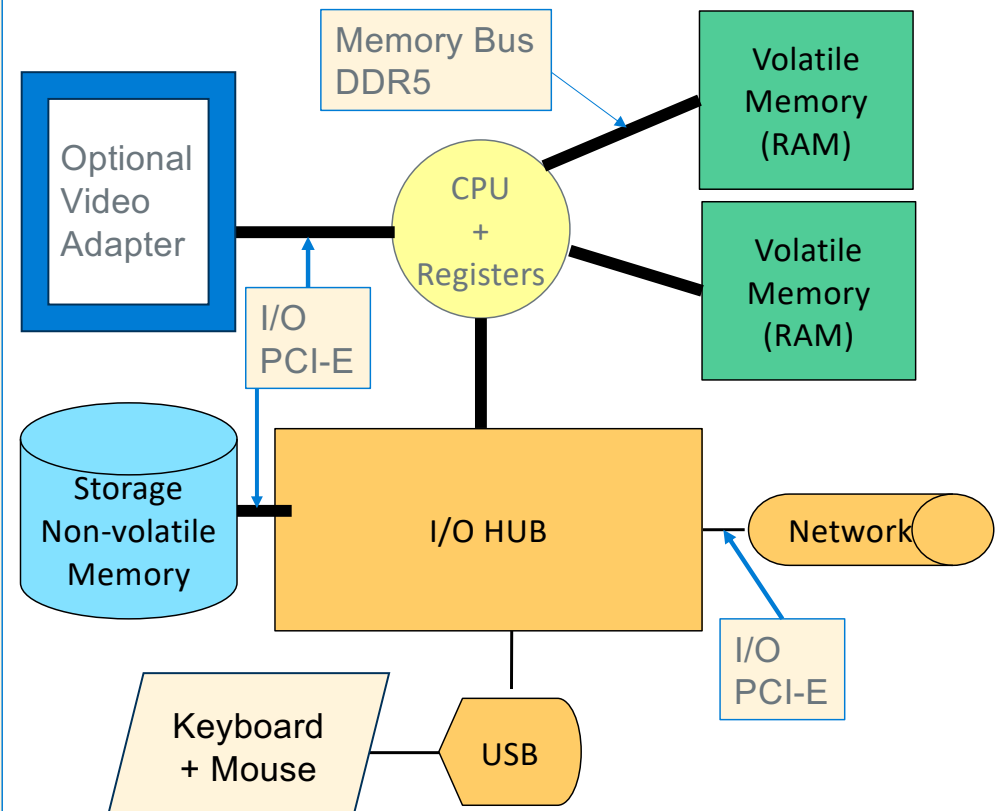
1110 0000 1000 0001 0000 0000 0011

- 0000 - AND
- 0001 - EOR
- 0010 - SUB
- 0011 - RSB
- 0100 - ADD
- 0101 - ADC
- 0110 - SBC
- 0111 - RSC
- 1000 - TST
- 1001 - TEQ
- 1010 - CMP
- 1011 - CMN
- 1100 - ORR
- 1101 - MOV
- 1110 - BIC
- 1111 - MVN

List of Different operations for this type of instruction

Review: Machine Organization – Von Neuman


1. CPU executes a machine code program
 - Machine code is specific to a particular CPU Instruction set Architecture (ISA)
2. **Memory** contains **both data and programs**
3. **I/O (input/Output)**: Connects the CPU and memory to the external world
 - An I/O operation is where data (including machine code) is copied between persistent storage (like an SSD) and ram memory
 - **Volatile (non-persistent) memory**
 - contents lost when power is removed
 - Memory dimms (memory bus)
 - CPU registers (memory inside the CPU)
 - **Non-volatile (persistent) memory**
 - contents preserved when power is removed
 - SSD (I/O bus attached)
 - NVDIMM (memory bus attached)



Memory Triangle: Hardware Cost/Performance/Capacity Tiers

Assume each instruction takes 1 clock cycle

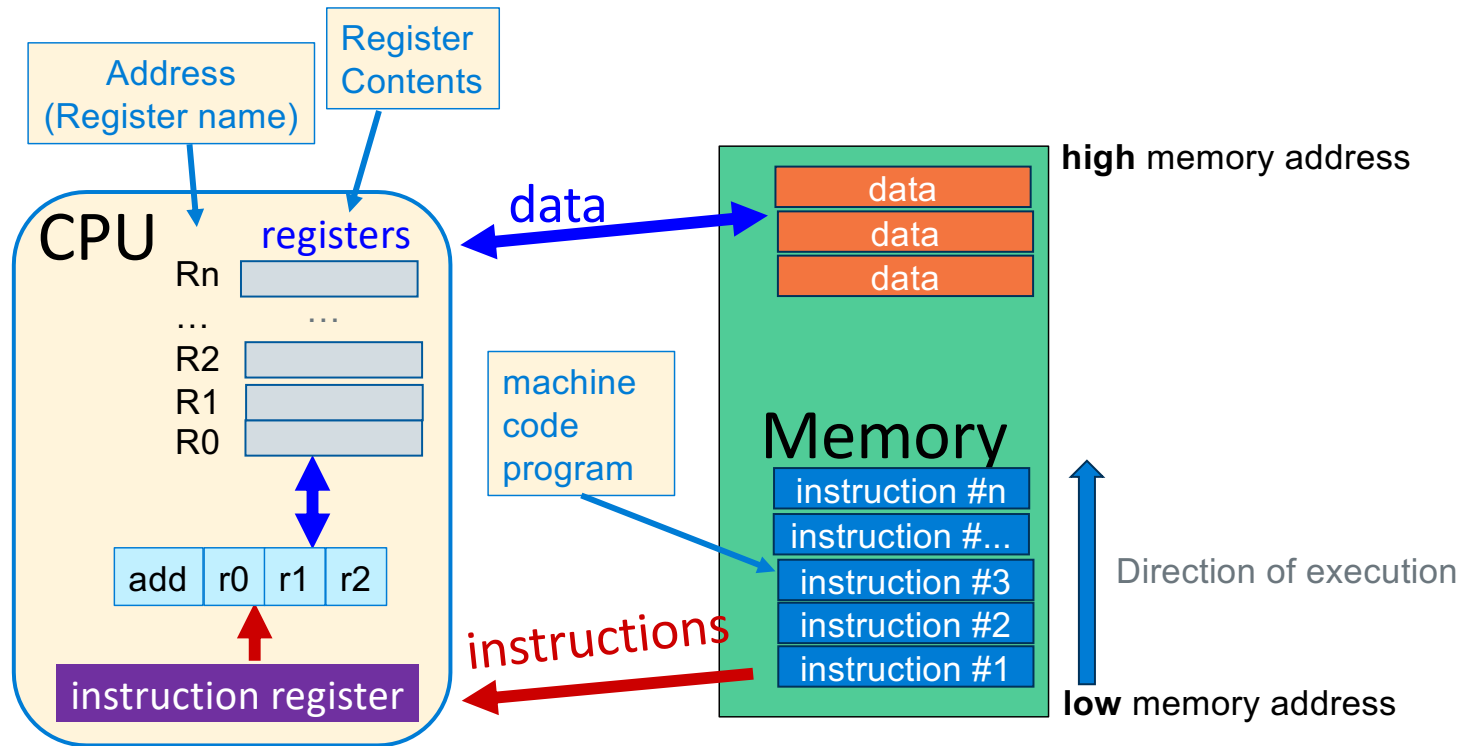
Clock cycle \approx time to access; larger is slower



Design Goal: best performance at the lowest (or specific) cost


Other goals: performance/energy (operating cost), expandability, high margin (price/cost)

Machine code execution order



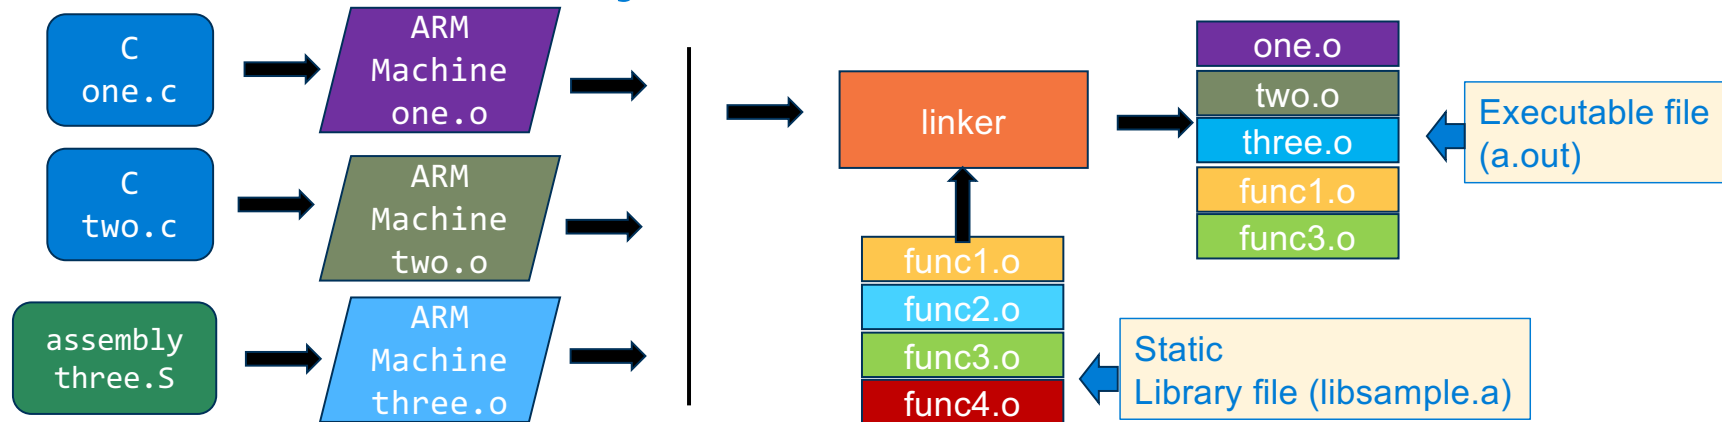
- **Execution order:** Programs execute from instructions located in **low address memory** to **high address memory** stepping **one machine instruction at a time** (called **execution order**) **unless there is a branch** (example: loop, if statement etc.)

From Source to Machine code




- The **granularity** of **compilation and assembly** is a **single text file** (called a **translation unit**)
 - **.c file** is a C **source** file (file.c)
 - **.S file** (upper case S) is a **human written assembly source** file (file.S)
 - **.s file** (lower case s) is a **compiler generated assemble source** file (file.s)
 - **.o file** is a **machine code binary (object)** file (file.o)
- Multiple **.o** files are **combined** (linked) into an **executable file**

Linker: Combines object files to create an executable file



- Each source file (**Translation unit**) is compiled (or assembled) independently to an object file
 - When we **modify a single file** in a **multi-source file program**, we want to only **recompile the file that changed** and **combine** it with the **other already compiled object files**
- **Library file** (libXX.a – where XX is the library name) is an **aggregation of distinct object (.o) files**
- **Linker combines** all the **listed object files together** plus **just those object files in libraries** whose **contents are referenced**
 - **Example:** one.c and two.c call functions contained in func1.o and func3.o (no calls to func2.o or func4.o)
- **Important:** **Object files created from C and assembly source can be linked** (call each other) into a working executable when certain rules are followed (**we will be doing a lot of this later this quarter**)

From Source to Execution: Different ISA



Pi-cluster system (all CSE30 PA's)

ieng6

x

From Source code to Execution

```
$ cat test.c
#include <stdlib.h>
#include <stdio.h>
int main (void)
{
    printf("Hello!\n");
    return EXIT_SUCCESS;
}
```

```
$ gcc -Wall -Wextra -Werror -c -S test.c
```

compile

```
$ ls -ls
```

```
total 8
```

```
4 -rw-r--r-- 1 kmuller kmuller 109 Mar 14 15:57 test.c
```

```
4 -rw-r--r-- 1 kmuller kmuller 725 Mar 14 15:58 test.s
```

```
$ gcc test.s
```

```
$ ls -ls
```

```
total 16
```

```
8 -rwxr-xr-x 1 kmuller kmuller 7708 Mar 14 15:58 a.out
```

```
4 -rw-r--r-- 1 kmuller kmuller 109 Mar 14 15:57 test.c
```

```
4 -rw-r--r-- 1 kmuller kmuller 725 Mar 14 15:58 test.s
```

```
$ ./a.out
```

```
Hello!
```


load and then execute

Source to Execution Steps

1. Compile (c source to assembler)
2. Assemble (assembler source to object)
3. Link (Combine object files to executable)
4. Load (Copy executable from into memory)
5. Execute (OS runs the code)

assemble and link

gcc automatically calls the assembler with .S or .s files



Equivalent Code: C -> Assembly -> Machine

```
#include <stdlib.h>
#include <stdio.h>
int main(void)
{
    printf("Hello!\n");
    return EXIT_SUCCESS;
}
```

C
source

```
mesg: .section .rodata
      .string "Hello!\n"
      .text
      .global main
      .type main, %function
      .equ FP_OFF, 4
      .equ EXIT_SUCCESS, 0
main:  push {fp, lr}
      add fp, sp, FP_OFF
      ldr r0, L1
      bl printf
      mov r0, EXIT_SUCCESS
      sub sp, fp, FP_OFF
      pop {fp, lr}
      bx lr
L1:   .word mesg
```

ARM-32 assembly

address of mesg

Code aka
TEXT

memory address high low bytes contents corresponding assembly

```
00010408 <main>:
  10408: e92d4800      push {fp, lr}
  1040c: e28db004      add fp, sp, 4
  10410: e59f0010      ldr r0, [pc, 16]
//10428 <L1>
  10414: ebffffb3      bl 102e8 <printf@plt>
  10418: e3a00000      mov r0, 0
  1041c: e24bd004      sub sp, fp, 4
  10420: e8bd4800      pop {fp, lr}
  10424: e12fff1e      bx lr
00010428 <L1>:
  10428: 0001049c      ← Machine instructions
                                     ← address of mesg
0001049c <mesg>:
  1049c: 6c6c6548      // 'l', 'l', 'e', 'h'
  104a0: 000a216f      // '\0', '\n', '!', 'o'
```

Data