


Version 2.00

UCSD CSE 30

Computer Organization and Systems Programming


Lecture 5

Keith Muller





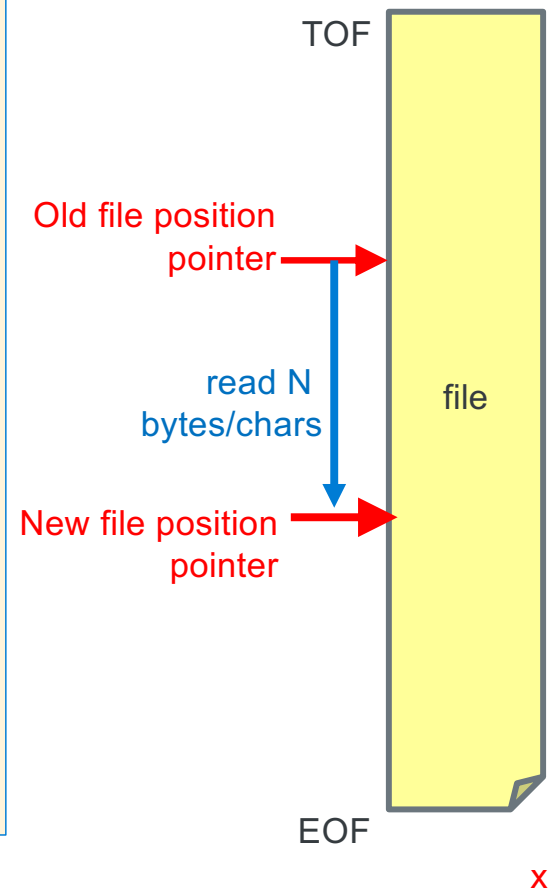
Compiling Multi-File Programs (assembly steps not shown)



1. compile each .c file independently to a .o object file this requires you use the `-c` flag to gcc to only compile and assemble and NOT to call the linker yet
`gcc -Wall -Wextra -Werror -c file.c # creates file.o`
`gcc -Wall -Wextra -Werror -c myprog.c # creates myprog.o`
2. link all the .o objects files and libraries (aggregation of multiple .o files) to produce an executable file (gcc calls ld, the linker)
 - The .o's in the libraries are automatically linked in as needed to produce an executable file`gcc -Wall -Wextra myprog.o file.o -o myprog`

C standard I/O Library (stdio) File I/O File Position Pointer and EOF

- Read/write functions in the standard I/O library *advances* the **file position pointer** from the **top of a file** (before the 1st byte if any) *towards* the **end of the file** after each call to a read/write function
 - **Side effect of call:** file position pointer moves towards the **end of file** by number of bytes read/written
- **standard I/O File position pointer** indicates where in the file (byte distance from the top of the file) the next read/write I/O will occur
- Performing a sequence of read/write operations (without using any other stdio functions to move the file pointer between the read/write calls) performs what is called **Sequential I/O** (sequential read & sequential write)
- EOF condition state may be set after a **read operation**
 - After the last byte is read in a file, additional reads results in a **function return value** of EOF
 - **EOF signals** no more data is available to be read
 - **EOF is NOT a character in the file**, but a condition state on the stream
 - EOF is usually a **#define EOF -1** macro located in the file stdio.h (later in course)



C Library Function API : Simple Character I/O – Used in PA3

Operation	Usage Examples
Write a char	<pre>int status; int c; status = putchar(c);</pre> <i>/* Writes to screen stdout */</i>
Read a char	<pre>int c; c = getchar();</pre> <i>/* Reads from keyboard stdin */</i>

```
#include <stdio.h> // import the public interface
```

```
int putchar(int c);
```

- writes c (demoted to a char) to **stdout**
- **returns** either: **c** on success **OR EOF** (a macro often defined as -1) on failure
- see % man 3 putchar

```
int getchar(void);
```

- **returns** the next input character (if present) **promoted to an int** read from **stdin**
- see % man 3 getchar
- Make sure you use **int variables** with **putchar()** and **getchar()**
- **Both functions return an int** because they must be able to return both valid chars and indicate the **EOF condition** (-1) which is outside the range of valid characters

Why is character I/O using an int?

Answer: Needs to indicate an EOF (-1) condition that is not a valid char

Character I/O (Also the Primary loop in PA3)

```
// copy stdin to stdout one char at a time
```

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
int main(void)
```

```
{
```

```
    int c;
```

```
    while ((c = getchar()) != EOF) {  
        (void)putchar(c);    // ignore return value
```

```
    }
```

```
    return EXIT_SUCCESS;
```

```
}
```

Always check return code to handle EOF
EOF is a macro integer in stdio.h

Always check return codes unless you do not need it

Sometimes you may see a (void) cast which indicates **ignoring the return value is deliberate** this is often required by many coding standards

```
% ./a.out
```

```
thIS is a TeSt
```

```
thIS is a TeSt
```

```
^d
```

```
%
```

```
%. /a.out < a > b
```

Typed on keyboard

Printed by program

Typed on keyboard

Copies file a to file b


Make sure you use int variable with getchar() and putchar()!

stdio File I/O – Working with a Keyboard

PROCESS

```
010000111001
0100001110011111000111
000111000111
```

KEYBOARD

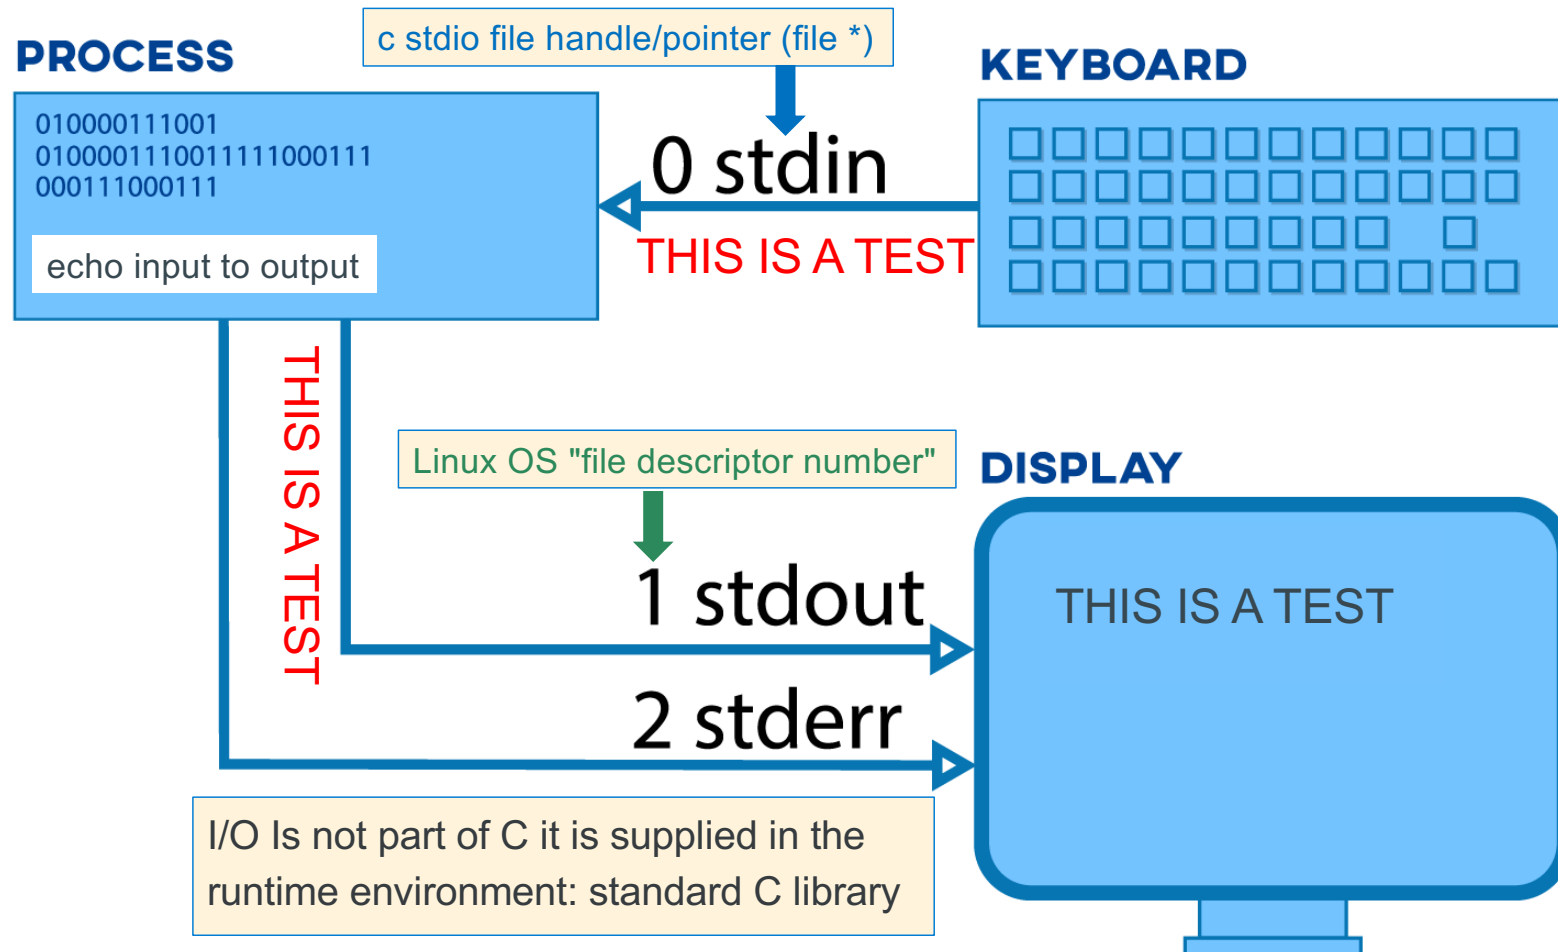


← 0 stdin

How do I
signal EOF?

- How can you have an **EOF** when reading from a keyboard?
- stdio I/O library functions **designed** to work primarily on **files**
 - With keyboard devices the semantics of *file operations* needs to be "*simulated*"
- **Example:** when a program (or a shell) is reading the keyboard and is blocked waiting for input it is waiting for you to type a line
 - **This is NOT an EOF condition**
- To **set** an **EOF condition from the keyboard**, **type** on an input line all by itself:
two key combination (ctrl key and the d key at same time), followed by a return/enter
ctrl-d *often shown in slides etc. as ^d*

Linux/Unix Process and Standard I/O (CSE 15L)



C Library Function: Simple Formatted Printing

Task	Example Function Calls
Write formatted data	<pre>int status; status = fprintf(stderr, "%d\n", i); status = printf("%d\n", i); /* Writes to stdout */</pre>

```
#include <stdio.h> // import the public interface
```

```
int fprintf(FILE *file, const char *format, ...);
```

- Write chars to the file identified by **file** (**stdout**, **stderr** are already open)
- Convert values to chars, as directed by **format**
- Return count of chars successfully written
- **Format** is the output specifications enclosed in a "string"
- Returns a negative value if an error occurs

```
int printf(const char *format, ...); // *format - later in course
```

- Equivalent to `fprintf(stdout, format, ...);`
- Type `% man 3 printf` for more information on **format**

Some Formatted Output Conversion Examples

- Conversion specifications example
 - **%d** conversion specifier for **int** variables
 - **%c** conversion specifier for **char** variables
 - many more conversion specifiers (online manual: `% man printf` and the textbooks)

```
int i = 10;
char z = 'i';
char a[] = " Hello\n";

printf("%c = %d,%s", z, i, a); // write to stdout
fprintf(stderr, "This is an error message to stderr\n");
```

- Output

```
i = 10, Hello
This is an error message to stderr
```

Conditional Statements (if, while, do...while, for)

- C conditional test expressions: 0 (NULL) is FALSE, any non-0 value is TRUE
- C comparison operators (==, !=, >, etc.) evaluate to either 0 (false) or 1 (true)
- Legal in Java and in C:

```
i = 0;  
if (i == 5)  
    statement1;  
else  
    statement2;
```

Which statement is executed after the if statement test?

- Illegal in Java, but legal in C (often a typo!):

```
i = 0;  
if (i = 5)  
    statement1;  
else  
    statement2;
```

Assignment operators evaluate to the value that is assigned, so.... Which statement is executed after the if statement test?

Program Flow – Short Circuit or Minimal Evaluation

- In evaluation of conditional guard expressions, C uses what is called **short circuit** or **minimal** evaluation

```
if ((x == 5) || (y > 3)) // if x == 5 then y > 3 is not evaluated
```



- Each** expression argument is evaluated **in sequence** from **left to right** including any **side effects** (modified using parenthesis), **before** (optionally) evaluating the next expression argument
- If after evaluating an argument, the **value of the entire expression can be determined**, then the **remaining arguments are NOT evaluated** (*for performance*)

Program Flow – Short Circuit or Minimal Evaluation

```
if ((a != 0) && func(b))    // if a is 0, func(b) is not called
    do_something();
```

```
// if ((x > 0) && (c == 'Q')) evaluates to non zero (true)
// then (b == 3) is not tested

while (((x > 0) && (c == 'Q')) || (b == 3)) { // c short circuit
    x = x / 2;
    if (x == 0) {
        return 0;
    }
}
```

Be Careful with the comma , sequence operator

- Sequence Operator ,
expr1, expr2
- Evaluates *expr1* first and then *expr2* evaluates to or returns *expr2*

```
for (i = 0, j = 0; i < 10; i++, j++)  
...
```

- Unexpected results with , operator (some compilers will warn)

```
i = 64, 323;           // i = 64 (assigns first)  
i = (64, 323);        // i = 323 (value of expression)
```

Review: Binary Numbering

- Binary is base 2
 - *adjective*: being in a state of one of two **mutually exclusive** conditions such as **on** or off, **true** or **false**, **molten** or **frozen**, **presence** or **absence** of a signal
 - From Late Latin *bīnārius* (“consisting of two”)
- **Two** symbols:
0 1
- Numbers in C that start with **0b** are binary
- Example: What is **0b110** in base 10?
 - $0b110 = 110_2 = (1 \times 2^2) + (1 \times 2^1) + (0 \times 2^0) = 6_{10}$
- A **bit** is a single binary digit
- A **byte** is an 8-bit value



powers of two

$$\text{Unsigned binary Number} = \sum_{i=0}^{n-1} b_i \times 2^i = b_{n-1}2^{N-1} + b_{n-2}2^{N-2} + \dots + b_12^1 + b_02^0$$

Review: Hexadecimal Numbering

- hexadecimal is base 16
 - From “hexa” (Ancient Greek ἑξά-) \Rightarrow six
 - and from “decem” (Latin) \Rightarrow ten

- Sixteen** symbols

0 1 2 3 4 5 6 7 8 9 a b c d e f



- Numbers in C that start with **0x** are hexadecimal numbers
 - $16_{10} = 0x10_{16}$
- Example: What is **0xa5** in base 10?
 - $0xa5 = a5_{16} = (10 \times 16^1) + (5 \times 16^0) = 165_{10}$
- Hexadecimal numbers are **very commonly used** in programming to express binary values
 - Imagine the difficulty in correctly expressing a 64-bit binary value in your code

$$\text{Unsigned Hex Number} = \sum_{i=0}^{n-1} b_i \times 16^i = b_{n-1}16^{n-1} + b_{n-2}16^{n-2} + \dots + b_116^1 + b_016^0$$

Binary <---> Hexadecimal Equivalences

- **Hex → Binary:** $16^1 = 2^4$ 1 digit hex = 4 digits binary
 1. Replace hex digits with binary digits
 2. Drop **leading zeros**
 - Example: 0x2d to binary
 - 0x2 is 0b0010, 0xd is 0b1101
 - Drop two leading zeros, answer is 0b101101
- **Binary → Hex:** $2^4 = 16^1$
 1. **Pad** with enough **leading zeros** until number of digits is a multiple of 4
 2. **Replace** each **group of 4** with the **HEX equivalent**
 - Example: 0b101101
 - **Pad on the left** to: 0b 0010 1101
 - Replace to get: 0x2d

Number Base Overview (as written in C)

- Decimal is base 10 and Hexadecimal is base 16,
- **Hex digits** have 16 values 0 - 9 a - f (written in C as 0x0 – 0xf)
- No standard prefix in C for binary (most use **hex**)
 - gcc (compiler) allows **0b** prefix **others might not**

Hex digit	0x0	0x1	0x2	0x3	0x4	0x5	0x6	0x7
Decimal value	0	1	2	3	4	5	6	7
Binary value	0b0000	0b0001	0b0010	0b0011	0b0100	0b0101	0b0110	0b0111

Hex digit	0x8	0x9	0xa	0xb	0xc	0xd	0xe	0xf
Decimal value	8	9	10	11	12	13	14	15
Binary value	0b1000	0b1001	0b1010	0b1011	0b1100	0b1101	0b1110	0b1111

Hex to Binary (group 4 bits per digit from the right)

- Each Hex digit is 4 bits in base 2 $16^1 = 2^4$

0x f a 5 3

1111 1010 0101 0011

0b1111101001010011

↑ binary start with a 0b in C

Binary to Hex (group 4 bits per digit from the right)

- 4 binary bits is one Hex digit $2^4 = 16^1$

0b 0110 1010 0011 1111

6 a 3 f

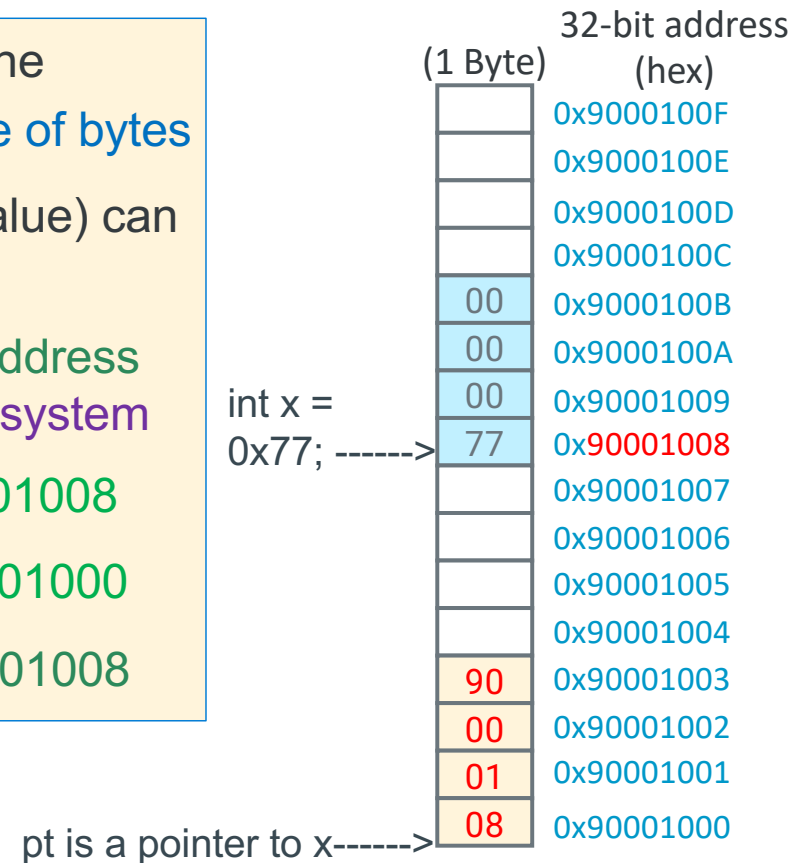
0x6a3f

hex start with 0x in C




Address and Pointers

- An **address** refers to a location in memory, the **lowest** or **first byte** in a **contiguous sequence of bytes**
- A **pointer** is a **variable** whose **contents** (or value) can be properly used as an **address**
 - The **value in a pointer** *should* be a **valid address allocated to the process** by the **operating system**
- The **variable x** is at **memory address 0x90001008**
- The **variable pt** is at **memory location 0x90001000**
- The **contents** of **pt** is the **address of x 0x90001008**



Variables in Memory: Size and Address

- The number of contiguous bytes a variable uses is based on the *type* of the variable
 - Different variable types require different numbers of contiguous bytes
- **Variable names** map to a starting address in memory
- **Example Below:** Variables all starting at address 0x80, each box is a byte



Variables: Size

- Integer types

- `char`, `int`

- Floating Point

- `float`, `double`

- Modifiers for each base type

- `short` [int]
- `long` [int, double]
- `signed` [char, int]
- `unsigned` [char, int]
- `const`: variable read only

- char type

- One byte in a byte addressable memory
- **Signed** vs **Unsigned** Char implementations
- **Be careful** `char` is unsigned on arm and signed on other HW like intel

C Data Type	AArch-32 contiguous Bytes	AArch-64 contiguous Bytes	printf specification
char (arm unsigned)	1	1	%c
short int	2	2	%hd
unsigned short int	2	2	%hu
int	4	4	%d / %i
unsigned int	4	4	%u
long int	4	8	%ld
long long int	8	8	%lld
float	4	4	%f
double	8	8	%lf
long double	8	16	%Lf
pointer *	4	8	%p

size of a pointer is the word size

sizeof(): Variable Size (number of bytes) Operator

```
#include <stddef.h>
/* size_t type may vary by system but is always unsigned */
```

sizeof() operator returns a value of type **size_t**:

the number of bytes used to store a variable or variable type

```
size_t size = sizeof(variable_type);
```

or

```
size_t size = sizeof(variable_name); // preferred!
```

- The argument to sizeof() is often an expression:

```
size = sizeof(int * 10);
```

- reads as:

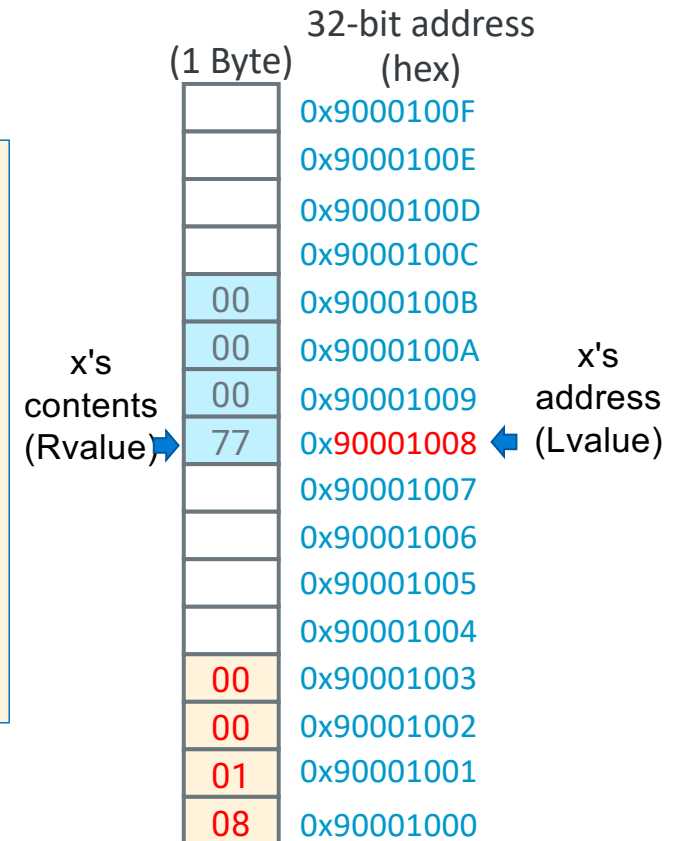
- number of bytes required to store **10 integers (an array of [10])**

Memory Addresses & Memory Content

Variable names in a C statement evaluation


```
x = x + 1;    // Lvalue = Rvalue
```

- **Lvalue:** when on the left side (Lside or Left value) of the = sign
 - address where it is stored in memory – a constant
 - Address assigned to a variable cannot be changed at runtime
 - Does not require a memory read
- **Rvalue:** when on the right side (Rside or Right value) of an = sign
 - contents or value stored in the variable (at its memory address)
 - requires a memory read to obtain contents



Memory Addresses & Memory Content

`y = 42;` One memory write required
`x = y;` One memory read required
 // Lvalue = Rvalue



- `x` on left side (**Lside**) of the assignment operator = evaluates to:
 - **Address** of the memory assigned to the `x` – this is `x`'s **Lvalue**
- `y` on right side (**Rside**) of the assignment operator = evaluates to:
 - **Contents** of the memory assigned to the variable `y` (type determines length – number of bytes) - this is `y`'s **Rvalue**
- So, `x = y;` is:

Read memory at `y` (**Rvalue**); write it to memory at `x`'s address (**Lvalue**)

Introduction: Address Operator: &

- Unary **address operator (&)** produces the **address** of where an **identifier** is in memory
 - Assigned address to **g**
- **Example** this might print:
value of g is: 42
address of g is: 0x71a0a0
(the address will vary)
- **Tip:** printf() format specifier to display an address/pointer (in hex) is "%p"

```
int main(void)
{
    int g = 42;

    printf("value of g is: %d\n", g);
    printf("address of g is: %p\n", &g);
    return EXIT_SUCCESS;
}
```

Introduction: Address Operator: &

- Requirement: **identifier must have a Lvalue**
 - Cannot be used with **constants** (e.g., 12) or **expressions** (e.g., x + y)
 - Example: **&12** does not have an *Lvalue*,
 - so, **&12** is not a legal expression
- How can I get an **address for use on the Rside**?
 - **&var** (any variable identifier or name)
 - **function_name** (name of a **function**, not func());
 - **&func_name** is equivalent
 - **array_name** (name of the **array** like array_name[5]);
 - **&array_name** is equivalent

Pointer Variables

- In C, there is a *variable type* for **storing an address**: a *pointer*
 - **Contents** of a pointer is an unsigned (positive numbers) memory address

```
type *name; // defines a pointer; name contains address of a variable of type
```

- A *pointer* is defined by placing a *star* (or *asterisk*) (*) before the identifier (name)
- You also must specify the *type of variable* to which the pointer points
- **Pointers are typed!** Why?
 - The compiler needs to know the *size* (sizeof()) of the data **you are pointing at** (number of consecutive bytes to access) to use the pointer
- When the **Rside** of a **variable** contains a **memory address**, (it **evaluates** to an **address**) the variable is called a **pointer variable**

Pointer Variables - 2

- A pointer cannot point at itself, why?

```
int *p = &p; /* is not legal - type mismatch */
```

- `p` is defined as `(int *)`, a pointer to an `int`, but
- the type of `&p` is `(int **)`, a pointer to a pointer to an `int`
- Pointer variables all use the **same amount of memory** no matter what they point at

```
int *iptr;  
char *cptr;  
  
printf("iptr(%u) cptr(%u)\n", sizeof(iptr), sizeof(cptr));
```

- Above prints on a 32-raspberry pi

```
% ./example  
iptr(4) cptr(4)
```

Defining Pointer Variables

- Assigning a value to a pointer:

```
int *p = &i;  /* p points at i (assign address i to p) */
```

- Is the same as writing the following definition and assignment statements



```
int *p;      /* p is defined (not initialized) */  
p = &i;      /* p points at i (assign address i to p) */
```

- The ***** is part of the definition of **p** and is not part of the variable name
 - The name of the variable is simply **p**, not ***p**
- C mostly ignores whitespace, so these three definitions are equivalent

```
int  *p = &i;    /* Style A */  
int *  p = &i;    /* Style B */  
int*  p = &i;    /* Style C */
```

Using Pointer Variables and the Address Operator & - 1

```
int i = 42;  
int *p; /* p contains the address of an integer */  
p = &i; /* p "points at" i (assign address of i to p) */
```



- **Recommended:** be careful when defining multiple pointers on the same line:

`int *p1, p2;` is not the same as: `int *p1, *p2;`


Use instead:

```
int *p1;  
int *p2;
```


Using Pointer Variables and the Address Operator & - 2

- As with any variable, its value can be changed

`p = &j;` */* p now points at j */*




`p = &i;` */* p now points at i */*



Indirection (or dereference) Operator: *

- The **indirection operator** (*) or the *dereference operator to a variable* is the **inverse** of the *address operator* (&)
- **address operator** (&) can be thought of as:

"get the address of this box"




- **indirection operator** (*) can be thought of as:
- "follow the arrow to the next box and get its contents"*
- **Indirection operator causes an additional read to occur**, when on either the Rside or Lside of a statement

Rside Indirection (or dereference) Operator: *

- Performs the following steps when the * is on the Rside:
 1. read the contents of the variable to get an address
 2. read and return the contents at that address
 - (requires two reads of memory on the Rside)

```
z = *x; // copy the contents of memory pointed at by x to z
```

Two reads here
(1) read to get an address
(2) read the address to get the value



Rside Indirection (or dereference) Operator: *

*Contents of **p** is the address of **i***
(p points at i)


```
int i = 42;  
int *p;  
p = &i;
```

No reads here

```
printf("*p is %d\n", *p);
```

```
% ./a.out  
*p is 42
```

Two reads here
(1) read to get an address
(2) read the address to get the value




Lside Indirection Operator

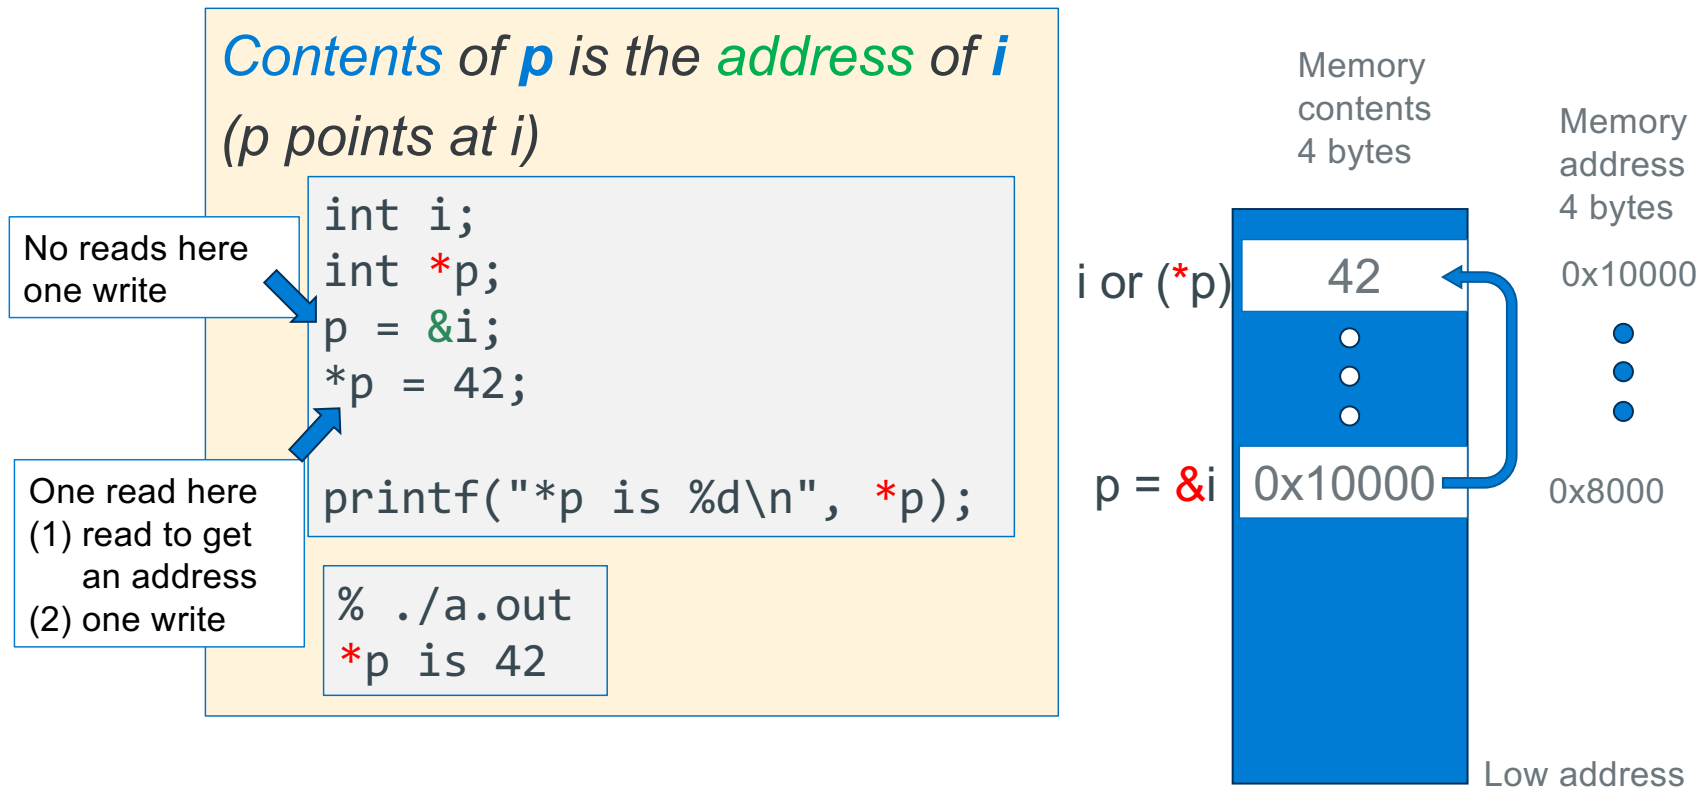
Performs the following steps when the ***** is on the Lside:

1. **read** the **contents** of the **variable** to get **an address**
2. **write** the evaluation of the Rside expression to that address
 - (requires **one read of memory and one write of memory on the Lside**)

```
*p = x; // copy the value of x to the memory pointed at by p
```





Left Side Indirection (or dereference) Operator: *




Each use of a * operator results in one additional read -1

RULE: Each * when used as a dereference operator in a **statement** (either **Lside** or **Rside**) generates an additional read

```
int z = 2, y = 1;  
z = y; // one read
```





```
int z = 2, y = 1;  
int *p;  
p = &y;  
z = *x; // two reads on rside
```



Aside: `z = *(&x);` // same as `z = x`

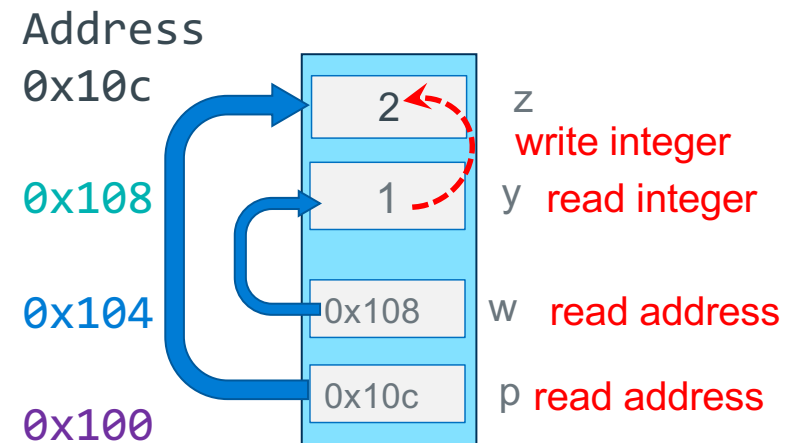

Each use of a * operator results in one additional read -2

```
int z = 2, y = 1;  
int *x;  
p = &z;  
*p = y;    // one read on lside
```



Each use of a * operator results in one additional read -2

```
int z = 2, y = 1;  
int *w;  
int *p;  
p = &z;  
w = &y;  
*p = *w;
```



Pointer to Pointers (Double Indirection)

- Define a pointer to a pointer (p2 below)


```
int i = 2;  
int *p1;  
int **p2;  
p1 = &i;  
p2 = &p1;  
printf("%d\n", (**p2) * (**p2));
```

- C allows any number of pointer indirections
 - more than two levels is very uncommon in real applications as it reduces readability and generates a lot of memory reads
- RULE (important):** number of ***** in the definition tells you how many reads it takes to get to the base type

#reads to base type = number ***** + 1



- Example:

```
int **p2;    // requires 3 reads to get to the int
```



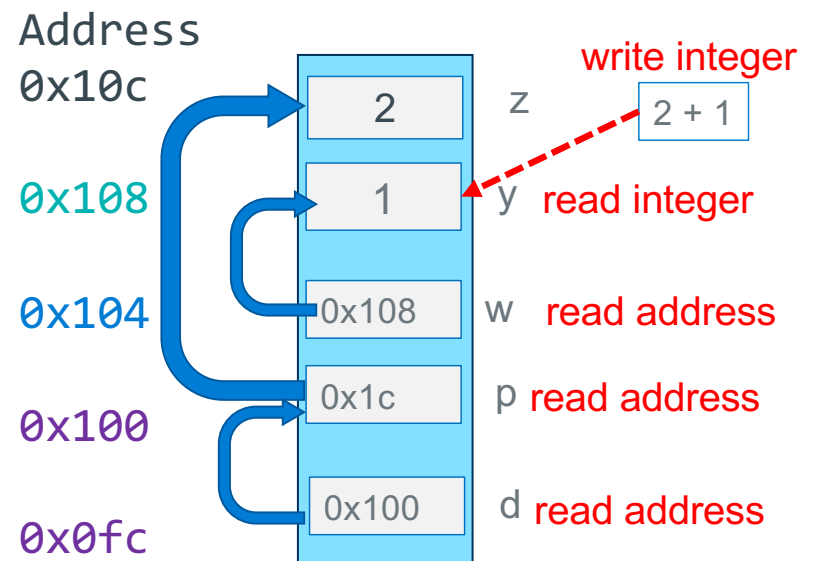
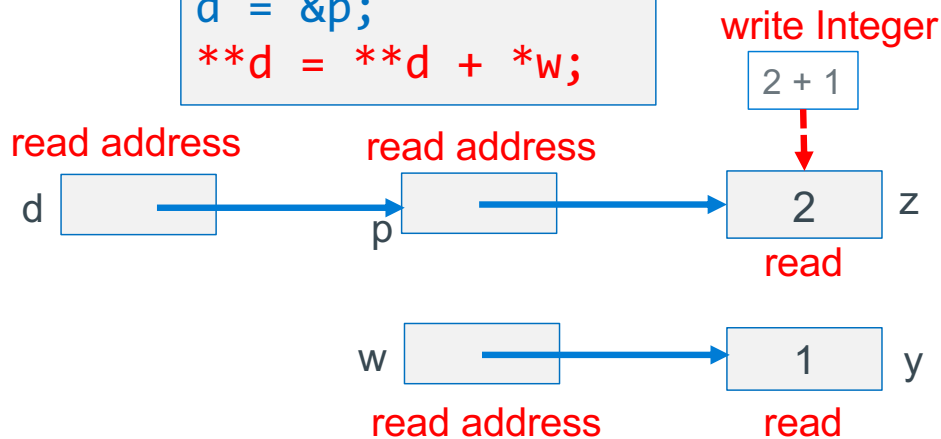
Double Indirection

```
int z = 2, y = 1;  
int *w;  
int *p;  
int **d;  
p = &z;  
w = &y;  
d = &p;  
**d = *w;
```



Double Indirection

```
int z = 2, y = 1;
int *w;
int *p;
int **d;
p = &z;
w = &y;
d = &p;
**d = **d + *w;
```



Important Observe


**d on Lside is two reads
 **d on Rside is three reads

What is Aliasing?

- **Two or more** variables are **aliases** of each other when they all reference the same memory (so different names, same memory location)
- **Example:** When one pointer is copied to another pointer it *creates an alias*
- **Side effect:** Changing one variables value (content) changes the value for other variables
 - **Multiple variables** all **read and write** the **same** memory location
 - Aliases occur either by **accident** (coding errors) or **deliberate** (careful: readability)

```
int i = 5;  
int *p;  
int *q;  
p = &i;
```


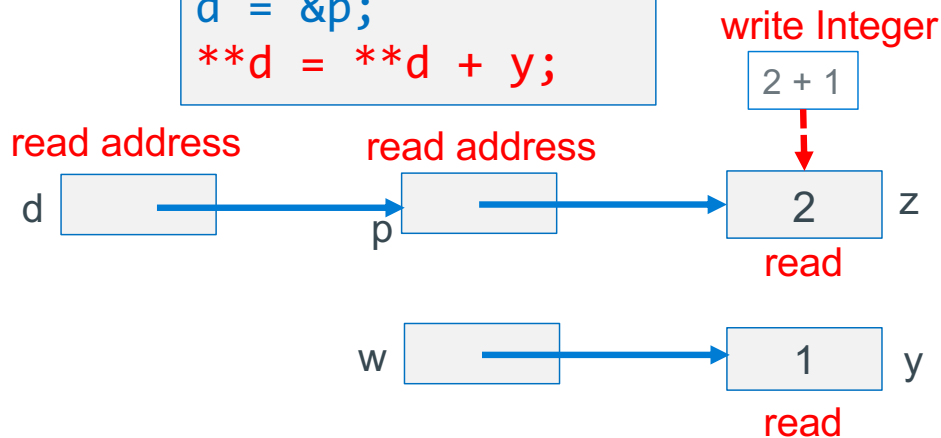
```
q = p;    // *p & *q are aliases  
*q = 4;   // changes i
```



Result *p, *q and i all have the value of 4

Double Indirection Aliases

```
int z = 2, y = 1;
int *w;
int *p;
int **d;
p = &z;
w = &y;
d = &p;
**d = **d + y;
```



Important Observe

`**d`, `*p` and `y` are all aliases

The NULL Constant and Pointers

- **NULL is a constant** that **evaluates to zero (0)**
- You **assign a pointer variable to contain NULL** to **indicate that the pointer does not point at anything**
- A **pointer variable** with a **value of NULL** is called a “**NULL pointer**” (invalid address!)
- Memory location 0 (address is 0) is not a valid memory address in any C program
- Dereferencing NULL at runtime will cause a program fault (segmentation fault)!

```
p = NULL;  
i = *p;           /* segmentation fault! */  
*(int *)900000 = 25; /* cast 900000 to a pointer */  
                  /* if writeable address space, it works */  
                  /* that memory location just changed */
```

Using the NULL Pointer

- Many functions return NULL to indicate an error has occurred

```
/* these are all equivalent */  
int *p = NULL;  
int *p = (int *)0;    // cast 0 to a pointer type  
int *p = (void *)0;   // automatically gets converted to the correct type
```

- NULL is considered “false” when used in a Boolean context
 - **Remember: false expressions** in C are defined to be zero or NULL
- The following two are equivalent (the second one is preferred for readability):

```
if (p) ...  
if (p != NULL) ...
```

Defining Arrays

Definition: `type name[count]`

- **"Compound"** data type where each value in an array is an element of `type`
- Allocates **name** with a *fixed* `count` array elements of type `type`
- Allocates $(\text{count} * \text{sizeof}(\text{type}))$ bytes of *contiguous memory*
- Common usage is to specify a compile-time constant for `count`

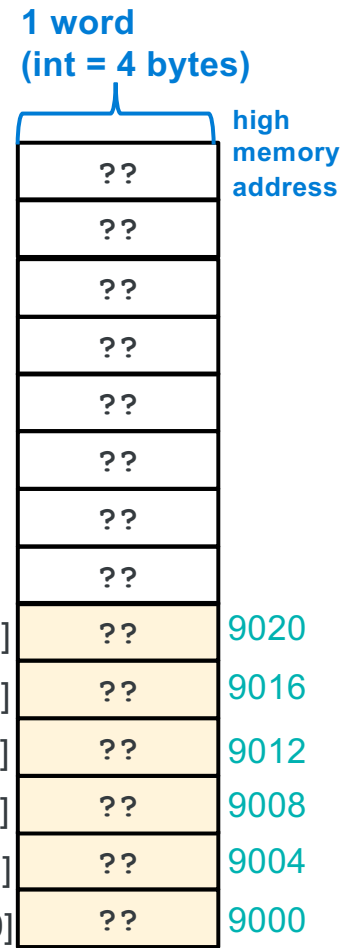
```
#define BSZ 6  
int b[BSZ];
```

BSZ is a macro replaced by the C preprocessor at compile time

- Array **names are constants** and *cannot be assigned* (the name cannot appear on the Lside by itself)

```
a = b;           // invalid does not copy the array  
                // copy arrays element by element
```

```
int b[6];
```



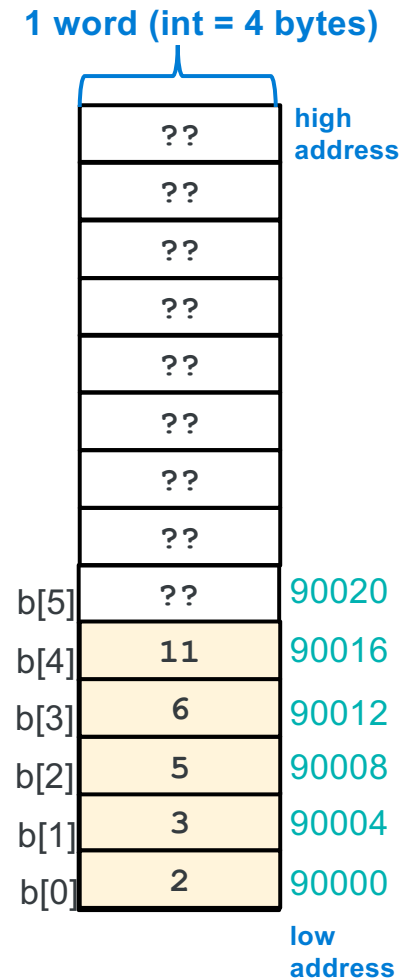
Array Initialization

- Initialization: `type name[count] = {val0,...,valN};`
 - `{ }` (*optional*) initialization list can only be used at **time of definition**
 - If no `count` supplied, `count` is determined by compiler using the number of array initializers
no initialization values given; then elements are initialized to 0
 - `int block[20] = {};` //only works with constant size arrays
 - defines an **array of 20 integers** each element filled with zeros
 - Performance comment: do not zero automatic arrays unless really needed!
 - When a `count` is given:
 - extra initialization values are **ignored**
 - missing initialization values are set to **zero**

```
int block[5] = {2, 3, 5, 6, 11, 13};
```

not needed and if used **may** truncate initialization list

6 initialization values given, **only 5 are used**



X

Accessing Arrays Using Indexing

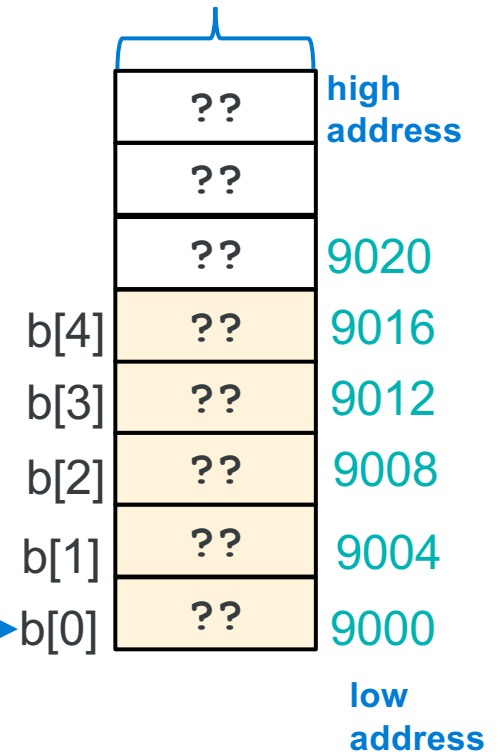
- **name** [**index**] selects the **index** element of the array
 - index **should be** unsigned
 - Elements range from: 0 to count – 1 (int x[count];)
- **name** [**index**] can be used as an assignment target or as a value in an expression
- Array name (by itself with no []) on the **Rside** evaluates to the address of the first element of the array

```
int a[5];  
int b[5];
```

```
int b[5];  
int *p = b;
```

p 9000

1 word
(int = 4 bytes)



How many elements are in an array?

- The number of elements of space allocated to an array (called **element count**) and indirectly the total size in bytes of an array is not stored anywhere!!!!!!
- An **array name** is just the **address of the first element in a block of contiguous memory**
 - So an array does not know its own size!

```
#define SZ 6
int block[SZ];      // you specify the array has SZ elements
int indx;           // use when SZ is defined

for (indx = 0; indx < SZ; indx++)
    block[indx] = 0;
```

```
int b[6];
```

1 word
(int = 4 bytes)

high
memory
address

	??	
	??	
	??	
	??	
	??	
	??	
	??	
	??	
b[5]	??	90020
b[4]	??	90016
b[3]	??	90012
b[2]	??	90008
b[1]	??	90004
b[0]	??	90000

Determining Element Count for a compiler calculated array

- Programmatically determining the element count in a compiler calculated array

`sizeof(array) / sizeof(of just one element in the array)`

- `sizeof(array)` only works when used in the SAME **scope** as where the array variable was **defined**

```
#include <stddef.h>

int block[] = {2, 3, 5, 6, 11, 13};    // automatic: compiler calculates array size

int cnt = (int)(sizeof(block) / sizeof(block[0])); // in this case cnt = 6

for (int indx = 0; indx < cnt; indx++)
    block[indx] = 0;
```

Pointers and Arrays - 1

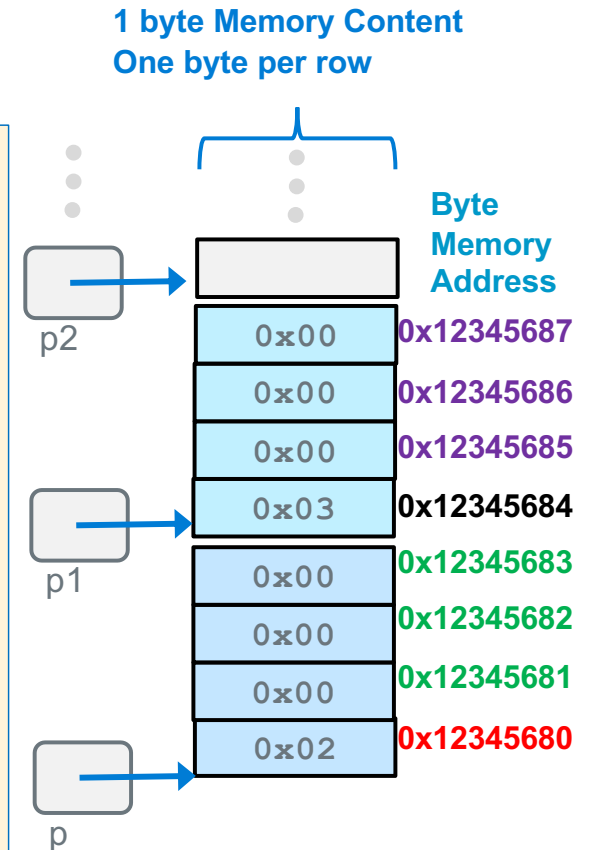
- A few slides back we stated: **Array name** (by itself) on the Rside evaluates to the **address of the first element of the array**

```
int buf[] = {2, 3, 5, 6, 11};
```

- Array indexing syntax (`[]`) an operator that performs *pointer arithmetic*
- buf** and **&buf[0]** on the **Rside** are **equivalent**, **both evaluate** to the address of the first array element

```
int *p = buf;           // or int *p = &buf[0];
int *p1 = &buf[1];
int *p2 = &buf[2];
int *p3 = &buf[3];

*p = *p + 10;
*p1 = *p1 + 10;         // {12, 13, 5, 6, 11}
```



Pointers and Arrays - 2


When `p` is a pointer, the actual value of `(p+1)` depends on the type that pointer `p` points at

- `(p+1)` adds `1 x sizeof(what p points at)` bytes to `p`
 - `++p` is equivalent to `p = p + 1`
- Using pointer arithmetic to find array elements:
 - Address of the second element `&buf[1]` is `(buf + 1)`
 - It can be referenced as `*(buf + 1)` or `buf[1]`

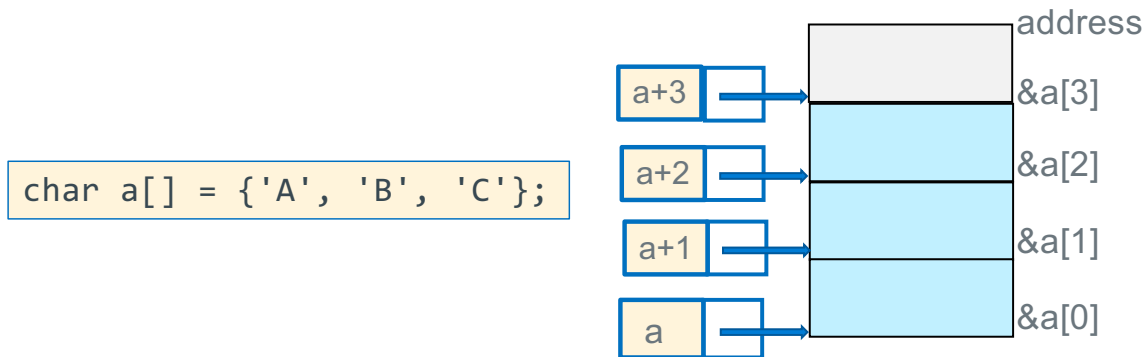
```
int buf[] = {2, 3, 5, 6, 11};
int *p;
p = buf;

*p = *p + 10;
*(p + 1) = *(p + 1) + 10; // {12, 13, 5, 6, 11}
```

1 byte Memory Content
One byte per row



Pointer Arithmetic In Use – C's Performance Focus



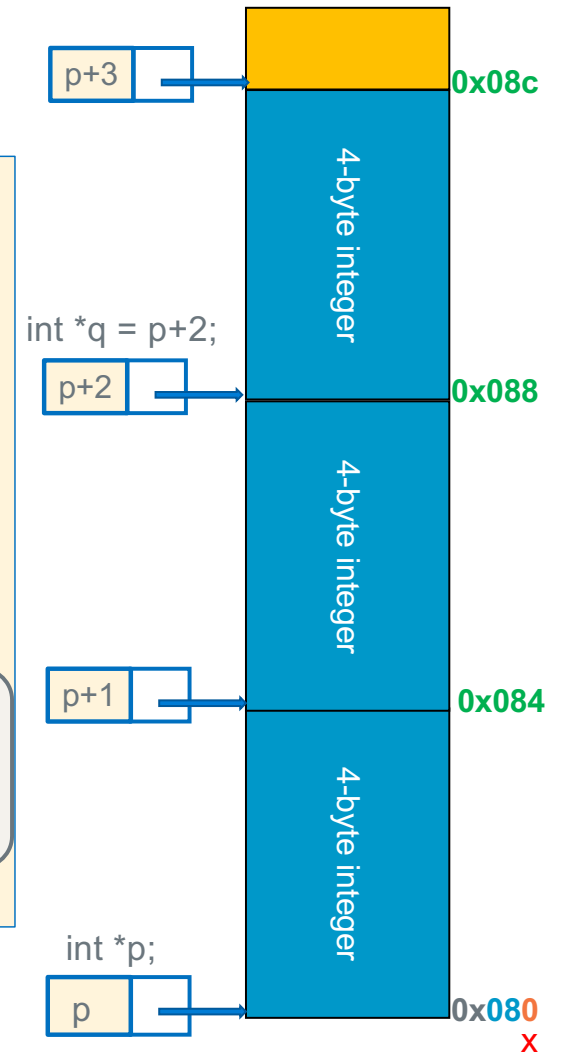
- **Alert!**: C performance focus does not perform any array “bounds checking”
- **Performance by Design**: *bound checking slows down execution of a properly written program*
- **Example**: array `a` of length `i`, C does not verify that `a[j]` or `*(a + j)` is valid (does not check: $0 \leq j < i$)
 - C simply “*translates*” and accesses the memory specified from: `a[j]` to be `*(a + j)` which may be *outside the bounds* of the array
 - OS only **“faults”** for an incorrect access to memory (read-only or not assigned to your process)
 - It does not fault for out of bound indexes or out of scope
- **lack of bound checking** is a **common source of errors and bugs** and is a common criticism of C

Pointer Arithmetic

- You cannot add two pointers (*what is the reason?*)
- A pointer *q* can be subtracted from another pointer *p* when the pointers are the same type – **best done only within arrays!**
- The value of $(p - q)$ is the number of **elements between** the two pointers
 - Using memory address arithmetic (*p* and *q* are both **byte addresses**):

distance in elements = $(p - q) / \text{sizeof}(*p)$

$(p + 3) - p = 3 = (0x08c - 0x080) / 4 = 3$



Pointer Comparisons

- Pointers (**same type**) can be compared with the comparison operators:

<, <=, ==, !=, >=, >

```
int numb[] = {9, 8, 1, 9, 5};
int *end;
int *a;
end = numb + (int) (sizeof(numb)/sizeof(*numb));
a = numb;
while (a < end) // compares two pointers (address)
    /* rest of code including doing an a++ */
```

- Invalid, Undefined, or **risky** pointer arithmetic (some examples)
 - Add, multiply, divide on two pointers
 - Subtract two pointers of different types or pointing at different arrays
 - Compare two pointers of different types
 - Subtract a pointer from an integer

Using Pointers to Traverse an array

```
int x[] = {0xd4c3b2a1, 0xd4c3b200, 0x12345684};
int cnt = (int)(sizeof(x) / sizeof(*x));

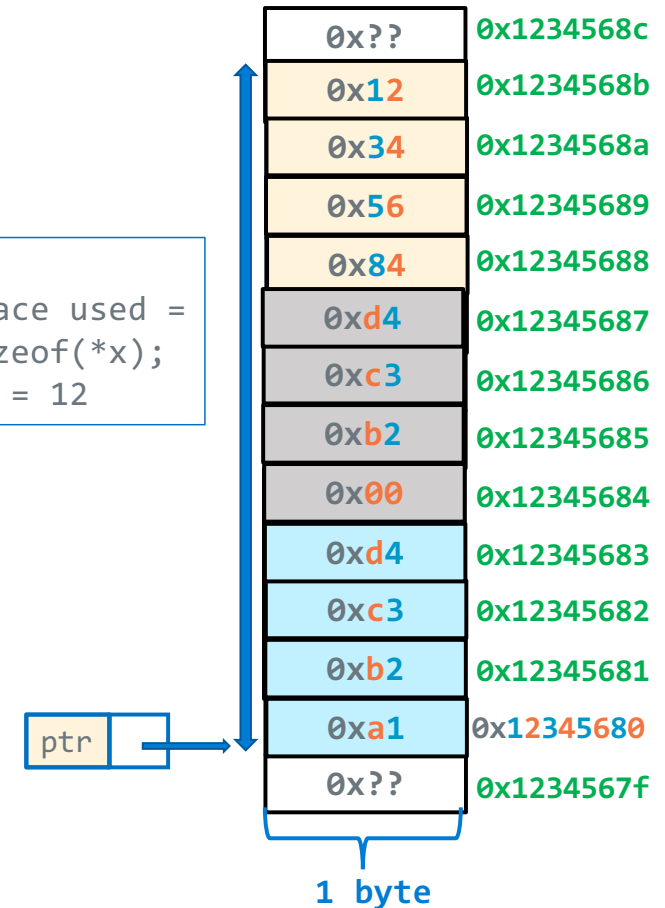
for (int j = 0; j < cnt; j++)
    printf("%#x\n", x[j]);
}
```

```
cnt = 3;
actual space used =
cnt * sizeof(*x);
= 12
```

```
int x[] = {0xd4c3b2a1, 0xd4c3b200, 0x12345684};
int cnt = (int)(sizeof(x) / sizeof(*x));
int *ptr = x;           // or &x[0]

for (int j = 0; j < cnt; j++)
    printf("%#x\n", *(ptr + j));
}
```

Brute force translation to pointers



Fast Ways to Traverse an Array: Use a Limit Pointer

```
int x[] = {0xd4c3b2a1, 0xd4c3b200, 0x12345684};
int cnt = (int)(sizeof(x) / sizeof(*x));
```

```
int *ptr;
int *xptr;
ptr = x; //or &x[0]
xptr = ptr + cnt;
```

xptr is a **loop limit pointer**
it points **1 element past**
the end of the array

```
while (ptr < xptr) {
    printf("%#x\n", *ptr);
    ptr++;
}
```

```
% ./a.out
0xd4c3b2a1
0xd4c3b200
0x12345684
```

cnt = 3;
actual space used =
cnt * sizeof(*x);
= 12

