

Version 2.17

UCSD CSE 30

Computer Organization and Systems Programming

Lecture - 9

Keith Muller

Vax 11/780 1980

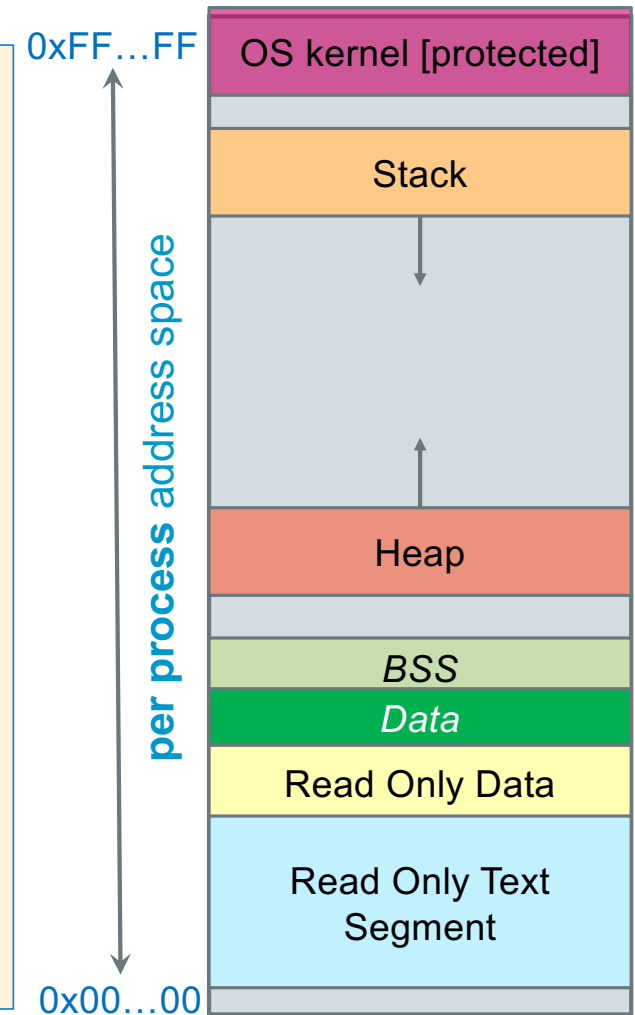




Thursday office hours moved to
Section B ZOOM Midterm Review
Thursday 7:30 PM – 8:30 PM
(My office hours zoom #)
(see canvas for number)

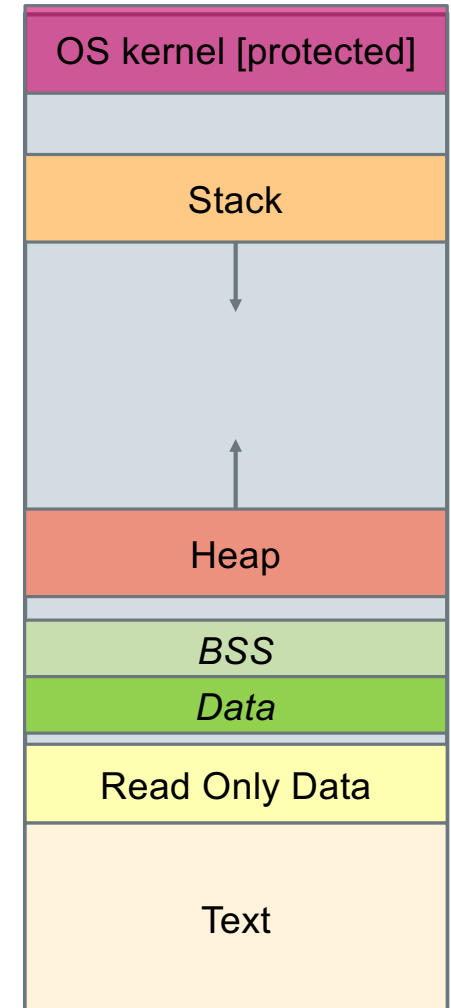
Process Memory Under Linux

- When your **program is running** it has been **loaded into memory** and is **called a process**
- **Stack segment:** Stores **Local** variables
 - Allocated and freed at function call entry & exit
- **Data segment + BSS:** Stores **Global** and **static** variables
 - **Allocated/freed** when the process **starts/exits**
 - **BSS** - Static variables with an implicit initial value
 - **Static Data** - Initialized with an explicit initial value
- **Heap segment:** Stores **dynamically-allocated** variables
 - Allocated with a function call
 - Managed by the stdio library malloc() routines
- **Read Only Data:** Stores **immutable** Literals
- **Text:** Stores your code in machine language + libraries



The Heap Memory Segment

- **Heap**: "pool" of memory that is available to a program
 - Managed by C runtime library and linked to your code; **not managed by the OS**
- Heap memory is **dynamically** *"borrowed"* or *"allocated"* by calling a library function
- When heap memory is no longer needed, it is *"returned"* or *deallocated* for **reuse**
- Heap memory has a lifetime from allocation until it is deallocated
 - Lifetime is independent of the scope it is allocated in (it is like a static variable)
- If too much memory has already been allocated, the library will attempt to borrow additional memory from the OS and will fail, returning a NULL



Heap Dynamic Memory Allocation Library Functions

<code>#include <stdlib.h></code>	args	Clears memory at runtime
<code>void *malloc(...)</code>	<code>size_t size</code>	no
<code>void *calloc(...)</code>	<code>size_t nmem, size_t memsize</code>	yes
<code>void free(...)</code>	<code>void *ptr</code>	no

- **void *** means these library functions return a pointer to **generic (untyped) memory**
 - Be careful with **void *** pointers and **pointer math** as void * points **at untyped memory**
 - When **assigned to a typed pointer**, it "**converts**" it from a **void *** to the **type of the pointer variable**
- **size_t** is an **unsigned integer data type**, the result of a **sizeof()** operator

```
int *ptr = malloc(sizeof(*ptr) * 100); // allocate an array of 100 ints
```

- please read: `% man 3 malloc`

Use of Malloc

```
void *malloc(size_t size)
```

- Returns a pointer to a **contiguous** block of `size` bytes of **uninitialized memory** from the heap
 - The block is **aligned to an 8-byte (arm32) or 16-byte (64-bit arm/intel) boundary**
 - returns **NULL** if allocation failed (also sets `errno`) **always CHECK for NULL RETURN!**
- Blocks returned on **different calls to malloc()** are **not necessarily adjacent**
- `void *` is implicitly cast into **any pointer type on assignment to a pointer variable**


```
char *bufptr;
```

```
/* ALWAYS CHECK THE RETURN VALUE FROM MALLOC!!!! */
```

```
if ((bufptr = malloc(cnt * sizeof(*bufptr))) == NULL) {  
    fprintf(stderr, "Unable to malloc memory");  
    return NULL;  
}
```

```
// allocates a character array with 10 elements
```

bufptr



Calloc()

```
void *calloc(size_t elementCnt, size_t elementSize)
```

- calloc() variant of malloc() but zeros out every byte of memory during program execution before returning a pointer to it (so this has a runtime cost!)
 - First parameter is the number of elements you would like to allocate space for
 - Second parameter is the size of each element

```
// allocate 10-element array of pointers to char, zero filled  
char **arr;  
arr = calloc(10, sizeof(*arr));  
if (arr == NULL) // handle the error
```



- Originally designed to allocate arrays but works for any memory allocation
 - calloc() multiplies the two parameters together for the total size
- calloc() is more expensive at runtime (uses both cpu and memory bandwidth) than malloc() because it must zero out memory it allocates at runtime
- Use calloc() only when you need the buffer to be zero filled prior to FIRST use

Using and Freeing Heap Memory

- void **free**(void *p)
 - Deallocates the **whole block pointed to by p** to the pool of available memory
 - **Freed memory is used in future allocations** (**expect the contents to change after freed**)
 - **Pointer p must be the same address as originally returned** by one of the heap allocation routines malloc(), calloc(), realloc()
 - **Pointer argument to free() is not changed by the call to free()**
- **Defensive programming: set the pointer to NULL** after passing it to free()

```
char *bufptr;

if ((bufptr = malloc(cnt * sizeof(*bufptr))) == NULL) {
    fprintf(stderr, "Unable to malloc memory");
    return NULL;
}
// other code
free(bufptr);           // returns memory to the heap
bufptr = NULL;          // set bufptr to NULL
```


Mis-Use of Free() - 1

- Call `free()`
 - With the same address that you obtained with `malloc()` (or other allocators)
 - It is NOT an error to pass `free()` a pointer to NULL

```
char *bytes;
if ((bytes = malloc(1024 * sizeof(*bytes)) != NULL) {
    /* some code */
    free(bytes + 5); // Program aborts free(): invalid pointer
```

- **Freeing unallocated memory:** Only call `free()` to free memory address that you obtain from one of the allocators (`malloc()`, `calloc()`, etc.)

```
char *ptr = "cse30";
...
/* some code */
free(ptr); // Program aborts free(): invalid pointer
```

Mis-Use of Free() - 2

- Continuing to write to memory after you `free()` it is likely to corrupt the heap or return changed values
 - Later calls to heap routines (`malloc()`, `calloc()`, `strdup()`) may fail or seg fault

```
char *bytes;
if ((bytes = malloc(1024 * sizeof(*bytes)) != NULL) {
    /* some code */
    free(bytes);
    strcpy(bytes, "cse30");    // INVALID! used after free
    .....
}
```

- Double Free:** Freeing allocated memory more than once will cause your program to abort (terminate)

```
char *bytes;
if ((bytes = malloc(1024 * sizeof(*bytes)) != NULL) {
    /* some code */
    free(bytes);
    free(bytes); // Program abort double free detected....
    .....
}
```

More Dangling Pointers: Continuing to use "freed" memory

- Review: **Dangling pointer** points to a memory location that is no longer "valid"
- Really hard to debug as the use of the return pointers may not generate a seg fault

```
char *dangling_freed_heap(void)
{
    char *buff = malloc(BLKSZ * sizeof(*buff));
    ...
    free(buff);    // memory pointed at buf may be reused
    return buff;   // but it is returned to the caller anyway - bad
}
```

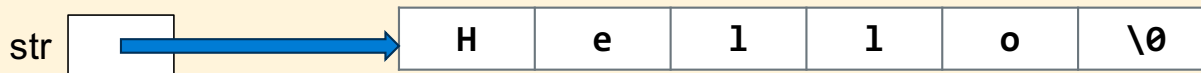
- `dangling_freed_heap()` **may cause** the allocators (`malloc()` and friends) to **seg fault when called later to allocate memory**
 - Why? Because it corrupts data structures the heap code uses to manage the memory pool (it often stores meta-data in the freed memory)

strdup(): Allocate Space and Copy a String

```
char *strdup(char *s);
```

- **strdup** is a function that has a **side effect** of returning a **null-terminated**, heap-allocated string copy of the provided text
- Alternative: **malloc** and copy the string with **strncpy()**;
- The caller is responsible for freeing this memory when no longer needed

```
char *str = strdup("Hello");  
*str = 'h'; // str points at a mutable string  
  
free(str); // caller correctly frees up space allocated by strdup()  
str = NULL;
```



Heap Memory "Leaks"

- A **memory leak** is when you **allocate memory** on the heap, **but never free it**

```
void
leaky_memory (void)
{
    char *buf = malloc(BLKSZ * sizeof(*bytes));
    ...
    /* code that never calls free() to deallocates the memory */
    return; // you lose the address in buf when leaving scope
}
```

- **Best practice:** free up memory **you allocated** when you no longer need it
 - If you keep allocating memory, you may run out of memory in the heap!
- **Memory leaks** may cause **long running programs to fault** when they **exhaust OS memory limits**
- **Valgrind** is a tool for finding memory leaks (not pre-installed in all linux distributions though!)

Valgrind – Finding Buffer Overflows and Memory leaks

```
1 #define SZ 50
2 #include <stdlib.h>
3 int main(void)
4 {
5     char *buf;
6     if ((buf = malloc(SZ * sizeof(*buf))) == NULL)
7         return EXIT_FAILURE;
8     *(buf + SZ) = 'A';
9     // free(buf);
10    return EXIT_SUCCESS;
11 }
```

```
% valgrind -q --leak-check=full --leak-resolution=med -s ./valgexample
==651== Invalid write of size 1
==651==    at 0x10444: main (valg.c:8)
==651== Address 0x49d305a is 0 bytes after a block of size 50 alloc'd
==651==    at 0x484A760: malloc (vg_replace_malloc.c:381)
==651==    by 0x1041B: main (valg.c:6)
==651==
==651== 50 bytes in 1 blocks are definitely lost in loss record 1 of 1
==651==    at 0x484A760: malloc (vg_replace_malloc.c:381)
==651==    by 0x1041B: main (valg.c:6)
==651==
==651== ERROR SUMMARY: 2 errors from 2 contexts (suppressed: 0 from 0)
```

Writing outside of allocated
buffer space

Memory not freed

Introduction to Structs – An Aggregate Data Type

- **Structs** are a **collection (or aggregation) of values** grouped under a **single name**
 - Each **variable** in a **struct** is called a **member** (sometimes **field** is used)
 - Each **member** is identified with a **name**
 - Each **member** can be (and quite often are) **different types**, include other structs
 - Like a Java class, but no associated methods or constructors with a struct
- Structure definition **does not** define a variable instance, nor does it allocate memory:
 - It creates a **new variable type** uniquely identified by its **tagname**:
"struct tagname" includes the **keyword struct** and the **tagname** for this type


Easy to forget
semicolon!

```
struct tagname {  
    type1 member1;  
    ...  
    typeN memberN;  
};
```

```
struct vehicle {  
    char *state;  
    char *plate;  
    char *make;  
    int year;  
};
```

Struct Variable Definitions

```
struct vehicle {  
    char *state;  
    char *plate;  
    char *make;  
    int year;  
};  
struct vehicle name1;
```



- Variable definitions like any other data type:

```
struct vehicle name1, *pn, ar[3];
```

↑
type: "struct vehicle"

↑
single variable instance

↑
pointer

↑
array

Accessing members of a struct

- Like arrays, struct variables are aggregated contiguous objects in memory
- The `.` structure operator *"selects"* the specified field or member

```
struct date {           // defining struct type
    int month;          // member month
    int day;            // member date
};
```

day	
month	

struct date type definition

```
struct date bday; // define a struct instance
bday.month = 1;
bday.day = 24;

// alternative initializer syntax
struct date new_years_eve = {12, 31};
struct date final = {.day= 24, .month= 1};
```

day	24
month	1

bday definition

Accessing members of a struct with pointers

```
struct date {      // defining struct type
    int month;     // member month
    int day;       // member date
};
```

day	
month	

- Define a *pointer* to a struct

```
struct date *ptr = &bday;
```

- Two ways to reference a member via a struct pointer (. is higher precedence than *):

1. Use * and . operators: (*ptr).month = 11;

2. Use -> operator for shorthand: ptr->month = 11;

Operator	Description	Associativity
() [] . -> ++ --	Parentheses or function call Brackets or array subscript Dot or Member selection operator Arrow operator Postfix increment/decrement	left to right
++ -- + - ! ~ (type) * & sizeof	Prefix increment/decrement Unary plus and minus not operator and bitwise complement type cast Indirection or dereference operator Address of operator Determine size in bytes	right to left

Accessing members of a struct

```
struct date {           // defining struct type
    int month;          // member month
    int day;            // member date
};
```

- You can create an array of structs and initialize them

```
struct date quarter[] =
    { {1,2}, {3,4}, {5,6}, {7,8}, {9,10} };
int cnt = sizeof(quarter)/sizeof(*quarter); // = 5
```

	High address	
quarter[4]	day	10
	month	9
quarter[3]	day	8
	month	7
quarter[2]	day	6
	month	5
quarter[1]	day	4
	month	3
quarter[0]	day	2
	month	1
	Low address	

Accessing members of a struct

```
struct date quarter[3];
struct date *ptr;

ptr = quarter + 1;      // array name = address
ptr->month = 2;
ptr->day = 21;           // or (*ptr).day = 21;

(ptr-1)->month = 1;      // or (*(ptr-1)).month = 4;
(ptr-1)->day = 7;

(++ptr)->month = 3;
ptr->day = 5;
```

