

Version 2.08


UCSD CSE 30

Computer Organization and Systems Programming

Lecture - 15

Keith Muller

Colossus MkII - 1944





Masking Summary - 1

Select a field: Use **and** with a mask of one's surrounded by zeros to select the bits that have a 1 in the mask, all other bits will be set to zero

selects this field when used with and

0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0	0 0 1 1 1 1 0 0
-----------------	-----------------	-----------------	-----------------

selection mask 0x3c

Clear a field: Use **and** with a mask of zero's surrounded by ones to select the bits that have a 1 in the mask, all other bits will be set to zero

clears this field when used with and

1 1 1 1 1 1 1 1	1 1 1 1 1 1 1 1	1 1 1 1 1 1 1 1	1 1 0 0 0 0 1 1
-----------------	-----------------	-----------------	-----------------

clear a field mask 0xfffffc3

Isolate a field: Use **lsl**, **lsl** to get a field surrounded by zeros

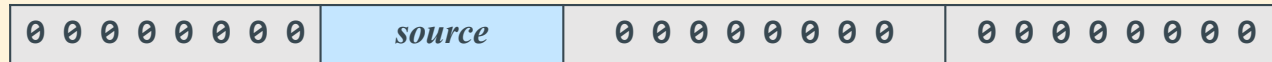
0 0 0 0 0 0 0 0	<i>source</i>	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
-----------------	---------------	---------------------------------

lsl this edge to bit 31 (left edge)
then lsr to move back

lsr this edge to bit 0 (right edge)
then lsl to move back

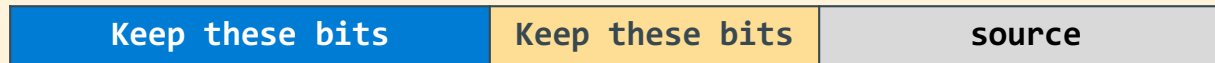
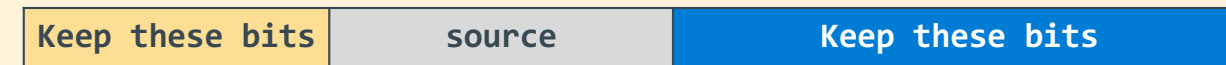
Masking Summary - 2

Isolate a field: Use **and** to get a **field surrounded by zeros**



selection mask 0x00ff0000

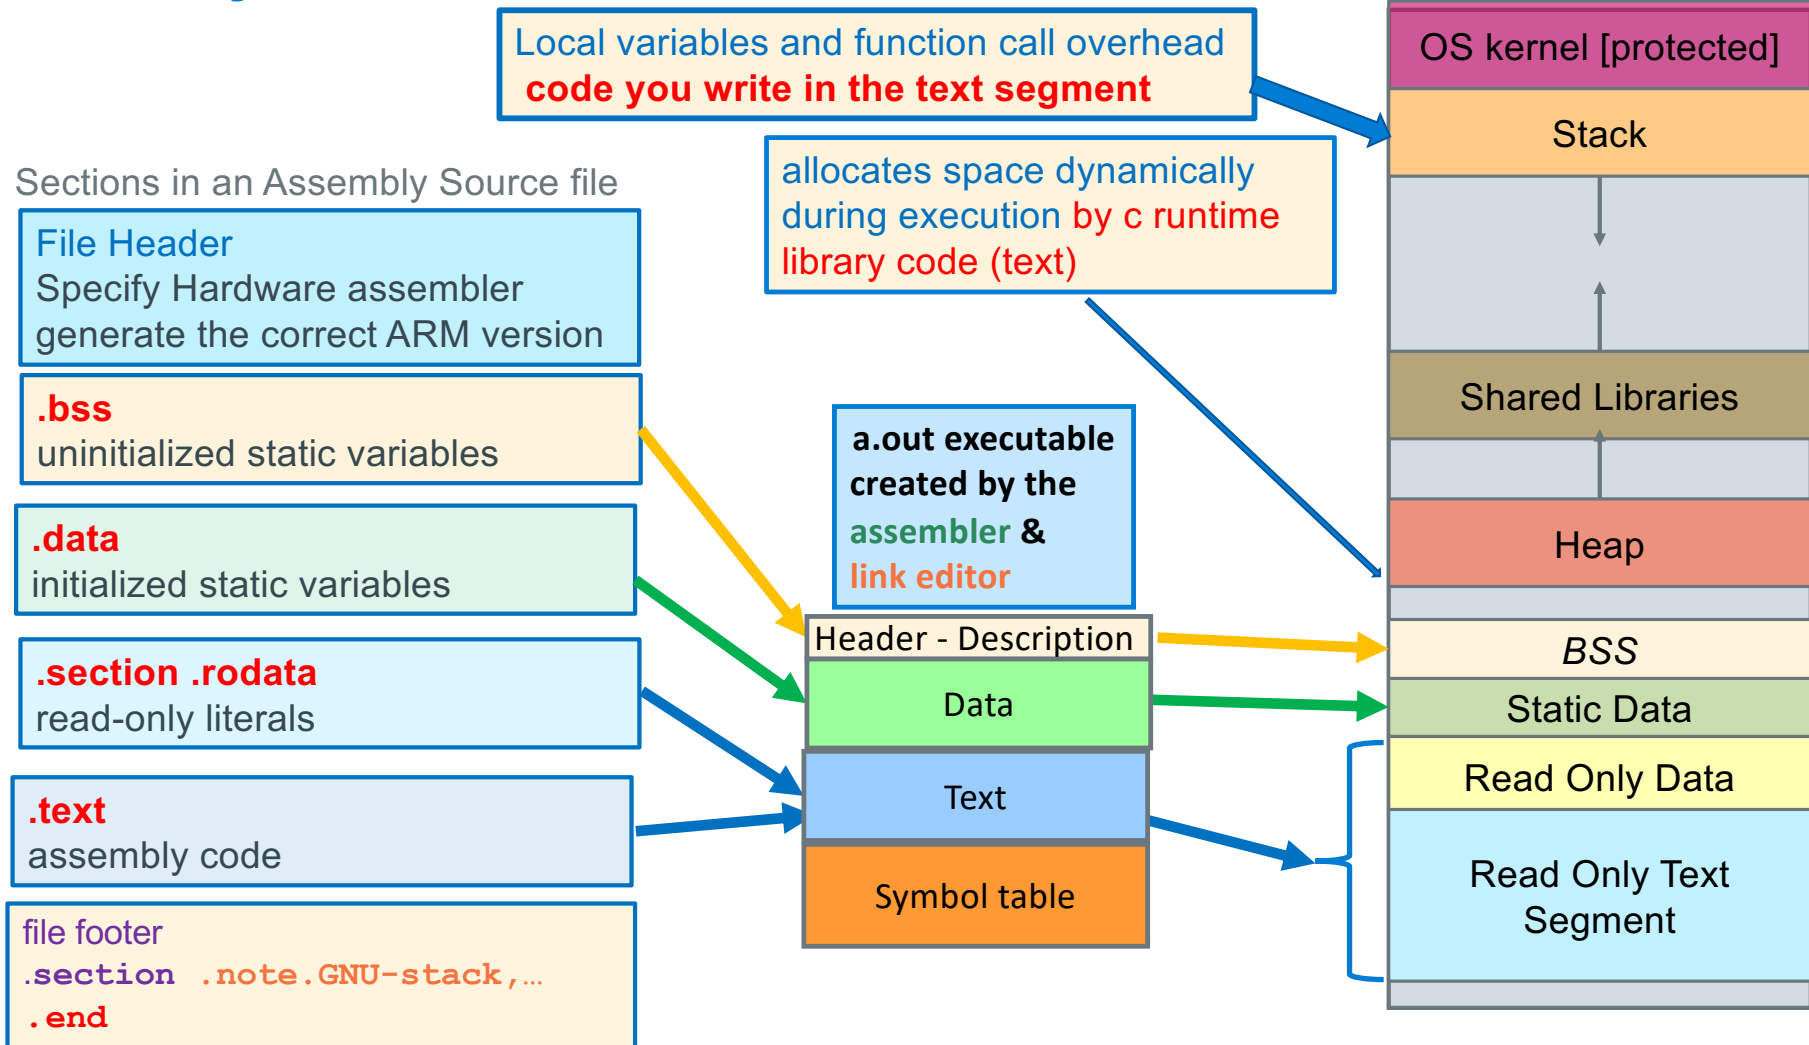
rotate a field: Use **ror** to move a field without changing other bits



Insert a field: Use **orr** with fields surrounded by zeros



Assembly Source File to Executable to Linux Memory



Assembly Source File Template

```
// File Header
.arch armv6                // armv6 architecture instructions
.arm                      // arm 32-bit instruction set
.fpu vfp                  // floating point co-processor
.syntax unified           // modern syntax

// BSS Segment (only when you have uninitialized globals)
.bss

// Data Segment (only when you have initialized globals)
.data

// Read-Only Data (only when you have literals)
.section .rodata

// Text Segment - your code
.text

// Function Header
.type    main, %function   // define main to be a function
.global main              // export function name
main:
// function prologue       // stack frame setup
    // your code for this function here
// function epilogue       //stack frame teardown

// function footer
.size    main, (. - main)

// File Footer
.section .note.GNU-stack,"",%progbits // stack/data non-exec
.end
```

- assembly programs end in **.S**
 - That is a **capital .S**
 - **example:** test.S
- Always use gcc to assemble
 - **_start()** and C runtime
- File has a complete program
gcc file.S
- File has a partial program
gcc -c file.S
- Link files together
gcc file.o cprog.o

Assembler Directives: .equ and .equiv

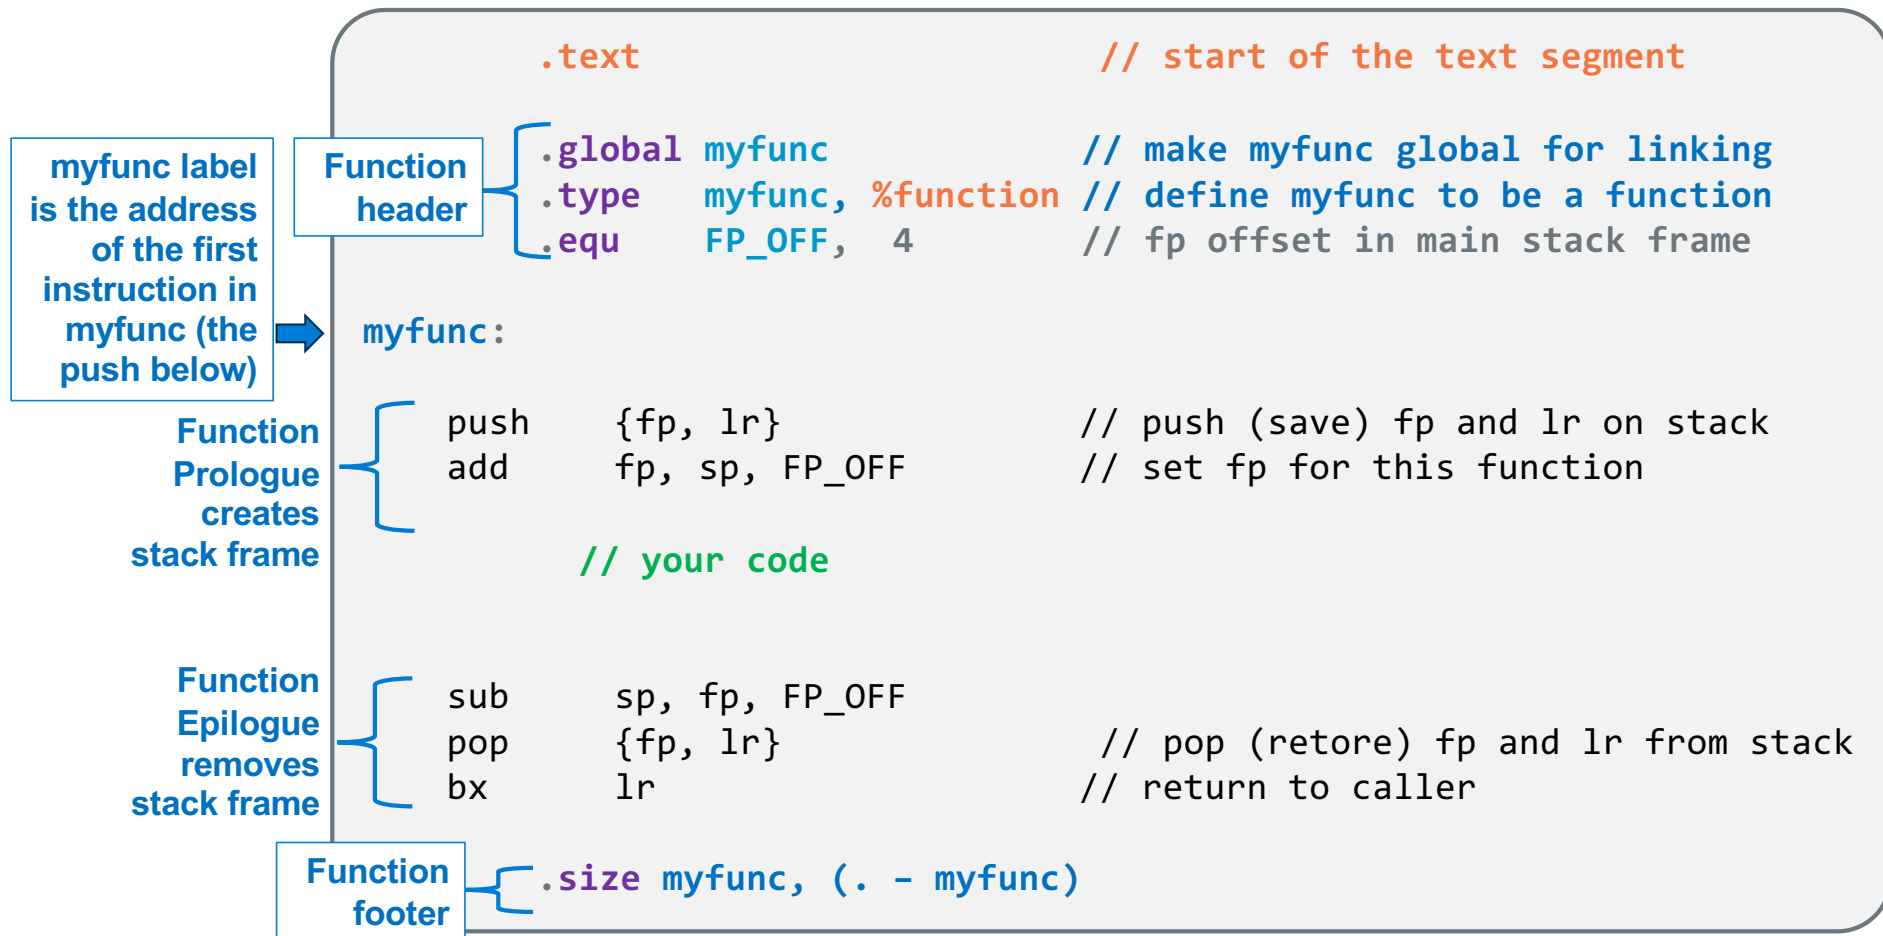
```
.equ    BLKSZ, 10240    // buffer size in bytes
.equ    BUFCNT, 100*4   // buffer for 100 ints
.equ    BLKSZ, STRSZ * 4 // redefine BLKSZ from here
```

.equ <symbol>, <expression>

- Defines and sets the value of a **symbol** to the **evaluation** of the **expression**
- Used for specifying constants, like a **#define** in C
- You can **(re)**set a symbol many times in the file, **last one seen applies**

```
.equ    BLKSZ, 10240    // buffer size in bytes
// other lines
.equ    BLKSZ, 1024     // buffer size in bytes
```

Function Template



Preview: Return Value and Passing Parameters to Functions

(Four parameters or less)

Register	Function Call Use
r0	1 st parameter
r1	2 nd parameter
r2	3 rd parameter
r3	4 th parameter

Register	Function Return Value Use
r0	8, 16 or 32-bit result, 32-bit address or least-significant half of a 64-bit result
r1	most-significant half of a 64-bit result

- Where **r0**, **r1**, **r2**, **r3** are arm registers, the function declaration is (first four arguments):

```
r0 = function(r0, r1, r2, r3)           // 32-bit return
```

```
r0, r1 = function(r0, r1, r2, r3)      // 64-bit return - long long
```
- Each **parameter and return value is limited to data that can fit in 4 bytes or less**
- You receive **up to the first four parameters in these four registers**
- You copy up to the first four parameters into these four registers before calling a function
- For parameter values using more than 4 bytes, a pointer to the parameter is passed (we will cover this later)
- You MUST ALWAYS assume** that the called function will **alter the contents of all four registers: r0-r3**
 - In terms of C runtime support, these registers contain the copies given to the called function
 - C allows the copies to be changed in any way by the called function

Preview: Writing an ARM32 function

```
#include <stdlib.h>
#include <stdio.h>
#include "sum4.h"
int main()
{
    int reslt;

    reslt = sum4(1,2,3,4);

    printf("%d\n", reslt);
    return EXIT_SUCCESS;
}
```

```
#ifndef SUM4_H
#define SUM4_H

#ifdef __ASSEMBLER__
int sum4(int, int, int, int);
#else
.extern sum4
#endif

#endif
```

```
#include "sum4.h"
.arch armv6
.arm
.fpu vfp
.syntax unified
.global sum4
.type sum4, %function
.equ FP_OFF, 28
// r0 = sum4(r0, r1, r2, r3)
sum4:
    push    {r4-r9, fp, lr}
    add     fp, sp, FP_OFF

    add     r0, r0, r1
    add     r0, r0, r2
    add     r0, r0, r3

    sub     sp, fp, FP_OFF
    pop     {r4-r9, fp, lr}
    bx      lr
    .size sum4, (. - sum4)
    .section .note.GNU-stack,"",%progbits
.end
```

```
$ gcc -Wall -Wextra -c main.c
$ gcc -c sum4.S
$ gcc sum4.o main.o
$ ./a.out
10
```

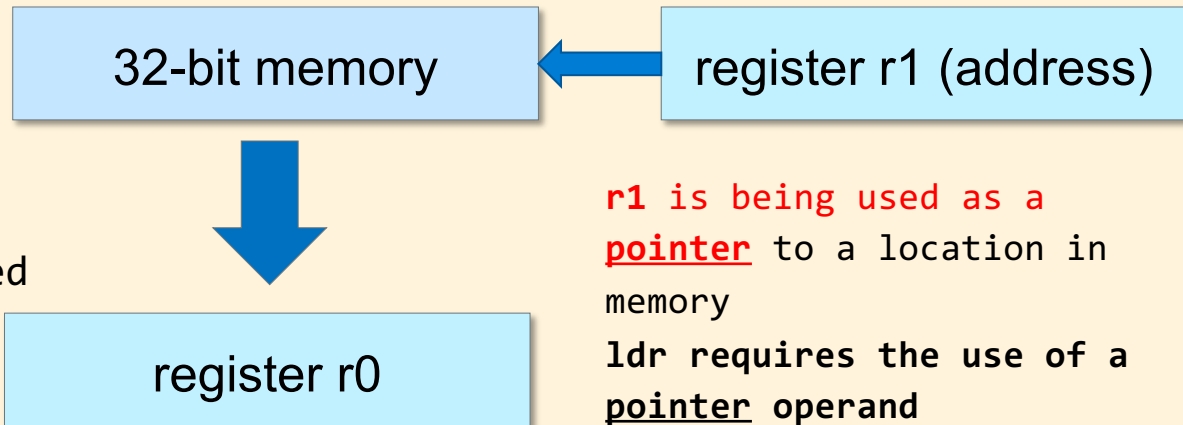
We will cover this
when we do stack frames

We will cover this
when we do stack frames

Load/Store: Register Base Addressing


ldr r0, [r1]

Copies a 32-bit word from the memory location whose address is contained in r1 (r1 is a pointer) into register r0




str r0, [r1]

Copies all 32 bits of the value held in register r0 to the 32-bit memory location contained in register r1 (r1 pointer)



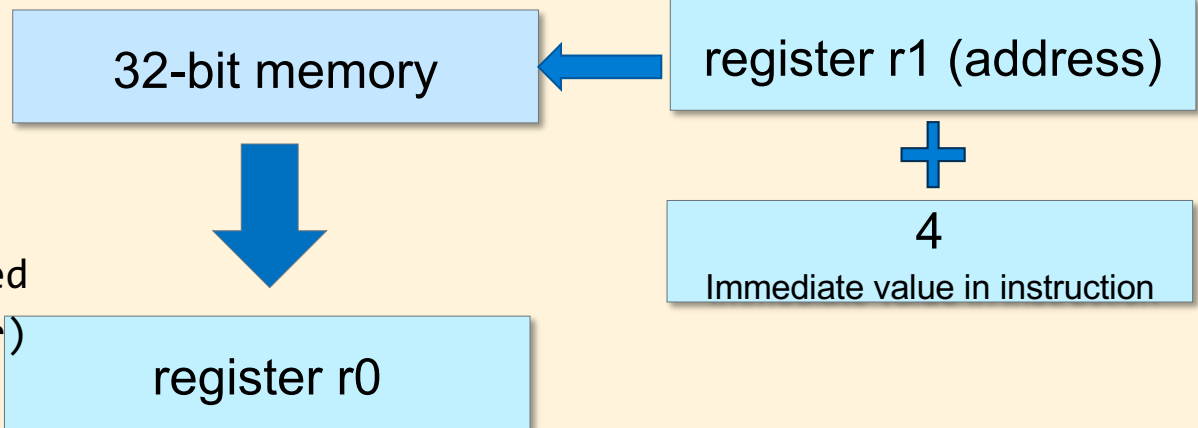
Example Base Register Addressing Load – Modify – Store



Load/Store: Register Base Addressing + Immediate

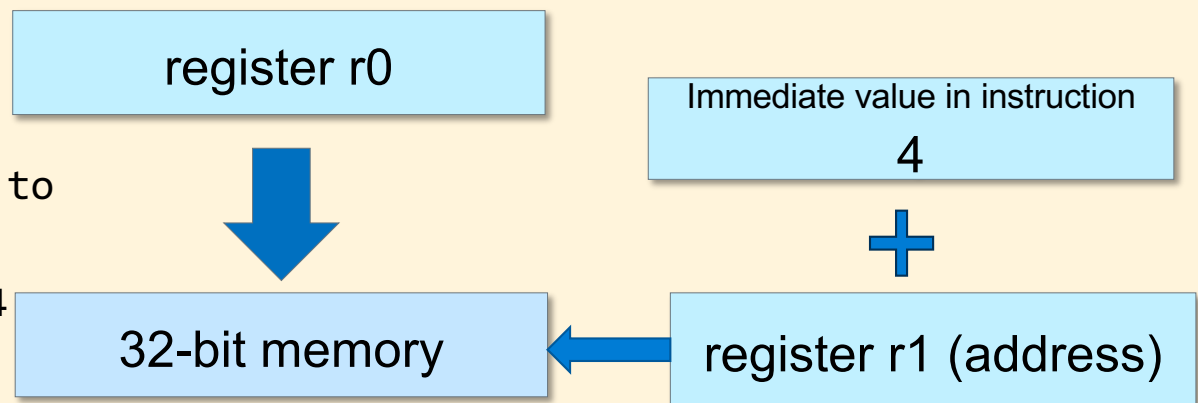
ldr r0, [r1, 4]

Copies a 32-bit word from the memory location whose address is contained in r1 +4 (r1 is a pointer) into register r0

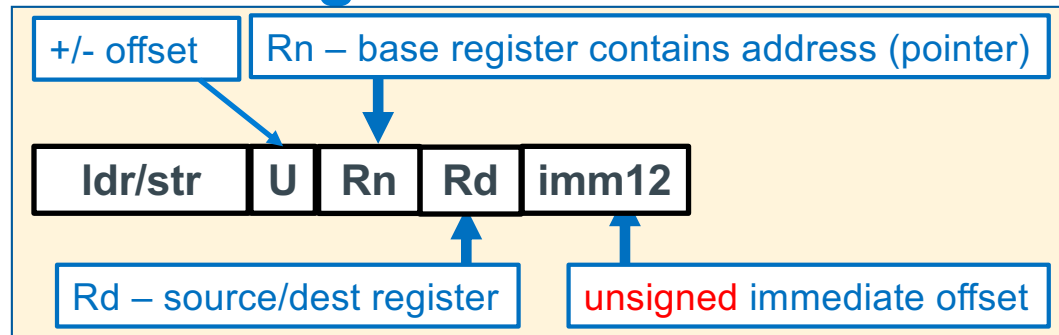


str r0, [r1, 4]

Copies all 32 bits of the value held in register r0 to the 32-bit memory location contained in register r1+4 (r1 pointer)



LDR/STR – Base Register + Immediate Offset Addressing



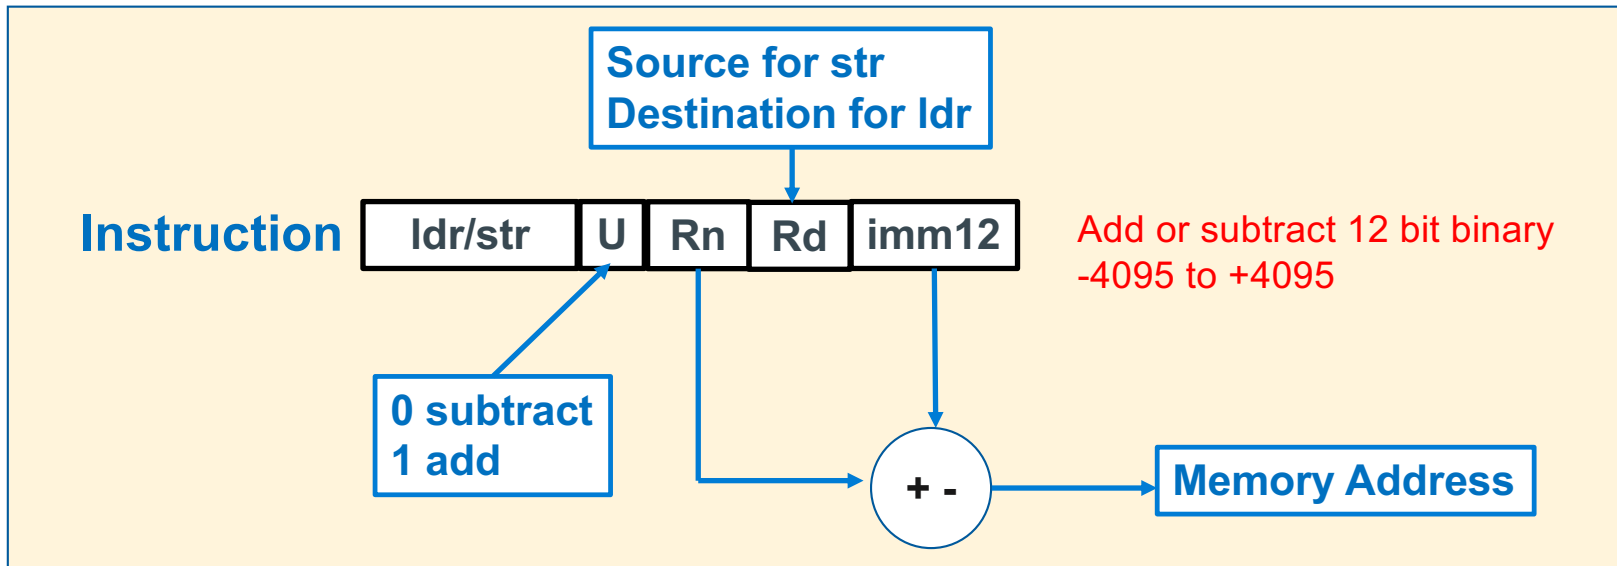
- **Register Base Addressing:**

- **Pointer Address:** Rn; **source/destination data:** Rd
- **Unsigned pointer address** is stored in the **base register**

- **Register Base + immediate offset Addressing:**

- **Pointer Address** = register content + immediate offset $-4095 \leq \text{imm12} \leq 4095$ (bytes)
- **Unsigned offset integer immediate value (bytes)** is **added or subtracted** (**U bit above says to add or subtract**) from the **pointer address** in the **base register**
- **Often used to address struct members**
 - Address of struct is address of the first member and subsequent members are a fixed offset from the first based on their size of the preceding members

ldr/str Register Base + Immediate Offset Addressing




Syntax	Address	Examples
<code>ldr/str Rd, [Rn, +/- constant]</code> constant is in bytes <code>ldr/str Rd, [Rn]</code>	<code>Rn + or - constant</code> same \longrightarrow	<code>ldr r0, [r5,100]</code> <code>str r1, [r5, 0]</code> <code>str r1, [r5]</code>

Loading and Storing: Variations List


- Load and store have **variations** that move 8-bits, 16-bits and 32-bits
- Load into a register with less than 32-bits will **set the upper bits not filled from memory differently** depending on which **variation of the load instruction** is used
- Store will only select the lower 8-bit, lower 16-bits or all 32-bits of the register to copy to memory, **register contents are not altered**

Instruction	Meaning	Sign Extension	Memory Address Requirement
ldrsb	load signed byte	sign extension	none (any byte)
ldrb	load unsigned byte	zero fill (extension)	none (any byte)
ldrsh	load signed halfword	sign extension	halfword (2-byte aligned)
ldrh	load unsigned halfword	zero fill (extension)	halfword (2-byte aligned)
ldr	load word	---	word (4-byte aligned)
strb	store low byte (bits 0-7)	---	none (any byte)
strh	store halfword (bits 0-15)	---	halfword (2-byte aligned)
str	store word (bits 0-31)	---	word (4-byte aligned)

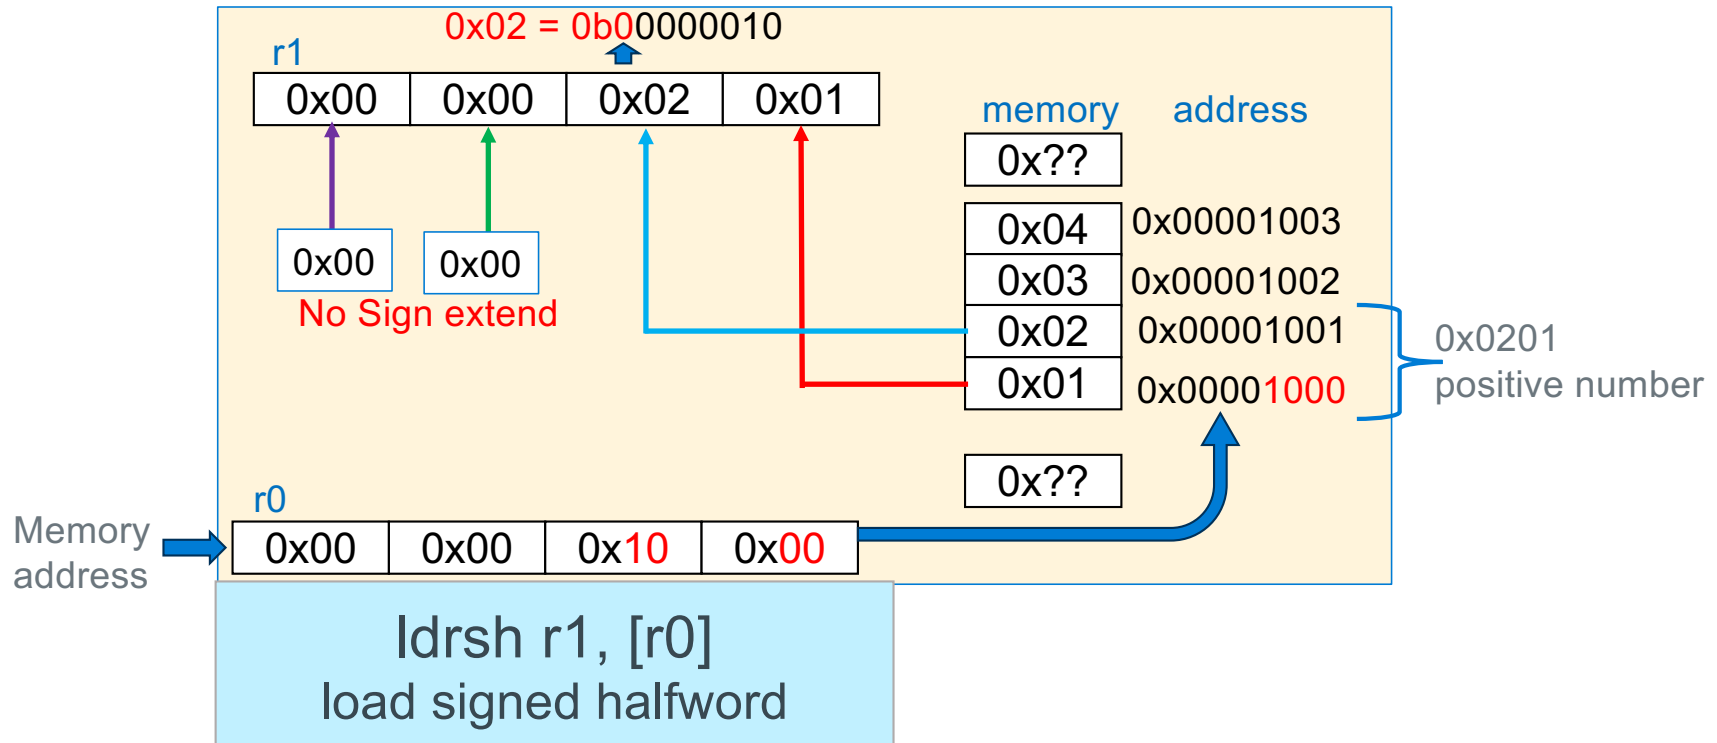
Loading 32-bit Registers From Memory, 32-bit



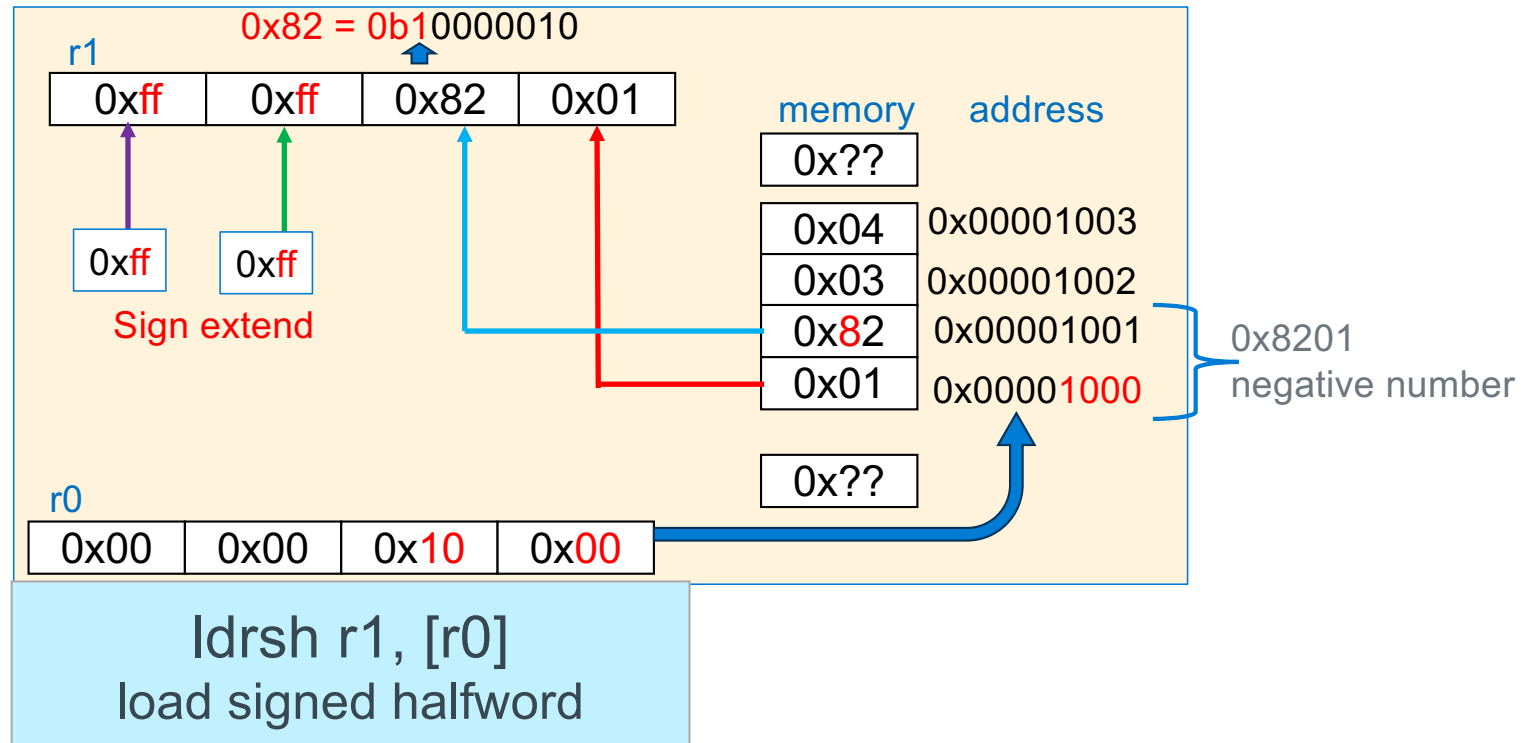
Loading 32-bit Registers From Memory, 16-bit




Loading 32-bit Registers From Memory, 16-bit



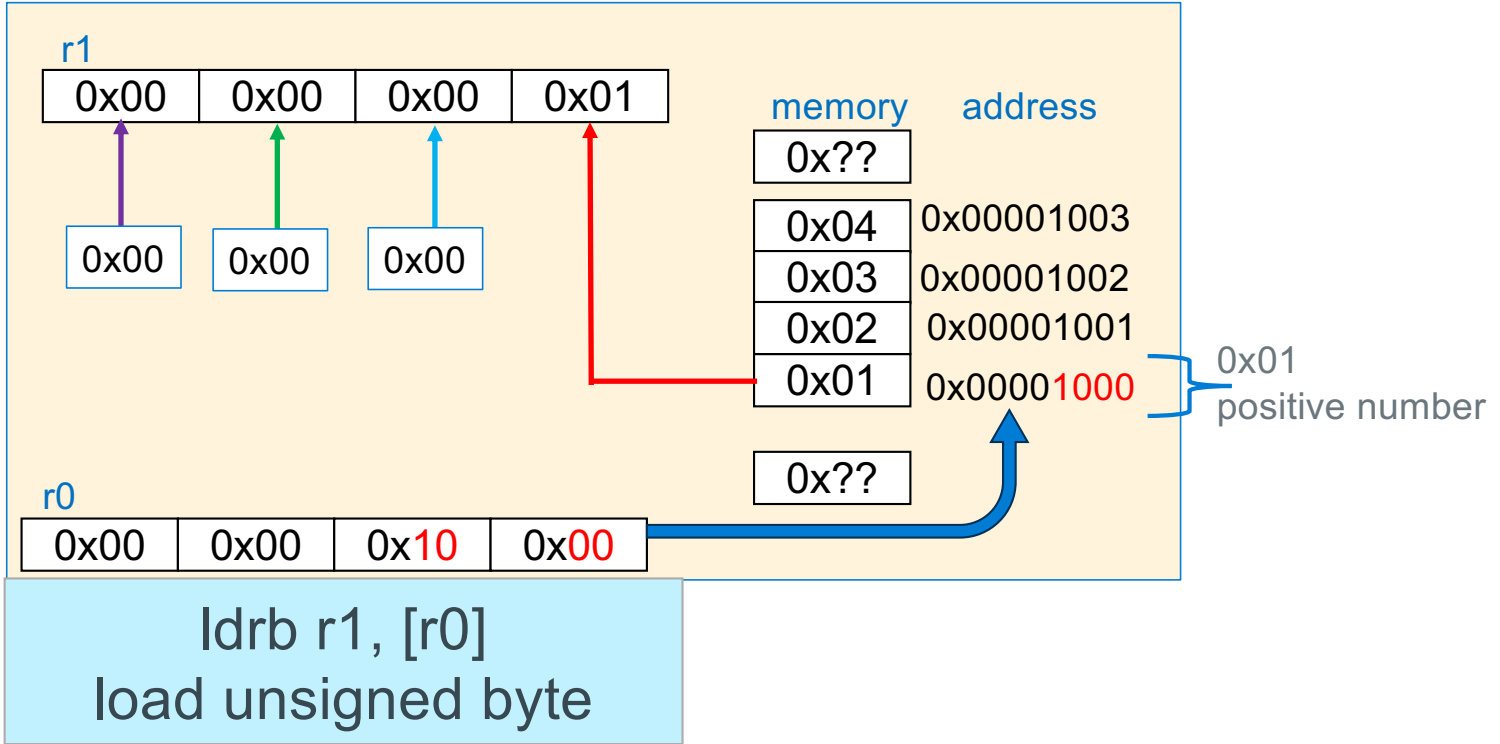
Loading 32-bit Registers From Memory, 16-bit Signed




Loading 32-bit Registers From Memory, 16-bit Unsigned




Loading 32-bit Registers From Memory, 8-bit



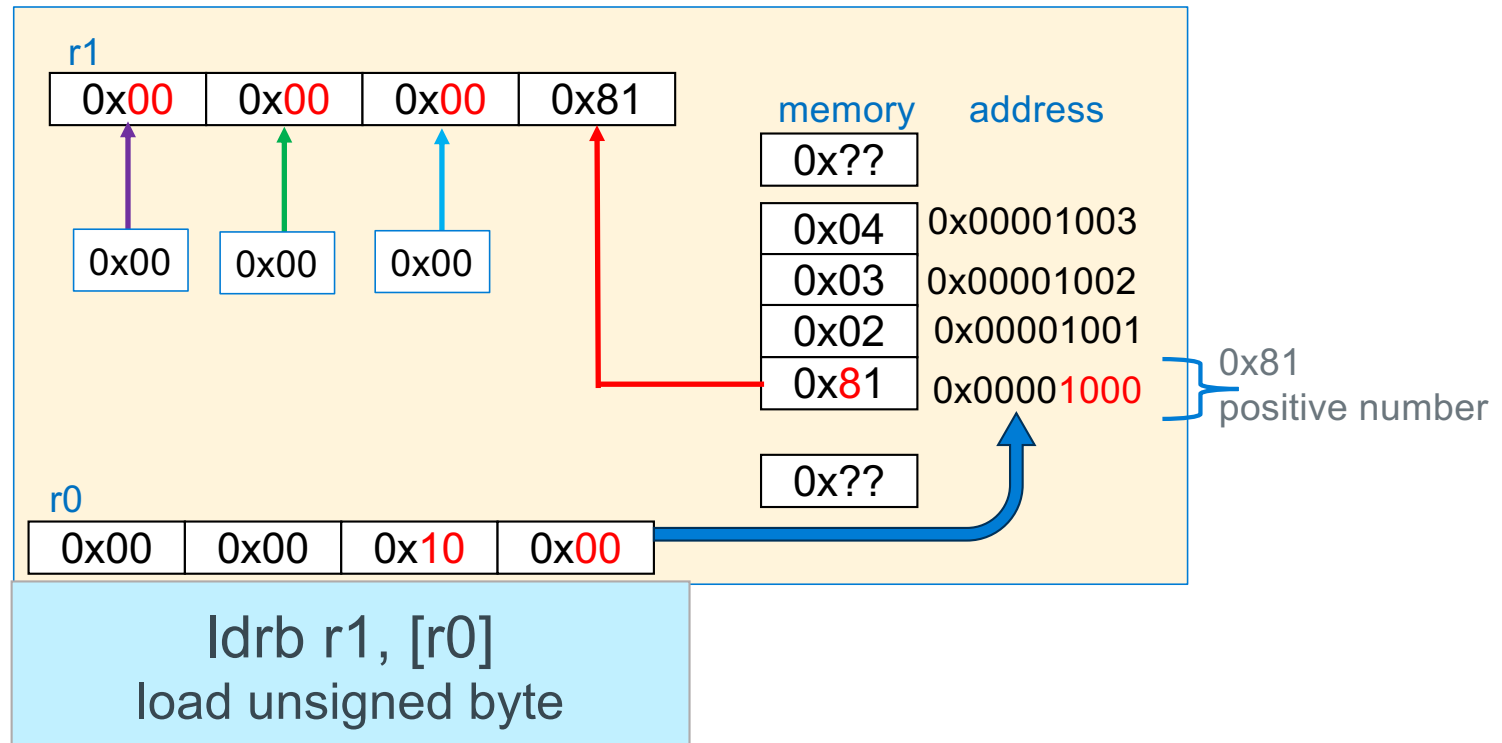
Loading 32-bit Registers From Memory, 8-bit



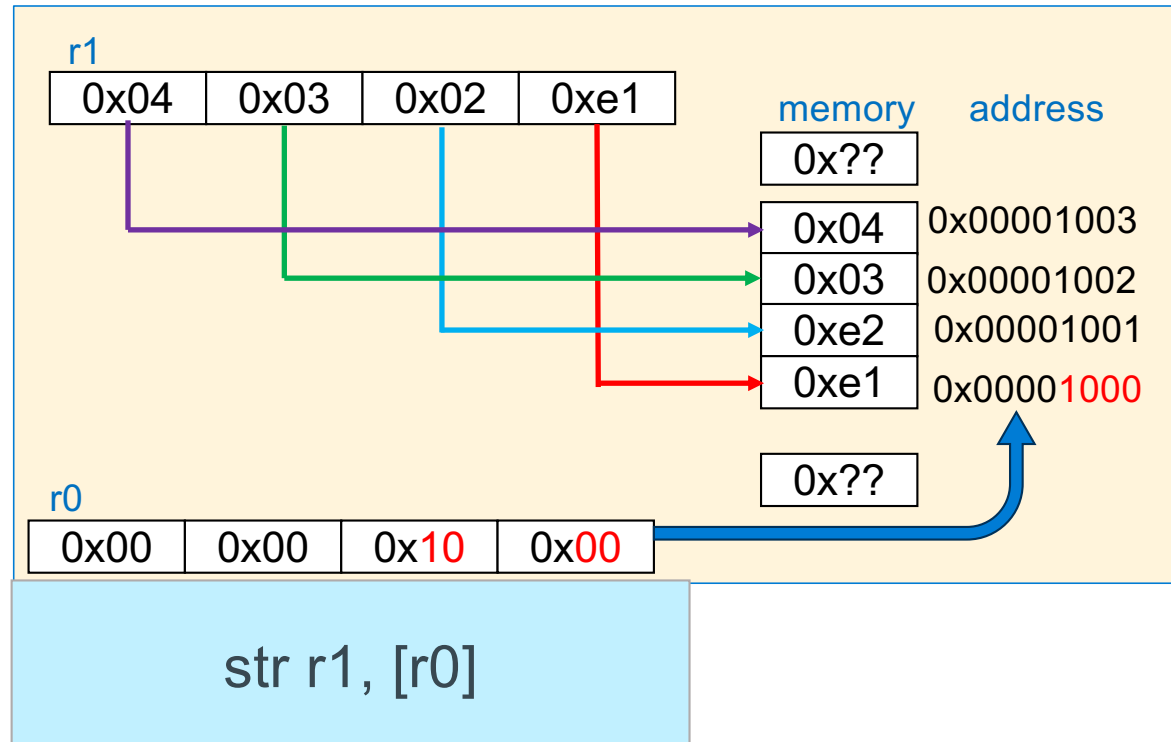
Loading 32-bit Registers From Memory, 8-bit Signed




Loading 32-bit Registers From Memory, 8-bit Signed



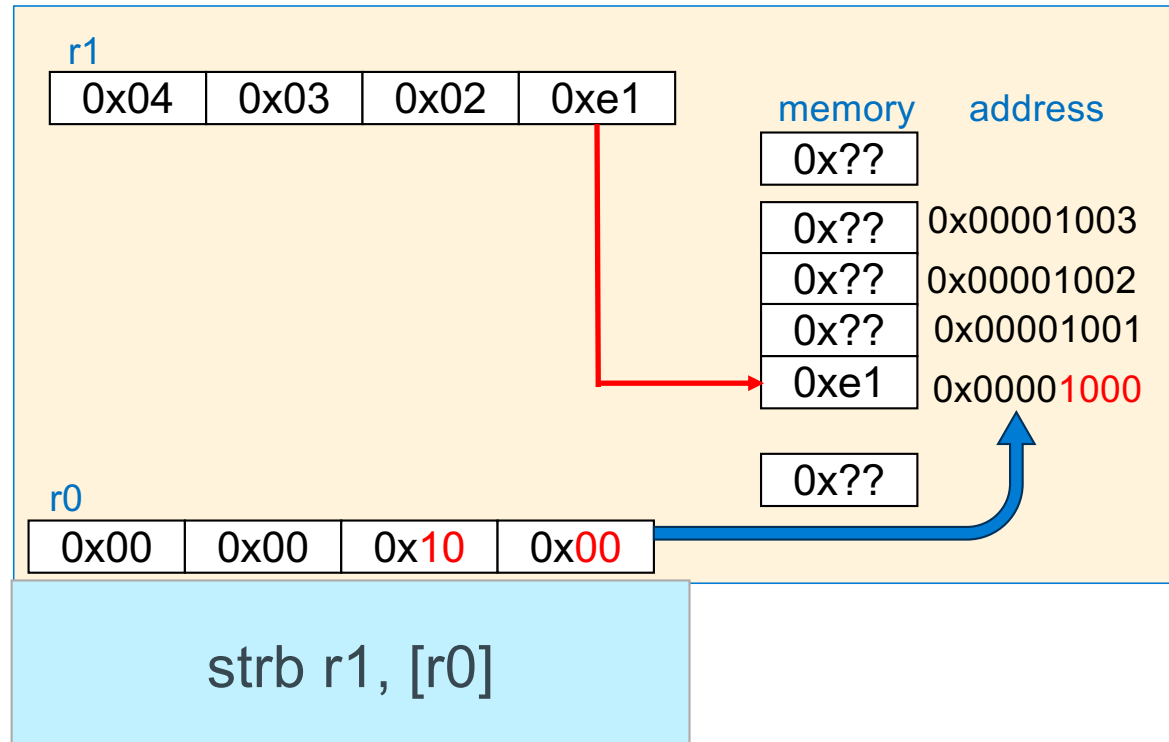
Storing 32-bit Registers To Memory, 32-bit




Storing 32-bit Registers To Memory, 16-bit



Storing 32-bit Registers To Memory, 8-bit

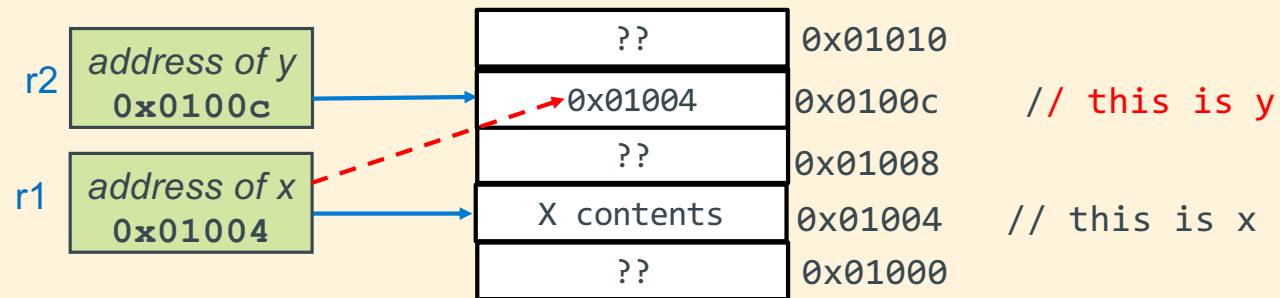


Storing 32-bit Registers To Memory, 8-bit – Storing different byte



ldr/str practice - 1

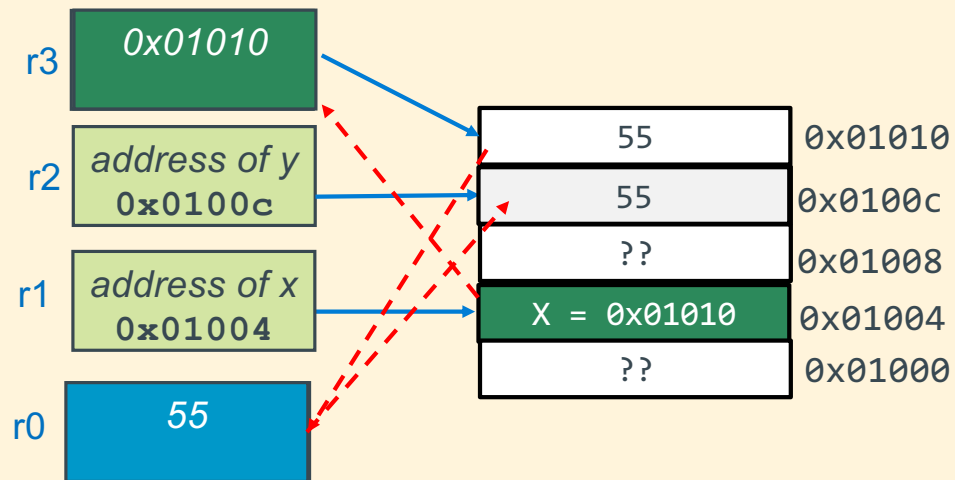
r1 contains the Address of X (defined as int X) in memory; r1 points at X
r2 contains the Address of Y (defined as int *Y) in memory; r2 points at Y
write Y = &X;



```
str    r1, [r2]    // y ← &x
```

ldr/str practice - 2

r1 contains the Address of X (defined as `int *X`) in memory r1 points at X
r2 contains the Address of Y (defined as `int Y`) in memory; r2 points at Y
write `Y = *X;`



```
ldr    r3, [r1]    // r3 ← x (read 1)
ldr    r0, [r3]    // r0 ← *x (read 2)
str    r0, [r2]    // y ← *x
```

using ldr/str: array copy

```
#include <stdio.h>
#include <stdlib.h>
#define SZ 6

void icpy(int *, int *, int);

int main(void)
{
    int  src[SZ] = {1, 2, 3, 4, 5, 6};
    int  dst[SZ];

    icpy(src, dst, SZ);
    for (int i = 0; i < SZ; i++)
        printf("%d\n", *(dst + i));

    return EXIT_SUCCESS;
}
```

```
void icpy(int *src, int *dst, int cnt)
{
    int *end = src + end;

    if (cnt <= 0)
        return;
    do {
        *dst++ = *src++;
    } while (src < end);
    return;
}
```

Base Register version

```
.arch armv6
.arm
.fpu vfp
.syntax unified
.text
.global icpy
.type icpy, %function
.equ FP_OFF, 12

// r0 contains int *src
// r1 contains int *dst
// r2 contains int cnt
// r3 use as loop term pointer
// r4 use as temp

icpy:
    push    {r4, r5, fp, lr}
    add     fp, sp, FP_OFF
    // see right ->
    sub     sp, fp, FP_OFF
    pop     {r4, r5, fp, lr}
    bx      lr
.size icpy, (. - icpy)
.end
```

```
    cmp     r2, 0
    ble     .Ldone      pre loop guard

    lsl     r2, r2, 2    //convert cnt to int size
    add     r3, r0, r2  // loop term pointer

.Ldo:
    ldr     r4, [r0]     // load from src
    str     r4, [r1]     // store to dest

    add     r0, r0, 4    // src++
    add     r1, r1, 4    // dst++

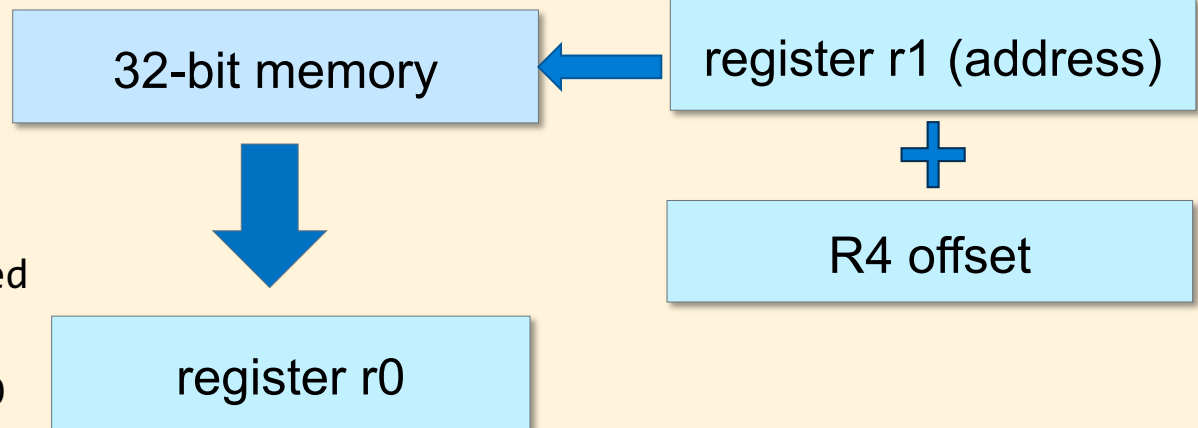
    cmp     r0, r3       // src < term pointer?
    blt     .Ldo
                                loop guard

.Ldone:
```

Load/Store: Register Base Addressing + Register Offset

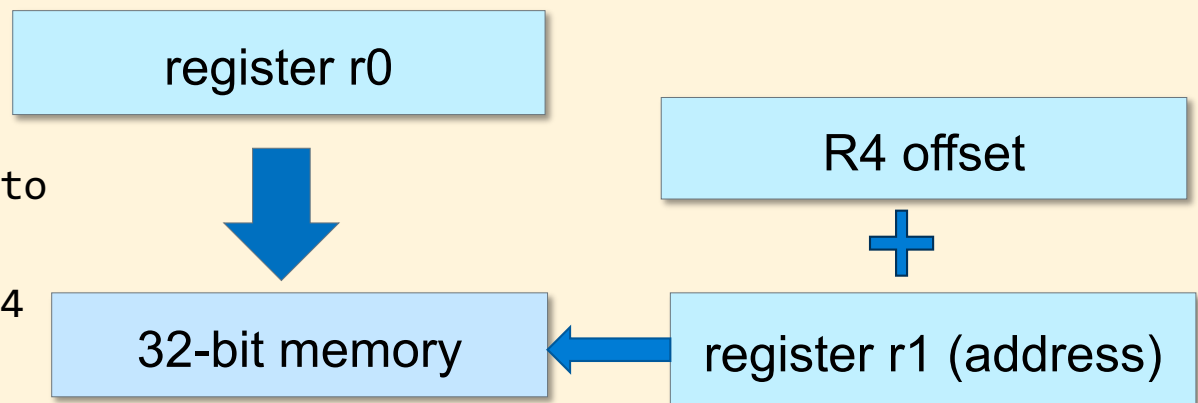
ldr r0, [r1, r4]

Copies a 32-bit word from the memory location whose address is contained in $r1 + r4$ ($r1$ is a pointer) into register $r0$

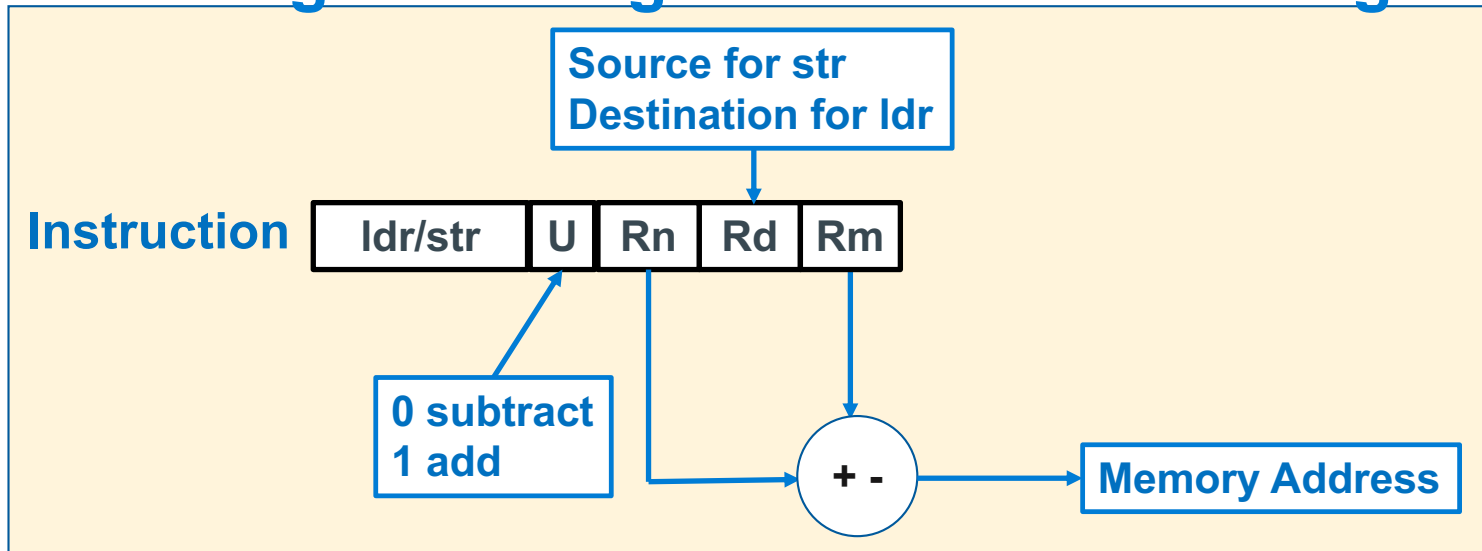


str r0, [r1, r4]

Copies all 32 bits of the value held in register $r0$ to the 32-bit memory location contained in register $r1+r4$ ($r1$ pointer)



ldr/str Base Register + Register Offset Addressing



Pointer Address = Base Register + Register Offset

- **Unsigned** offset integer **in a register (bytes)** is either added/subtracted from the **pointer address** in the **base register**

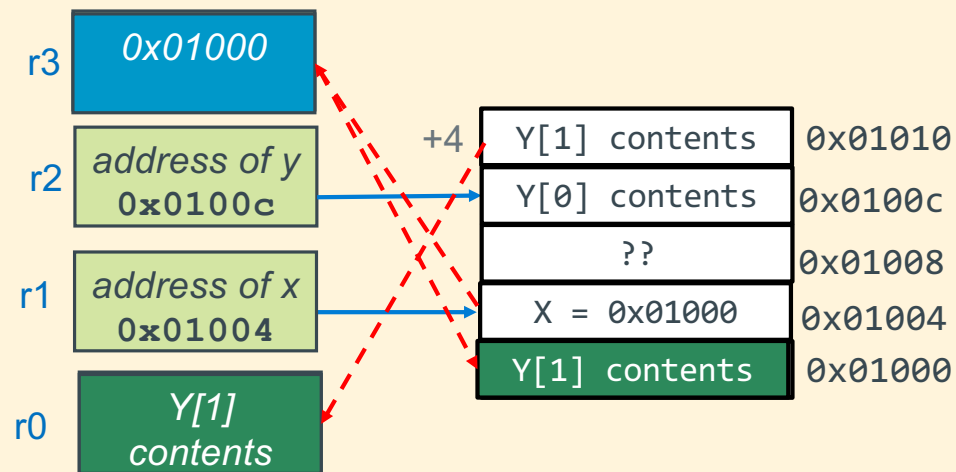
Syntax	Address	Examples
<code>ldr/str Rd, [Rn +/- Rm]</code>	$Rn + \text{ or } - Rm$	<code>ldr r0, [r5, r4]</code> <code>str r1, [r5, r4]</code>

ldr/str practice - 3

r1 contains Address of X (defined as `int *X`) in memory; r1 points at X

r2 contains Address of Y (defined as `int Y[2]`) in memory; r2 points at `&(Y[0])`

`write *X = Y[1];`



```
ldr    r0, [r2, 4]    // r0 ← y[1]
ldr    r3, [r1]        // r3 ← x
str     r0, [r3]       // *x ← y[1]
```

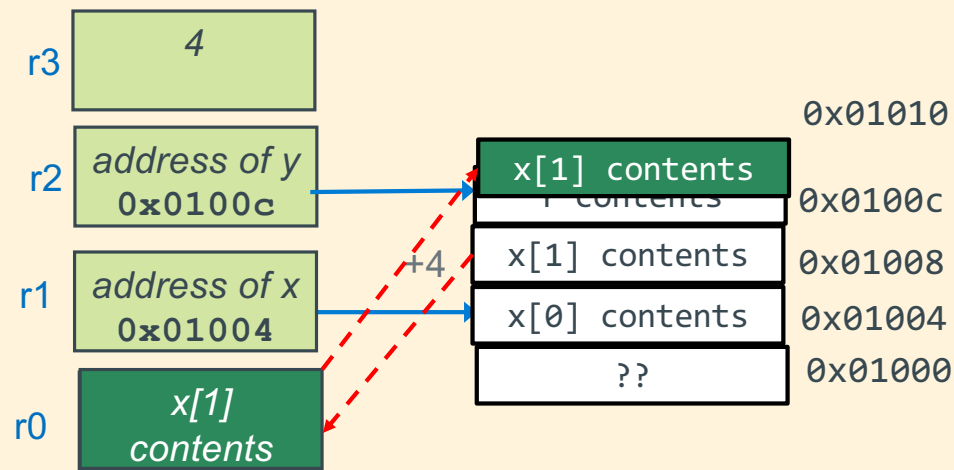

ldr/str practice - 4

r1 contains Address of X (defined as `int X[2]`) in memory; r1 points at `&(x[0])`

r2 contains Address of Y (defined as `int Y`) in memory; r2 points at Y

r3 contains a 4

write `Y = X[1];`



```
ldr    r0, [r1, r3]  // r0 ← x[1]
```

```
str    r0, [r2]      // y ← x[1]
```

Base Register + Register Offset Version

```
.arch armv6
.arm
.fpu vfp
.syntax unified
.text
.global icpy
.type icpy, %function
.equ FP_OFF, 12
// r0 contains int *src
// r1 contains int *dst
// r2 contains int cnt
// r3 use as loop counter
// r4 use as temp
```

```
icpy:
    push    {r4, r5, fp, lr}
    add     fp, sp, FP_OFF
    // see right ->
    sub     sp, fp, FP_OFF
    pop     {r4, r5, fp, lr}
    bx      lr
    .size icpy, (. - cpy)
.end
```

```
    cmp     r2, 0
    ble     .Ldone
    lsl     r2, r2, 2
    mov     r3, 0
.Ldo:
    ldr     r4, [r0, r3]
    str     r4, [r1, r3]
    add     r3, r3, 4
    cmp     r3, r2
    blt     .Ldo
```

pre loop guard

loop guard

.Ldone:

one increment
covers both arrays

Base Register + Register Offset With chars

```
#include <stdio.h>
#include <stdlib.h>
#define SZ 6
void cpy(char *, char *, int);
int main(void)
{
    char src[SZ] =
        {'a', 'b', 'c', 'd', 'e', '\0'};
    char dst[SZ];

    cpy(src, dst, SZ);
    printf("%s\n", dst);
    return EXIT_SUCCESS;
}
```

```
    cmp    r2, 0
    ble    .Ldone

    mov    r3, 0           // initialize counter
.Ldo:
    ldrb   r4, [r0, r3]    // load from src
    strb   r4, [r1, r3]    // store to dest
    add    r3, r3, 1       // counter++
    cmp    r3, r2          // count < r3
    blt    .Ldo

.Ldone:
```

Reference: Addressing Mode Summary for use in CSE30

index Type	Example	Description
Pre-index immediate	<code>ldr r1, [r0]</code>	$r1 \leftarrow \text{memory}[r0]$ $r0$ is unchanged
Pre-index immediate	<code>ldr r1, [r0, 4]</code>	$r1 \leftarrow \text{memory}[r0 + 4]$ $r0$ is unchanged
Pre-index immediate	<code>str r1, [r0]</code>	$\text{memory}[r0] \leftarrow r1$ $r0$ is unchanged
Pre-index immediate	<code>str r1, [r0, 4]</code>	$\text{memory}[r0 + 4] \leftarrow r1$ $r0$ is unchanged
Pre-index register	<code>ldr r1, [r0, +-r2]</code>	$r1 \leftarrow \text{memory}[r0 \pm r2]$ $r0$ is unchanged
Pre-index register	<code>str r1, [r0, +-r2]</code>	$\text{memory}[r0 \pm r2] \leftarrow r1$ $r0$ is unchanged

Base Register Addressing + Offset register

```
#include <stdio.h>
#include <stdlib.h>
int count(char *, int);
int main(void)
{
    char msg[] = "Hello CSE30! We Are CountinG UpPER cASe letters!";

    printf("%d\n", count(msg, sizeof(msg)/sizeof(*msg)));
    return EXIT_SUCCESS;
}
```

```
int count(char *ptr, int len)
{
    int cnt = 0;
    int i;

    for (i = 0; i < len; i++) {
        if ((ptr[i] >= 'A') && (ptr[i] <= 'Z'))
            cnt++;
    }
    return cnt;
}
```

Base Register + Offset register

```
.arch armv6
.arm
.fpu vfp
.syntax unified
.text
.global count
.type count, %function
.equ FP_OFF, 12
// r0 contains char *ptr
// r1 contains int len
// r2 contains int cnt
// r3 contains int i
// r4 contains char

count:
    push    {r4, r5, fp, lr}
    add     fp, sp, FP_OFF
    // see right ->
    sub     sp, fp, FP_OFF
    pop     {r4, r5, fp, lr}
    bx      lr
    .size count, (. - count)
.end
```

byte array
Also use ldrb here
offsets are 0,1,2,...

```
count:
    push    {r4, r5, fp, lr}
    add     fp, sp, FP_OFF

    mov     r2, 0
    cmp     r1, 0
    ble     .Ldone
    mov     r3, 0

.Lfor:
    cmp     r3, r1
    bge     .Ldone

    ldrb     r4, [r0, r3]
    cmp     r4, 'A'
    blt     .Lendif
    cmp     r4, 'Z'
    bgt     .Lendif
    add     r2, r2, 1



.Lendif:
    add     r3, r3, 1
    b       .Lfor

.Ldone:
    mov     r0, r2
```

loop guard

Variable Alignment In Memory and Performance

Accessing **address aligned** memory on many systems **based on data type** has **the best performance** (due to hardware implementation)

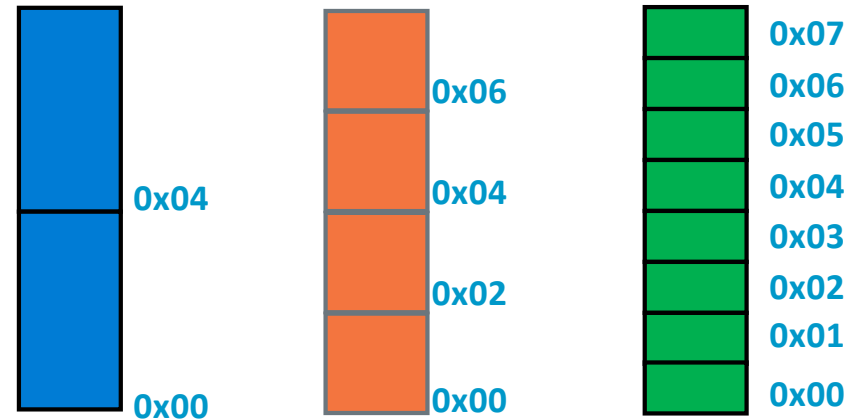


Defining Static Variables: Allocation and Initialization

Variable SIZE	Directive	.align	C static variable Definition	Assembler static variable Definition
8-bit char (1 byte)	.byte		char chx = 'A'; char string[] = {'A','B','C', 0};	chx: .byte 'A' string: .byte 'A','B',0x42,0
16-bit int (2 bytes)	.short	.align 1	short length = 0x55aa;	length: .short 0x55aa
32-bit int (4 bytes)	.word .long	.align 2	int dist = 5; int *distptr = &dist; unsigned int mask = 0xaa55aa55; int array[] = {12,~0x1,0xCD,-1};	dist: .word 5 distptr: .word dist mask: .word 0xff array: .word 12,~0x1,0xCD,-3
string with '\0'	.string		char class[] = "cse30";	class: .string "cse30"

Rule: Place the align above the variable
 .align 1
 len: .short 0x55aa

SIZE	Address ends in	Align
8-bit char -1 byte	0b..0 or 0b..1	
16-bit int -2 bytes	0b..0	.align 1
32-bit int -4 bytes and all arrays	0b..00	.align 2



Defining Static Variables: Allocation and Initialization

```
.data
chx: .byte 'A'
cnt: .word 4
```

See cnt is not word aligned



```
<chx>: 12028 00000041
<cnt>: 12029 00000004
```

```
.data
chx: .byte 'A'
    .align 2
cnt: .word 4
```

See cnt is word aligned



```
<chx>: 12028 00000041
<cnt>: 1202c 00000004
```

Sets alignment for
next variable

```
int num;           //4 bytes
int *ptr = &num;    //4 bytes
char *lit = "456";  //4 bytes, "456" string literal
char msg[] = "123"; //4 bytes - array
```

```
.bss
    .align 2
num:    .word 0
.data
    .align 2
ptr:    .word num
    .align 2
lit:    .word .Lmsg
    .align 2
msg:    .string "123"
.section .rodata
.Lmsg:  .string "456"
```

initializes
a pointer

read-only
string literal

Defining Static Array Variables

```
Label:    .size_directive expression, ... expression
```

```
In C:      int int_buf[100];  
            int array[] = {1, 2, 3, 4, 5};  
            char buffer[100];
```

```
.bss  
int_buf:    .space 400    // convert 100 to 400 bytes  
            .align 2  
char_buf:   .space 100  
.data  
array:      .word 1, 2, 3, 4, 5  
            .align 2  
one_buf:    .space 100, 1    // 100 bytes each byte filled with 1
```

.space size, fill

- Allocates **size** bytes, each of which contain the value **fill**
- Both **size** and **fill** are absolute expressions
- If the comma and **fill** are **omitted**, **fill** is assumed to be **zero**
- **.bss section**: Must be used **without a specified fill**

Loading Static variables into a register

- Tell the assembler load the address (Lvalue) of a label into a register:

`ldr/str Rd, =Label // Rd = address`

- Example to the right: $y = x$;*

two step to **load** a **memory** variable

- load the pointer to the memory
- read (load) from *pointer

two steps **store** to a **memory** variable

- load the pointer to the memory
- write (store) to *pointer

```
.bss
y: .space 4
```

```
.data
x: .word 200
```

```
.text
// function header
main:

// load the address, then contents
// using r2
ldr r2, =x      // int *r2 = &x
ldr r2, [r2]    // r2 = *r2;

// &x was only needed once above
// Note: r2 was a pointer then an int
// no "type" checking in assembly!

// store the contents of r2
ldr r1, =y      // int *r1 = &y
str r2, [r1]    // *r1 = r2
...
```

Loading large constants into a register:

Error: invalid constant (3ff) after fixup

- In data processing instructions, the field **imm8 + rotate 4 bits** is too small to store the immediate value, how do you get larger immediate values into a register?



fails



```
mov    r0, 1023
```

xxx.s:24: Error: invalid constant (3ff) after fixup

replacement

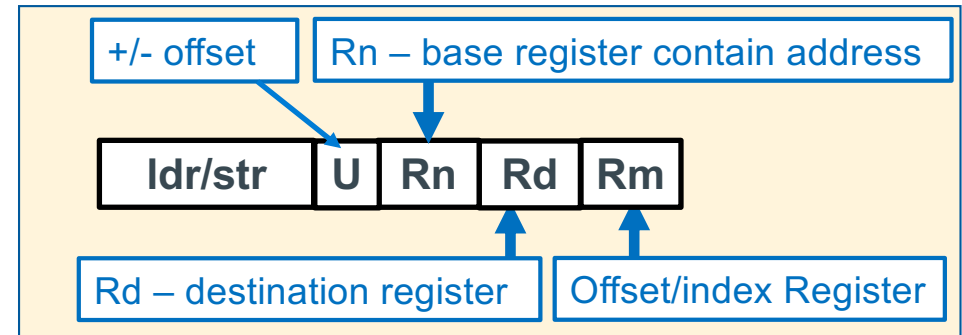
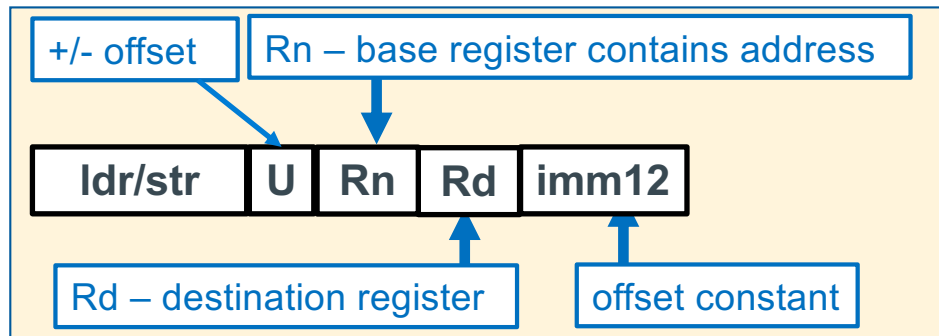


```
ldr    r0, =1023
```

- Answer: use **ldr** instruction with the constant as an operand: **=constant**
- Assembler creates a **literal table entry** with the **constant**

```
ldr    Rd, =constant    // =constant
ldr    r1, =0x2468abcd   // loads the constant 0x246abcd into r1
```

Reference: LDR/STR – Register To/From Memory Copy

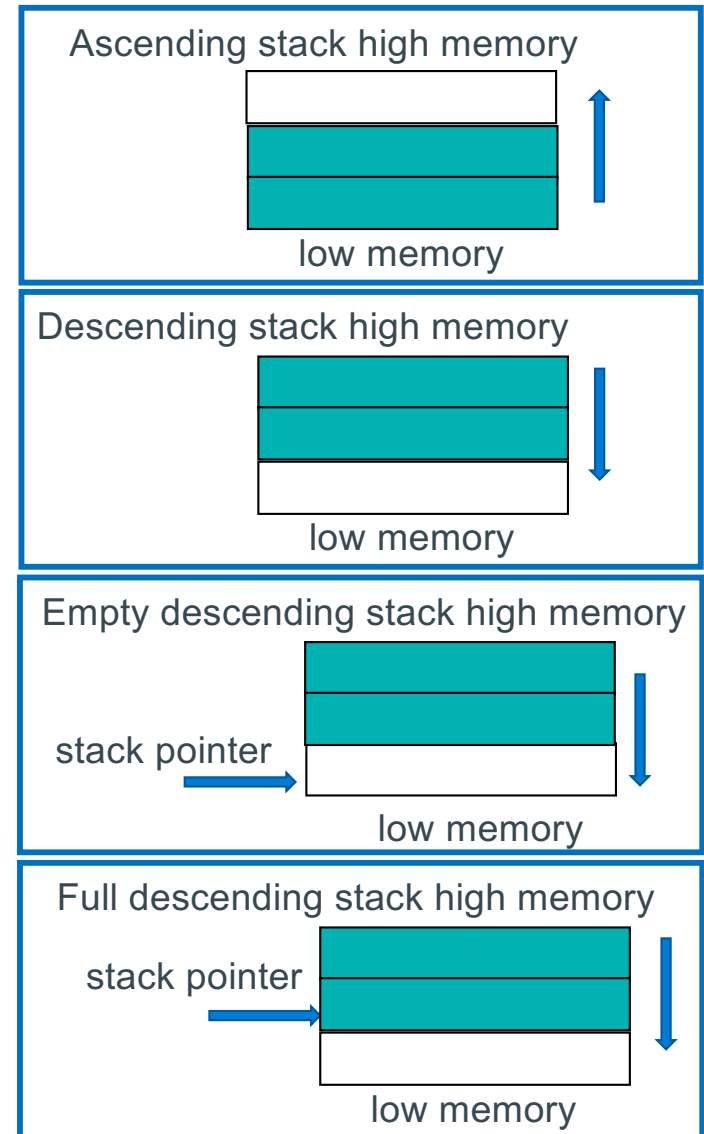


```
ldr/str Rd, [Rn, +/- imm12] // base register pointer + offset imm12 in bytes
                             -4095 <= imm12 <= 4095 (bytes)
ldr/str Rd, [Rn]             // base register pointer + 0 (imm12 is 0)
ldr/str Rd, [Rn, +/- Rm]     // base register pointer +/- offset register
```

```
ldr    r1, =var_x           // r1 = &var_x
str    r1, =mylabel+4       // *(mylabel+4) = r1
ldr    r1, =0x246abcd       // load an immediate into r1
ldr    r1, [r3]             // y = *r3 (4 bytes)
str    r1, [r0]             // *r0 = r1
ldr    r1, [r3, -4]         // y = *(r3 - 4) (4 bytes)
str    r1, [r0, r2]         // *(r0 + r2) = r1
```

Stack types

- A Stack Implements a **last-in first-out** (LIFO) protocol
- Each time a **function is called**, a **stack frame is activated**
 - space is allocated by moving the stack pointer
 - push adds space, pop removes space
- Stack growth direction
 - **Ascending stack**: grows from low memory towards high memory (adding to the sp to allocate memory)
 - **Descending stack**: grows from high memory towards low memory (subtracting from the sp to allocate memory)
- Full versus empty stacks
 - **Empty stack**: **stack pointer** (sp) points at the **next word address** after the last item pushed on the stack
 - **Full stack**: **stack pointer** (sp) points at the **last item pushed on the stack**
- ARM on Linux uses a **full descending stack**



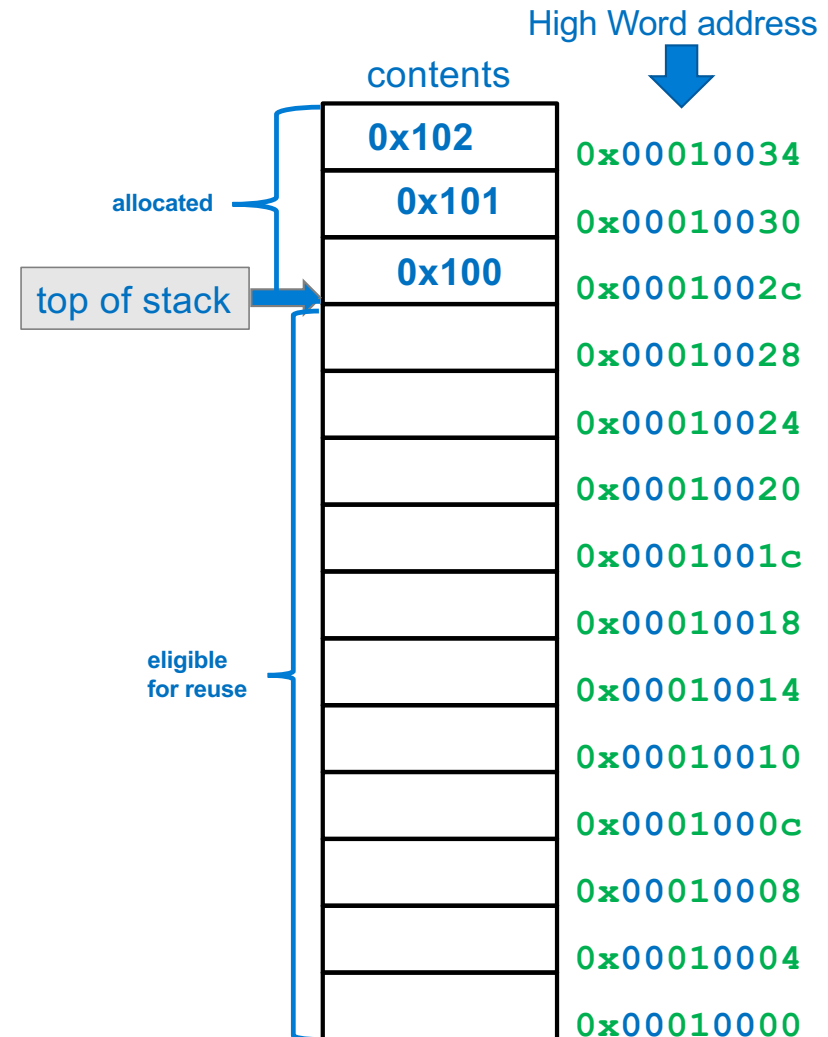
Arm: Stack Operation

- **Stack** is expandable and grows downward from high memory address towards low memory address
- **Stack pointer (sp)** always points at the **top of stack**
 - contains the starting address of the top element
- New items are **pushed** (*added*) onto the **top of the stack** by **subtracting** from the stack pointer the **size of the element** and then writing the element

push (sp - element size) & write

- Existing items are **popped** (*removed*) from the top of the stack by **adding** to the stack pointer the **size of the element** (leaving the **old contents unchanged**)

pop (sp + element size)



Function Calls, Parameters and Locals: Requirements

```
int
main(int argc, char *argv[])
{
    int x, z = 4;

    x = a(z);
    z = b(z);
    return EXIT_SUCCESS;
}

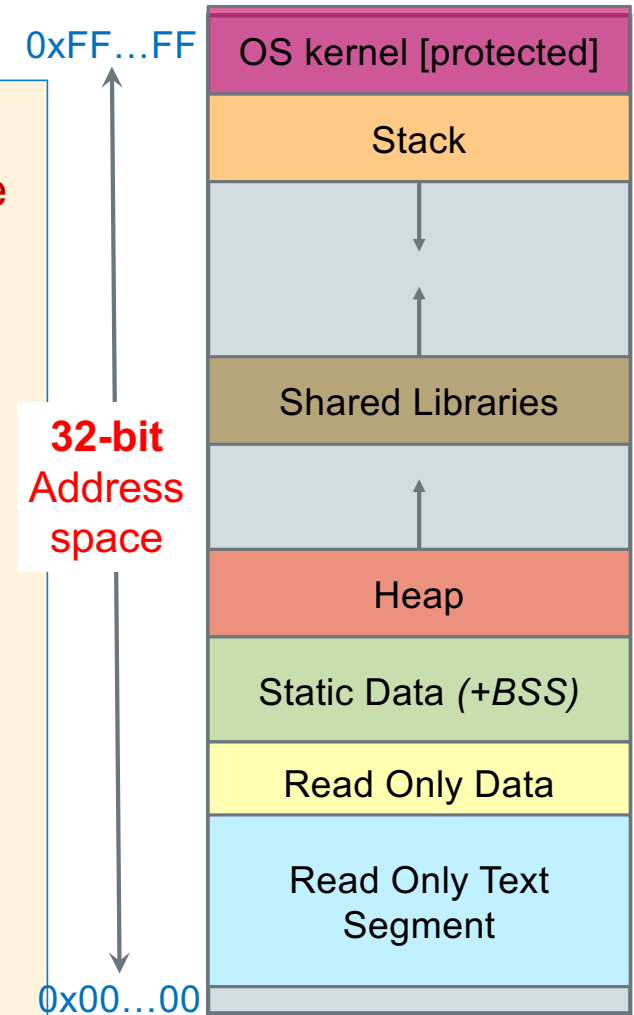
int
a(int n)
{
    int i = 0;
    if (n == 1)
        i = b(n);
    return i;
}

int
b(int m)
{
    return m+1;
}
/* the return cannot be done with a
branch */
```

- Since **b()** is called both by main and a() how does the **return m+1** statement in b() know where to return to? (Obviously, it cannot be a branch)
- Where are the parameters (args) to a function stored so the function has a copy that it can alter?
- Where is the return value from a function call stored?
- How are Automatic variables *lifetime* and *scope* implemented?
 - When you enter a variables scope: memory is allocated for the variables
 - When you leave a variable scope: memory lifetime is ended (memory can be reused -- deallocated) – contents are **no longer valid**

Stack Segment: Support of Functions

- The stack consists of a series of "*stack frames*" or "*activation frames*", one is **created** each time a function is called at runtime
- Each frame represents a function that is currently being executed and has not yet completed (why activation frame)
- A function's stack "frame" goes away when the function returns
- Specifically, a new stack frame is
 - allocated (**pushed** on the stack) for each function call (**contents are not implicitly zeroed**)
 - deallocated (**popped** from the stack) on function return
- **Stack frame** contains:
 - Local variables, parameters of function called
 - Where to return to which caller when the function completes (the return address)



ARM Assembly Source File: Header and Footer

File Header

At the top of every
ARM source file

```
.arch    armv6           // armv6 architecture
.arm     // arm 32-bit instruction set
.fpu     vfp             // floating point co-processor
.syntax  unified         // modern syntax
```

```
// Contents of the other memory segment include .text (your code)
```

File Footer

At the bottom of every
ARM source file

```
.section .note.GNU-stack,"",%progbits // set stack/data non-exec
.end
// everything past the .end is ignored!
// Debugging notes etc
```

`.syntax unified`

- use the standard ARM assembly language syntax called *Unified Assembler Language (UAL)*

`.section .note.GNU-stack,"",%progbits`

- tells the linker to **make the stack and all data segments not-executable** (no instructions in those sections) – security measure

`.end`

- at the end of the source file, everything written after the `.end` is ignored

Function Header and Footer Assembler Directives

function entry point
address of the first
instruction in the function
Must not be a local label
(does not start with .L)

```
        .text
Function Header {
    .global myfunc           // make myfunc global for linking
    .type   myfunc, %function // define myfunc to be a function
    .equ    FP_OFF, 4        // fp offset in main stack frame
myfunc:
    // function prologue, stack frame setup
    // your code
    // function epilogue, stack frame teardown
Function Footer {
    .size myfunc, (. - myfunc)
```

.global function_name

- Exports the function name to other files. Required for main function, optional for others

.type name, %function

- The **.type** directive sets the **type of a symbol/label name**
- %function** specifies that **name** is a function (name is the address of the first instruction)

equ FP_OFF, 4

- Used for basic stack frame setup; the number 4 will change – later slides

.size name, bytes

- The **.size** directive is used to **set the size associated with a symbol**
- Used by the linker to exclude unneeded code and/or data when creating an executable file
- It is also used by the **debugger** gdb
- bytes is best calculated as an expression: (period is the current address in a memory segment)**

.size name, (. - name)

Example: Assembler Directive and Instructions

assembler directive `.equ` does not allocate any memory (NULL = 0)

Regular label `main` is associated with memory location 0x3000

Local label `.Lloop` is associated with memory location 0x3004

space.S

```
10  .equ NULL, 0
11 main:
12 3000 0310A0E1 mov r1, r3
13 .Lloop:
14 3004 043083E2 add r3, r3, 4
15 3008 001093E5 ldr r1, [r3]
16 300c 000051E3 cmp r1, NULL
17 3010 FBFFFF1A bne .Lloop
```

output generated with
`gcc -c -Wa,-ahlns space.S`
partial output is shown

Memory Contents

Warning contents shown in "reverse" byte order: Lsb – Msb

Instruction Memory Addresses (lowest 2-bits are always are 00)
Notice alignment and how addresses increase by 4 (32-bit instructions)

Preview: Return Value and Passing Parameters to Functions

(Four parameters or less)

Register	Function Call Use
r0	1 st parameter
r1	2 nd parameter
r2	3 rd parameter
r3	4 th parameter

Register	Function Return Value Use
r0	8, 16 or 32-bit result, 32-bit address or least-significant half of a 64-bit result
r1	most-significant half of a 64-bit result

- Where **r0**, **r1**, **r2**, **r3** are arm registers, the function declaration is (first four arguments):

```
r0 = function(r0, r1, r2, r3)           // 32-bit return
```

```
r0, r1 = function(r0, r1, r2, r3)      // 64-bit return - long long
```
- Each **parameter** and **return value** is limited to data that **can fit in 4 bytes or less**
- You receive **up to the first four parameters in these four registers**
- You copy up to the first four parameters into these four registers before calling a function
- For parameter values using more than 4 bytes, a pointer to the parameter is passed (we will cover this later)
- You MUST ALWAYS assume** that the called function will **alter the contents of all four registers: r0-r3**
 - In terms of C runtime support, these registers contain the copies given to the called function
 - C allows the copies to be changed in any way by the called function

Assembler Directives: Label Scope Control (Normal Labels only)

```
.extern printf
.extern fgets
.extern strcpy
.global fbuf
```

.extern <label>

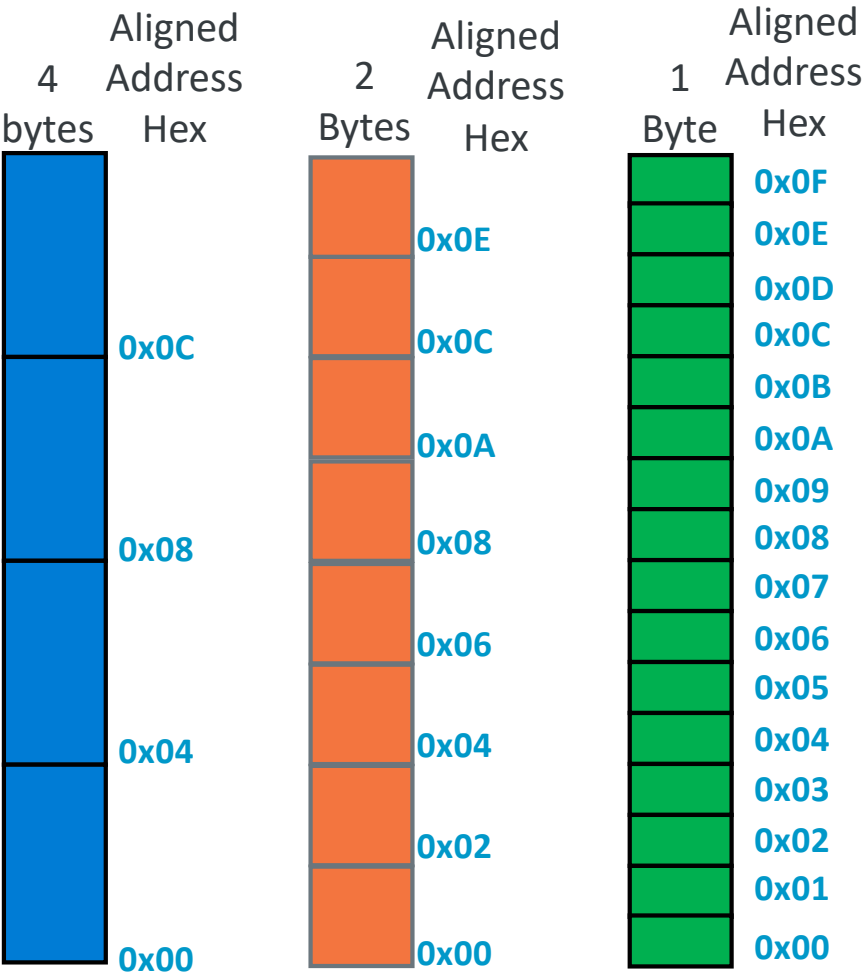
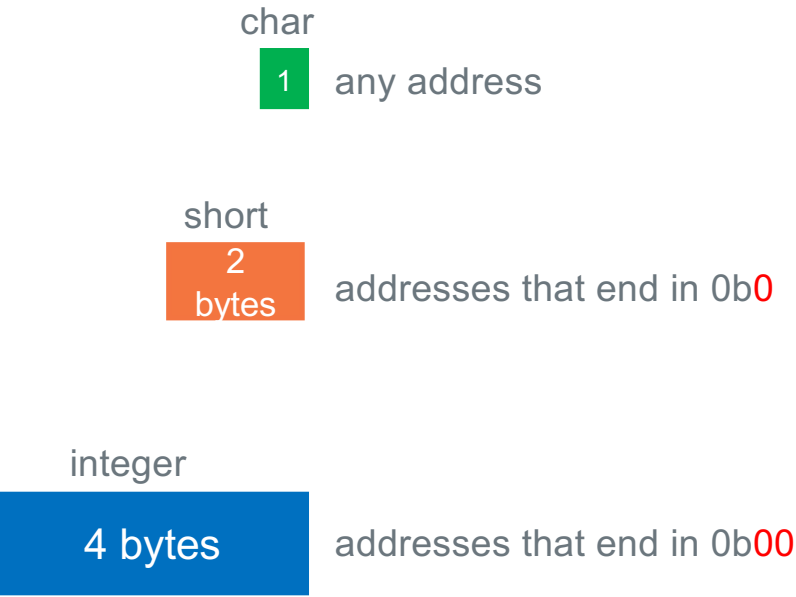
- **Imports** label (function name, symbol or a static variable name);
- An address associated with the label from another file can be used by code in this file

.global <label>

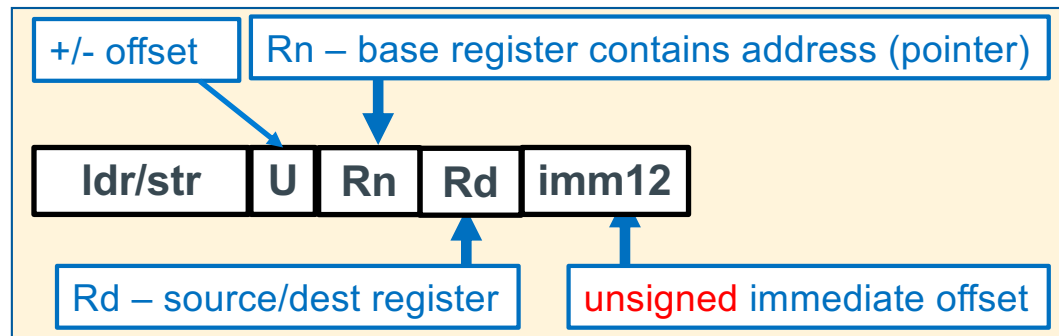
- **Exports** label (or symbol) to be visible outside the source file boundary (other assembly or c source)
 - label is either a function name or a global variable name
 - Only use with function names or static variables
- **Without** .global, labels are usually **local to the file** from the point where they are defined

Variable Alignment In Memory and Performance

Accessing **address aligned** memory on many systems **based on data type** has **the best performance** (due to hardware implementation)

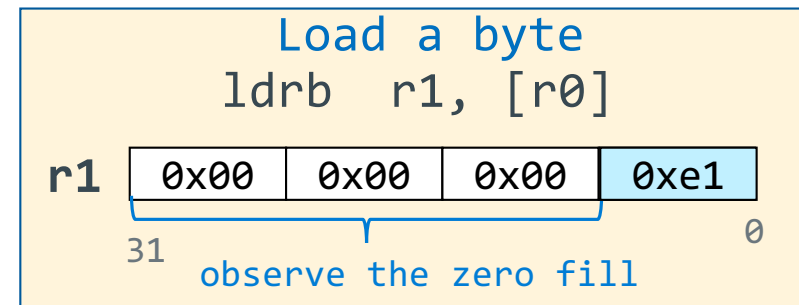
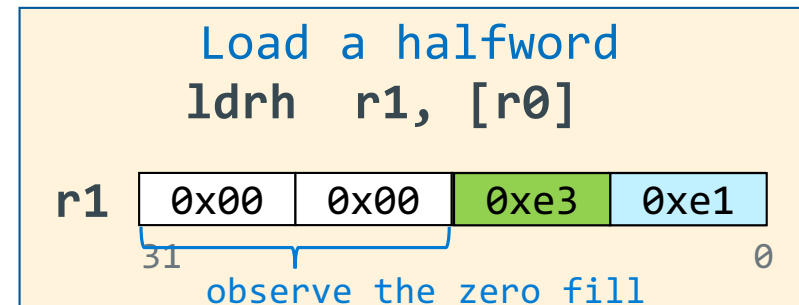
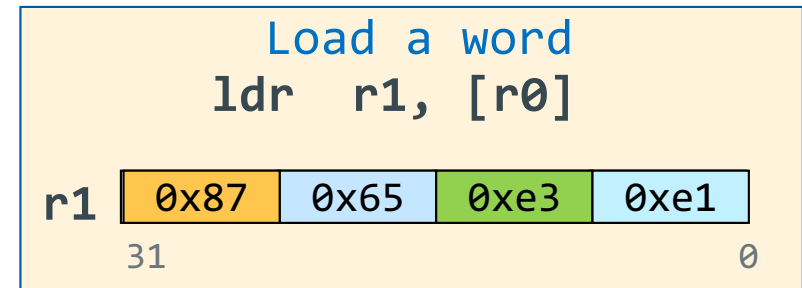
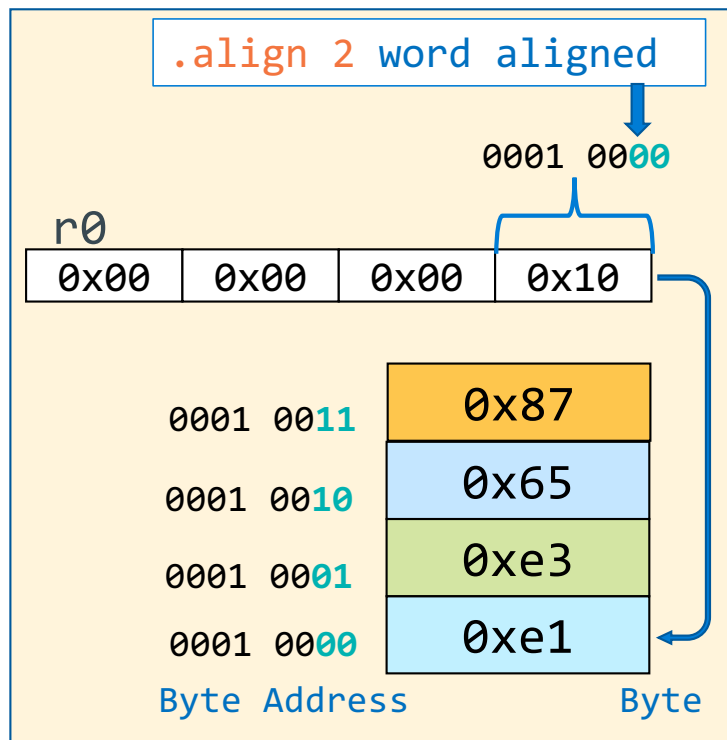


LDR/STR – Base Register + Immediate Offset Addressing

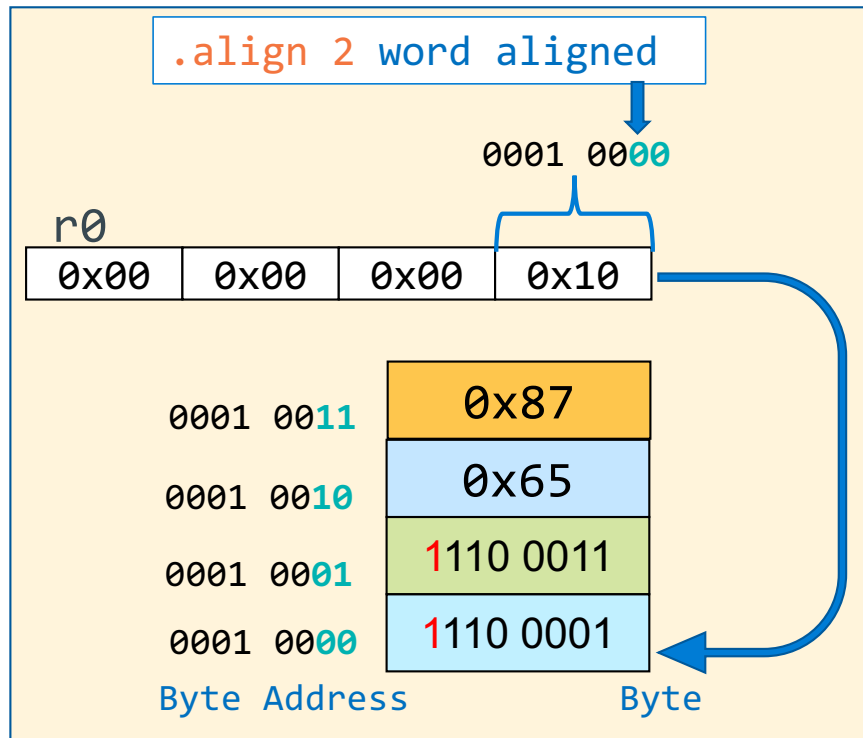


- **Register Base Addressing:**
 - Pointer Address: **Rn**; **source/destination data: Rd**
 - **Unsigned pointer address** is stored in the **base register**
- **Register Base + immediate offset Addressing:**
 - Pointer Address = register content + immediate offset $-4095 \leq \text{imm12} \leq 4095$ (bytes)
 - Unsigned offset integer **immediate value (bytes)** is added or subtracted (**U bit above says to add or subtract**) from the **pointer address** in the **base register**

Load a Byte, Half-word, Word




Signed Load a Byte, Half-word, Word




Load a word (no change)

```
ldr r1, [r0]
```



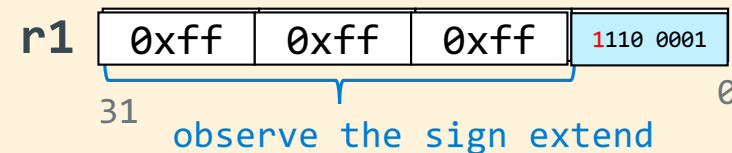
Load a halfword

```
ldrsh r1, [r0]
```

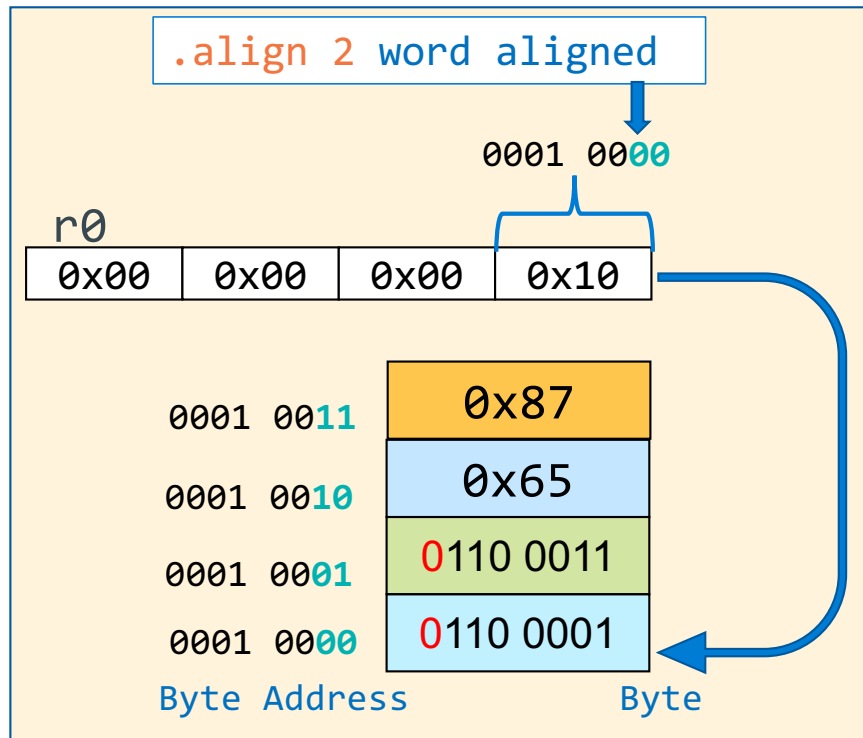


Load a byte

```
ldrsb r1, [r0]
```




Signed Load a Byte, Half-word, Word




Load a word (no change)

```
ldr r1, [r0]
```




Load a halfword

```
ldrsh r1, [r0]
```



Load a byte

```
ldrsb r1, [r0]
```



Store a Byte, Half-word, Word

initial value in r0

0x20	0x00	0x00	0x00
------	------	------	------

Store a byte
`strb r1, [r0]`

r1:

0x87	0x65	0xe3	0xe1
------	------	------	------

 31 0

Byte Address Byte

0x20000003	0x33
0x20000002	0x22
0x20000001	0x11
0x20000000	0xe1

observe other bytes NOT altered

Store a halfword
`strh r1, [r0]`

r1:

0x87	0x65	0xe3	0xe1
------	------	------	------

 31 0

Byte Address Byte

0x20000003	0x33
0x20000002	0x22
0x20000001	0xe3
0x20000000	0xe1

Store a word
`str r1, [r0]`

r1:

0x87	0x65	0xe3	0xe1
------	------	------	------

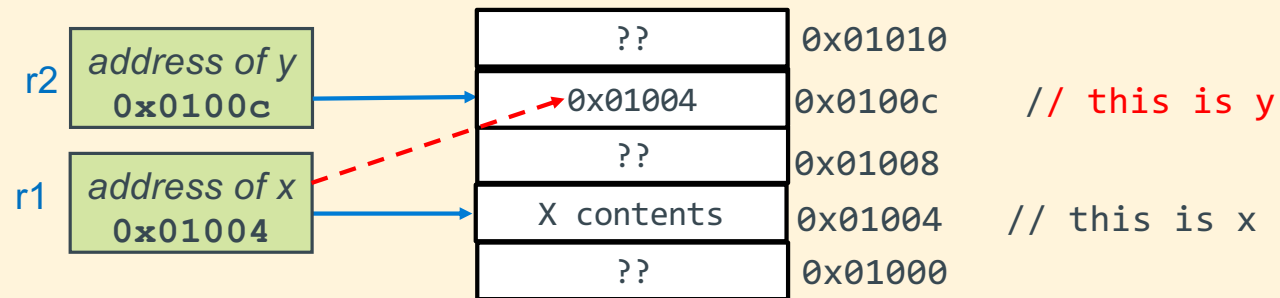
 31 0

Byte Address Byte

0x20000003	0x87
0x20000002	0x65
0x20000001	0xe3
0x20000000	0xe1

ldr/str practice - 1


r1 contains the Address of X (defined as int X) in memory; r1 points at X
r2 contains the Address of Y (defined as int *Y) in memory; r2 points at Y
write Y = &X;



```
str    r1, [r2]    // y ← &x
```

ldr/str practice - 2

r1 contains the Address of X (defined as `int *X`) in memory r1 points at X
r2 contains the Address of Y (defined as `int Y`) in memory; r2 points at Y
write `Y = *X;`



```
ldr    r3, [r1]    // r3 ← x (read 1)
ldr    r0, [r3]    // r0 ← *x (read 2)
str    r0, [r2]    // y ← *x
```