

Version 2.14


UCSD CSE 30

Computer Organization and Systems Programming

Lecture - 16

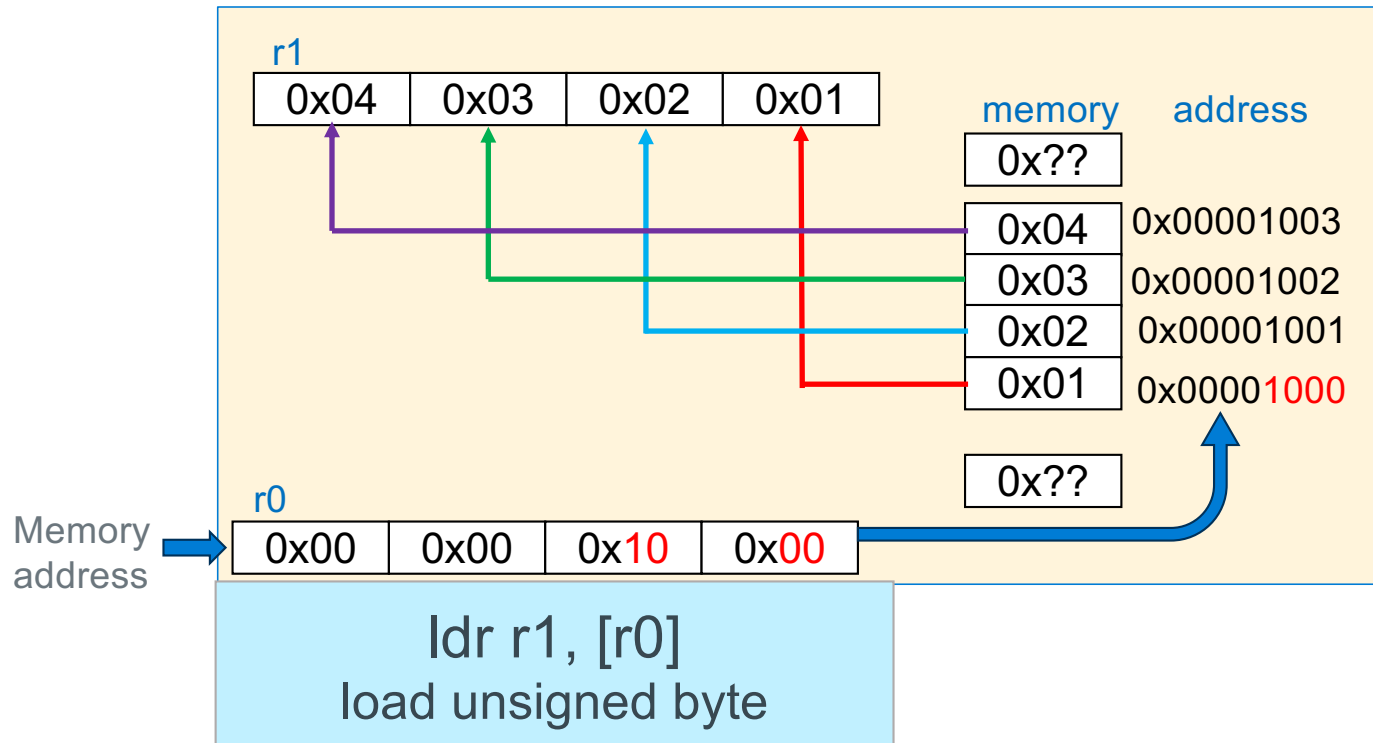
Keith Muller

Colossus MkII - 1944

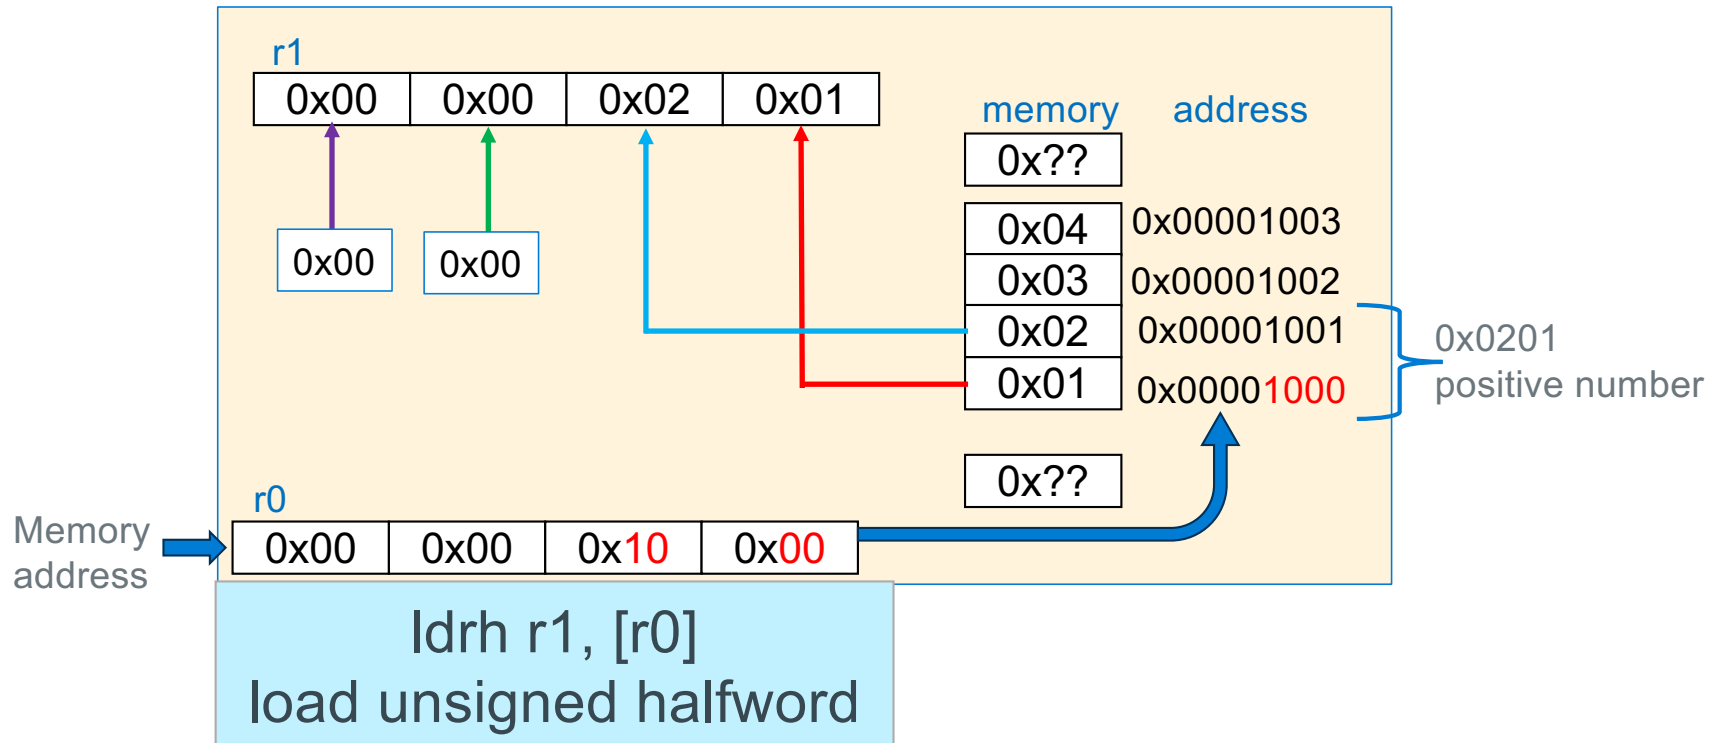




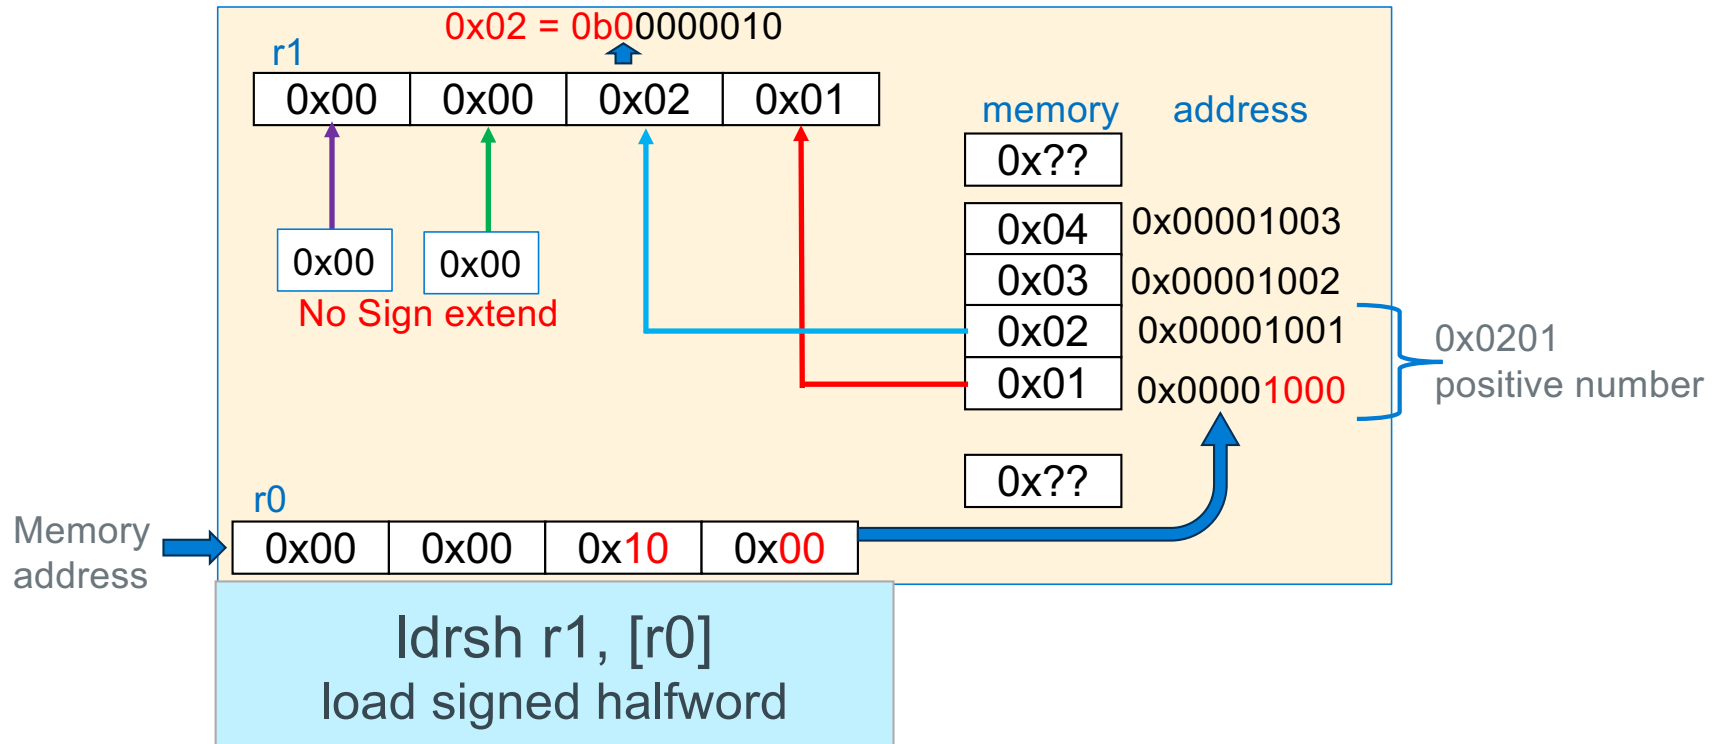
Loading 32-bit Registers From Memory, 32-bit



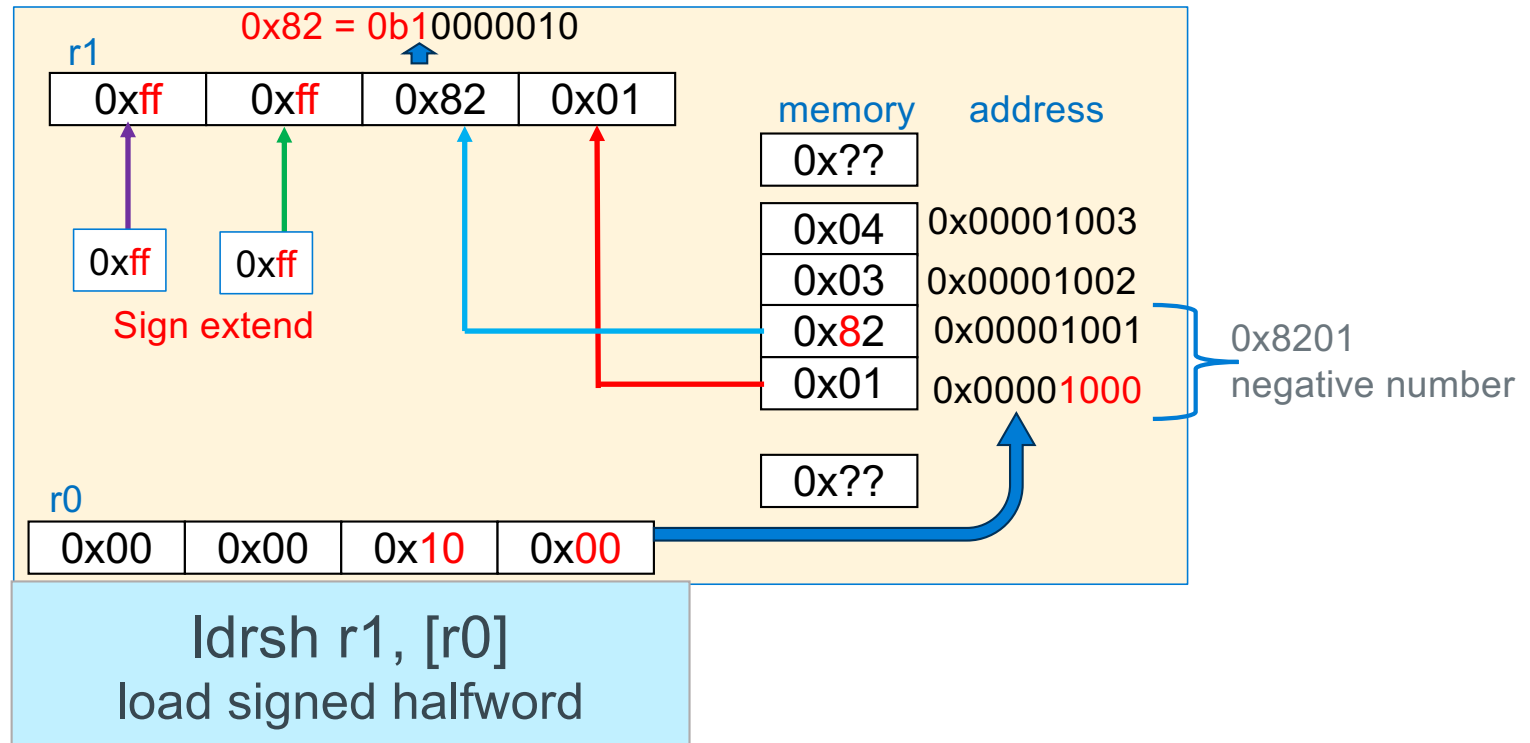
Loading 32-bit Registers From Memory, 16-bit



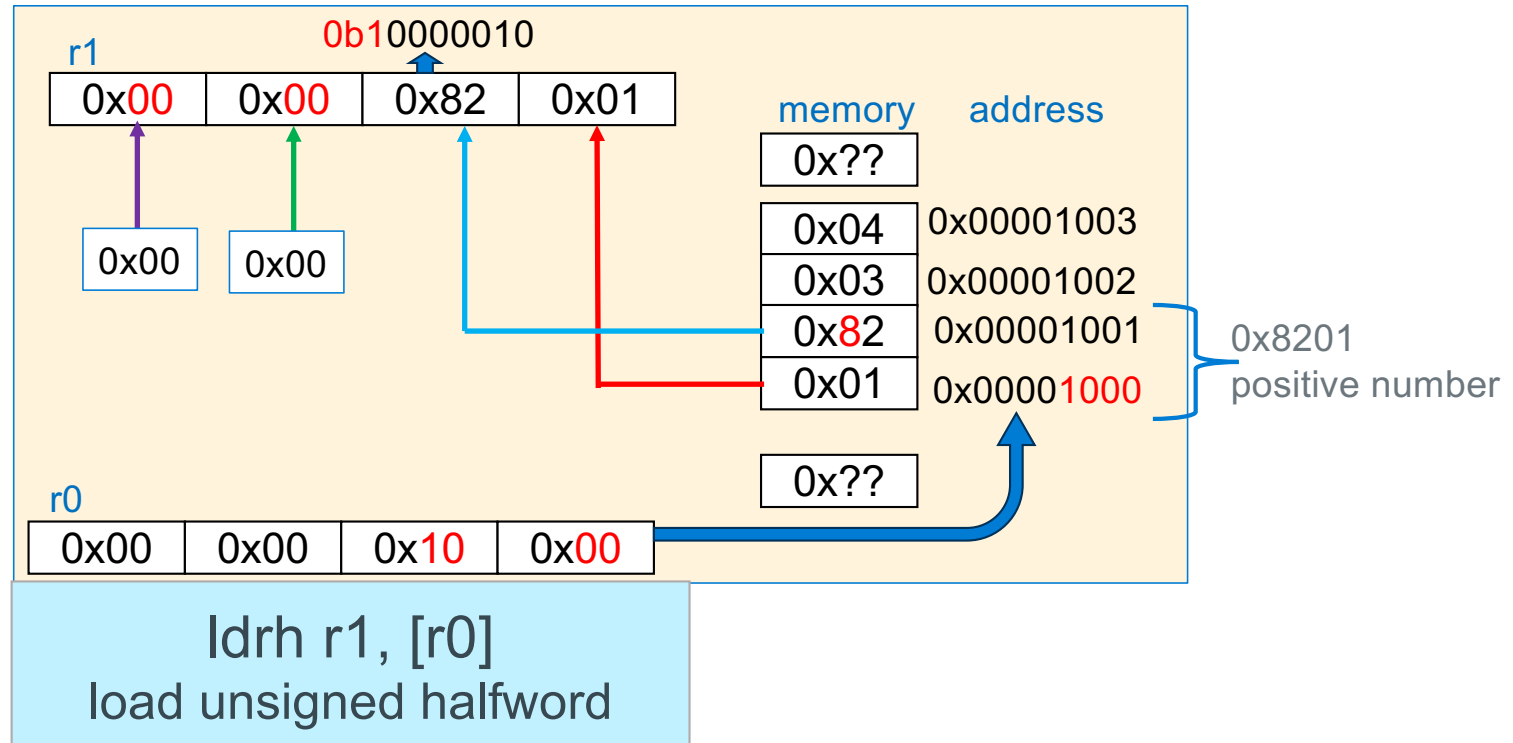
Loading 32-bit Registers From Memory, 16-bit



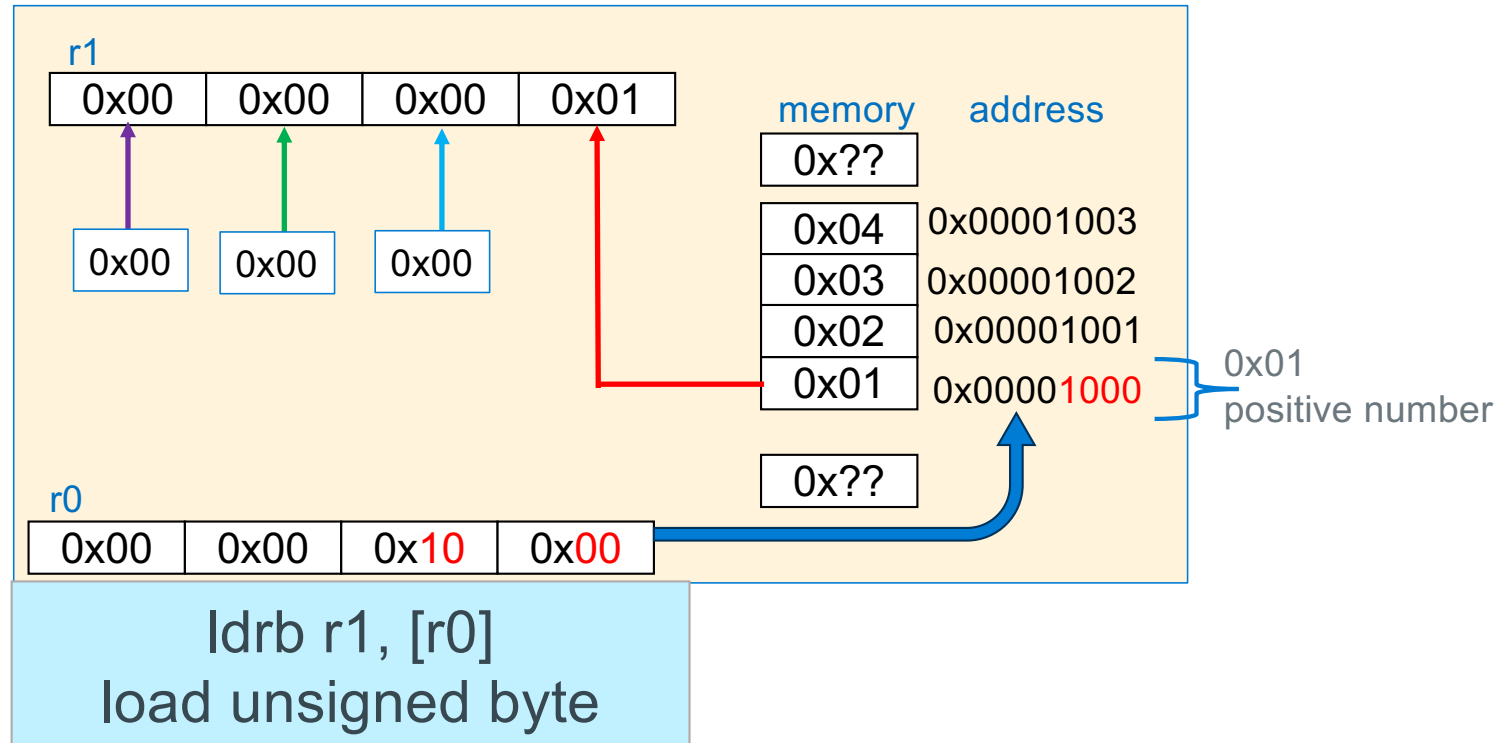
Loading 32-bit Registers From Memory, 16-bit Signed



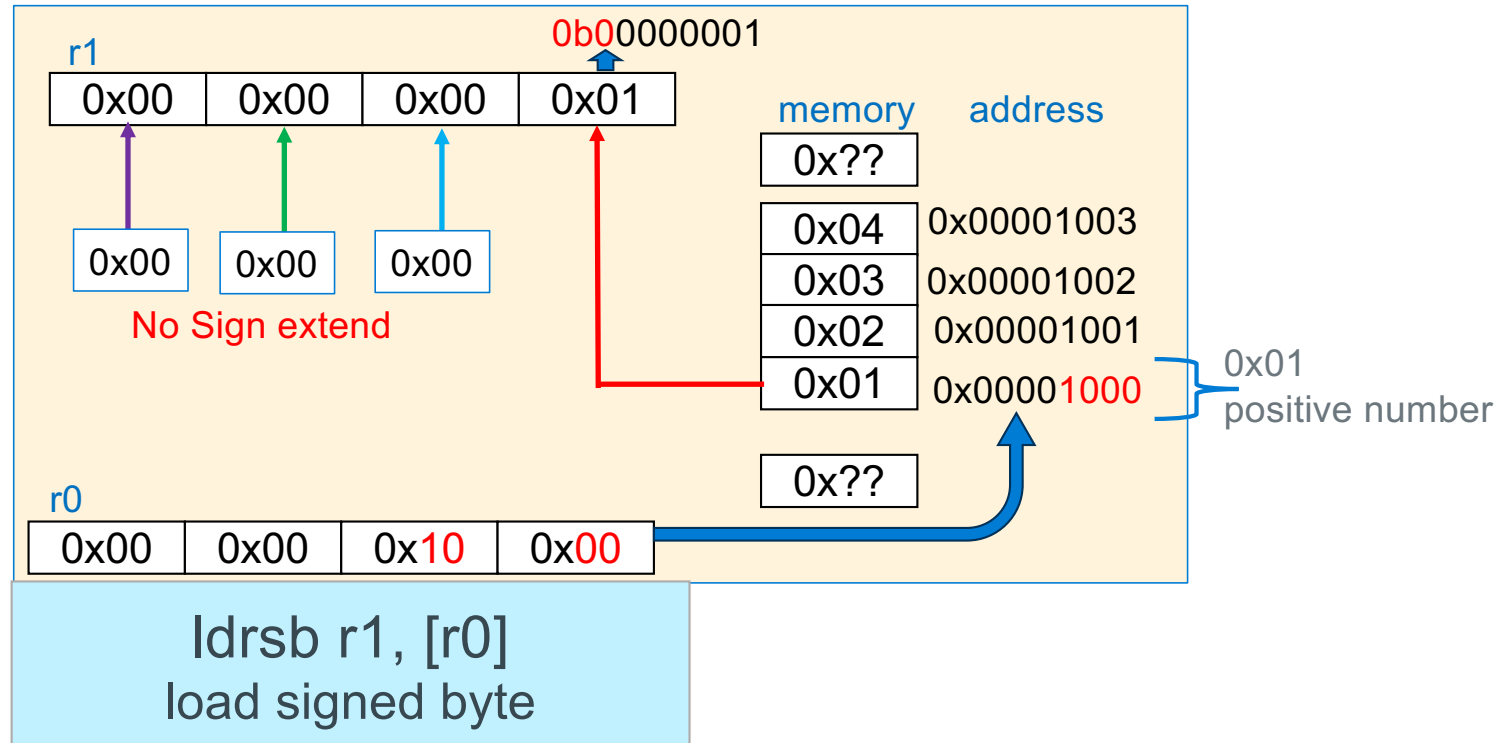
Loading 32-bit Registers From Memory, 16-bit Unsigned



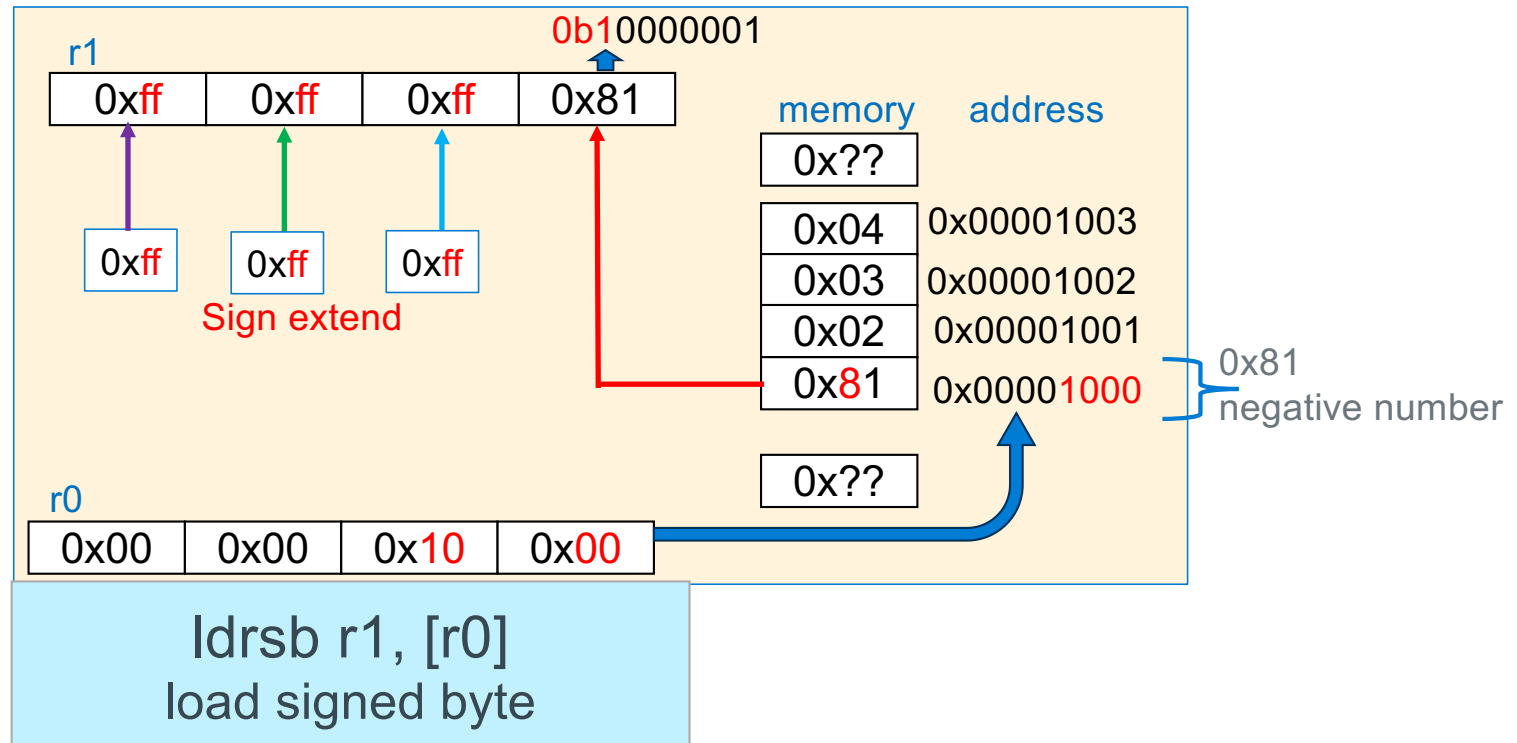
Loading 32-bit Registers From Memory, 8-bit



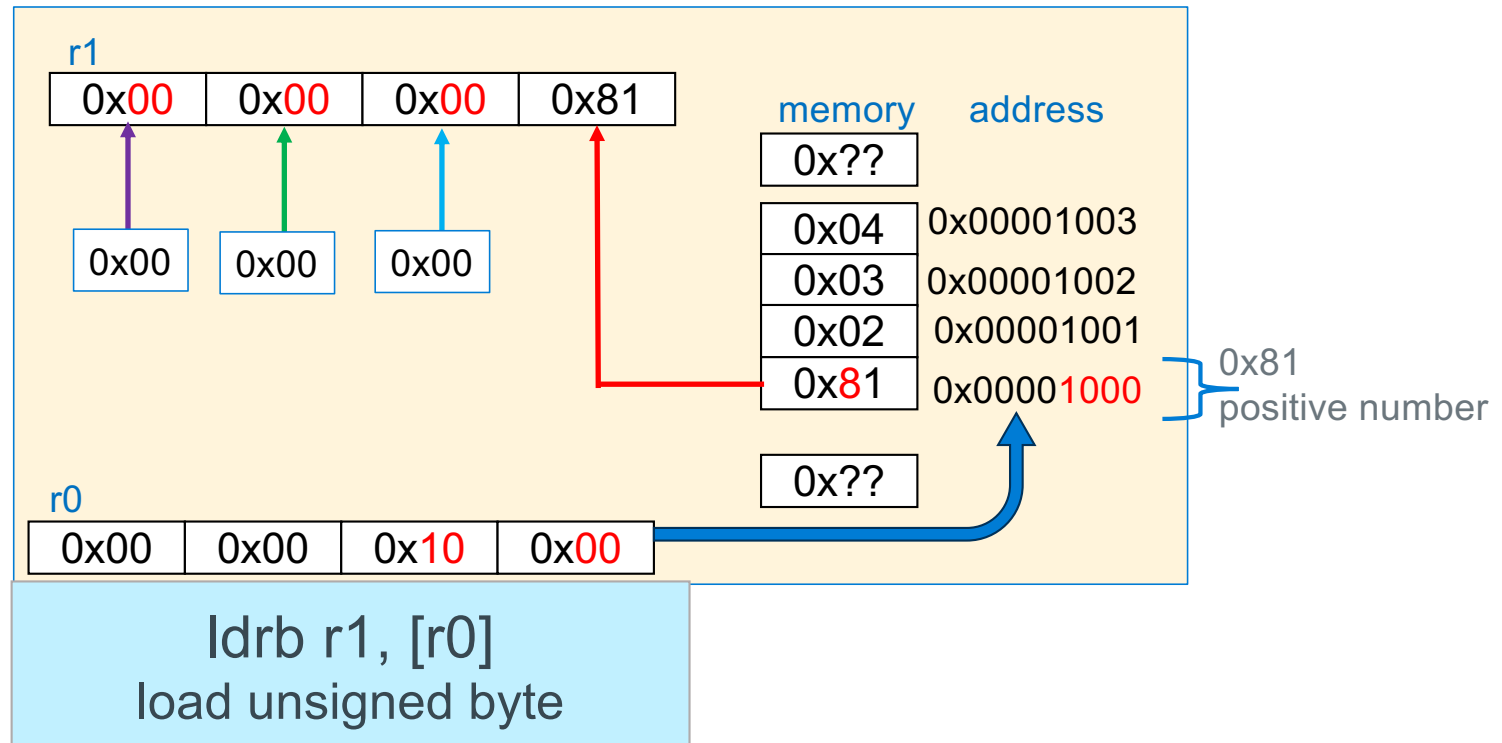
Loading 32-bit Registers From Memory, 8-bit



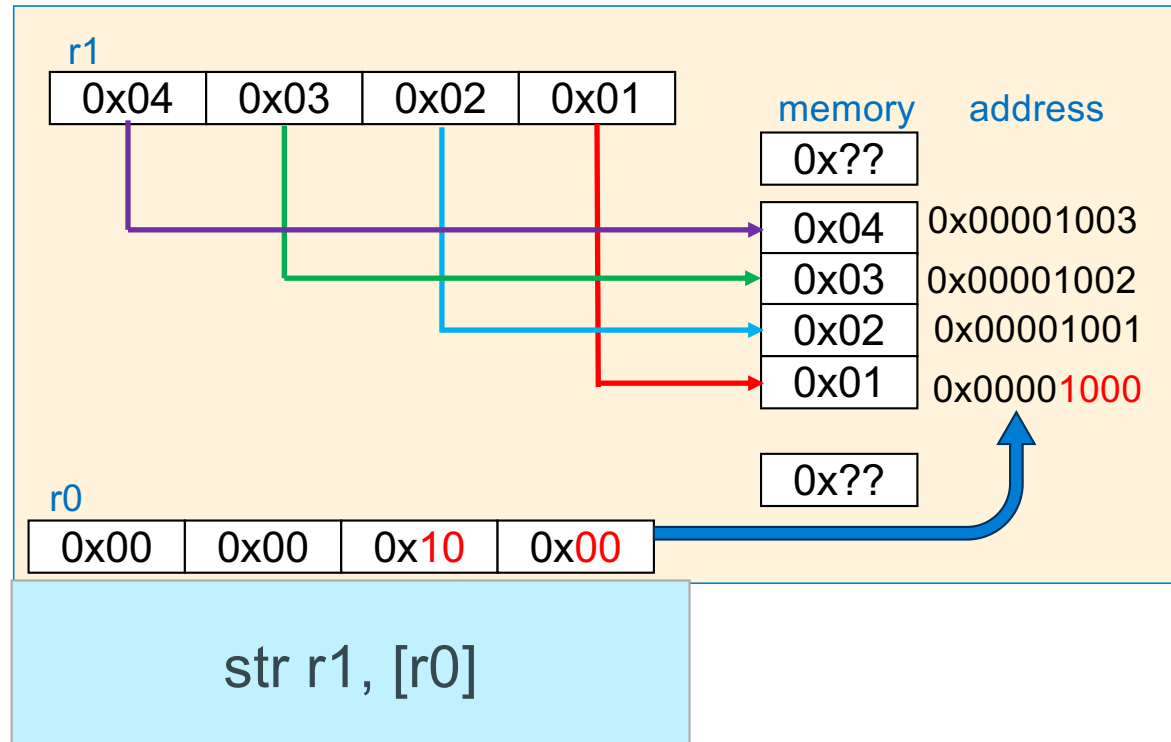
Loading 32-bit Registers From Memory, 8-bit Signed



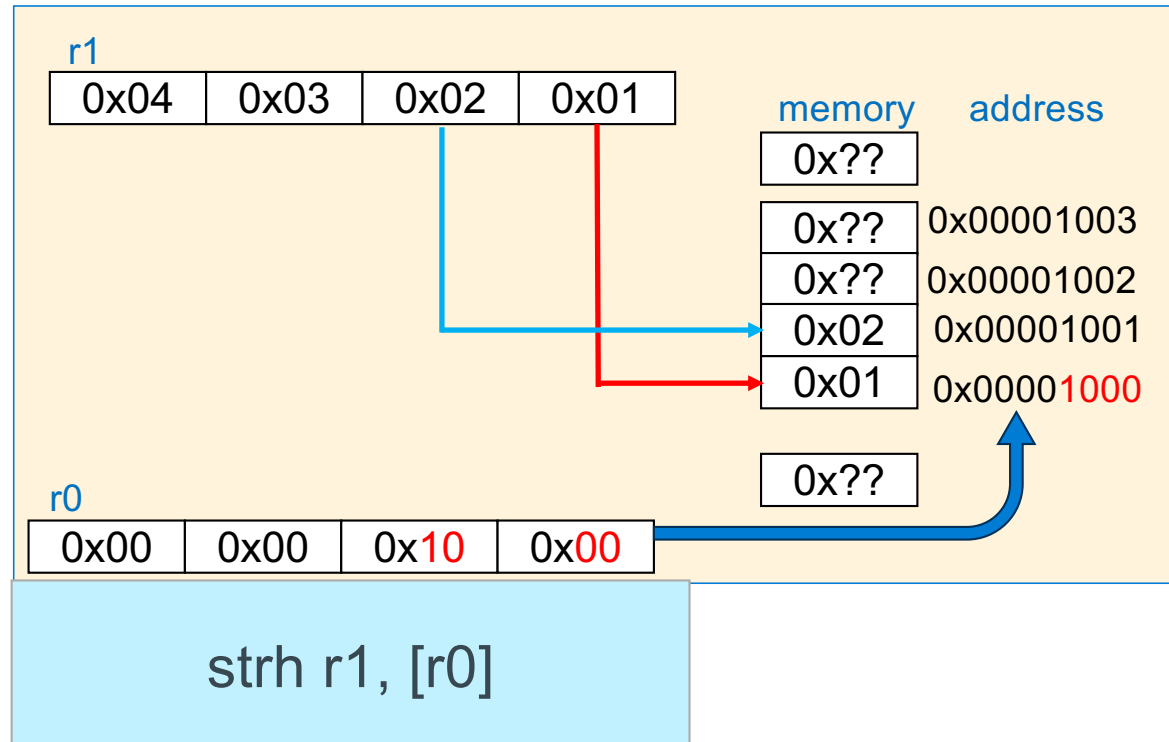
Loading 32-bit Registers From Memory, 8-bit Signed



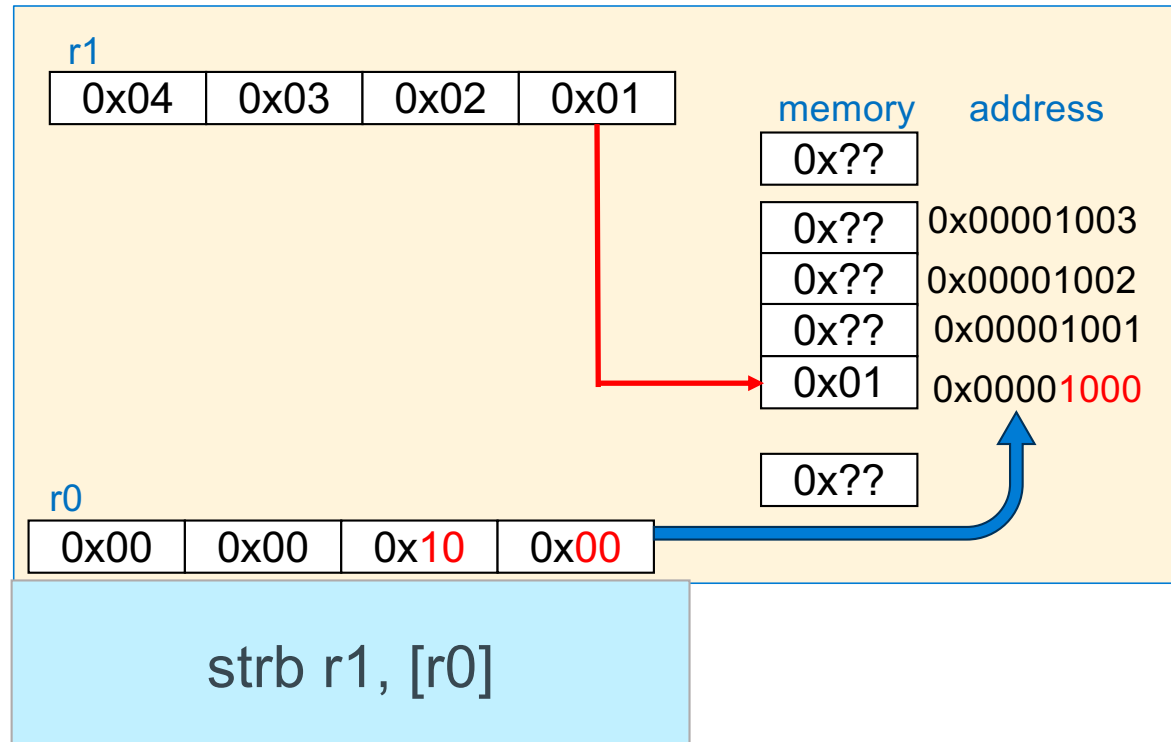
Storing 32-bit Registers To Memory, 32-bit



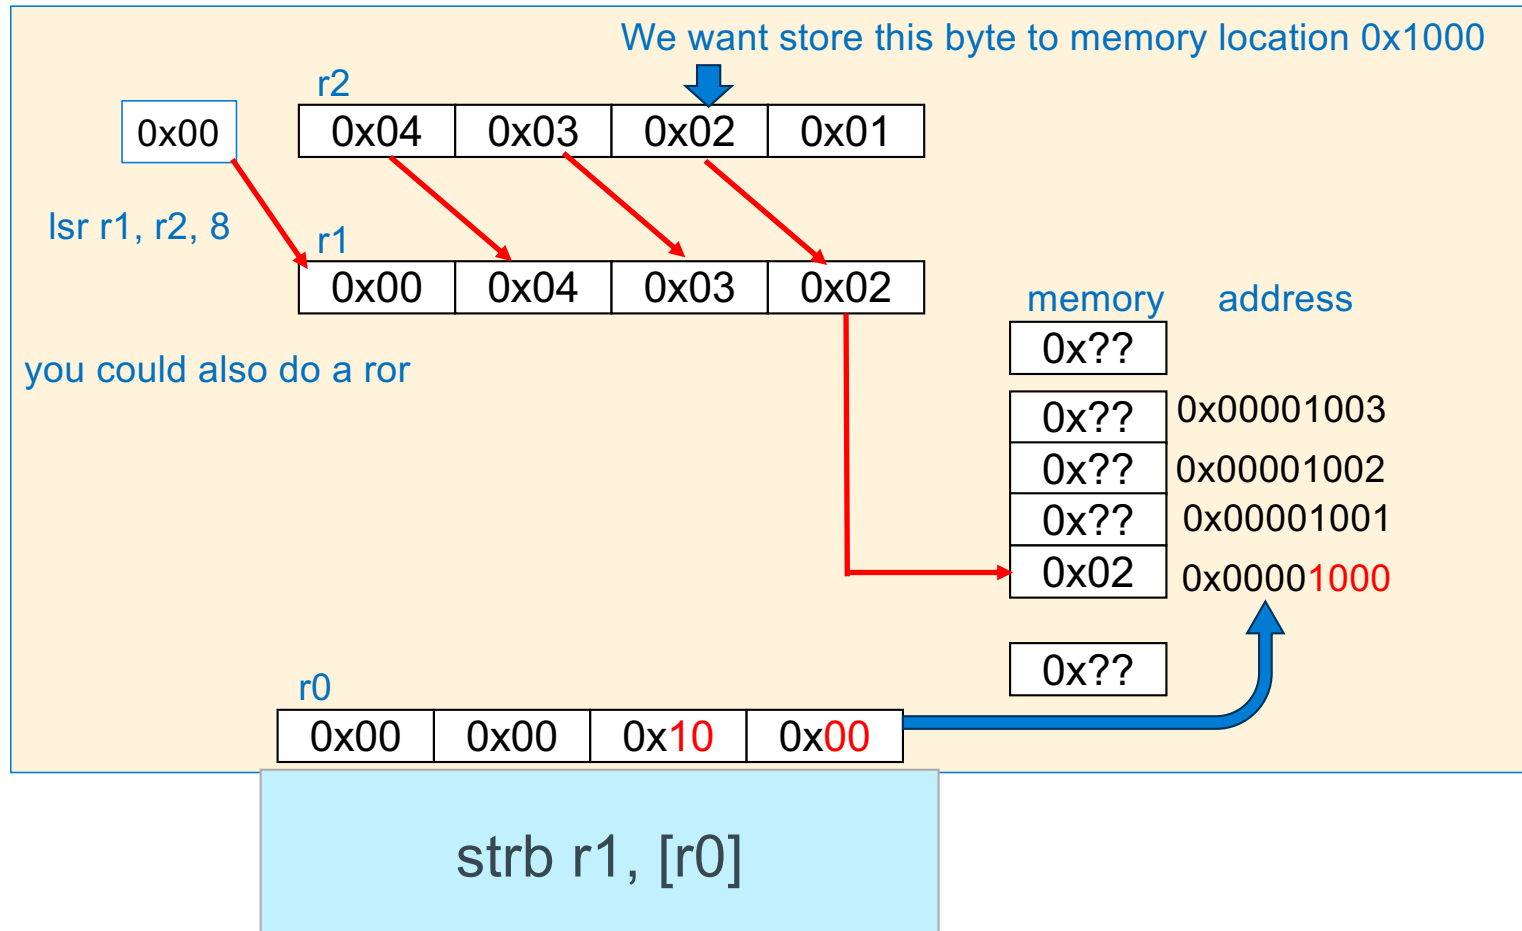
Storing 32-bit Registers To Memory, 16-bit



Storing 32-bit Registers To Memory, 8-bit



Storing 32-bit Registers To Memory, 8-bit – Storing different byte



using ldr/str: array copy

```
#include <stdio.h>
#include <stdlib.h>
#define SZ 6

void icpy(int *, int *, int);

int main(void)
{
    int src[SZ] = {1, 2, 3, 4, 5, 6};
    int dst[SZ];

    icpy(src, dst, SZ);
    for (int i = 0; i < SZ; i++)
        printf("%d\n", *(dst + i));

    return EXIT_SUCCESS;
}
```

```
void icpy(int *src, int *dst, int cnt)
{
    int *end = src + cnt;

    if (cnt <= 0)
        return;
    do {
        *dst++ = *src++;
    } while (src < end);
    return;
}
```


Base Register version

```
.arch armv6
.arm
.fpu vfp
.syntax unified
.text
.global icpy
.type icpy, %function
.equ FP_OFF, 12

// r0 contains int *src
// r1 contains int *dst
// r2 contains int cnt
// r3 use as loop term pointer
// r4 use as temp

icpy:
    push    {r4, r5, fp, lr}
    add     fp, sp, FP_OFF
    // see right ->
    sub     sp, fp, FP_OFF
    pop     {r4, r5, fp, lr}
    bx      lr
.size icpy, (. - icpy)
.end
```

```
    cmp     r2, 0
    ble     .Ldone
    lsl     r2, r2, 2 //convert cnt to int size
    add     r3, r0, r2 // loop term pointer

.Ldo:
    ldr     r4, [r0] // load from src
    str     r4, [r1] // store to dest

    add     r0, r0, 4 // src++
    add     r1, r1, 4 // dst++

    cmp     r0, r3 // src < term pointer?
    blt     .Ldo
    .Ldone:
```

pre loop guard

loop guard

Base Register + Register Offset Version

```
.arch armv6
.arm
.fpu vfp
.syntax unified
.text
.global icpy
.type icpy, %function
.equ FP_OFF, 12
// r0 contains int *src
// r1 contains int *dst
// r2 contains int cnt
// r3 use as loop counter
// r4 use as temp
```

```
icpy:
    push    {r4, r5, fp, lr}
    add     fp, sp, FP_OFF
    // see right ->
    sub     sp, fp, FP_OFF
    pop     {r4, r5, fp, lr}
    bx      lr
    .size icpy, (. - cpy)
.end
```

```
    cmp     r2, 0
    ble     .Ldone
    lsl     r2, r2, 2
    mov     r3, 0
    .Ldo:
    ldr     r4, [r0, r3]
    str     r4, [r1, r3]
    add     r3, r3, 4
    cmp     r3, r2
    blt     .Ldo
    .Ldone:
```

pre loop guard

loop guard

one increment
covers both arrays

Base Register + Register Offset With chars

```
#include <stdio.h>
#include <stdlib.h>
#define SZ 6
void cpy(char *, char *, int);
int main(void)
{
    char src[SZ] =
        {'a', 'b', 'c', 'd', 'e', '\0'};
    char dst[SZ];

    cpy(src, dst, SZ);
    printf("%s\n", dst);
    return EXIT_SUCCESS;
}
```

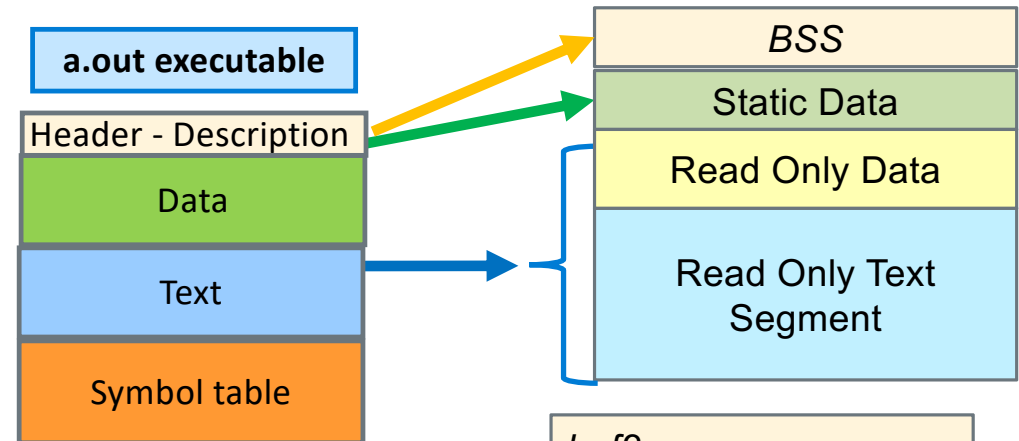
```
    cmp    r2, 0
    ble    .Ldone

    mov    r3, 0           // initialize counter
.Ldo:
    ldrb   r4, [r0, r3]    // load from src
    strb   r4, [r1, r3]    // store to dest
    add    r3, r3, 1       // counter++
    cmp    r3, r2          // count < r3
    blt    .Ldo

.Ldone:
```

What is the conceptual difference between .bss and .data?

- All static variables that do not specify an initial value default to an initial value of 0 and are placed in .bss segment
- To save file system space in the executable file (the a.out file) the assembler collapses these .bss variables to a location and size "table"
- .data segment variables use the same space in the executable file as they have in memory
- .section .rodata is handled the same as .data



```
// these are .bss variables
int buf1[4096];
int buf2[4096];
```

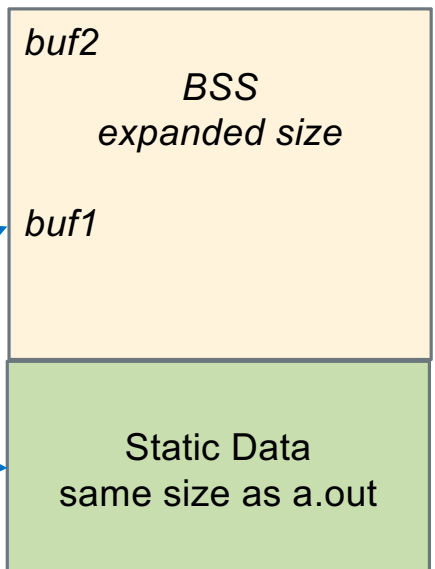
just big enough for address, size

```
// these are .data variables
int table[] = {1,2};
char string[] = "CSE30!!";
```

same size as specified

| | | | |
|-------------------|-----|-----|-----|
| buf2 address size | | | |
| buf1 address size | | | |
| 2 | | | |
| 1 | | | |
| '\0' | '!' | '!' | '0' |
| '3' | 'E' | 'S' | 'c' |

executable file



low main mameory

Variable Alignment In .data, .bss and .section .rodata


Use .align directive to force the assembler to align the address of the next variable defined after the .align

char **1** any address

short **2 bytes** addresses that end in 0b**0**

integer **4 bytes** addresses that end in 0b**00**

| SIZE Alignment Requirements | Starting Address must end in | Align Directive |
|---|--------------------------------|-----------------|
| 8-bit char -1 byte | 0b.. 0 or 0b.. 1 | |
| 16-bit int -2 bytes | 0b.. 0 | .align 1 |
| 32-bit int -4 bytes pointers, all arrays | 0b.. 00 | .align 2 |



Defining Static Variables: Allocation and Initialization

| Variable SIZE | Directive | .align | C static variable Definition | Assembler static variable Definition |
|-------------------------|----------------|-------------|---|--|
| 8-bit char (1 byte) | .byte | | char chx = 'A'; char string[] = {'A','B','C', 0}; | chx: .byte 'A' string: .byte 'A','B',0x42,0 |
| 16-bit int (2 bytes) | .short | .align 1 | short length = 0x55aa; | length: .short 0x55aa |
| 32-bit int (4 bytes) | .word .long | .align 2 | int dist = 5; int *distptr = &dist; unsigned int mask = 0xaa55; int array[] = {12,~0x1,0xCD,-1}; | dist: .word 5 distptr: .word dist mask: .word 0xaa55 array: .word 12,~0x1,0xCD,-3 |
| string with '\0' | .string | | char class[] = "cse30"; | class: .string "cse30" |

Rule: Place the .align above the variable

.align 1

len: .short 0x55aa

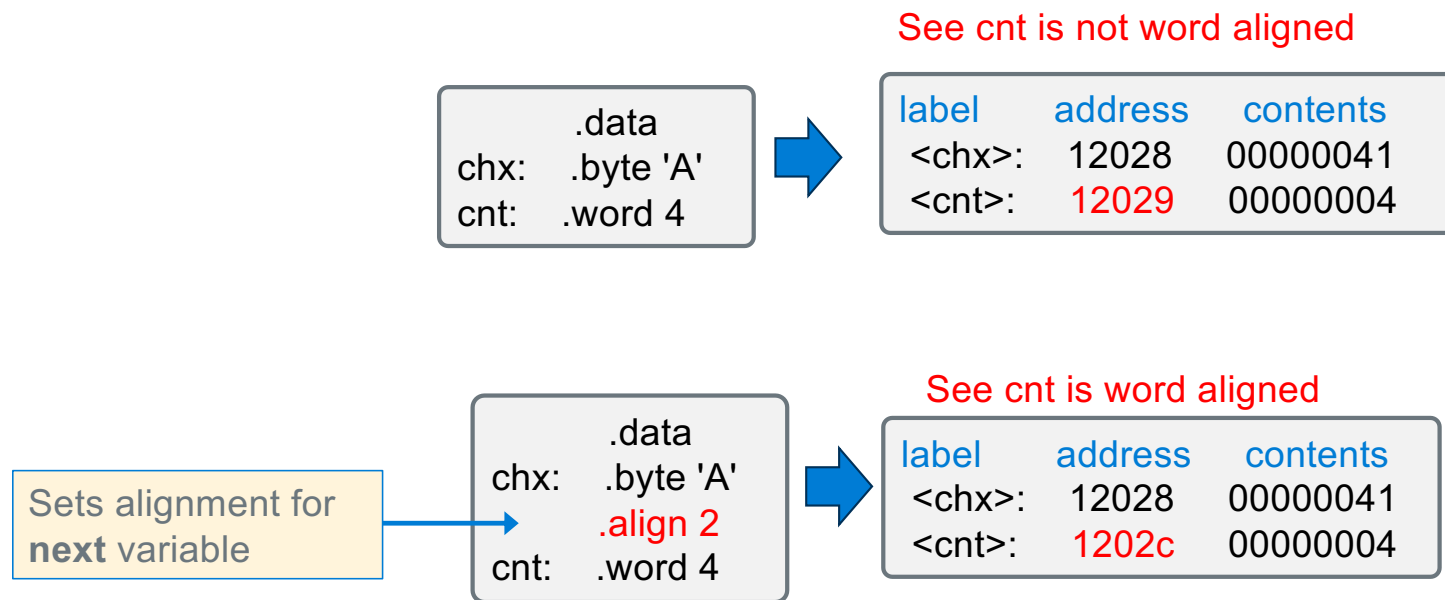
Rule: use .align 2 before every array regardless of type

Rule: place variables with explicit initialized values in a .data segment

Rule: place variables with no explicit initiali value (default to 0) in .bss segment

Rule: place string literals in .section .rodata and use a local label (.Llabel:)

Defining Static Variables: Why the .align?



Defining Static variables

```
// format: <var_name> is the address, <value> is the initial value
    var_name: <directive> <value>, <value>, ...

// Use regular labels for all <var_name> if anonymous use local labels .Llabel
```

```
.bss
// put all static variables without an explicit initial value here
// until another section directive is seen everything from this point is in .bss
// format: the value field if specified must be zero in .bss
.align 2
count: .word 0
buf:   .size 400      // int buf[100];

.data
put all static variables with an explicit initial value here
.align 2
array: .word 1, 2, 3, 4 // int array[] = {1, 2, 3, 4};

.section .rodata
// put all immutable string literals here variables
.align 2
.Lmess: .string "count is %d size is %d\n" // for a printf
```


Defining Static Array Variables (large Arrays)

Label: `.space <size>, <fill>`

`.space size, fill`

- Allocates `size` bytes, each of which contain the value `fill`
- If the comma and `fill` are **omitted**, `fill` is assumed to be **zero**
- if used in `.bss` section: Must be used **without a specified fill**

```
.bss
int_buf:  .space 400    // int int_buf[100];
          .align 2
char_buf: .space 100    // char char_buf[100];
          .data
one_buf:  .space 100, 1 // 100 bytes each byte filled with 1
```

Loading Static variables into a register

- Tell the assembler load the address (Lvalue) of a label into a register:
`ldr Rd, =Label // Rd = address`
- Tell the assembler load the contents into a register
- `ldr R0, [Rd] // Rd = address`
- *Example to the right: $y = x$;*

load a static **memory** variable

1. load the pointer to the memory
2. read (load) from *pointer

store to a static **memory** variable

1. load the pointer to the memory
2. write (store) to *pointer

```
.bss
y: .space 4

.data
x: .word 200

.text
// function header
main:

// load the address, then contents
// using r2

ldr r2, =x      // int *r2 = &x
ldr r2, [r2]    // r2 = *r2;

// &x was only needed once above
// Note: r2 was a pointer then an int
// no "type" checking in assembly!

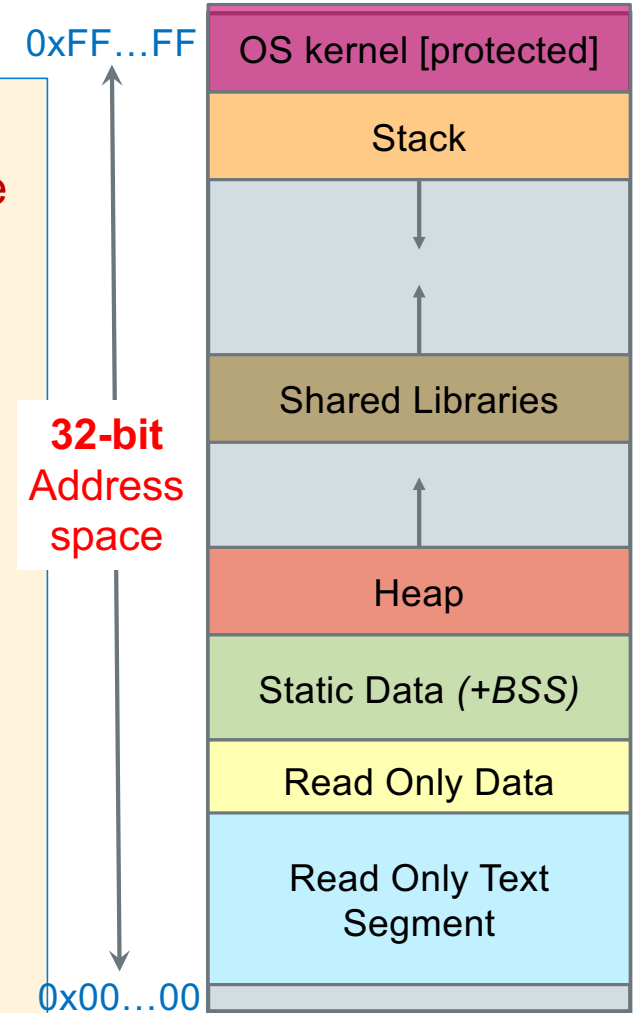
// store the contents of r2

ldr r1, =y      // int *r1 = &y
str r2, [r1]    // *r1 = r2
```

x

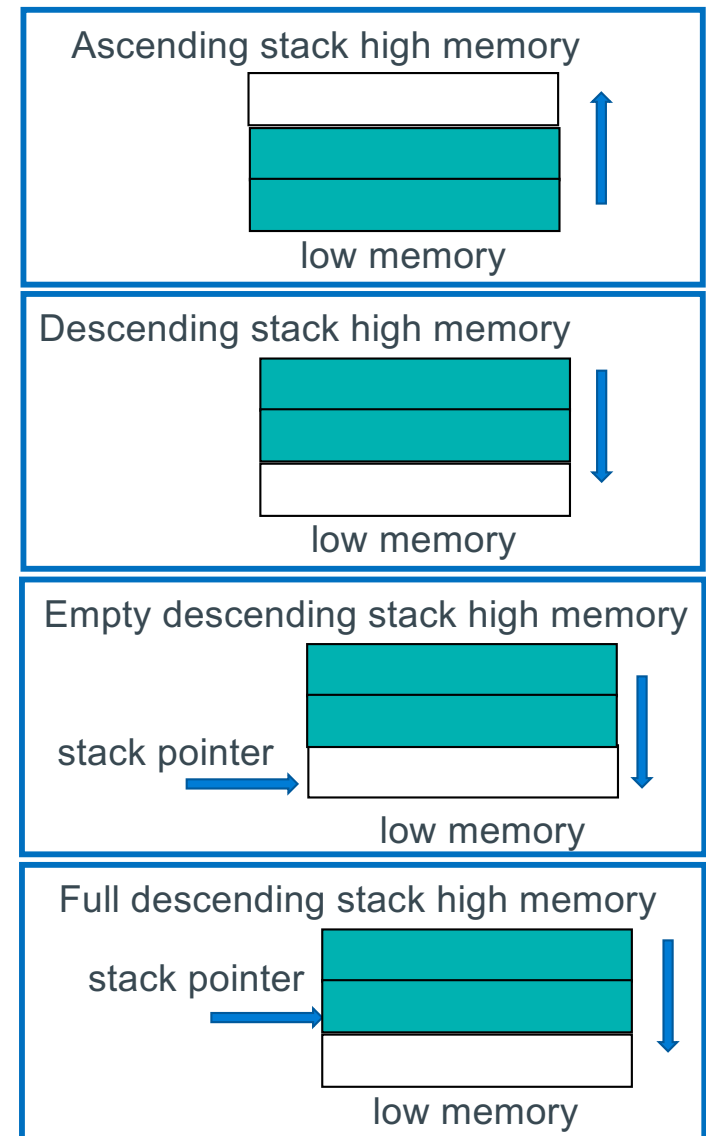
Stack Segment: Support of Functions

- The stack consists of a series of "*stack frames*" or "*activation frames*", one is **created** each time a function is called at runtime
- Each frame represents a function that is currently being executed and has not yet completed (why activation frame)
- A function's stack "frame" goes away when the function returns
- Specifically, a new stack frame is
 - allocated (**pushed** on the stack) for each function call (**contents are not implicitly zeroed**)
 - deallocated (**popped** from the stack) on function return
- **Stack frame** contains:
 - Local variables, parameters of function called
 - Where to return to which caller when the function completes (the return address)



Stack types

- A Stack Implements a **last-in first-out** (LIFO) protocol
- Each time a **function is called**, a **stack frame is activated**
 - space is allocated by moving the stack pointer
 - push adds space, pop removes space
- Stack growth direction
 - **Ascending stack**: grows from low memory towards high memory (adding to the sp to allocate memory)
 - **Descending stack**: grows from high memory towards low memory (subtracting from the sp to allocate memory)
- Full versus empty stacks
 - **Empty stack**: **stack pointer** (sp) points at the **next word address** after the last item pushed on the stack
 - **Full stack**: **stack pointer** (sp) points at the **last item pushed on the stack**
- ARM on Linux uses a **full descending stack**



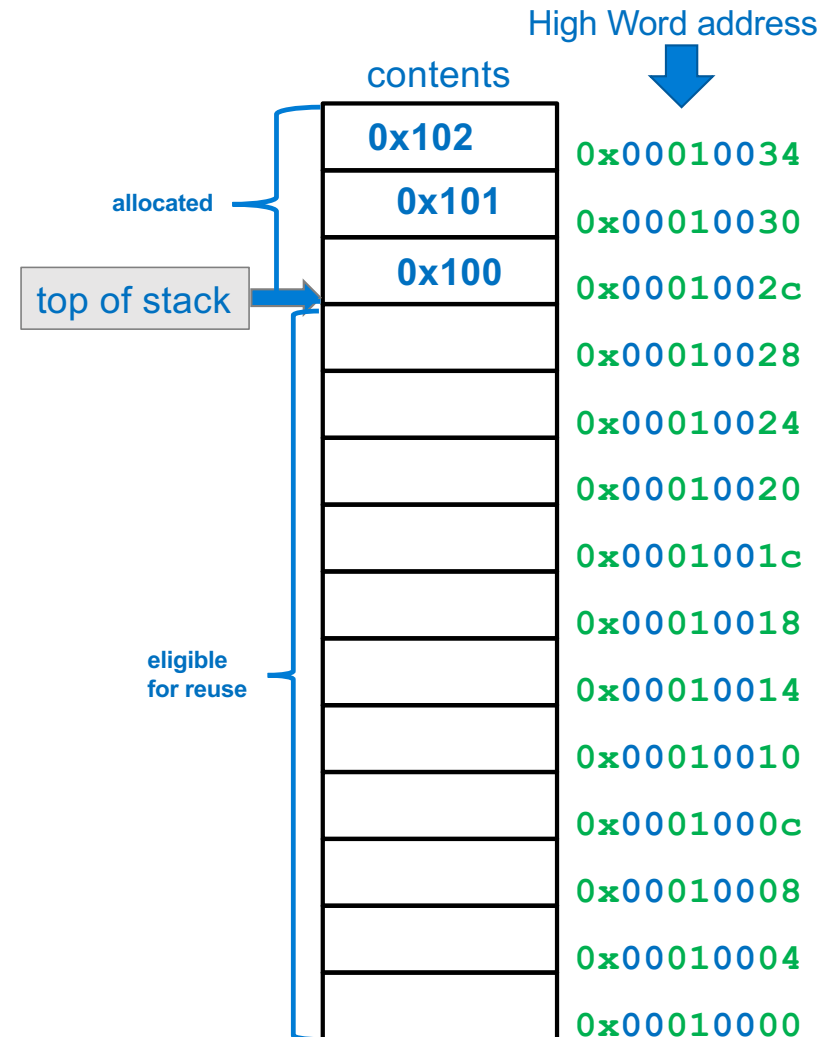
Arm: Stack Operation

- **Stack** is expandable and grows downward from high memory address towards low memory address
- **Stack pointer (sp)** always points at the **top of stack**
 - contains the starting address of the top element
- New items are **pushed** (*added*) onto the **top of the stack** by **subtracting** from the stack pointer the **size of the element** and then writing the element

push (sp - element size) & write

- Existing items are **popped** (*removed*) from the top of the stack by **adding** to the stack pointer the **size of the element** (leaving the **old contents unchanged**)


pop (sp + element size)



The Stack

Each function **call** has its own *stack frame* for its own copy of variables

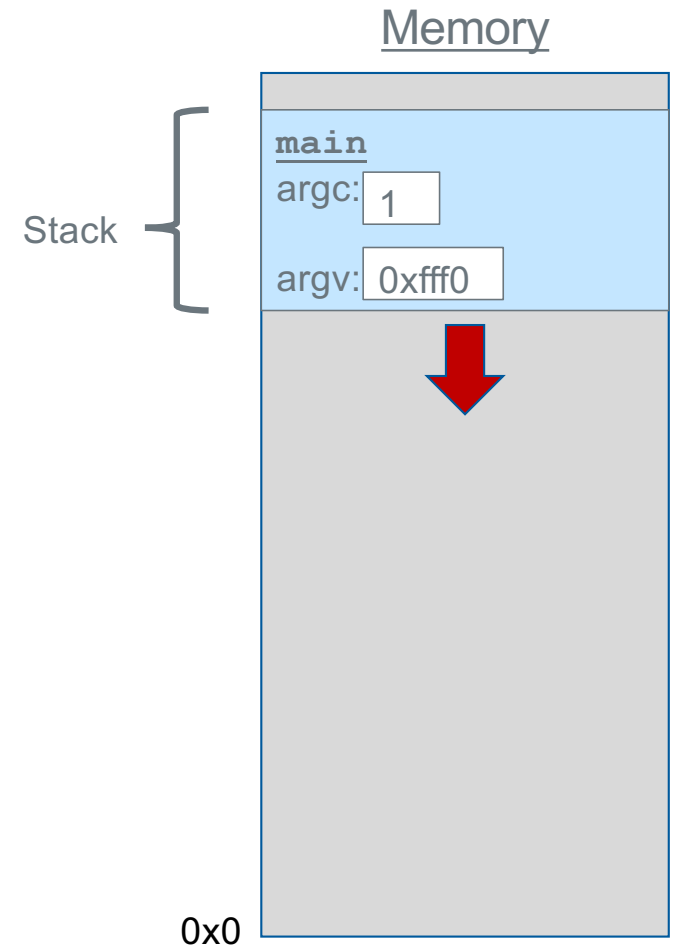
```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.


```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.

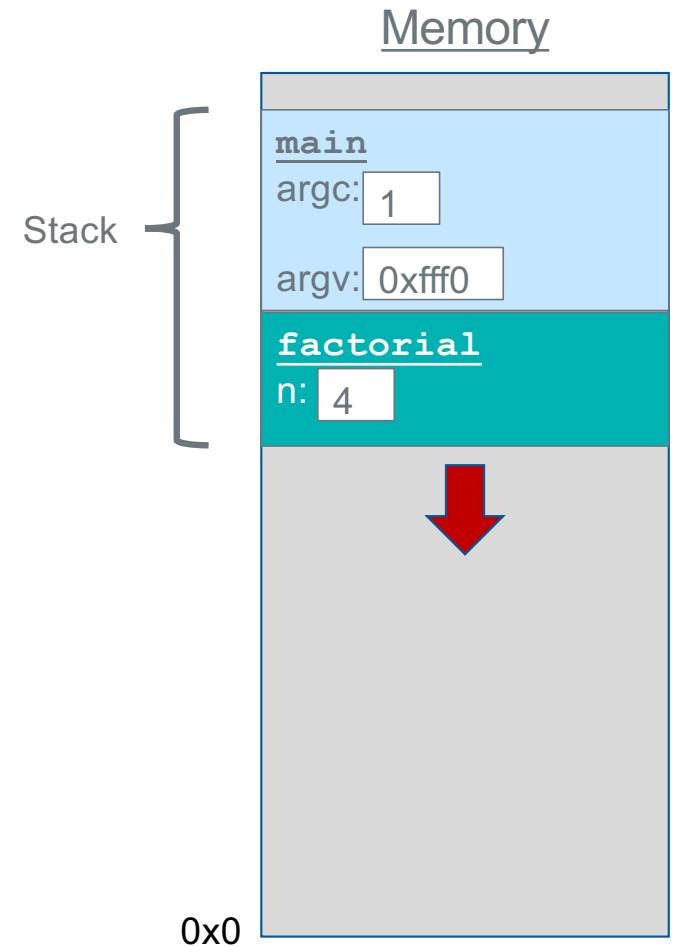
```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.


```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.

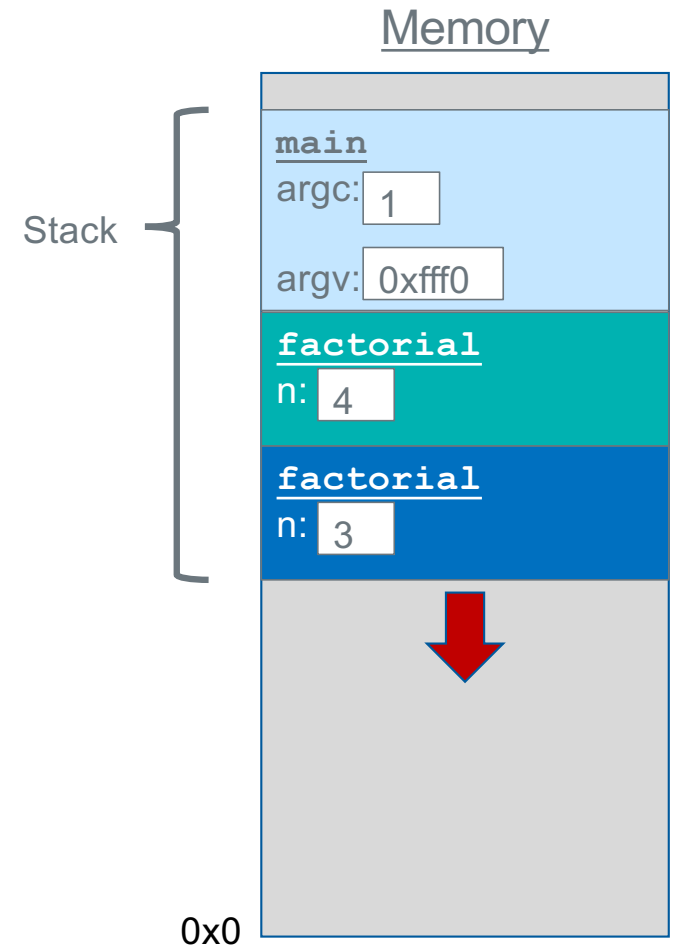
```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.

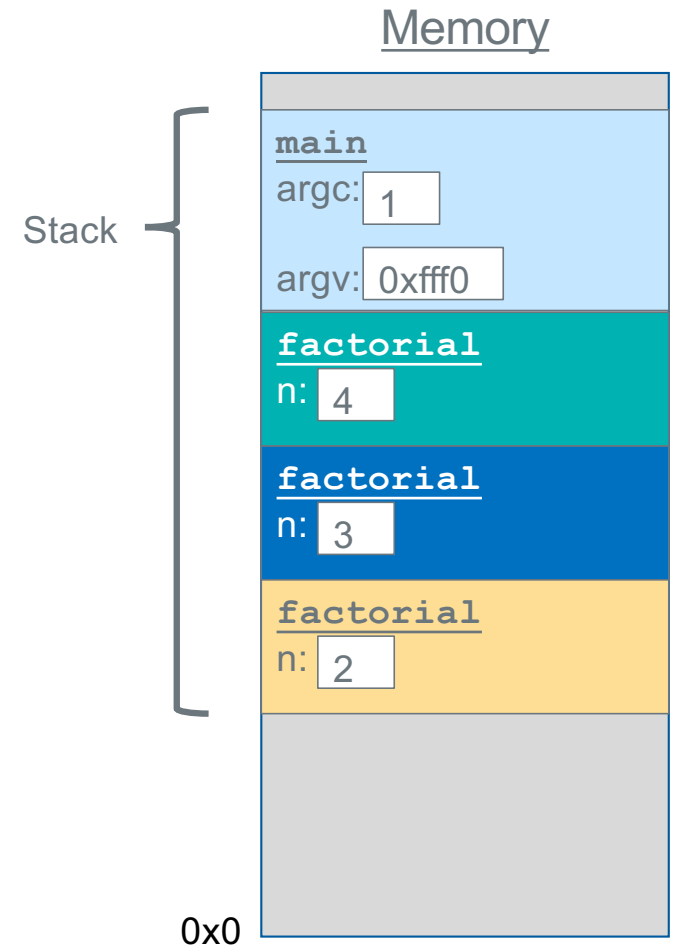
```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.

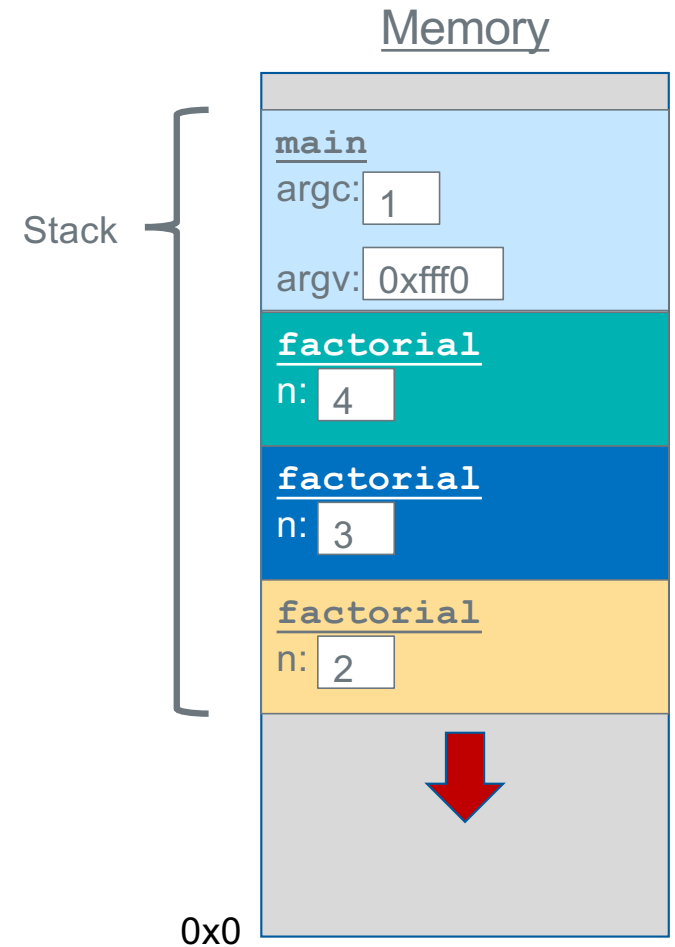
```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.


```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.

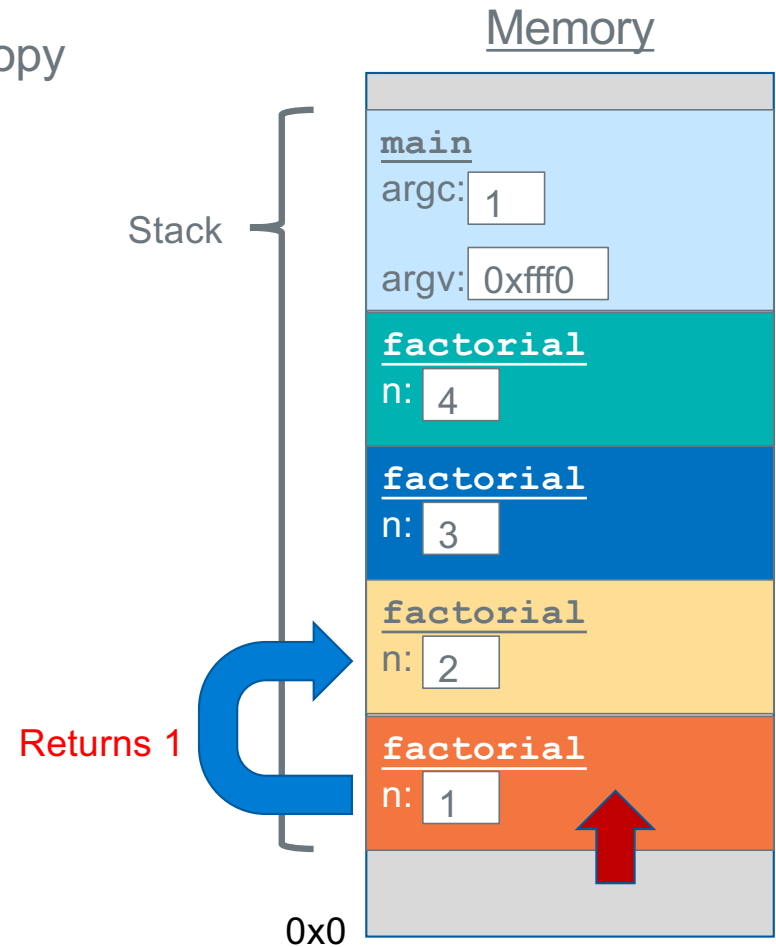
```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.

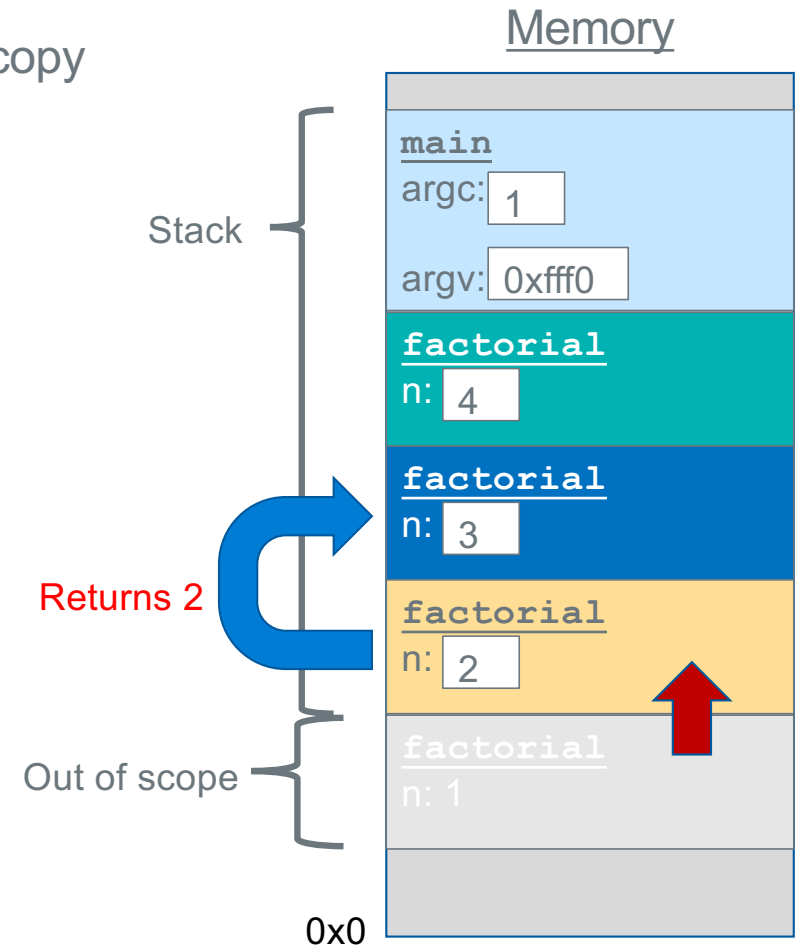
```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.

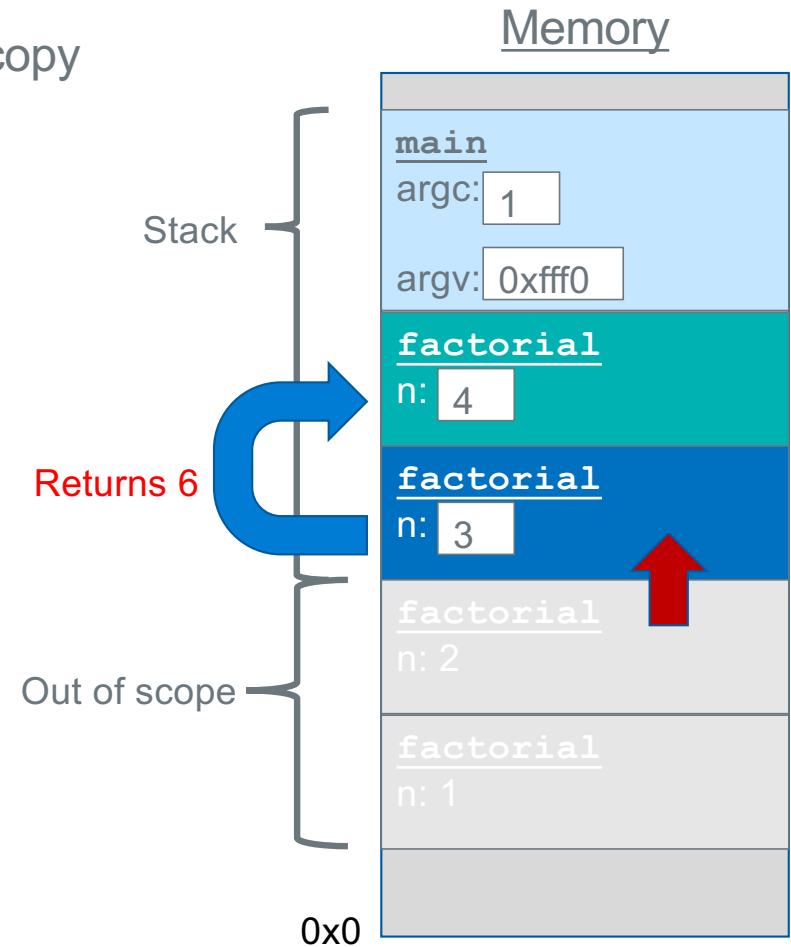
```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.

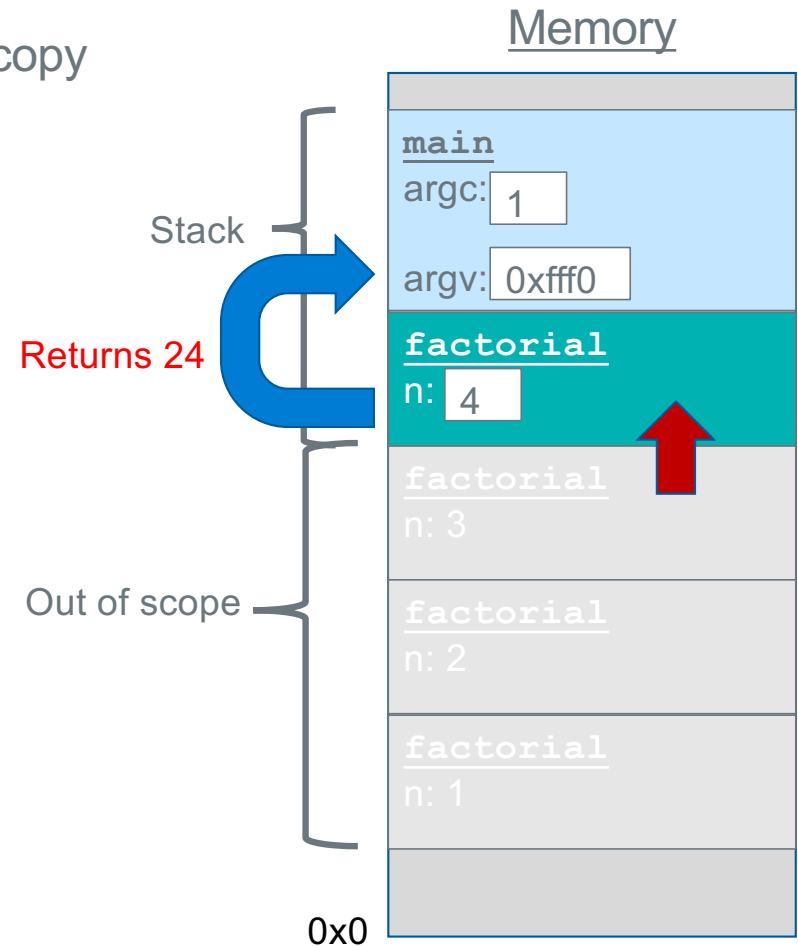
```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.

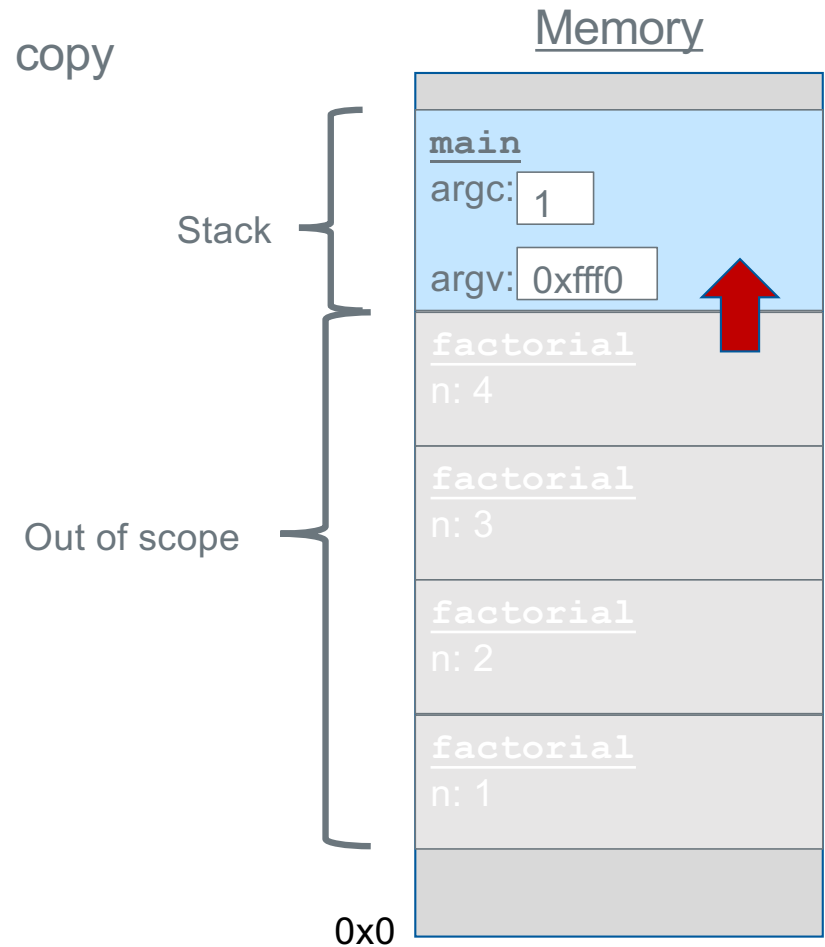
```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.

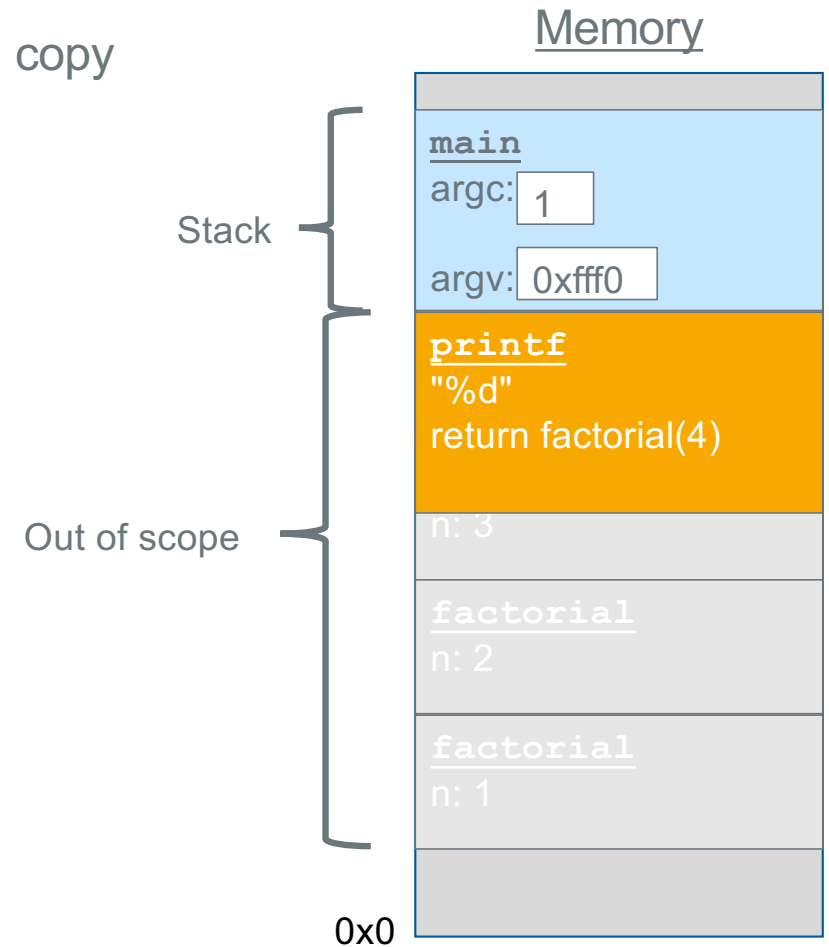
```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



The Stack

Each function **call** has its own *stack frame* for its own copy of variables.

```
int factorial(int n) {  
    if (n == 1) {  
        return 1;  
    } else {  
        return n * factorial(n - 1);  
    }  
}  
  
int main(int argc, char *argv[]) {  
    printf("%d", factorial(4));  
    return 0;  
}
```



Function Calls

Branch with Link (function call) instruction

`bl label`

`bl`

`imm24`

- Function call to the instruction with the address `label` (no local labels for functions)
 - `imm24` number of instructions from pc+8 (24-bits)
 - `label` any function label in the current file, any function label that is defined as `.global` in any file that it is linked to, any C function that is not static

Branch with Link Indirect (function call) instruction

`blx Rm`

`blx`

`Rm`

- Function call to the instruction whose address is stored in Rm (Rm is a function pointer)
- `bl` and `blx` both save the address of the instruction immediately following the `bl` or `blx` instruction in register `lr` (link register is also known as r14)
- The contents of the link register is the return address in the calling function

- Branch to the instruction with the label f1
- copies the address of the instruction AFTER the `bl` in `lr`

main:

•

`bl f1` → `f1:`

•

•

Function Call Return

Branch & exchange (function return) instruction

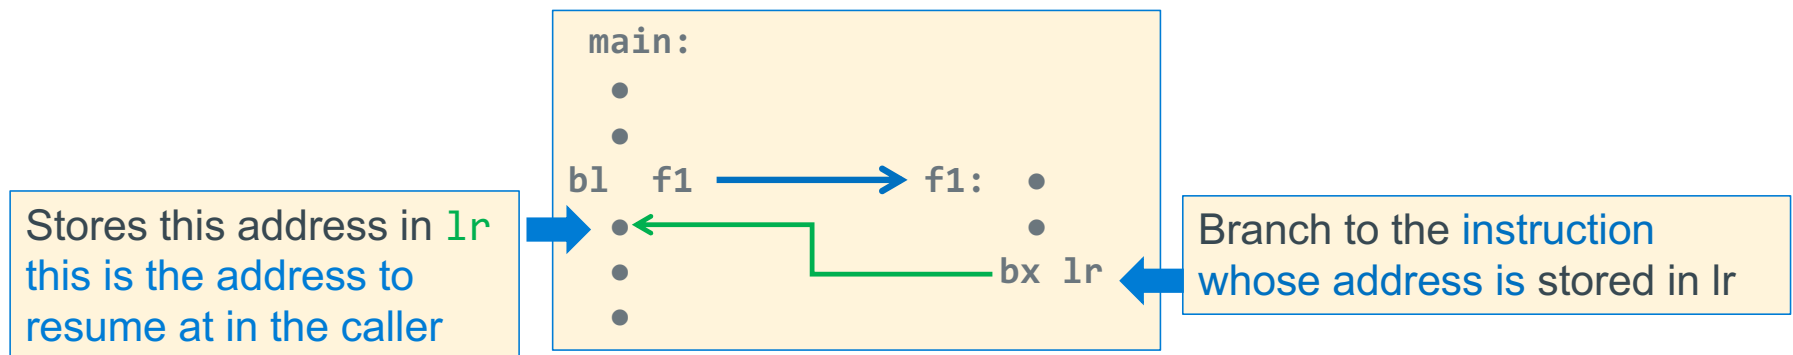
`bx lr`

`bx`

`Rn`

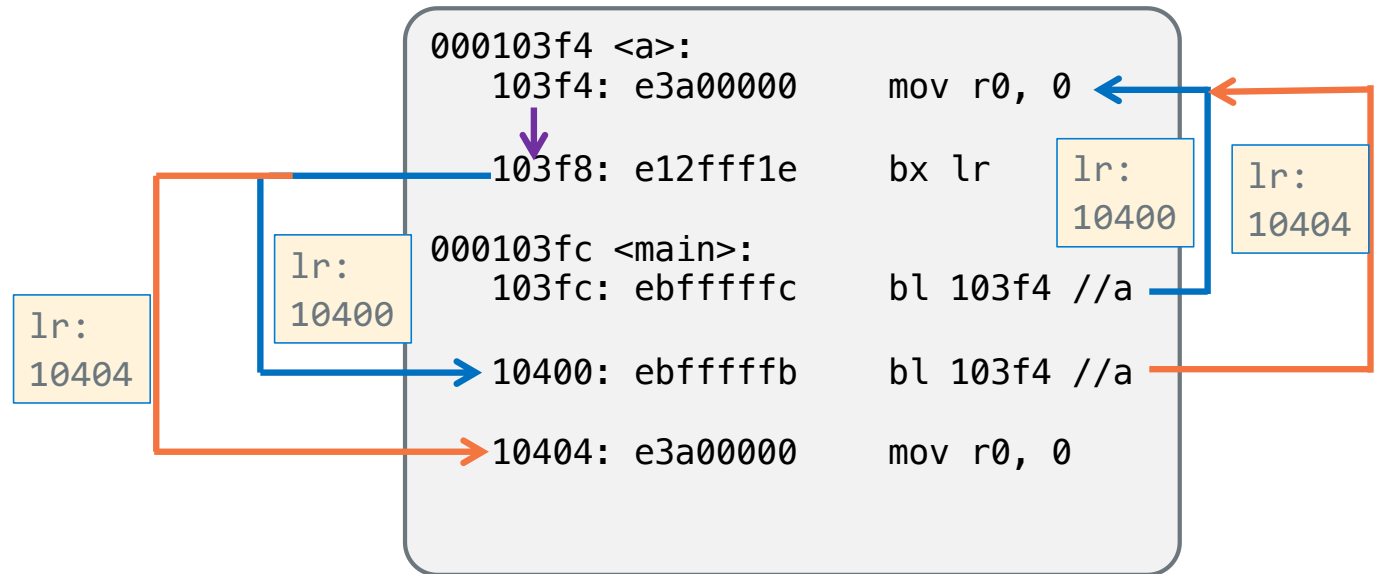
// we will always use `lr`

- Causes a branch to the instruction whose address is stored in register `<lr>`
 - It copies `lr` to the PC
- This is often used to implement a return from a function call (exactly like a C return) when the function is called using either `bl label`, or `blx Rm`



Understanding bl and bx - 1

```
int a(void)
{
    return 0;
}
int main(void)
{
    a();
    a();
    // not shown
}
```

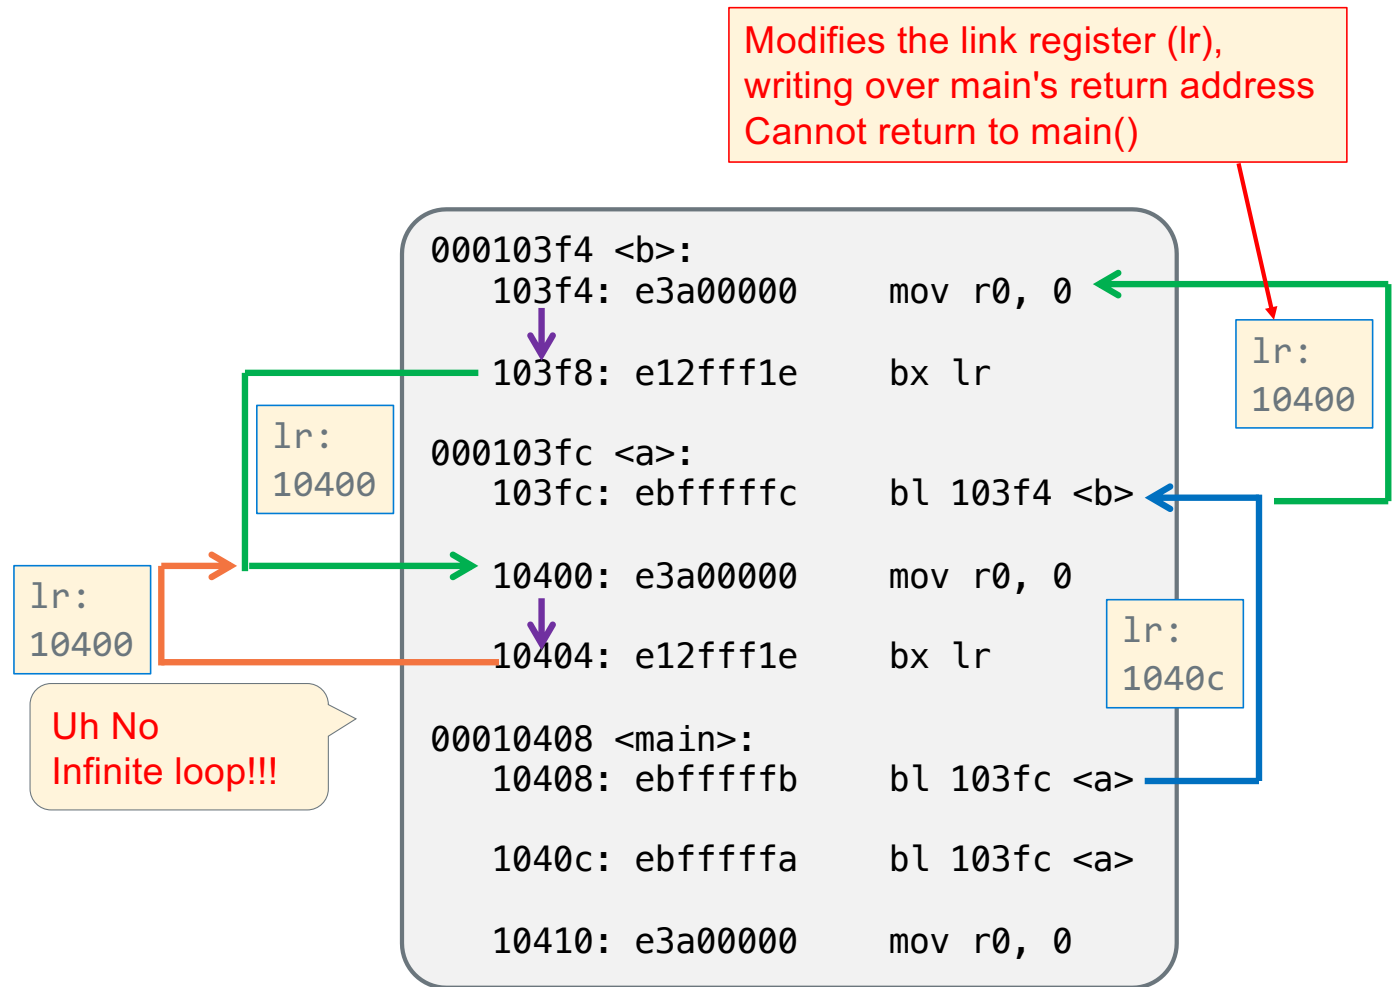


But there is a problem we must address here – next slide

Understanding bl and bx - 2

```
int b(void)
{
    return 0;
}
int a(void)
{
    b();
    return 0;
}
int main(void)
{
    a();
    a();
    // not shown
}
```

We need to preserve the lr!



Understanding bl and blx - 3

```
int a(void)
{
    return 0;
}

int (*func)() = a;

int main(void)
{
    (*func)();
    // not shown
}
```

But this has the same infinite loop problem when main() returns!

```
.data
func: .word a // func initialized with address of a()

.text
.global a
.type a, %function
.equ FP_OFF, 4

a:
    mov     r0, 0
    bx      lr
    .size a, (. - a)

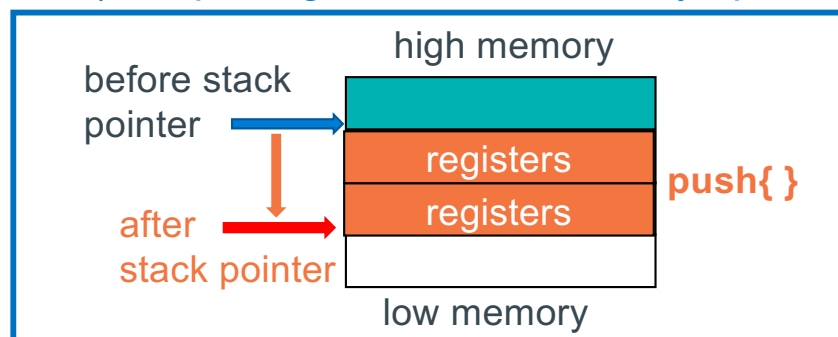
.global main
.type main, %function
.equ FP_OFF, 4

main:
    ldr     r4, =func // load address of func in r4
    ldr     r4, [r4]   // load contents of func in r4
    blx     r4         // we lose the lr for main!
    // not shown
    bx      lr         // infinite loop!
```


Preserving and Restoring Registers on the stack - 1

| Operation | Pseudo Instruction | Operation |
|----------------------------------|------------------------------|--|
| Push registers Function Entry | <code>push {reg list}</code> | $sp = sp - 4 \times \text{\#registers}$ Copy registers to <code>mem[sp]</code> |
| Pop registers Function Exit | <code>pop {reg list}</code> | Copy <code>mem[sp]</code> to registers, $sp = sp + 4 \times \text{\#registers}$ |

push (multiple register **str** to memory operation)



push (multiple register **ldr** from memory operation)

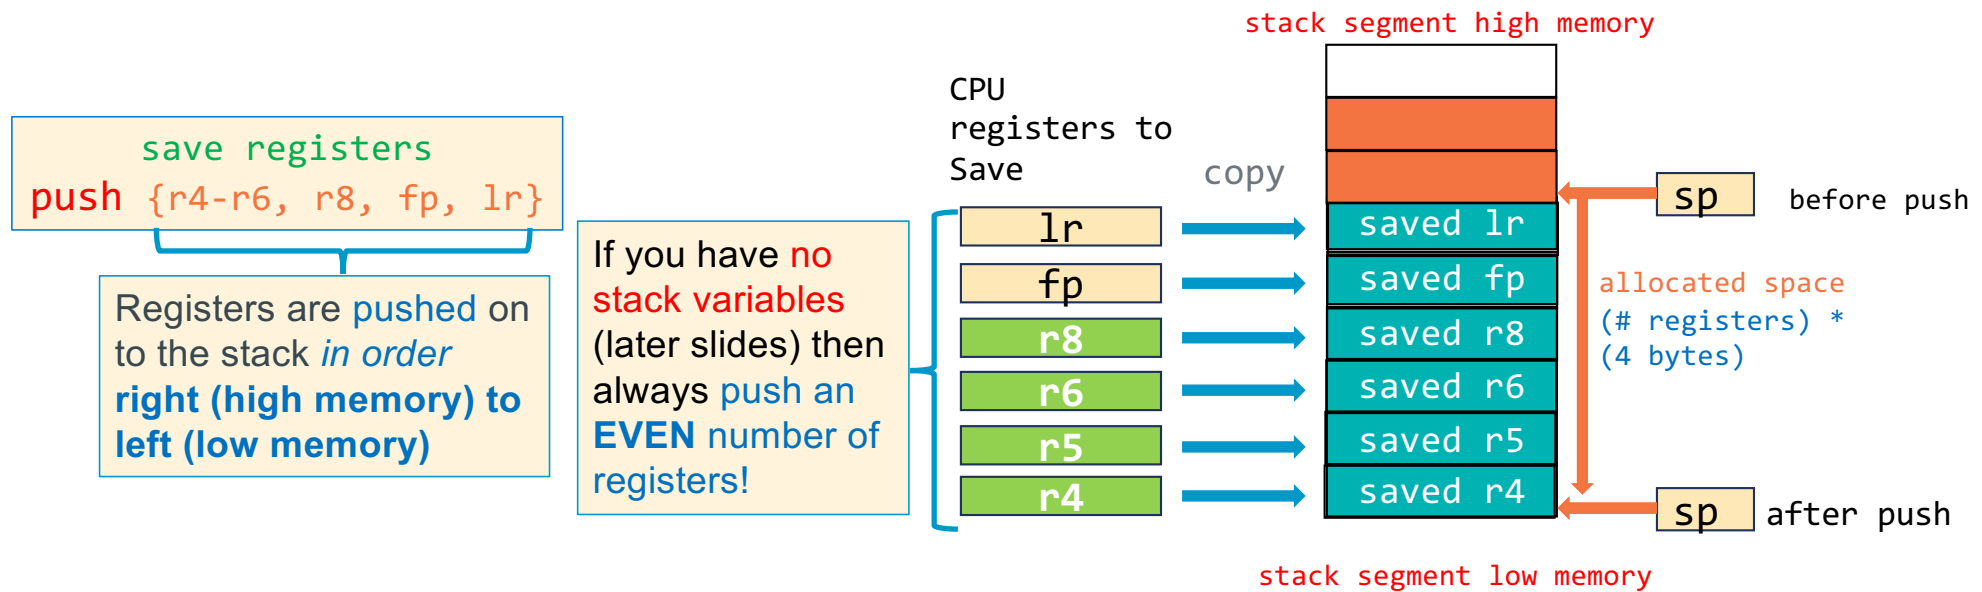


Preserving and Restoring Registers on the Stack - 2

| Operation | Pseudo Instruction | Operation |
|----------------------------------|------------------------------|--|
| Push registers Function Entry | <code>push {reg list}</code> | $sp = sp - 4 \times \text{\#registers}$ Copy registers to <code>mem[sp]</code> |
| Pop registers Function Exit | <code>pop {reg list}</code> | Copy <code>mem[sp]</code> to registers, $sp = sp + 4 \times \text{\#registers}$ |

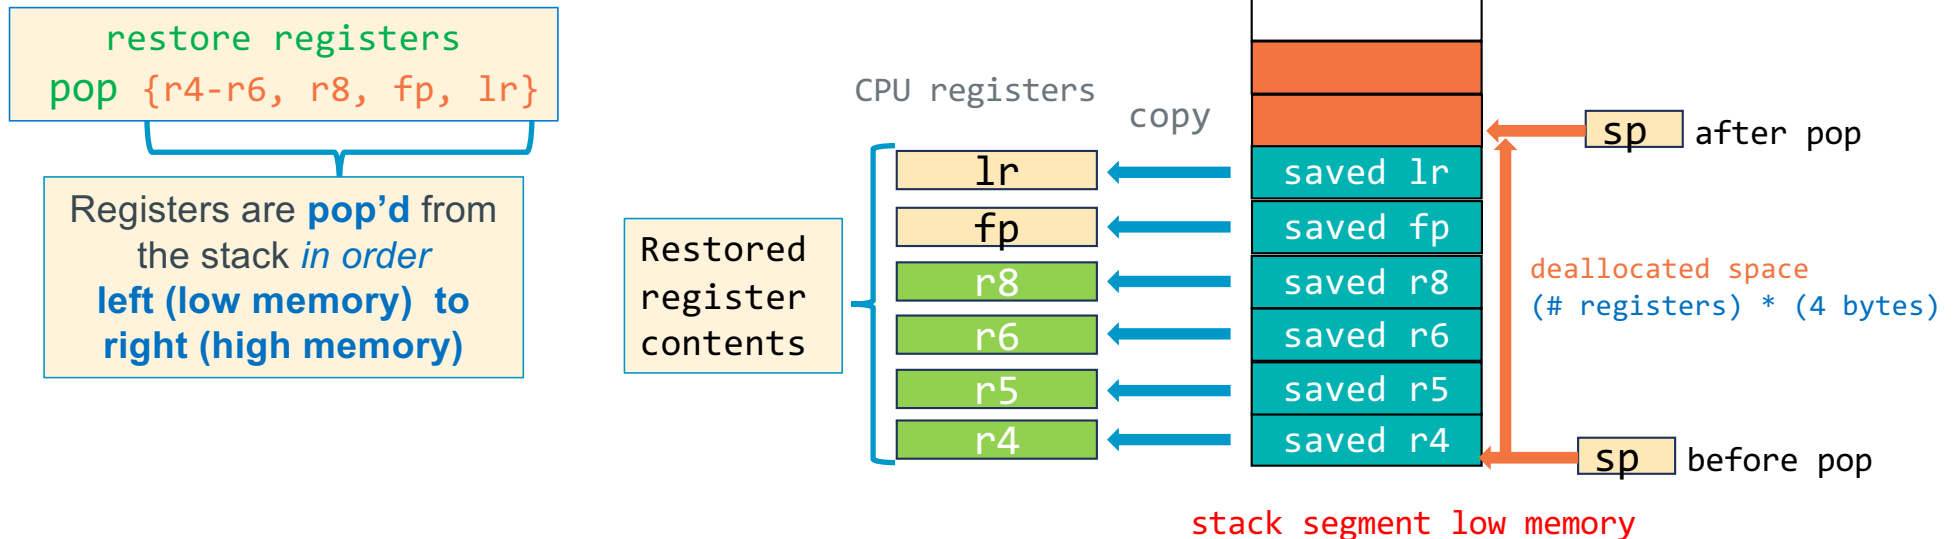
- `{reg list}` is a **list of registers in numerically increasing order, left to right**
`push {r4-r10, fp, lr}` // *fp is r11, lr is r14*
- Registers **cannot be**:
 1. duplicated in the list
 2. listed out of increasing numeric order (left to right)
- Register ranges can be specified `{r4, r5, r8-r10, fp, lr}`
- **Never!** push/pop `r12, r13, or r15`
 - the top two registers on the stack must always be `fp, lr` // ARM function spec – later slides

push: Multiple Register Save to the stack



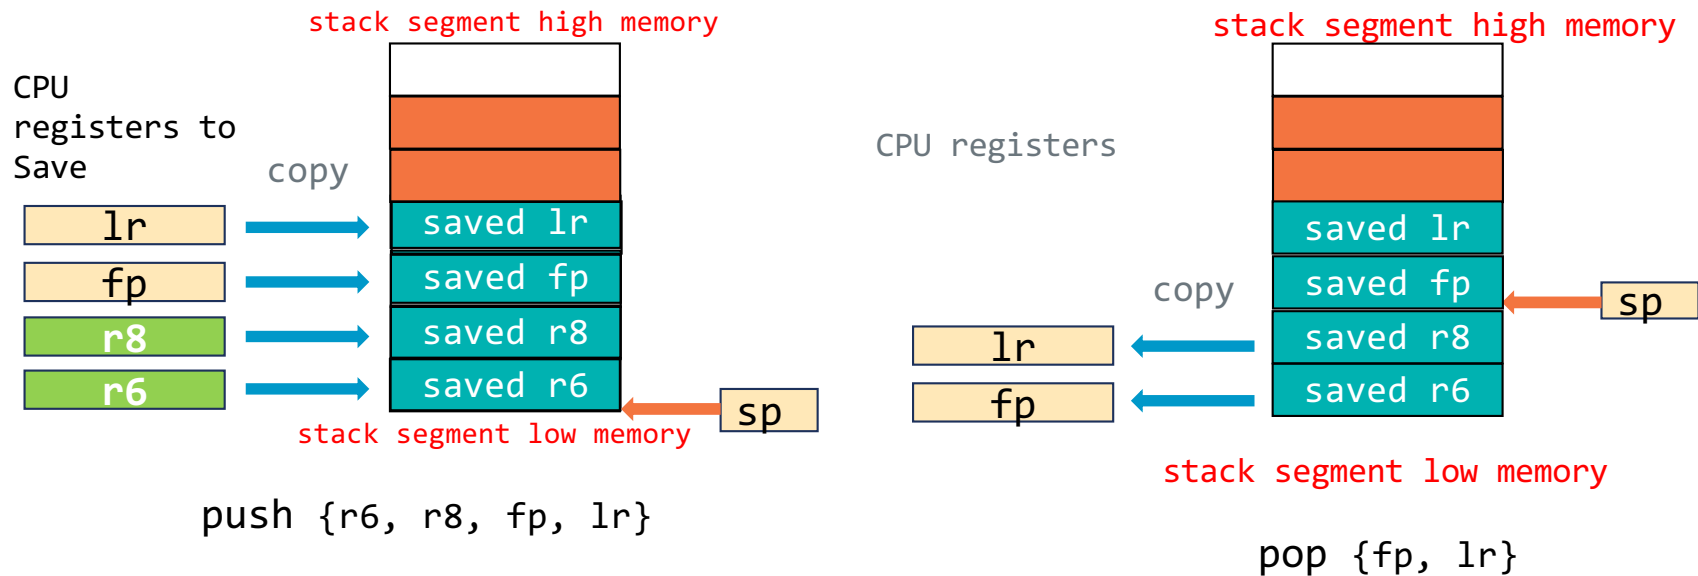
- **push** copies the contents of the **{reg list}** to stack segment memory
- **push** subtracts $(\text{\# of registers saved}) * (4 \text{ bytes})$ from the **sp** to **allocate** space on the stack
 - $sp = sp - (\text{\# registers_saved} * 4)$
- this must always be true: **$sp \% 8 == 0$**

pop: Multiple Register Restore from the stack



- **pop** copies the contents of stack segment memory to the **{reg list}**
- **pop adds:** $(\# \text{ of registers restored}) * (4 \text{ bytes})$ to **sp** to **deallocate** space on the stack
 - $sp = sp + (\# \text{ registers restored} * 4)$
- **Remember:** **{reg list}** must be the same in both the **push** and the corresponding **pop**

Consequences of inconsistent push and pop operands

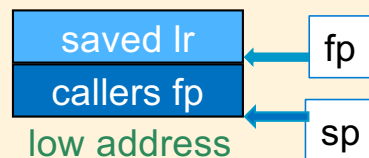


- lr gets an address on the stack, likely segmentation fault

Minimum Stack Frame (Arm Arch32 Procedure Call Standards)

- **Minimal frame: allocating at function entry: `push {fp, lr}`**

Minimum stack frame



- **sp** always points at top element in the stack (lowest byte address)
- **fp** always points at the bottom element in the stack
 - Bottom element is always the saved **lr** (contains the return address of caller)
 - A saved copy of **callers fp** is always the next element below the **lr**
 - **fp** will be used later when referencing stack variables
- **Minimal frame: deallocating at function exit: `pop {fp, lr}`**
- **On function entry:** **sp** must be 8-byte aligned (`sp % 8 == 0`)

Minimum Stack Frame (Arm Arch32 Procedure Call Standards)

- **Function entry (Function Prologue):**

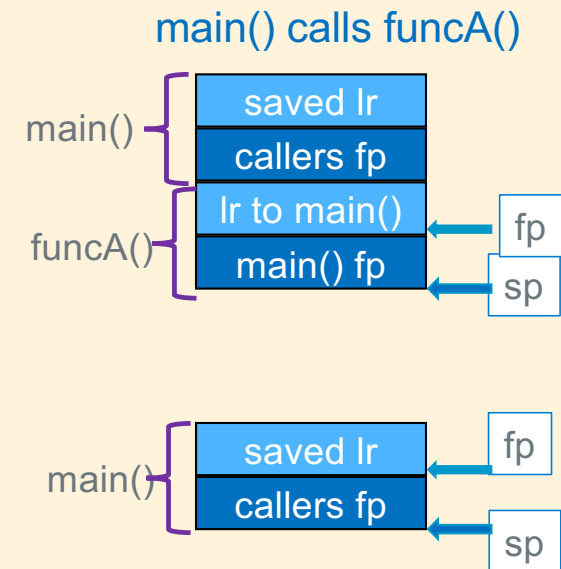
1. create (activate) frame
2. save preserved registers
3. allocate space for locals

allocate stack space
 $SP = SP - \text{"space"}$
grows "down"

- **Function return (Function Epilogue):**

1. deallocate space for locals
2. restores preserved registers
3. removes the frame

deallocate stack space
 $SP = SP + \text{"space"}$
shrinks "up"



How to set the FP – Minimum Activation Frame

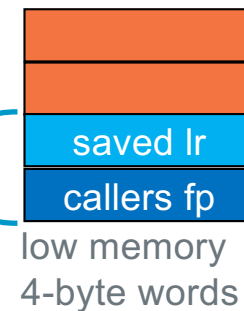
```
// other code
.equ    FP_OFF, 4
main:
  push   {fp, lr}
  add    fp, sp, FP_OFF
  .....
  sub    sp, fp, FP_OFF
  pop    {fp, lr}
  bx     lr
```

Function Prologue
always at top of function
push saves regs and
allocates space by
subtracting from sp and
sets fp with the add

Function Epilogue
always at bottom of
function pop restores
regs fp, lr
and deallocates space
by adding to sp

main()
Stack
Frame

after push {fp,lr}
add fp, sp, FP_OFF



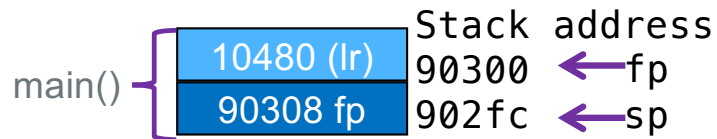
fp = sp + 4 bytes

IMPORTANT: FP_OFF has two uses:

1. Where to set fp after prologue push (remember sp position)
2. Restore sp (deallocates locals) right before epilogue pop

Using Minimal Stack Frames

```
int b(void)
{
    return 0;
}
int a(void)
{
    b();
    return 0;
}
int main(void)
{
    a();
    a();
}
```



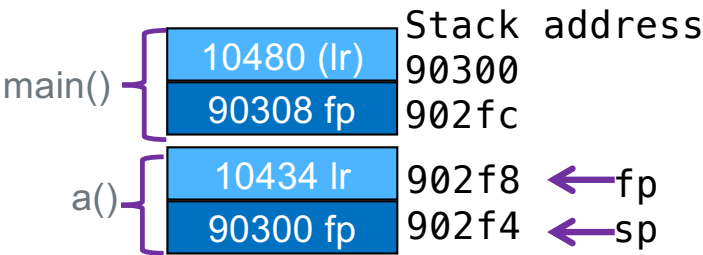
```
000103f4 <b>:
    103f4: e92d4800    push {fp, lr}
    103f8: e28db004    add fp, sp, 4
    103fc: e3a00000    mov r0, 0
    10400: e24bd004    sub sp, fp, 4
    10404: e8bd4800    pop {fp, lr}
    10408: e12fff1e    bx lr
```

```
0001040c <a>:
    1040c: e92d4800    push {fp, lr}
    10410: e28db004    add fp, sp, 4
    10414: ebfffff6    bl 103f4 <b>
    10418: e3a00000    mov r0, 0
    1041c: e24bd004    sub sp, fp, 4
    10420: e8bd4800    pop {fp, lr}
    10424: e12fff1e    bx lr
```

```
00010428 <main>:
    10428: e92d4800    push {fp, lr}
    1042c: e28db004    add fp, sp, 4
    10430: ebfffff5    bl 1040c <a>
    10434: ebfffff4    bl 1040c <a>
    // not shown
```

Using Minimal Stack Frames

```
int b(void)
{
    return 0;
}
int a(void)
{
    b();
    return 0;
}
int main(void)
{
    a();
    a();
}
```



```
000103f4 <b>:
    103f4: e92d4800    push {fp, lr}
    103f8: e28db004    add fp, sp, 4
    103fc: e3a00000    mov r0, 0
    10400: e24bd004    sub sp, fp, 4
    10404: e8bd4800    pop {fp, lr}
    10408: e12fff1e    bx lr


0001040c <a>:
    1040c: e92d4800    push {fp, lr}
    10410: e28db004    add fp, sp, 4
    10414: ebfffff6    bl 103f4 <b>
    10418: e3a00000    mov r0, 0
    1041c: e24bd004    sub sp, fp, 4
    10420: e8bd4800    pop {fp, lr}
    10424: e12fff1e    bx lr

00010428 <main>:
    10428: e92d4800    push {fp, lr}
    1042c: e28db004    add fp, sp, 4
    10430: ebfffff5    bl 1040c <a>
    10434: ebfffff4    bl 1040c <a>
// not shown
```

lr:
10434

Using Minimal Stack Frames

```
int b(void)
{
    return 0;
}
int a(void)
{
    b();
    return 0;
}
int main(void)
{
    a();
    a();
}
```



```
000103f4 <b>:
  103f4: e92d4800
  103f8: e28db004
  103fc: e3a00000
  10400: e24bd004
  10404: e8bd4800
  10408: e12fff1e
```

```
push {fp, lr}
add fp, sp, 4
mov r0, 0
sub sp, fp, 4
pop {fp, lr}
bx lr
```

lr:
10418

```
0001040c <a>:
  1040c: e92d4800
  10410: e28db004
  10414: ebfffff6
  10418: e3a00000
  1041c: e24bd004
  10420: e8bd4800
  10424: e12fff1e
```

```
push {fp, lr}
add fp, sp, 4
bl 103f4 <b>
mov r0, 0
sub sp, fp, 4
pop {fp, lr}
bx lr
```

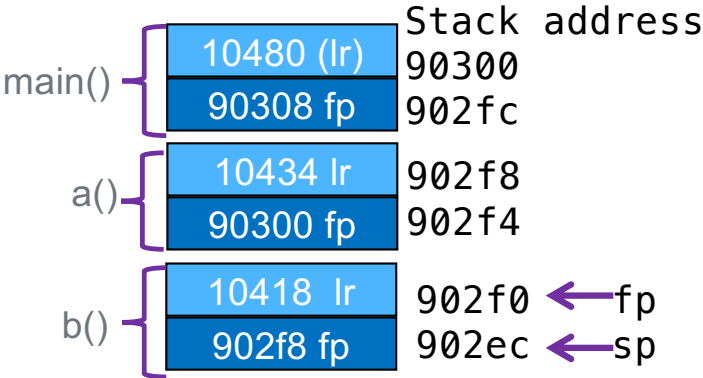
lr:
10434

```
00010428 <main>:
  10428: e92d4800
  1042c: e28db004
  10430: ebfffff5
  10434: ebfffff4
// not shown
```

```
push {fp, lr}
add fp, sp, 4
bl 1040c <a>
bl 1040c <a>
```

Using Minimal Stack Frames

```
int b(void)
{
    return 0;
}
int a(void)
{
    b();
    return 0;
}
int main(void)
{
    a();
    a();
}
```



lr:
10418

000103f4 :
103f4: e92d4800
103f8: e28db004
103fc: e3a00000
10400: e24bd004
10404: e8bd4800
10408: e12fff1e

push {fp, lr}
add fp, sp, 4
mov r0, 0
sub sp, fp, 4
pop {fp, lr}
bx lr

0001040c <a>:
1040c: e92d4800
10410: e28db004
10414: ebfffff6
10418: e3a00000
1041c: e24bd004
10420: e8bd4800
10424: e12fff1e

push {fp, lr}
add fp, sp, 4
bl 103f4
mov r0, 0
sub sp, fp, 4
pop {fp, lr}
bx lr

00010428 <main>:
10428: e92d4800
1042c: e28db004
10430: ebfffff5
10434: ebfffff4
// not shown

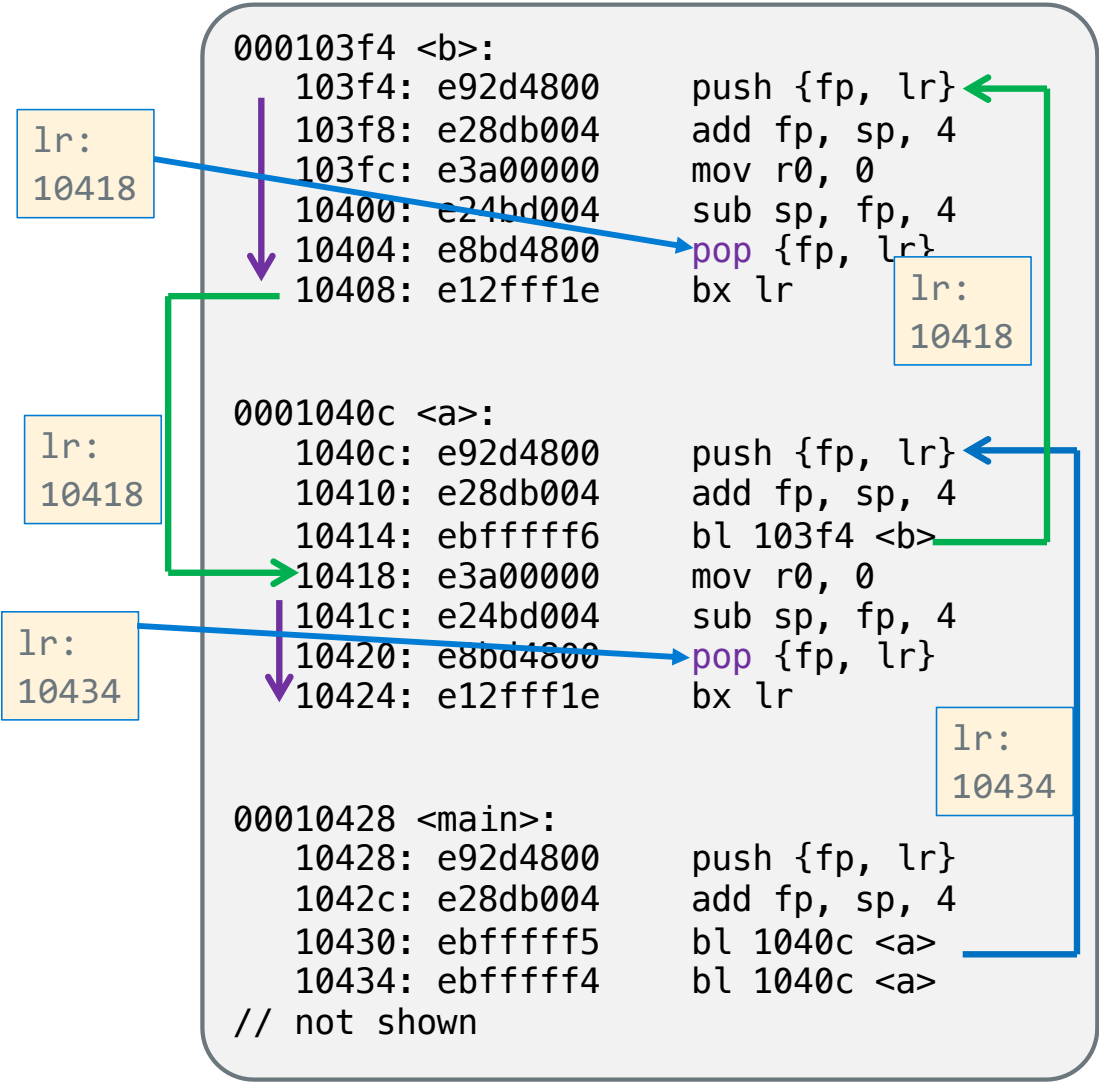

push {fp, lr}
add fp, sp, 4
bl 1040c <a>
bl 1040c <a>

lr:
10418

lr:
10434

Using Minimal Stack Frames

```
int b(void)
{
    return 0;
}
int a(void)
{
    b();
    return 0;
}
int main(void)
{
    a();
    a();
}
```



Using Minimal Stack Frames

```
int b(void)
{
    return 0;
}
int a(void)
{
    b();
    return 0;
}
int main(void)
{
    a();
    a();
}
```

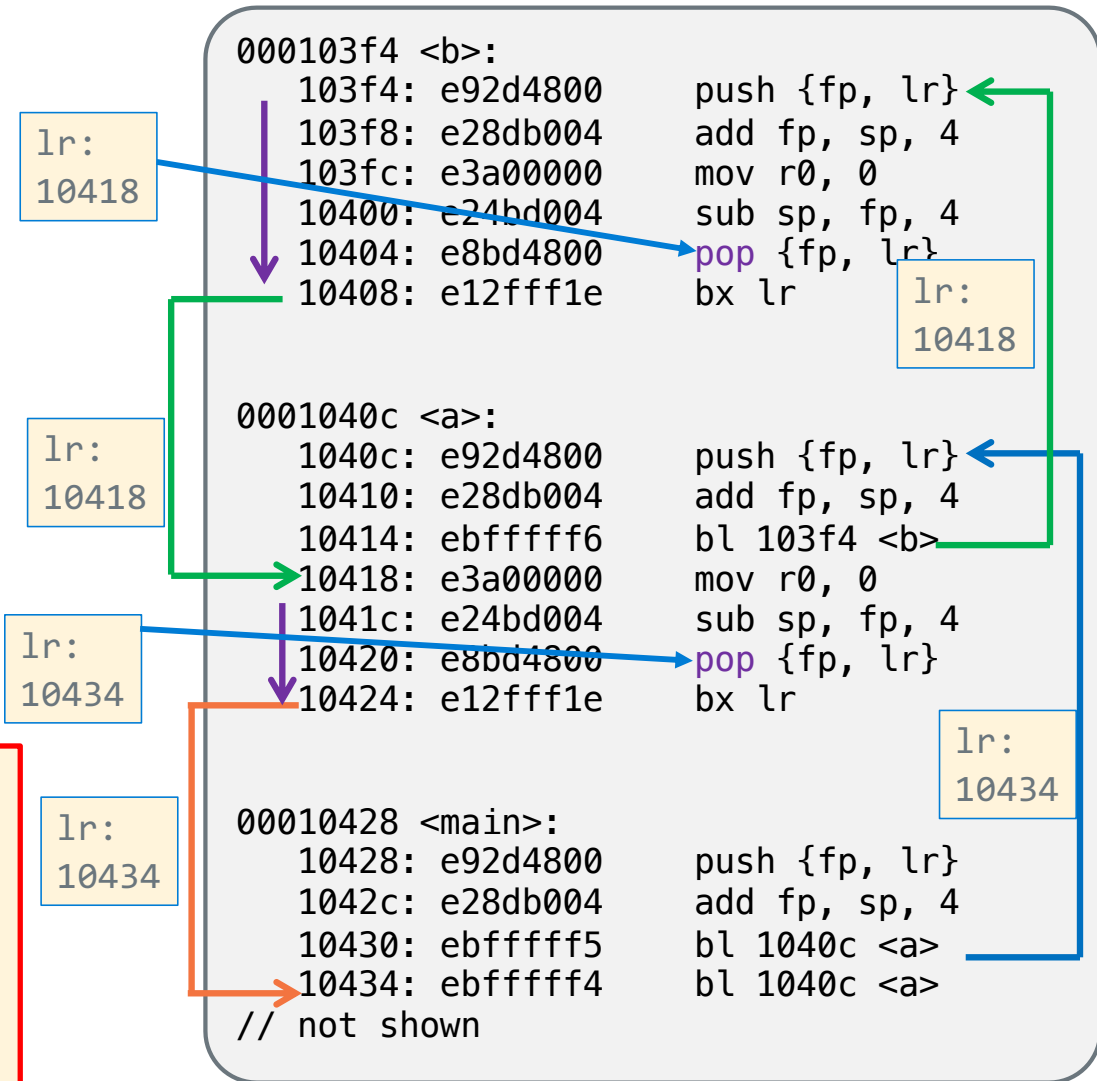
main() {

| | Stack address | |
|------------|---------------|------|
| 10480 (lr) | 90300 | ← fp |
| 90308 fp | 902fc | ← sp |

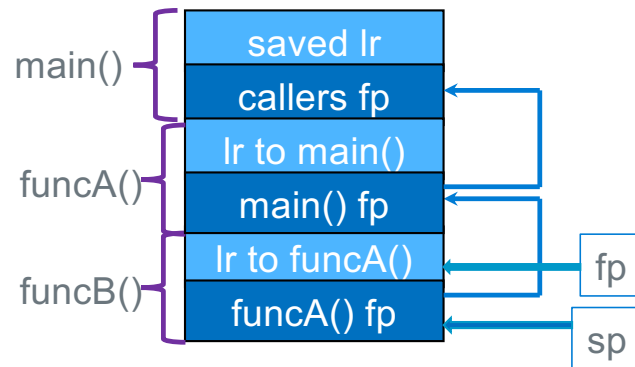
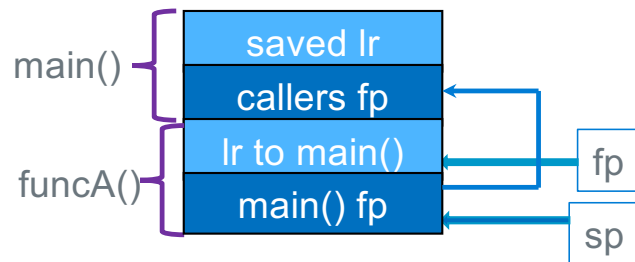
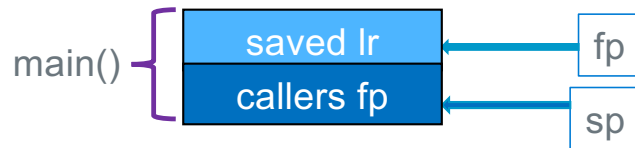
}

We are saving the lr on the stack on each function call and restoring it before returning.

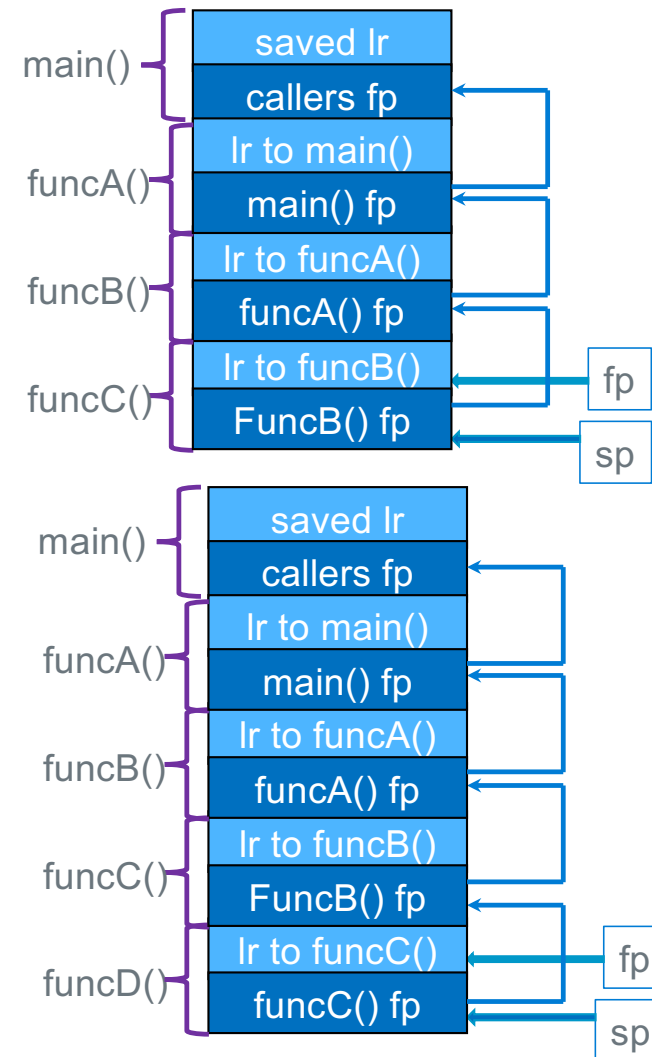
Result: NO infinite loop and we return to the correct instruction in the caller no matter how many functions we call.
Even recursion will work!



By following the saved fp, you can find each stack frame



How gdb finds stack frames



Registers: Requirements for Use

| <i>Register</i> | <i>Function Call Use</i> | <i>Function Body Use</i> | <i>Save before use Restore before return</i> |
|-----------------|--|--|--|
| r0 | arg1 and return value | scratch registers | No |
| r1-r3 | arg2 to arg4 | scratch registers | No |
| r4-r10 | preserved registers | contents preserved across function calls | Yes |
| r11/fp | stack frame pointer | Use to locate variables on the stack | Yes |
| r12/ip | may used by assembler with large text file | can be used as a scratch if really needed | No |
| r13/sp | stack pointer | stack space allocation | Yes |
| r14/lr | link register | contains return address for function calls | Yes |
| r15 | Do not use | Do not use | No |

- Any value you have in a **preserved register before a function call will still be there after the function returns**
- Contents are “preserved” across function calls

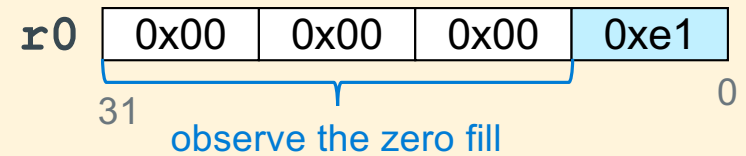
If the function wants to use a preserved register it must:

- Save** the **value contained in the register** at **function entry**
- Use the register in the body of the function
- Restore** the **original saved value** to the register at **function exit** (before returning to the caller)

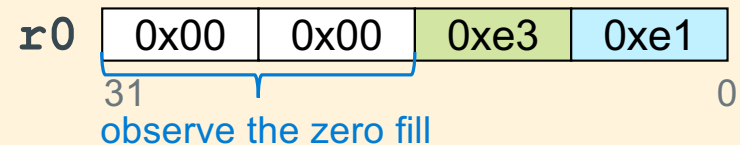
Argument and Return Value Requirements

- When passing or returning values from a function you must do the following:
 - Make sure that the values in the registers r0-r3 are in their **properly aligned position in the register based on data type**
 - Upper bytes in byte and halfword values in registers r0-r3 when passing arguments and returning values **are zero filled**

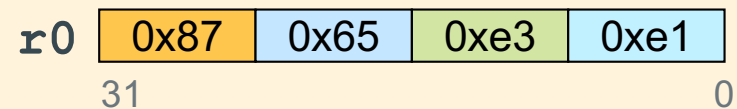
Single Byte (char)



Single Halfword (short)



Full Word (int or pointer)



Global Variable access

| var | global variable address into r0 (lside) | global variable contents into r0 (rside) | contents of r0 into global variable |
|--------|--|--|---|
| x | ldr r0, =x | ldr r0, =x ldr r0, [r0] | ldr r1, =x str r0, [r1] |
| *x | ldr r0, =x ldr r0, [r0] | ldr r0, =x ldr r0, [r0] ldr r0, [r0] | ldr r1, =x ldr r1, [r1] str r0, [r1] |
| **x | ldr r0, =x ldr r0, [r0] ldr r0, [r0] | ldr r0, =x ldr r0, [r0] ldr r0, [r0] ldr r0, [r0] | ldr r1, =x ldr r1, [r1] ldr r1, [r1] str r0, [r1] |
| stderr | ldr r0, =stderr | ldr r0, =stderr ldr r0, [r0] | <do not write unless you really know what you are doing> |
| .Lstr | ldr r0, =.Lstr | ldr r0, =.Lstr ldrb r0, [r0] | <read only> |

```
.bss // from libc
stderr:.space 4 // FILE *
```

```
.data
x: .data y //x = &y
```

```
.section .rodata
.Lstr: .string "HI\n"
```

stdin, stdout and stderr are global variables