

Version 2.00


UCSD CSE 30

Computer Organization and Systems Programming

Numbers Data and Memory

Keith Muller

Cray XMP-1982



Positive Number (unsigned) in 4 bits

- Real hardware has a fixed number of bits to store numbers (pi-cluster is 32 bits)
- There are only 2^n distinct values in n bits
- This limits the range of positive number to be 0 (unsigned min) to $2^n - 1$ (unsigned max)

Hex digit	0x0	0x1	0x2	0x3	0x4	0x5	0x6	0x7
Decimal value	0	1	2	3	4	5	6	7
Binary value	0b0000	0b0001	0b0010	0b0011	0b0100	0b0101	0b0110	0b0111


umin

Hex digit	0x8	0x9	0xa	0xb	0xc	0xd	0xe	0xf
Decimal value	8	9	10	11	12	13	14	15
Binary value	0b1000	0b1001	0b1010	0b1011	0b1100	0b1101	0b1110	0b1111

umax

Unsigned Integers (positive numbers) with Fixed # of Bits

- 4 bits is $2^4 = \text{ONLY } 16$ distinct values
- **Modular** (C operator: `%`) or **clock math**
 - Numbers start at 0 and “wrap around” after 15 and go back to 0
- Keep **adding** 1
 - wraps (**clockwise**)
 - 0000 \rightarrow 0001 ... \rightarrow 1111 \rightarrow 0000
- Keep **subtracting** 1
 - wraps (**counter-clockwise**)
 - 1111 \rightarrow 1110 ... \rightarrow 0000 \rightarrow 1111
- Addition and subtraction use normal “**carry**” and “**borrow**” rules, just operate in binary



Unsigned Binary Number: Addition in 4 bits

Be Aware in Binary

1 + 1 = 10

base 10: (1 + 1 = 2)

1 + 1 + 1 = 11

base10: (1 + 1 + 1 = 3)

Carry Bit


carries 0 1 1 1

+

0 0 0 1

0 1 1 1

sum 1 0 0 0 = 8




Unsigned Binary Number: Subtraction in 4 bits

Be Aware in Binary

1 - 1 = 0 base 10: (1 - 1 = 0)
10 - 1 = 1 base10: (2 - 1 = 1)

Borrows

$$\begin{array}{r}
 0 1 0 1 \\
 - 0 0 1 1 \\
 \hline
 \text{sum } 0 0 1 0 = 2
 \end{array}$$



Unsigned Binary Number: Addition in 4 bits – Overflow!

Be Aware in Binary

$$1 + 1 = 10$$

$$\text{base 10: } (1 + 1 = 2)$$

$$1 + 1 + 1 = 11$$

$$\text{base10: } (1 + 1 + 1 = 3)$$

Carry Bit

carries

1

1

1

1

0

1

0

10

+

0

1

1

1

7

sum

0


0

0

1

≠

17



Unsigned Binary Number: Subtraction in 4 bits – Overflow!


Be Aware in Binary

1 - 1 = 0 base 10: (1 - 1 = 0)
10 - 1 = 1 base 10: (2 - 1 = 1)


Borrows

$$\begin{array}{r}
 \text{11101} \\
 - 0111 \\
 \hline
 \text{sum } 1110 \neq -2
 \end{array}$$

The diagram shows a binary subtraction of 5 (101) from 7 (0111). The result is 1110, which is -2 in two's complement. The text "sum 1110 ≠ -2" indicates an overflow error.



Overflow: Going Past the Boundary Between umax and umin



Overflow with unsigned numbers:


Occurs when an arithmetic result (from addition or subtraction for example) is **more than min** or **max** limits

C (and Java) ignore overflow exceptions

- You end up with a bad value in your program and absolutely no warning or indication... **happy debugging!....**

Problem: How to Encode Both Positive and Negative Integers

- How do we represent the negative numbers within a fixed number of bits?
 - Allocate some bit patterns to negative and others to positive numbers (and zero)
- 2^n distinct bit patterns to encode positive and negative values
- Unsigned values:** $0 \dots 2^n - 1$ ← -1 comes from counting 0 as a "positive" number
- Signed values:** $-2^{n-1} \dots 2^{n-1} - 1$ (dividing the range in ~ half including 0)
- On a number line (below):** 8-bit integers – signed and unsigned (e.g., `char` in C)




Same "width" (same number of encodings), just shifted in value

Negative Integer Numbers: Sign + Magnitude Method



- Use the **M**ost **S**ignificant **B**it as a sign bit
 - 0 as the MSB represents positive numbers
 - 1 as the MSB represents negative numbers
- **Two** (oops) representations for **zero**: 0000, 1000
- Tricky Math (must handle sign bit independently)
 - Positive and Negatives “*increment*” (+1) in the **opposite directions**!



Signed Magnitude Examples (Sign bit is always MSB)

0 110
positive 6

1 011
negative 3

Examples (4 bits):

1 000 = -0	0 000 = 0
1 001 = -1	0 001 = 1
1 010 = -2	0 010 = 2
1 011 = -3	0 011 = 3
1 100 = -4	0 100 = 4
1 101 = -5	0 101 = 5
1 110 = -6	0 110 = 6
1 111 = -7	0 111 = 7


0 00000000
positive 0

1 0001100
negative 12


Two's Complement: The MSB Has a *Negative Weight*

$$2's\ Comp = -b_{n-1}2^{n-1} + b_{n-2}2^{n-2} + \dots + b_12^1 + b_02^0$$

b_{n-1} weight is (-2^{n-1}) , all other bits: have positive weights $(+2^i)$




- 4-bit ($w = 4$) weight = $-2^{4-1} = -2^3 = -8$
 - 1010_2 **unsigned**:
 $1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 0 \times 2^0 = 10$
 - 1010_2 **two's complement**:
 $-1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 0 \times 2^0 = -8 + 2 = -6$
 - 8 in **two's complement**:
 $1000_2 = -2^3 + 0 = -8$
 - 1 in **two's complement**:
 $1111_2 = -2^3 + (2^3 - 1) = -8 + 7 = -1$



2's Complement Signed Integer Method

- Positive numbers encoded same as unsigned numbers
- All **negative values** have a **one in the leftmost bit**
- All **positive values** have a **zero in the leftmost bit**
 - This implies that 0 is a positive value
- **Only one zero**
- **For n bits, Number range is $-(2^{n-1})$ to $+(2^{n-1} - 1)$**
 - Negative values “go 1 further” than the positive values
- Example: the range for 8 bits:
-128, -127, .. 0, .. 126, +127
- Example the range for 32 bits:
-2147483648 .. 0, .. +2147483647
- *Arithmetic is the same as with unsigned binary!*



Summary: Min, Max Values: Unsigned and Two's Complement

Two's Complement → Unsigned for n bits

- **Unsigned Value Range**


$$\begin{aligned} \text{UMin} &= 0b00\dots00 \\ &= 0 \end{aligned}$$

$$\begin{aligned} \text{UMax} &= 0b11\dots11 \\ &= 2^n - 1 \end{aligned}$$

- **Two's Complement Range**

$$\begin{aligned} \text{SMin} &= 0b10\dots00 \\ &= -2^{n-1} \end{aligned}$$

$$\begin{aligned} \text{SMax} &= 0b01\dots11 \\ &= 2^{n-1} - 1 \end{aligned}$$



Negation Of a Two's Complement Number (Method 1)

$$\begin{array}{r}
 7 = 0111 \\
 \downarrow \downarrow \downarrow \downarrow \\
 \text{invert} = 1000 \\
 \text{add } 1 \quad + \quad \underline{1} \\
 -7 \quad \quad 1001
 \end{array}$$

$$\begin{array}{r}
 -7 = 1001 \\
 \downarrow \downarrow \downarrow \downarrow \\
 \text{invert} = 0110 \\
 \text{add } 1 \quad + \quad \underline{1} \\
 7 \quad \quad 0111
 \end{array}$$

$$-x == \sim x + 1;$$

$$\begin{array}{r}
 7 = \quad \quad 0111 \\
 -7 = \quad + \quad \underline{1001} \\
 \text{(discard carry)} \quad 0000
 \end{array}$$






$$\begin{array}{r}
 1 = 0001 \\
 \downarrow \downarrow \downarrow \downarrow \\
 \text{invert} = 1110 \\
 \text{add } 1 \quad + \quad \underline{1} \\
 -1 \quad \quad 1111
 \end{array}$$

$$\begin{array}{r}
 -1 = 1111 \\
 \downarrow \downarrow \downarrow \downarrow \\
 \text{invert} = 0000 \\
 \text{add } 1 \quad + \quad \underline{1} \\
 1 \quad \quad 0001
 \end{array}$$

$$\begin{array}{r}
 -8 = 1000 \\
 \downarrow \downarrow \downarrow \downarrow \\
 \text{invert} = \underline{0111} \\
 \text{add } 1 \quad + \quad \underline{1} \\
 -8 \quad \quad 1000 \text{ oops!}
 \end{array}$$

Negation of a Two's Complement Number (Method 2)

1. **copy unchanged** right most bit containing a 1 and all the 0's to its right
2. Invert all the bits to the left of the right-most 1



Signed Decimal to Two's Complement Conversion

	dividend -102	Quotient	Remainder	Bit Position
➡	102/2	51	➡ 0	b0
➡	51/2	25	➡ 1	b1
➡	25/2	12	➡ 1	b2
➡	12/2	6	➡ 0	b3
➡	6/2	3	➡ 0	b4
➡	3/2	1	➡ 1	b5
➡	1/2	0	➡ 1	b6
➡	0/2	0	➡ 0	b7

102(base 10) = $b_7 b_6 b_5 b_4 b_3 b_2 b_1 b_0$ = 0b0110 0110

Get the two complement of 01100110 is 10011010



Two's Complement to Signed Decimal Conversion - Positive

What is $b_7 \ b_6 \ b_5 \ b_4 \ b_3 \ b_2 \ b_1 \ b_0$
 0 1 1 0 0 1 0 1_(base 2) in decimal (N)?

Signed Bit Bias	Bit	Bit Position	Bias
$-2^{W-1} = -2^{8-1} = -128$	x 0	b7	0 ←
Product Shift Left	Addend	Bit Position	Product
0	+ 1	b6	1
2 x 1 = 2	+ 1	b5	3
2 x 3 = 6	+ 0	b4	6
2 x 6 = 12	+ 0	b3	12
2 x 12 = 24	+ 1	b2	25
2 x 25 = 50	+ 0	b1	50
2 x 50 = 100	+ 1	b0	SUM = 101
		Bias + SUM:	0 + 101 = 101

Two's Complement to Signed Decimal Conversion - Negative

What is $\overset{b_7}{1} \overset{b_6}{1} \overset{b_5}{1} \overset{b_4}{0} \overset{b_3}{0} \overset{b_2}{1} \overset{b_1}{0} \overset{b_0}{1}_{(\text{base } 2)}$ in decimal (N)?

Signed Bit Bias	Bit	Bit Position	Bias
$-2^{W-1} = -2^{8-1} = -128$	x 1	b7	-128
Product Shift Left	Addend	Bit Position	Product
0	+ 1	b6	1
2 x 1 = 2	+ 1	b5	3
2 x 3 = 6	+ 0	b4	6
2 x 6 = 12	+ 0	b3	12
2 x 12 = 24	+ 1	b2	25
2 x 25 = 50	+ 0	b1	50
2 x 50 = 100	+ 1	b0	SUM = 101
		Bias + SUM:	-128 + 101 = -27

Two's Complement Addition and Subtraction

- **Addition:** just add the two number directly
- **Subtraction:** you can convert to addition: **difference = minuend + 2's complement (subtrahend)**

	Count	0	0	0	0	0	0	1	1
x	=	0	1	0	1	0	0	1	1
y	=	0	0	0	0	1	0	1	1
x + y	=	0	1	0	1	1	1	1	0

$$\begin{array}{r} x = 0 \ 1 \ 0 \ 1 \ 0 \ 0 \ 1 \ 1 \\ y = 0 \ 0 \ 0 \ 0 \ 1 \ 0 \ 1 \ 1 \\ \hline x-y \end{array}$$



2's complement first and then add


$$\begin{array}{r} \mathbf{x} = 0 \ 1 \ 0 \ 1 \ 0 \ 0 \ 1 \ 1 \\ + \ (-\mathbf{y}) = 1 \ 1 \ 1 \ 1 \ 0 \ 1 \ 0 \ 1 \\ \hline \mathbf{x} - \mathbf{y} = \mathbf{x} + (-\mathbf{y}) = 0 \ 1 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0 \end{array}$$

Two's Complement Positive Overflow

- **4-bit** Two's complement numbers (positive overflow)

Cout	0	1	0	0	
	0	1	0	1	5
+	0	1	1	0	6
<hr/>					
	1	0	1	1	-5
					!= 11

signed numbers: overflow occurs if
operands have same sign and result's sign is different




Overflow: Occurs when an arithmetic result is beyond the min or max limits

Two's Complement Negative Overflow

- **4-bit** Two's complement numbers (negative overflow)

	Cout	1	0	1	1	
		1	0	0	1	-7
	+	1	0	1	1	-5
		0	1	0	0	+4
						!= -12

carry bit is dropped from result



signed numbers: overflow occurs if
operands have same sign and result's sign is different

Overflow: Occurs when an arithmetic result is beyond the min or max limits

Summary: When Does Overflow Occur

Operand 1
+ Operand 2
Result

Operand 1 Sign	Operand 2 Sign	Is overflow Possible?
+	+	YES
-	-	YES
+	-	NO
-	+	NO

Sign Extension in C: Type casts

- Convert from smaller to larger integral data types
- C and Java automatically performs sign extension
- Example (on pi-cluster with 32-bit int and 16-bit short)

```
#include <stdlib.h>
#include <stdio.h>
int main(void)
{
    signed char c = -1;
    signed int i = c;
    unsigned char d = 1;
    unsigned int j = d;
    printf("c decimal = %hd\n", c);
    printf("c = 0x%hhx\n", c);
    printf("i decimal = %d\n", i);
    printf("i = 0x%x\n", i);
    printf("\nd decimal = %hd\n", d);
    printf("d = 0x%hhx\n", d);
    printf("j decimal = %d\n", j);
    printf("j = 0x%x\n", j);
    return EXIT_SUCCESS;
}
```

```
./a.out
c decimal = -1
c = 0xff
i decimal = -1
i = 0xffffffff

d decimal = 1
d = 0x1
j decimal = 1
j = 0x1
```

Sign Extension (how type promotion works)

- Sometimes you need to work with integers encoded with different number of bits

8 bits (char) → (16 bits) short → (32 bits) int


- Sign extension increases the number of bits:** n -bit wide signed integer X , **EXPANDS** to a **wider** n -bit + k -bit signed integer X' where **both have the same value**

Unsigned

- Just add leading zeroes to the left side



Two's Complement Signed:

- If **positive**, add leading **zeroes on the left**
 - Observe: Positive stay positive
- If **negative**, add **leading ones on the left**
 - Observe: Negative stays negative



Example: Two's Complement Sign or bit Extension - 1

- Adding 0's in front of a positive number does not change its value



Example: Two's Complement Sign or bit Extension -2

- Adding 1's if front of a negative number does not change its value

7 = 0111
invert = 1000
add 1 + 1
-7 = 1001

-7 = 1001

extend to
8 bits

11111001

$$\begin{aligned} 1001 &= -8 + 1 = -7 \\ \text{1111}1001 &= \\ (-128 + 64 + 32 + 16 + 8) + 1 &= \\ = -8 + 1 &= -7 \end{aligned}$$

7 = 00000111


invert = 11111000

add 1 + 1


-7 11111001

Sign Extension Under Different representations

How to sign extend this bit pattern?




unsigned




0x00d6

signed
magnitude



0x8056

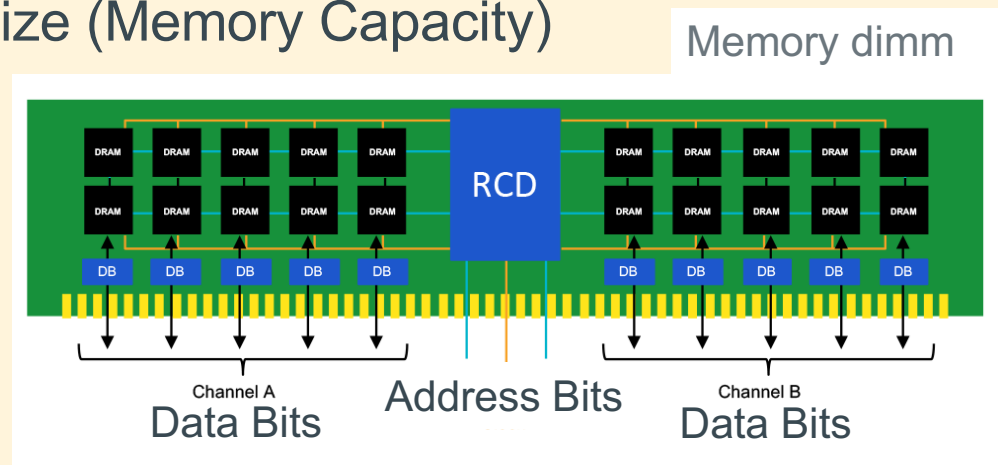
two's
complement



0xffd6

Memory Size

- Since memory addresses are implemented in hardware using binary
 - The **Size (number of byte sized cells)** of Memory is specified in **powers of 2**
- Memory size/capacity in **bytes** is specified by the “**Number of bits**” in an address
 - 32 bits of address = $2^{32} = 4,294,967,296$
 - Address Range is 0 to $2^{32} - 1$ (unsigned)
- Shorthand notation for address size (Memory Capacity)
 - KB = 2^{10} (K=1024) kilobyte
 - MB = 2^{20} megabyte
 - GB = 2^{30} gigabyte
 - TB = 2^{40} terabyte
 - PB = 2^{50} petabyte



Different Type of Numbers each have a Fixed # of Bits

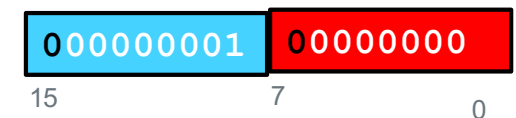
Spanning one or more contiguous bytes of memory

C Data Type	AArch-32 contiguous Bytes
char (arm unsigned)	1
short int	2
unsigned short int	2
int	4
unsigned int	4
long int	4
long long int	8
float	4
double	8
long double	8
pointer *	4

Byte 8-bit integer uses 1 byte



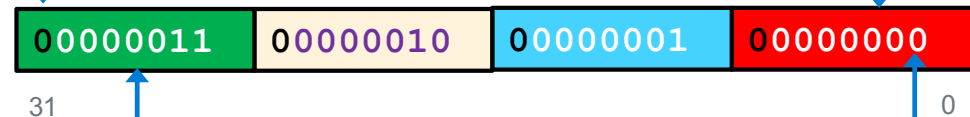
Half Word 16-bit integer uses 2 bytes



most significant bit (largest power of 2)

least significant byte

Word 32-bit integer uses 4 bytes

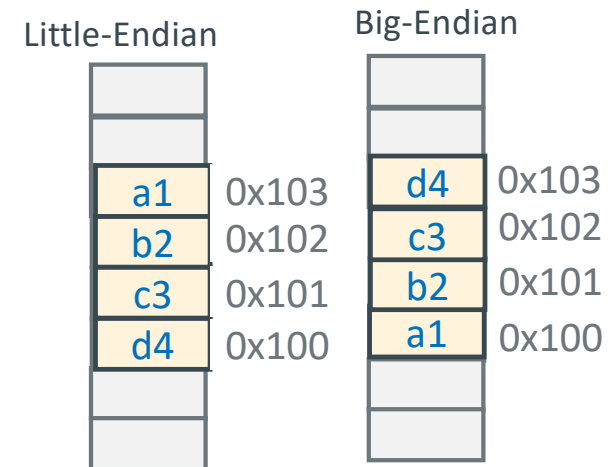
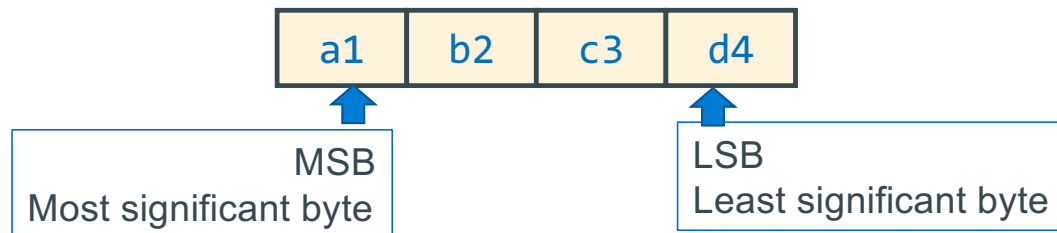


least significant bit (smallest power of 2)

most significant byte

Byte Ordering of Numbers In Memory: Endianness

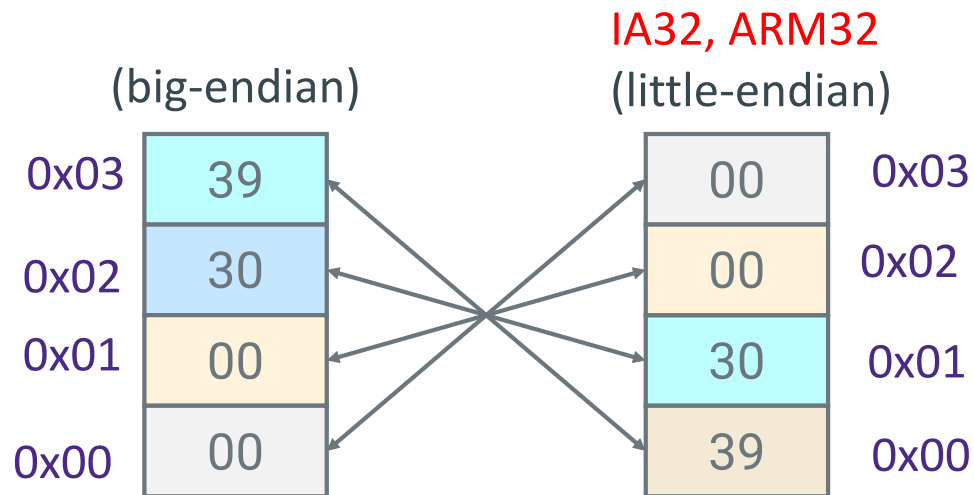
- Two different ways to place multi-byte integers in a **byte addressable** memory
- **Big-endian**: **Most** Significant Byte (“**big end**”) starts at the **lowest (starting)** address
- **Little-endian**: **Least** Significant Byte (“**little end**”) starts at the **lowest (starting)** address
- Example: 32-bit integer with 4-byte data



Byte Ordering Example

Decimal:	12345
Binary:	0011 0000 0011 1001
Hex:	3 0 3 9


```
int x = 12345;  
// or x = 0x00003039; // show all 32 bits
```



Byte Addressable Memory Shown as 32-bit words

1 byte Memory Content

One byte per row




Contents of Memory
One 32-bit (4 byte) word per row

MSByte · · · LSByte				Word Memory Address
				0x12345694
				0x12345690
				0x1234568C
				0x12345688
0x07	0x06	0x05	0x04	0x12345684
0x03	0x02	0x01	0x00	0x12345680
0x12345683	0x12345682	0x12345681	0x12345680	

Byte address

Observation
32-bit aligned addresses
rightmost 2 bits of the address are always 0

Scientific Notation Binary



- Computer hardware that supports this is called **floating point hardware** due to the “floating” of the binary point
- Declare such variable in C as `float` (or `double`)

Floating Point Representation

- Analogous to scientific notation
- In Decimal:
 - Not 12000000, but 1.2×10^7 In C: 1.2e7
 - Not 0.0000012, but 1.2×10^{-6} In C: 1.2e-6
- In Binary:
 - Not 11000.000, but 1.1×2^4
 - Not 0.000101, but 1.01×2^{-4}

Normalized Scientific Notation

- Convert from **scientific notation** to fixed **binary point**
- Perform the multiplication by shifting the decimal until the exponent disappears

Binary	Decimal
2^{-1}	0.5
2^{-2}	0.25
2^{-3}	0.125
2^{-4}	0.0625

- Example: $1.011_2 \times 2^4 = 10110_2 = 22_{10}$
- Example: $1.011_2 \times 2^{-2} = 0.01011_2 = 0.34375_{10}$
- Convert from **binary point** to **normalized scientific notation**
 - Distribute out exponents until binary point is to the right of a single digit
 - Example: $1101.001_2 = 1.101001_2 \times 2^3$

Encoding Fractions Observations

In Base 2:

$$10.1 \times 2^5 = 1.01 \times 2^6$$

$$1011.1 \times 2^5 = 1.0111 \times 2^8$$

$$0.110 \times 2^5 = 1.10 \times 2^4$$

Normalizing with base 2 :

adjust so there *always* a 1 to the **left of the decimal point!**

this 1 is **called the hidden bit** as we do not have use a bit to store it since it is there in every normalized mantissa

- Adjust x to always be in the format **1.XXXXXXXXXX...** (**fraction is normalized**)
- Fraction portion ONLY **encodes** what is *to the right* of the decimal point
- “Hidden bit” allows number to have **One additional digit for increased precision**

Fraction encoding is **1.[FRACTION BINARY DIGITS]**

Floating Point Numbers: Implementation Approach

- Supports a wide range of numbers
- Flexible “floating” decimal point
- Represent scientific notation numbers like 1.202×10^6

$$(-1)^S M 2^E$$



- **Sign bit** (a single bit): 0 positive, 1 negative
- **Exponent:** encoding of E above (it is NOT E directly represented in binary)
- **Fraction:** encoding of M above (it is NOT M directly represented in binary)

Excess Bias Encoding (As used in floating point numbers)


- Given a number in E bits, to divide the range in about $1/2$ the following is used:

$$\text{excess N bias} = (2^{E-1} - 1) \quad (\text{this is just one of many bias formulas})$$

- With this excess N Bias approach:** actual numbers range from most negative to most positive is: **-(bias) to bias+1**
- So, for a number that is limited to 4 bits (0 to 15 unsigned)**
 - Then excess N bias = $2^{4-1} - 1 = 2^3 - 1 =$ a bias of +7

actual	-7	-6	-5	-4	-3	-2	-1	0	1	2	3	4	5	6	7	8
bias	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7
bias encoded	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15


Floating Point Number in a Byte (Not A Real Format)



- **Mantissa encoding:** = 1.[xxxx] encoded as an unsigned value
- **Exponent encoding:** 3 bits encoded as an unsigned value using bias encoding
 - Bias encoding = $(2^{E-1} - 1)$
 - 3 bits for the bias we have $2^{3-1} - 1 = 2^2 - 1 =$ a bias of 3
 - **With a Bias of 3:** positive and negative numbers range: small to large is: 2^{-3} to 2^4

Actual	-3	-2	-1	0	1	2	3	4
Bias	+ 3	+ 3	+ 3	+ 3	+ 3	+ 3	+ 3	+3
Biased	0	1	2	3	4	5	6	7

Scientific Notation Decimal



- *Scientific Normalized form:*


exactly one digit (non-zero) to left of decimal point

- Alternatives to representing 1/1,000,000,000

- **Normalized:** 1.0×10^{-9}

- Not normalized: 0.1×10^{-8} , 10.0×10^{-10}

Floating Point Number in a Byte (Not A Real Format)



- **Mantissa encoding:** = 1.[xxxx] encoded as an unsigned value
- **Exponent encoding:** 3 bits encoded as an unsigned value using bias encoding
 - Bias encoding = $(2^{E-1} - 1)$
 - 3 bits for the bias we have $2^{3-1} - 1 = 2^2 - 1 =$ a bias of 3
 - **With a Bias of 3:** positive and negative numbers range: small to large is: 2^{-3} to 2^4

Actual	-3	-2	-1	0	1	2	3	4
Bias	+ 3	+ 3	+ 3	+ 3	+ 3	+ 3	+ 3	+3
Biased	0	1	2	3	4	5	6	7

Floating Point Number (8-bits) Number Range: 2^{-3} to 2^4



0.0 Special case in this simple model
we do not put back the “hidden bit”



Smallest Non-zero Positive
 $0.00\textcolor{blue}{1}0001 = \textcolor{blue}{1}/8 + 1/128 = 0.1328125$ base 10




Largest Positive/Negative
 $\textcolor{blue}{1}.\textcolor{blue}{1}111 \times 2^4 = \textcolor{blue}{1}1111 = 31$ base 10



Smallest (closest to zero) Number
 $\textcolor{blue}{1}.0000 \times 2^{-3} = 0.00\textcolor{blue}{1}000 = \textcolor{blue}{1}/8 = -0.125$ base 10

Note: Orange is hidden bit added back

Decimal to Float



Step 1: convert from base 10 to binary (absolute value)

$$-0.375 (\text{decimal}) = 0000.0110_2$$

Step 2: Find out how many places to shift to get the number into the normalized 1.xxxx mantissa format

$$0000.0110_2 = 1.1000 \times (2^{-2})_{\text{base } 10}$$

$$\text{exponent: } -2_{10} + \text{bias of } 3_{10} = 1_{10} = 0b001 \text{ for the exponent (after adding the bias)}$$

Step 3: Use as many digits that fit to the right of the decimal point in the fractional .xxxx part

$$1.1000$$


Step 4: Sign bit

positive sign bit is 0

negative sign bit is 1

s	exponent	fraction
1	0b001	0b1000
0x9		0x8

$$= 0x98$$



Step 1: Break into binary fields

0x45 =

Step 2: Extract the unbiased exponent

	0x4	0x5
s	exponent	fraction
0	0b100	0b0101

0b100 = 4_{base 10} - bias of 3₁₀ = 1₁₀ for the exponent (bias removed)

Step 3: Express the mantissa (restore the hidden bit)

1.0101

Step 4: Apply the unbiased exponent

1.0101_{base 2} × (2¹)_{base 10} = 10.101

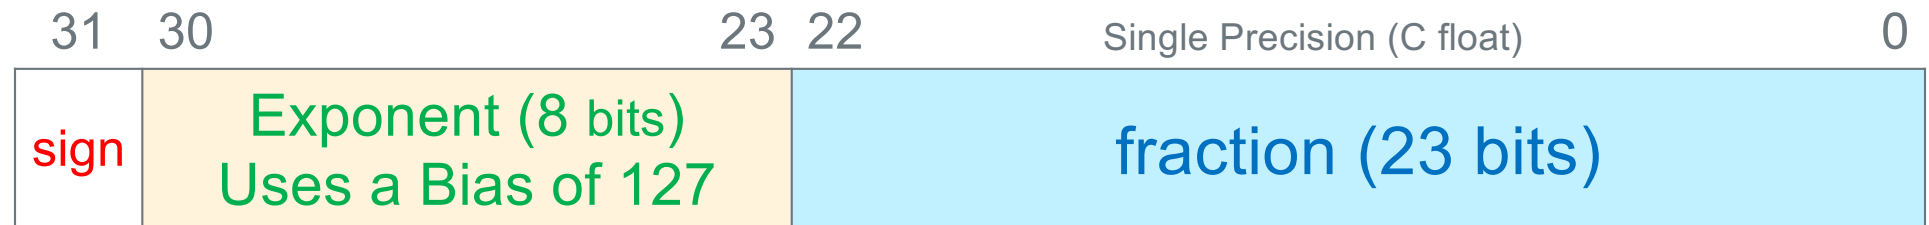
Step 5: Convert to decimal

10.101 = 2.625_{base 10}

Step 6: Apply the Sign


+ 2.625_{base 10}

IEEE “754” Floating Point Double and Single Precision



Bias is $(2^{8-1} - 1) = 127$

single precision floating point number = $(-1)^s \times 2^{E-127} \times 1.\text{fraction}$



bias is $(2^{11-1} - 1) = 1023$

double precision floating point number = $(-1)^s \times 2^{E-1023} \times 1.\text{fraction}$

Extra slides


Another Way to Look at 2's Complement Encoding

- A 2's complement value can be thought of as using a slightly different **bias encoding** for negative numbers only (more negative values): -2^{w-1}
- The **leftmost bit** is then interpreted as a **decision to apply the bias** (if **1**) or not (if **0**)
 - **1** apply the bias
 - **0** do not apply the bias
- For example, for a 4-bit number ($w = 4$), the negative number bias weight would be $= -2^{4-1} = -2^3 = -8$

2's	1000	1001	1010	1011	1100	1101	1110	1111	0000	0001	0010	0011	0100	0101	0110	0111
3 bit	000	001	010	011	100	101	110	111	000	001	010	011	100	101	110	111
decimal	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7
+Bias	-8	-8	-8	-8	-8	-8	-8	-8	0	0	0	0	0	0	0	0
Actual	-8	-7	-6	-5	-4	-3	-2	-1	0	1	2	3	4	5	6	7

Observe: adding +1 makes the number more positive for both negative and positive numbers

Decimal to IEEE Single Precision Float



Step 1: convert from base 10 to binary (absolute value)

$-13.375(\text{decimal}) = 1101.0110$

Step 2: Find out how many places to shift to get the number into the normalized 1.xxxx mantissa format

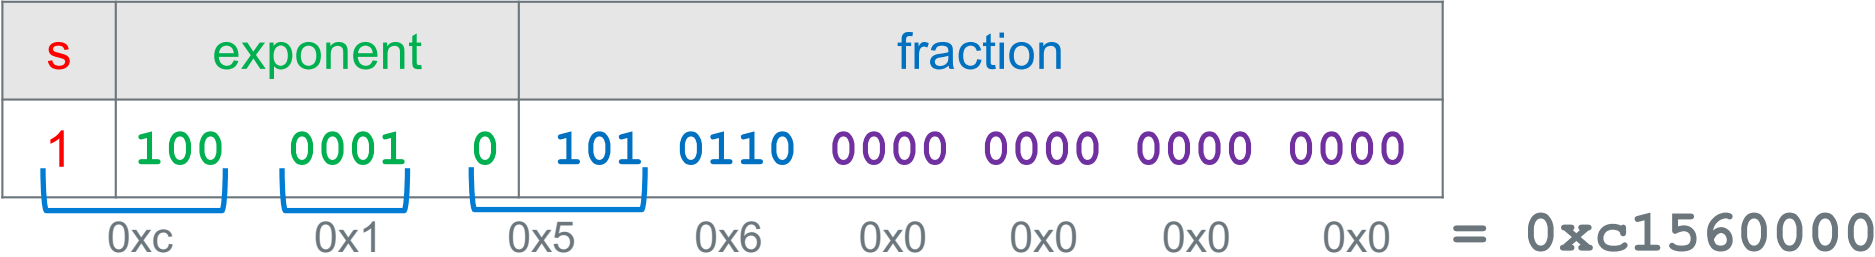
$1101.0110 = 1.1010110 \times (2^3)_{\text{base } 10}$

$3 + \text{bias of } 127 = 130 \text{ for the exponent} = 0b1000\ 0010$


Step 3: Use as many digits that fit to the right of the decimal point in the fractional .xxxx part (0 pad)

$1.1010110\ 0000\ 0000\ 0000\ 0000$

Step 4: If the sign is positive sign bit is 0, otherwise it is 1




IEEE Single Precision Float to Decimal



Step 1: Break into binary fields and **expand** as needed

$0xc0b00000 =$

Step 2: Find the exponent



$0b10000001 = 129_{\text{base } 10} - \text{bias of } 127_{10} = 2_{10}$ exponent with **bias added**

Step 3: Express the mantissa (restore the hidden **bit**)

1.0110

Step 4: Apply the exponent

$1.0110 \times (2^2)_{\text{base } 10} = 101.10$

Step 5: Convert to decimal

$101.10 = 5.5$

Step 6: Apply the Sign

-5.5